



MSR for Vulnerability Prediction

Mining Vulnerability-Contributing Commits

Emanuele Iannone

SeSa Lab @ University of Salerno, Italy

eiannone@unisa.it

Who are
you?









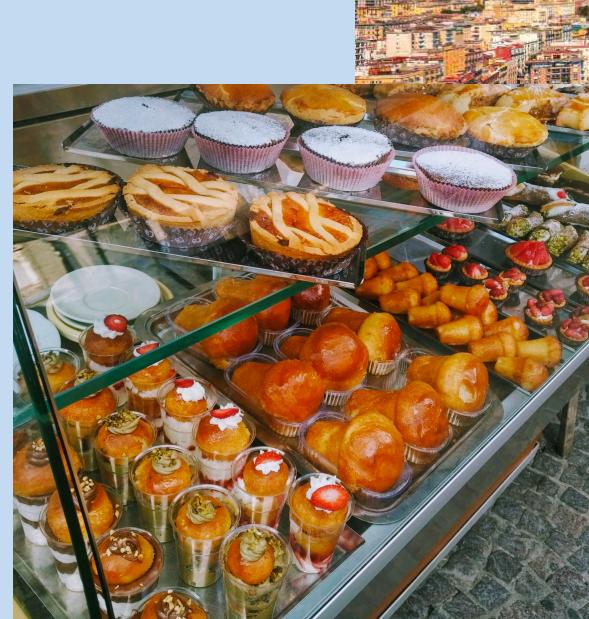








MAR
ADRIATICO











~\$ whoami



Nationality



Italian

Emanuele Iannone

✉ eiannone@unisa.it

🌐 <https://emaiannone.github.io/>

🐦 @EmanueleIannone3

~\$ whoami



Nationality  Italian

Age  26 y.o.

Emanuele Iannone

 eiannone@unisa.it

 <https://emaiannone.github.io/>

 @EmanueleIannone3

~\$ whoami



Nationality	 Italian
Age	 26 y.o.
Affiliation	 University of Salerno Department of Computer Science Software Engineering (SeSa) Lab

Emanuele Iannone

 eiannone@unisa.it

 <https://emaiannone.github.io/>

 @EmanueleIannone3

~\$ whoami



Emanuele Iannone

✉ eiannone@unisa.it

🌐 <https://emaiannone.github.io/>

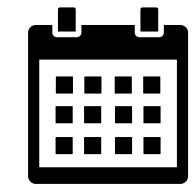
🐦 @EmanueleIannone3

Nationality



Italian

Age



26 y.o.

Affiliation



University of Salerno

Department of Computer Science

Software Engineering (SeSa) Lab

Curriculum



B.Sc. in Computer Science, thesis on
Refactoring for Android Energy Consumption

2015-2018

~\$ whoami



Emanuele Iannone

✉️ eiannone@unisa.it

🌐 <https://emaiannone.github.io/>

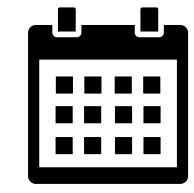
🐦 @EmanueleIannone3

Nationality



Italian

Age



26 y.o.

Affiliation



University of Salerno

Department of Computer Science

Software Engineering (SeSa) Lab

Curriculum



B.Sc. in Computer Science, thesis on
Refactoring for Android Energy Consumption

2015-2018

M.Sc. in Computer Science, thesis on
Test Generation for Third-party Vulnerabilities

2018-2020

~\$ whoami



Emanuele Iannone

✉️ eiannone@unisa.it

🌐 <https://emaiannone.github.io/>

🐦 [@EmanueleIannone3](https://twitter.com/EmanueleIannone)

Nationality



Italian

Age



26 y.o.

Affiliation



University of Salerno

Department of Computer Science

Software Engineering (SeSa) Lab

Curriculum



B.Sc. in Computer Science, thesis on
Refactoring for Android Energy Consumption

2015-2018

M.Sc. in Computer Science, thesis on
Test Generation for Third-party Vulnerabilities

2018-2020

Ph.D. in Computer Science, researching on
*Software Vulnerabilities Analysis,
Prediction, and Assessment*

2020-2023

~\$ whoami



Emanuele Iannone

✉️ eiannone@unisa.it

🌐 <https://emaiannone.github.io/>

🐦 @EmanueleIannone3

Teaching

⚙️ **[BSc] Software Engineering**

🧠 **[BSc] Fundamentals of AI**

🔨 **[MSc] Software Maintenance & Evolution**

🔨 **[MSc] Software Dependability**

~\$ whoami



Emanuele Iannone

✉️ eiannone@unisa.it

🌐 <https://emaiannone.github.io/>

🐦 @EmanueleIannone3

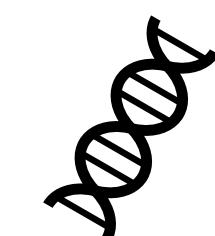
My Research



Mining Software Repositories for Security



Machine Learning for Vulnerability Prediction



Evolutionary Algorithms for Vulnerability Assessment

Other topics of Empirical Software Engineering:

- ▶ Source Code Refactoring
- ▶ Program Comprehension
- ▶ Energy Consumption of Mobile Apps

Can we
just start?

Vulnerability-contributing commit (VCC)

Vulnerability-contributing commit (VCC)



“A commit that contributed to the introduction of a post-release vulnerability.”

Andrew (Andy) Meneely

Vulnerability-contributing commit (VCC)



“A commit that contributed to the introduction of a post-release vulnerability.”

Andrew (Andy) Meneely

Vulnerability-contributing commit (VCC)



Andrew (Andy) Meneely

“A commit that contributed to the introduction of a post-release vulnerability.”

Code changes that “move” the code toward the state in which it contains the weakness.

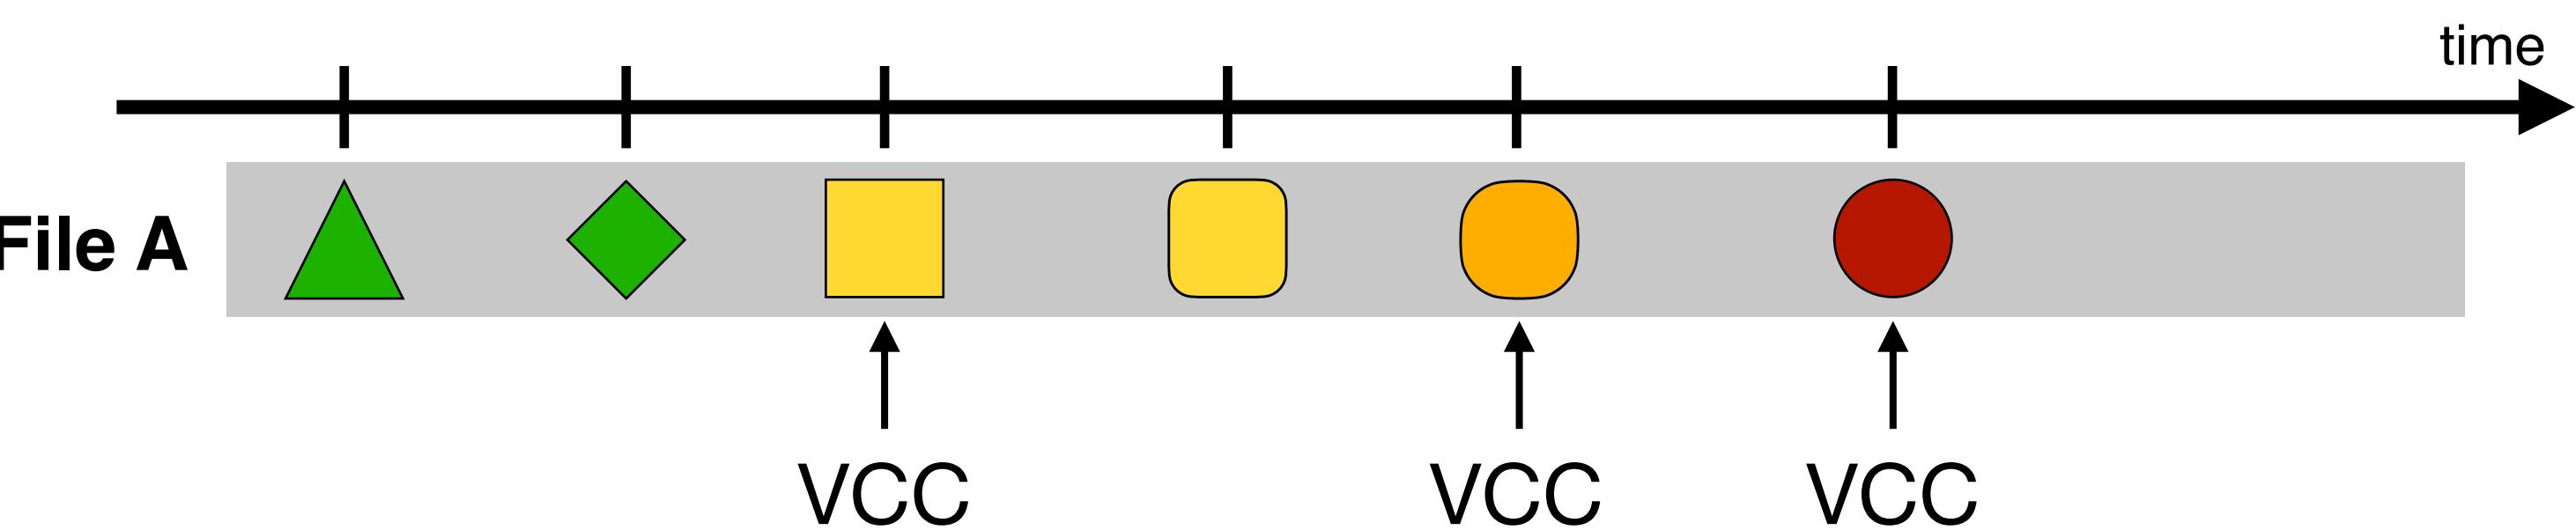
Vulnerability-contributing commit (VCC)



Andrew (Andy) Meneely

“A commit that contributed to the introduction of a post-release vulnerability.”

Code changes that “move” the code toward the state in which it contains the weakness.



Vulnerability-contributing commit (VCC)



“A commit that contributed to the introduction of a post-release vulnerability.”

Andrew (Andy) Meneely

Vulnerability-contributing commit (VCC)



Andrew (Andy) Meneely

“A commit that contributed to the introduction of a post-release vulnerability.”

The concept is irrelevant for pre-release vulnerabilities, fixed the exposure to externals.

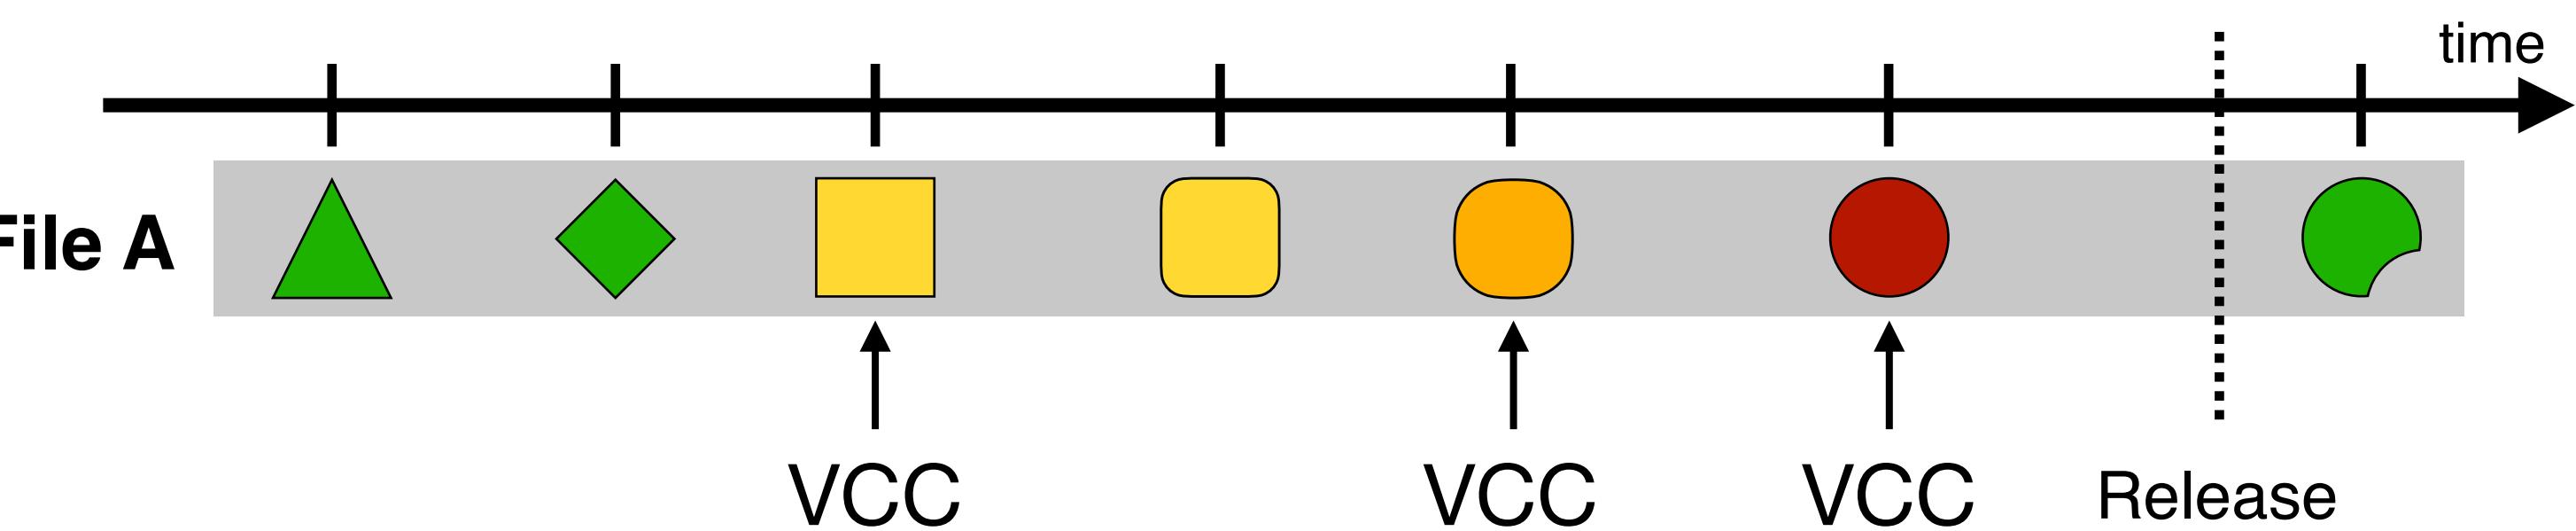
Vulnerability-contributing commit (VCC)



Andrew (Andy) Meneely

“A commit that contributed to the introduction of a post-release vulnerability.”

The concept is irrelevant for pre-release vulnerabilities, fixed the exposure to externals.



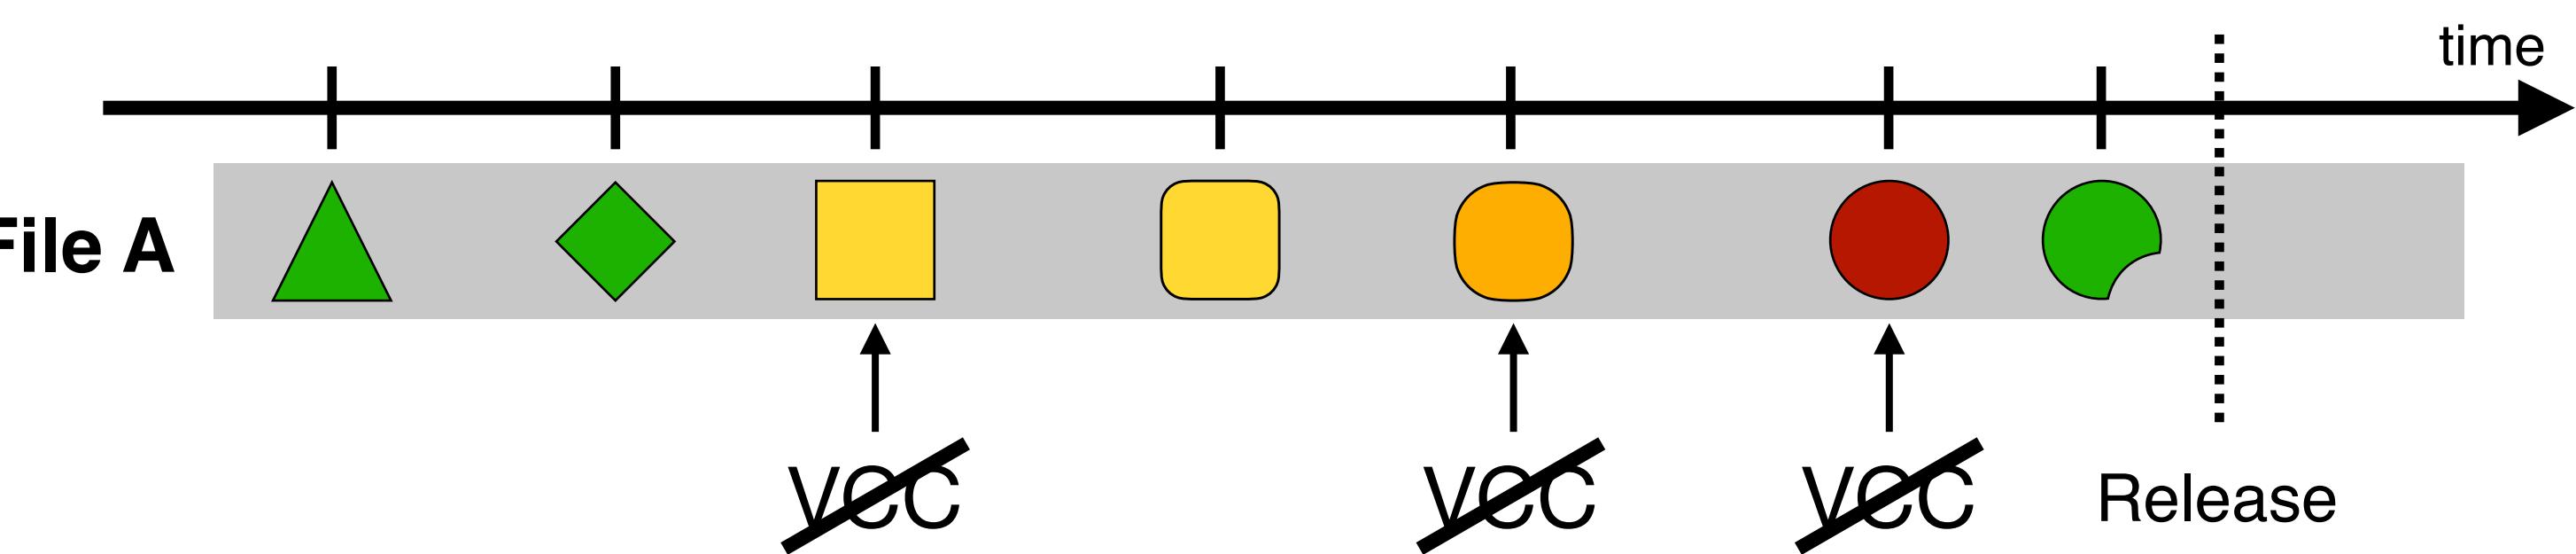
Vulnerability-contributing commit (VCC)



Andrew (Andy) Meneely

“A commit that contributed to the introduction of a post-release vulnerability.”

The concept is irrelevant for pre-release vulnerabilities, fixed the exposure to externals.



Example of a VCC

CVE-2019-11274

“Cloud Foundry UAA, versions prior to 74.0.0, is vulnerable to an XSS attack. A remote unauthenticated malicious attacker could craft a URL that contains a SCIM filter that contains malicious JavaScript, which older browsers may execute.”

Example of a VCC

CVE-2019-11274

“Cloud Foundry UAA, versions prior to 74.0.0, is vulnerable to an XSS attack. A remote unauthenticated malicious attacker could craft a URL that contains a SCIM filter that contains malicious JavaScript, which older browsers may execute.”

CWE-79: Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting')

Example of a VCC

CVE-2019-11274

“Cloud Foundry UAA, versions prior to 74.0.0, is vulnerable to an XSS attack. A remote unauthenticated malicious attacker could craft a URL that contains a SCIM filter that contains malicious JavaScript, which older browsers may execute.”

CWE-79: Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting')

We expect unescaped or unvalidated data supplied from the user via URL parameters that end up directly in the response.

Example of a VCC

CVE-2019-11274

Fix
a34f55fc

```
@RequestMapping(value = {"/Groups"}, method = RequestMethod.GET)
@ResponseBody
public SearchResults<?> listGroups(
    @RequestParam(value = "attributes", required = false) String attributesCommaSeparated,
    @RequestParam(required = false, defaultValue = "id pr") String filter,
    @RequestParam(required = false, defaultValue = "created") String sortBy,
    @RequestParam(required = false, defaultValue = "ascending") String sortOrder,
    @RequestParam(required = false, defaultValue = "1") int startIndex,
    @RequestParam(required = false, defaultValue = "100") int count) {
    if (count > groupMaxCount) {
        count = groupMaxCount;
    }
    List<ScimGroup> result;
    try {
        result = dao.query(filter, sortBy, "ascending".equalsIgnoreCase(sortOrder),
            identityZoneManager.getCurrentIdentityZoneId());
    } catch (IllegalArgumentException e) {
        throw new ScimException("Invalid filter expression: [" + filter + "]",
            HttpStatus.BAD_REQUEST);
        throw new ScimException("Invalid filter expression: [" + HtmlUtils.htmlEscape(filter) + "]",
            HttpStatus.BAD_REQUEST);
    }
    [...]
```

Example of a VCC

CVE-2019-11274

Fix
a34f55fc

```
@RequestMapping(value = {"/Groups"}, method = RequestMethod.GET)
@ResponseBody
public SearchResults<?> listGroups(
    @RequestParam(value = "attributes", required = false) String attributesCommaSeparated,
    @RequestParam(required = false, defaultValue = "id pr") String filter,
    @RequestParam(required = false, defaultValue = "created") String sortBy,
    @RequestParam(required = false, defaultValue = "ascending") String sortOrder,
    @RequestParam(re...
```

Essentially, the filter parameter is not sanitized and is placed directly in this ScimException. Then, this exception message is placed verbatim on an error page.

```
        identityZoneManager.getCurrentIdentityZoneId());
    } catch (IllegalArgumentException e) {
        throw new ScimException("Invalid filter expression: [" + filter + "]",
            HttpStatus.BAD_REQUEST);
        throw new ScimException("Invalid filter expression: [" + HtmlUtils.htmlEscape(filter) + "]",
            HttpStatus.BAD_REQUEST);
    }
    [...]
```

Example of a VCC

CVE-2019-11274

Fix
a34f55fc

```
@RequestMapping(value = {"/Groups"}, method = RequestMethod.GET)
@ResponseBody
public SearchResults<?> listGroups(
    @RequestParam(value = "attributes", required = false) String attributesCommaSeparated,
    @RequestParam(required = false, defaultValue = "id pr") String filter,
    @RequestParam(required = false, defaultValue = "created") String sortBy,
    @RequestParam(required = false, defaultValue = "ascending") String sortOrder,
    @RequestParam(required = false) Integer count,
    @RequestParam(required = false) Integer groupNumber,
    @RequestParam(required = false) String sortField,
    @RequestParam(required = false) String sortDirection) {
    if (count > groupNumber) {
        count = groupNumber;
    }
    List<ScimGroup> result = new ArrayList<ScimGroup>();
    try {
        result = dao.listGroups(attributesCommaSeparated, filter, sortBy, sortDirection, count, groupNumber);
    } catch (SQLException e) {
        throw new ScimException("Error occurred while listing groups: " + e.getMessage());
    }
    if (result.isEmpty()) {
        throw new ScimException("No groups found.", HttpStatus.NO_CONTENT);
    }
    return SearchResults.of(result);
}

private void validateFilter(String filter) {
    if (filter == null || filter.isEmpty() || filter.contains("<") || filter.contains(">")) {
        throw new ScimException("Invalid filter expression: [" + filter + "]", HttpStatus.BAD_REQUEST);
    }
}

private void validateSortBy(String sortBy) {
    if (sortBy == null || sortBy.isEmpty() || sortBy.equals("created")) {
        throw new ScimException("Invalid sort by field: [" + sortBy + "]", HttpStatus.BAD_REQUEST);
    }
}

private void validateSortOrder(String sortOrder) {
    if (sortOrder == null || sortOrder.isEmpty() || !sortOrder.equals("ascending")) {
        throw new ScimException("Invalid sort order: [" + sortOrder + "]", HttpStatus.BAD_REQUEST);
    }
}
```

Let's go back in time to find the commit
that contributed to this problem!

Example of a VCC

CVE-2019-11274

VCC
bb8ff8f4

```
@RequestMapping(value = { "/Groups/External/list" }, method = RequestMethod.GET)
@ResponseBody
public SearchResults<?> listExternalGroups(
    @RequestParam(required = false, defaultValue = "1") int startIndex,
    @RequestParam(required = false, defaultValue = "100") int count) {
    String filter = "";
    List<ScimGroupExternalMember> result;
    try {
        result = externalMembershipManager.query(filter);
    } catch (IllegalArgumentException e) {
        throw new ScimException("Invalid filter expression: [" + filter + "]",
            HttpStatus.BAD_REQUEST);
    }
    [...]
```

Example of a VCC

CVE-2019-11274

VCC
bb8ff8f4

```
@RequestMapping(value = { "/Groups/External/list" }, method = RequestMethod.GET)
@ResponseBody
public SearchResults<?> listExternalGroups(
    @RequestParam(required = false, defaultValue = "1") int startIndex,
    @RequestParam(required = false, defaultValue = "100") int count) {
    String filter = "";
    List<ScimGroupExternalMember> result;
    try {
        result = externalMembershipManager.query(filter);
    } catch (IllegalArgumentException e) {
        throw new ScimException("Invalid filter expression: [" + filter + "]",
            HttpStatus.BAD_REQUEST);
    }
    [...]
```

This was the first revision where the filter parameters was put inside the exception message: the vulnerability was there since the method (with a different name) was born.

Terminology

2013 ACM / IEEE International Symposium on Empirical Software Engineering and Measurement

When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits

Andrew Mencely, Harshavardhan Srinivasan, Ayemi Musa,
Alberto Rodríguez Tejeda, Matthew Mokary, Brian Spates
Department of Software Engineering
Rochester Institute of Technology
Rochester, NY, USA
andy@se.rit.edu, {hxs8839, ajm661, acr921, mxm6060, bxs4361}@rit.edu

Abstract—Security is a harsh reality for software teams today. Developers must engineer secure software by preventing vulnerabilities, which are design and coding mistakes that have security consequences. Even in open source projects, vulnerable source code can remain unnoticed for years. In this paper, we traced 68 vulnerabilities in the Apache HTTP server back to the version control commits that contributed the vulnerable code originally. We manually found 124 Vulnerability-Contributing Commits (VCCs), spanning 17 years. In this exploratory study, we analyzed these VCCs quantitatively and qualitatively with the over-arching question: “What could developers have looked for to identify security concerns in this commit?” Specifically, we examined the size of the commit via code churn metrics, the amount developers overwrite each others’ code via interactive churn metrics, exposure time between VCC and fix, and dissemination of the VCC to the development community via release notes and voting mechanisms. Our results show that VCCs are large: more than twice as much code churn on average than non-VCCs, even when normalized against lines of code. Furthermore, a commit was twice as likely to be a VCC when the author was a new developer to the source code. The insight from this study can help developers understand how vulnerabilities originate in a system so that security-related mistakes can be prevented or caught in the future.

Index Terms—vulnerability, churn, socio-technical, empirical.

I. INTRODUCTION

Security is a harsh reality for software teams today. Insecure software is not only expensive to maintain, but can cause immeasurable damage to a brand, or worse, to the livelihood of customers, patients, and citizens.

To software developers, the key to secure software lies in preventing vulnerabilities. Software vulnerabilities are special types of “faults that violate an [implicit or explicit] security policy” [1]. If developers want to find and fix vulnerabilities they must focus beyond making the system work as specified and prevent the system’s functionality from being abused. According to security experts [2]–[4], finding vulnerabilities requires expertise in both the specific product and in software security in general.

The field of engineering secure software has a plethora of security practices for finding vulnerabilities, such as threat modeling, penetration testing, code inspections, misuse and

abuse cases [5], and automated static analysis [2]–[4]. While these practices have been shown to be effective, they can also be inefficient. Development teams are then faced with the challenge of prioritizing their fortification efforts within the entire development process. Developers might know what is possible, but lack a firm grip on what is probable. As a result, an uninformed development team can easily focus on the wrong areas for fortification.

Fortunately, an historical, longitudinal analysis of how vulnerabilities originated in professional products can inform fortification prioritization. Understanding the specific trends of how vulnerabilities can arise in a software development product can help developers understand where to look and what to look for in their own product. Some of these trends have been quantified in vulnerability prediction [6]–[10] studies using metrics aggregated at the file level, but little has been done to explore the original coding mistakes that contributed the vulnerabilities in the first place. In this study, we have identified and analyzed original coding mistakes as Vulnerability-Contributing Commits (VCCs), or commits in the version control repository that contributed to the introduction of a post-release vulnerability.

A myriad of factors can lead to the introduction and lack of detection of vulnerabilities. A developer may make a single massive change to the system, leaving his peers with an overwhelmingly large review. Furthermore, a developer may make small, incremental changes, but his work might be affecting the work of many other developers. Or, a developer may forget to disseminate her work in the change notes and so the code may miss out on be reviewed entirely.

The objective of this research is to improve software security by analyzing the size, interactive churn, and community dissemination of VCCs. We conducted an empirical case study of the Apache HTTP Server project (HTTPD). Using a multi-researcher, cross-validating, semi-automated, semi-manual process, we identified the VCCs for each known post-release vulnerability in HTTPD. To explore commit size, we analyzed three code churn metrics. Interactive churn is a suite of five recently-developed [6] socio-technical variants of code churn metrics that measure the degree to which developers’ changes overwrite each others’ code at the line level. To explore community dissemination, we analyzed the

The core idea behind VCCs is not new to the MSR world, and stems from research on traditional bugs.

Terminology

2013 ACM / IEEE International Symposium on Empirical Software Engineering and Measurement

When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits

Andrew Mencely, Harshavardhan Srinivasan, Ayemi Musa,

Alberto Rodríguez Tejeda, Matthew Mokary, Brian Spates

Department of Software Engineering
Rochester Institute of Technology
Rochester, NY, USA

andy@se.rit.edu, {hxs8839, ajm661, acr921, mxm6060, bxs4361}@rit.edu

Abstract—Security is a harsh reality for software teams today. Developers must engineer secure software by preventing vulnerabilities, which are design and coding mistakes that have security consequences. Even in open source projects, vulnerable source code can remain unnoticed for years. In this paper, we traced 68 vulnerabilities in the Apache HTTP server back to the version control commits that contributed the vulnerable code originally. We manually found 124 Vulnerability-Contributing Commits (VCCs), spanning 17 years. In this exploratory study, we analyzed these VCCs quantitatively and qualitatively with the overarching question: “What could developers have looked for to identify security concerns in this commit?” Specifically, we examined the size of the commit via code churn metrics, the amount developers overwrite each others’ code via interactive churn metrics, exposure time between VCC and fix, and dissemination of the VCC to the development community via release notes and voting mechanisms. Our results show that VCCs are large: more than twice as much code churn on average than non-VCCs, even when normalized against lines of code. Furthermore, a commit was twice as likely to be a VCC when the author was a new developer to the source code. The insight from this study can help developers understand how vulnerabilities originate in a system so that security-related mistakes can be prevented or caught in the future.

Index Terms—vulnerability, churn, socio-technical, empirical.

I. INTRODUCTION

Security is a harsh reality for software teams today. Insecure software is not only expensive to maintain, but can cause immeasurable damage to a brand, or worse, to the livelihood of customers, patients, and citizens.

To software developers, the key to secure software lies in preventing vulnerabilities. Software vulnerabilities are special types of “faults that violate an [implicit or explicit] security policy” [1]. If developers want to find and fix vulnerabilities they must focus beyond making the system work as specified and prevent the system’s functionality from being abused. According to security experts [2]–[4], finding vulnerabilities requires expertise in both the specific product and in software security in general.

The field of engineering secure software has a plethora of security practices for finding vulnerabilities, such as threat modeling, penetration testing, code inspections, misuse and

abuse cases [5], and automated static analysis [2]–[4]. While these practices have been shown to be effective, they can also be inefficient. Development teams are then faced with the challenge of prioritizing their fortification efforts within the entire development process. Developers might know what is possible, but lack a firm grip on what is probable. As a result, an uninformed development team can easily focus on the wrong areas for fortification.

Fortunately, an historical, longitudinal analysis of how vulnerabilities originated in professional products can inform fortification prioritization. Understanding the specific trends of how vulnerabilities can arise in a software development product can help developers understand where to look and what to look for in their own product. Some of these trends have been quantified in vulnerability prediction [6]–[10] studies using metrics aggregated at the file level, but little has been done to explore the original coding mistakes that contributed the vulnerabilities in the first place. In this study, we have identified and analyzed original coding mistakes as Vulnerability-Contributing Commits (VCCs), or commits in the version control repository that contributed to the introduction of a post-release vulnerability.

A myriad of factors can lead to the introduction and lack of detection of vulnerabilities. A developer may make a single massive change to the system, leaving his peers with an overwhelmingly large review. Furthermore, a developer may make small, incremental changes, but his work might be affecting the work of many other developers. Or, a developer may forget to disseminate her work in the change notes and so the code may miss out on be reviewed entirely.

The objective of this research is to improve software security by analyzing the size, interactive churn, and community dissemination of VCCs. We conducted an empirical case study of the Apache HTTP Server project (HTTPD). Using a multi-researcher, cross-validating, semi-automated, semi-manual process, we identified the VCCs for each known post-release vulnerability in HTTPD. To explore commit size, we analyzed three code churn metrics. Interactive churn is a suite of five recently-developed [6] socio-technical variants of code churn metrics that measure the degree to which developers’ changes overwrite each others’ code at the line level. To explore community dissemination, we analyzed the

The core idea behind VCCs is not new to the MSR world, and stems from research on traditional bugs.

Fix-inducing Change

Bug-introducing Change

Bug-inducing Change

Bug-injecting Change

Terminology

2013 ACM / IEEE International Symposium on Empirical Software Engineering and Measurement

When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits

Andrew Meneely, Harshavardhan Srinivasan, Ayemi Musa,
Alberto Rodríguez Tejeda, Matthew Mokary, Brian Spates

Department of Software Engineering
Rochester Institute of Technology
Rochester, NY, USA

andy@se.rit.edu, {hxs8839, ajm661, acr921, mxm6060, bxs4361}@rit.edu

Abstract—Security is a harsh reality for software teams today. Developers must engineer secure software by preventing vulnerabilities, which are design and coding mistakes that have security consequences. Even in open source projects, vulnerable source code can remain unnoticed for years. In this paper, we traced 68 vulnerabilities in the Apache HTTP server back to the version control commits that contributed the vulnerable code originally. We manually found 124 Vulnerability-Contributing Commits (VCCs), spanning 17 years. In this exploratory study, we analyzed these VCCs quantitatively and qualitatively with the overarching question: “What could developers have looked for to identify security concerns in this commit?” Specifically, we examined the size of the commit via code churn metrics, the amount developers overwrite each others’ code via interactive churn metrics, exposure time between VCC and fix, and dissemination of the VCC to the development community via release notes and voting mechanisms. Our results show that VCCs are large: more than twice as much code churn on average than non-VCCs, even when normalized against lines of code. Furthermore, a commit was twice as likely to be a VCC when the author was a new developer to the source code. The insight from this study can help developers understand how vulnerabilities originate in a system so that security-related mistakes can be prevented or caught in the future.

Index Terms—vulnerability, churn, socio-technical, empirical.

I. INTRODUCTION

Security is a harsh reality for software teams today. Insecure software is not only expensive to maintain, but can cause immeasurable damage to a brand, or worse, to the livelihood of customers, patients, and citizens.

To software developers, the key to secure software lies in preventing vulnerabilities. Software vulnerabilities are special types of “faults that violate an [implicit or explicit] security policy” [1]. If developers want to find and fix vulnerabilities they must focus beyond making the system work as specified and prevent the system’s functionality from being abused. According to security experts [2]–[4], finding vulnerabilities requires expertise in both the specific product and in software security in general.

The field of engineering secure software has a plethora of security practices for finding vulnerabilities, such as threat modeling, penetration testing, code inspections, misuse and

abuse cases [5], and automated static analysis [2]–[4]. While these practices have been shown to be effective, they can also be inefficient. Development teams are then faced with the challenge of prioritizing their fortification efforts within the entire development process. Developers might know what is possible, but lack a firm grip on what is probable. As a result, an uninformative development team can easily focus on the wrong areas for fortification.

Fortunately, an historical, longitudinal analysis of how vulnerabilities originated in professional products can inform fortification prioritization. Understanding the specific trends of how vulnerabilities can arise in a software development product can help developers understand where to look and what to look for in their own product. Some of these trends have been quantified in vulnerability prediction [6]–[10] studies using metrics aggregated at the file level, but little has been done to explore the original coding mistakes that contributed the vulnerabilities in the first place. In this study, we have identified and analyzed original coding mistakes as Vulnerability-Contributing Commits (VCCs), or commits in the version control repository that contributed to the introduction of a post-release vulnerability.

A myriad of factors can lead to the introduction and lack of detection of vulnerabilities. A developer may make a single massive change to the system, leaving his peers with an overwhelmingly large review. Furthermore, a developer may make small, incremental changes, but his work might be affecting the work of many other developers. Or, a developer may forget to disseminate her work in the change notes and so the code may miss out on being reviewed entirely.

The objective of this research is to improve software security by analyzing the size, interactive churn, and community dissemination of VCCs. We conducted an empirical case study of the Apache HTTP Server project (HTTPD). Using a multi-researcher, cross-validating, semi-automated, semi-manual process, we identified the VCCs for each known post-release vulnerability in HTTPD. To explore commit size, we analyzed three code churn metrics. Interactive churn is a suite of five recently-developed [6] socio-technical variants of code churn metrics that measure the degree to which developers’ changes overwrite each others’ code at the line level. To explore community dissemination, we analyzed the

The core idea behind VCCs is not new to the MSR world, and stems from research on traditional bugs.

Fix-inducing Change

Bug-introducing Change

Bug-inducing Change

Bug-injecting Change

Meneely et al. argued about the term “fix-inducing”, which can be translated into “persuade to fix (the bug)”. In their view, a VCC does not persuade developers to fix the vulnerability... the vulnerability is fixed after its discovery, not because of a flawed commit!

Terminology

2013 ACM / IEEE International Symposium on Empirical Software Engineering and Measurement

When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits

Andrew Meneely, Harshavardhan Srinivasan, Ayemi Musa,
Alberto Rodríguez Tejeda, Matthew Mokary, Brian Spates

Department of Software Engineering
Rochester Institute of Technology
Rochester, NY, USA

andy@se.rit.edu, {hxs8839, ajm661, acr921, mxm6060, bxs4361}@rit.edu

Abstract—Security is a harsh reality for software teams today. Developers must engineer secure software by preventing vulnerabilities, which are design and coding mistakes that have security consequences. Even in open source projects, vulnerable source code can remain unnoticed for years. In this paper, we traced 68 vulnerabilities in the Apache HTTP server back to the version control commits that contributed the vulnerable code originally. We manually found 124 Vulnerability-Contributing Commits (VCCs), spanning 17 years. In this exploratory study, we analyzed these VCCs quantitatively and qualitatively with the overarching question: “What could developers have looked for to identify security concerns in this commit?” Specifically, we examined the size of the commit via code churn metrics, the amount developers overwrite each others’ code via interactive churn metrics, exposure time between VCC and fix, and dissemination of the VCC to the development community via release notes and voting mechanisms. Our results show that VCCs are large: more than twice as much code churn on average than non-VCCs, even when normalized against lines of code. Furthermore, a commit was twice as likely to be a VCC when the author was a new developer to the source code. The insight from this study can help developers understand how vulnerabilities originate in a system so that security-related mistakes can be prevented or caught in the future.

Index Terms—vulnerability, churn, socio-technical, empirical.

I. INTRODUCTION

Security is a harsh reality for software teams today. Insecure software is not only expensive to maintain, but can cause immeasurable damage to a brand, or worse, to the livelihood of customers, patients, and citizens.

To software developers, the key to secure software lies in preventing vulnerabilities. Software vulnerabilities are special types of “faults that violate an [implicit or explicit] security policy” [1]. If developers want to find and fix vulnerabilities they must focus beyond making the system work as specified and prevent the system’s functionality from being abused. According to security experts [2]–[4], finding vulnerabilities requires expertise in both the specific product and in software security in general.

The field of engineering secure software has a plethora of security practices for finding vulnerabilities, such as threat modeling, penetration testing, code inspections, misuse and

abuse cases [5], and automated static analysis [2]–[4]. While these practices have been shown to be effective, they can also be inefficient. Development teams are then faced with the challenge of prioritizing their fortification efforts within the entire development process. Developers might know what is possible, but lack a firm grip on what is probable. As a result, an uninformative development team can easily focus on the wrong areas for fortification.

Fortunately, an historical, longitudinal analysis of how vulnerabilities originated in professional products can inform fortification prioritization. Understanding the specific trends of how vulnerabilities can arise in a software development product can help developers understand where to look and what to look for in their own product. Some of these trends have been quantified in vulnerability prediction [6]–[10] studies using metrics aggregated at the file level, but little has been done to explore the original coding mistakes that contributed the vulnerabilities in the first place. In this study, we have identified and analyzed original coding mistakes as Vulnerability-Contributing Commits (VCCs), or commits in the version control repository that contributed to the introduction of a post-release vulnerability.

A myriad of factors can lead to the introduction and lack of detection of vulnerabilities. A developer may make a single massive change to the system, leaving his peers with an overwhelmingly large review. Furthermore, a developer may make small, incremental changes, but his work might be affecting the work of many other developers. Or, a developer may forget to disseminate her work in the change notes and so the code may miss out on being reviewed entirely.

The objective of this research is to improve software security by analyzing the size, interactive churn, and community dissemination of VCCs. We conducted an empirical case study of the Apache HTTP Server project (HTTPD). Using a multi-researcher, cross-validating, semi-automated, semi-manual process, we identified the VCCs for each known post-release vulnerability in HTTPD. To explore commit size, we analyzed three code churn metrics. Interactive churn is a suite of five recently-developed [6] socio-technical variants of code churn metrics that measure the degree to which developers’ changes overwrite each others’ code at the line level. To explore community dissemination, we analyzed the

The core idea behind VCCs is not new to the MSR world, and stems from research on traditional bugs.

Fix-inducing Change

Bug-introducing Change

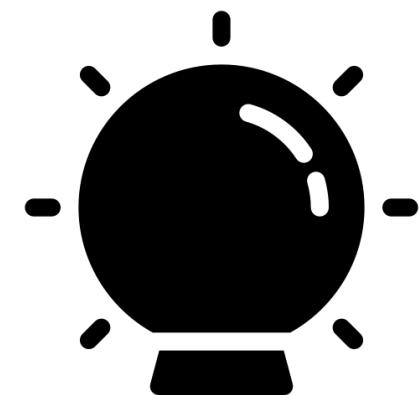
Bug-inducing Change

Bug-injecting Change

Meneely et al. argued about the term “fix-inducing”, which can be translated into “persuade to fix (the bug)”. In their view, a VCC does not persuade developers to fix the vulnerability... the vulnerability is fixed after its discovery, not because of a flawed commit!

Long story short: as long as we all agree, it makes no (real) difference.

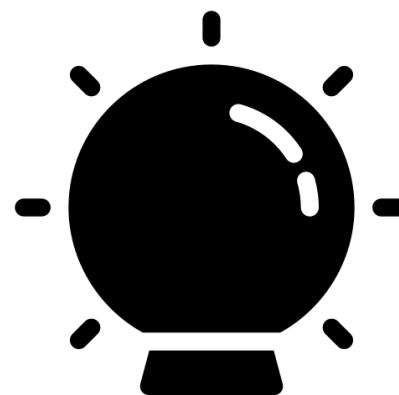
Main Uses of VCCs



**Train Vulnerability
Prediction Models**

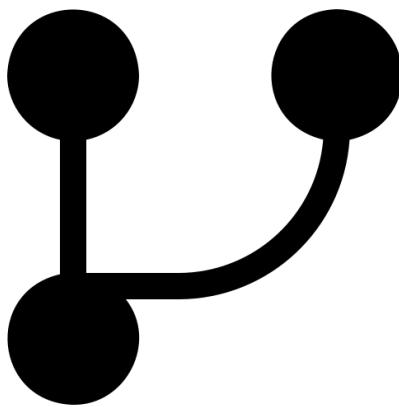
We can build a **just-in-time
vulnerability prediction model** if
the dataset is made of VCCs and
non-VCCs.

Main Uses of VCCs



Train Vulnerability Prediction Models

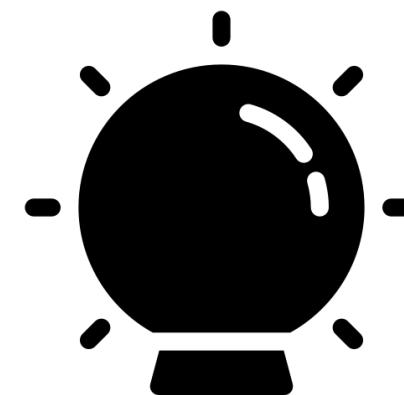
We can build a **just-in-time vulnerability prediction model** if the dataset is made of VCCs and non-VCCs.



Recover Vulnerable Versions/Releases

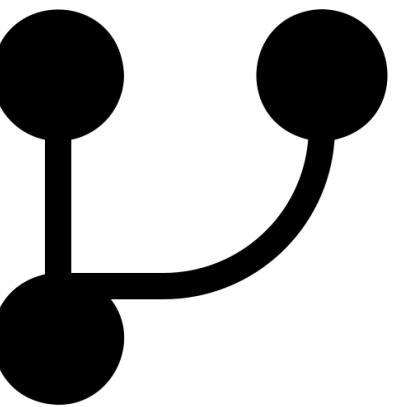
VCCs can help understand which project releases are affected by the vulnerability.

Main Uses of VCCs



Train Vulnerability Prediction Models

We can build a **just-in-time vulnerability prediction model** if the dataset is made of VCCs and non-VCCs.



Recover Vulnerable Versions/Releases

VCCs can help understand which project releases are affected by the vulnerability.



Expand the Knowledge on Vulnerabilities

Understand how vulnerabilities are progressively introduced in the code, drawing out interesting facts.

Key Characteristics of VCCs

VCCs vs non-VCCs

A case study on *Apache HTTP Server* with 68 post-release vulnerabilities and 124 VCCs.

2013 ACM / IEEE International Symposium on Empirical Software Engineering and Measurement

When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits

Andrew Meneely, Harshavardhan Srinivasan, Ayemni Musa,
Alberto Rodriguez Tejeda, Matthew Mokary, Brian Spates
Department of Software Engineering
Rochester Institute of Technology
Rochester, NY, USA
andy@se.rit.edu, {bxs839, ajm661, acr921, mxm6060, bxs4361}@rit.edu

Abstract—Security is a harsh reality for software teams today. Developers must engineer secure software by preventing vulnerabilities, which are design and coding mistakes that have security consequences. Even in open source projects, vulnerable code can be introduced to the system. In this paper, we traced 68 vulnerabilities in the Apache HTTP Server back to the version control commits that contributed the vulnerable code originally. We manually found 124 Vulnerability-Contributing Commits (VCCs), spanning 17 years. In this exploratory study, we analyzed these VCCs quantitatively and qualitatively with the over-arching question: “What could developers have looked for to identify security concerns in this commit?” Specifically, we examined the size of the commit via code churn metrics, the amount developers overwrite each others’ code via interactive churn metrics, and the time between VCC and fix. We also analyzed the VCCs to determine community via release notes and voting mechanisms. Our results show that VCCs are large; more than twice as much code churn on average than non-VCCs, even when normalized against lines of code. Furthermore, a commit was twice as likely to be a VCC when the author was a new developer to the source code. The insight from this study can help developers understand how vulnerabilities originate in a system so that security-related mistakes can be prevented or caught in the future.

Index Terms—vulnerability, churn, socio-technical, empirical.

I. INTRODUCTION

Security is a harsh reality for software teams today. Insecure software is not only expensive to maintain, but can cause immeasurable damage to a brand, or worse, to the livelihood of customers, patients, and citizens.

To software developers, the key to secure software lies in preventing vulnerabilities. Software vulnerabilities are special types of “faults that violate an [implicit or explicit] security policy” [1]. If developers want to find and fix vulnerabilities they must focus beyond making the system work as specified and prevent the system’s functionality from being abused. According to security experts [2]–[4], finding vulnerabilities requires expertise in both the specific product and in software security in general.

The field of engineering secure software has a plethora of security practices for finding vulnerabilities, such as threat modeling, penetration testing, code inspections, misuse and abuse cases [5], and automated static analysis [2]–[4]. While these practices have been shown to be effective, they can also be inefficient. Development teams are then faced with the challenge of prioritizing their fortification efforts within the entire development process. Developers might know what is possible, but lack a firm grip on what is probable. As a result, an uninformed development team can easily focus on the wrong areas for fortification.

Fortunately, an historical, longitudinal analysis of how vulnerabilities originated in professional products can inform fortification prioritization. Understanding the specific trends of how vulnerabilities can arise in a software development product can help developers understand where to look and what to look for in their own product. Some of these trends have been quantified in vulnerability prediction [6]–[10] studies using metrics aggregated at the file level, but little has been done to explore the original coding mistakes that contributed the vulnerabilities in the first place. In this study, we have identified and analyzed original coding mistakes as Vulnerability-Contributing Commits (VCCs), or commits in the version control repository that contributed to the introduction of a post-release vulnerability.

A myriad of factors can lead to the introduction and lack of detection of vulnerabilities. A developer may make a single massive change to the system, leaving his peers with an overwhelmingly large review. Furthermore, a developer may make small, incremental changes, but his work might be affecting the work of many other developers. Or, a developer may forget to disseminate her work in the change notes and so the code may miss out on a reviewed entirely.

The objective of this research is to improve software security by analyzing the size, interactive churn, and community dissemination of VCCs. We conducted an empirical case study of the Apache HTTP Server project (HTTPD). Using a multi-researcher, cross-validating, semi-automated, semi-manual process, we identified the VCCs for each known post-release vulnerability in HTTPD. To explore commit size, we analyzed three code churn metrics. Interactive churn is a suite of five recently-developed [6] socio-technical variants of code churn metrics that measure the degree to which developers’ changes overwrite each others’ code at the line level. To explore community dissemination, we analyzed the

978-0-7695-5056-5/13 \$26.00 © 2013 IEEE
DOI 10.1109/ESEM.2013.19

65

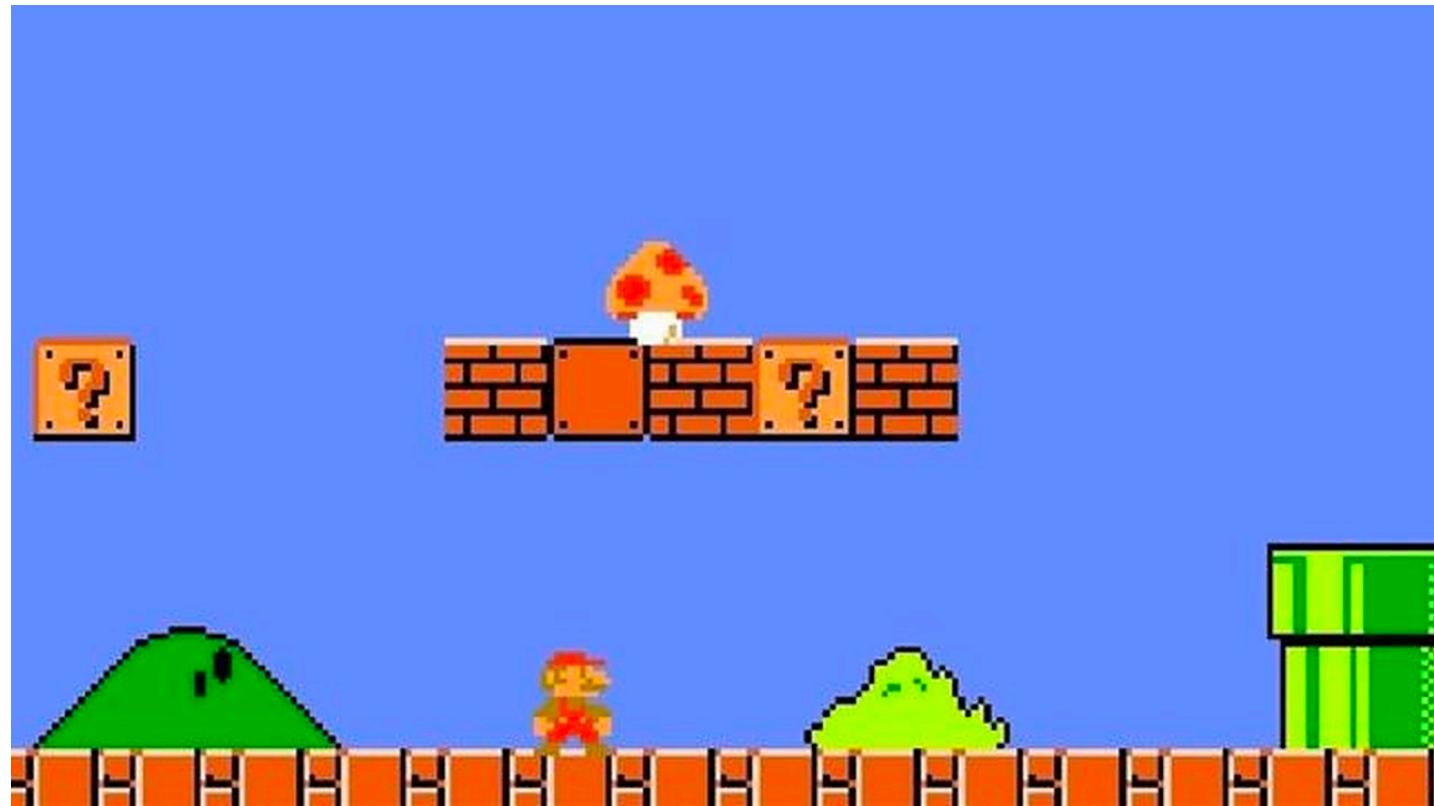
IEEE Computer Society

Authorized licensed use limited to: Università degli Studi di Salerno. Downloaded on May 04, 2023 at 13:28:42 UTC from IEEE Xplore. Restrictions apply.

Key Characteristics of VCCs

VCCs vs non-VCCs

A case study on *Apache HTTP Server* with 68 post-release vulnerabilities and 124 VCCs.



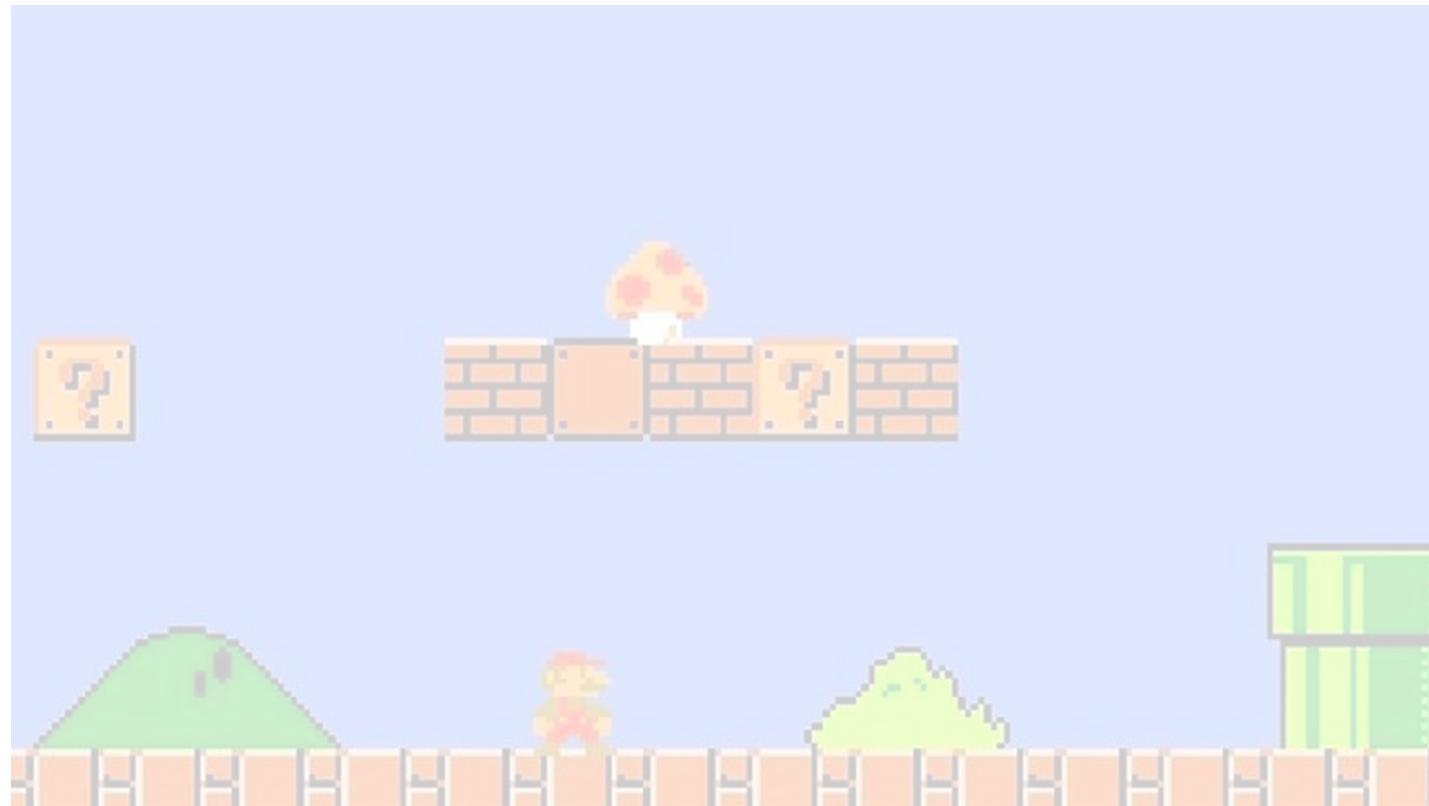
Size matters

VCCs change x10 more lines of code than non-VCCs.

Key Characteristics of VCCs

VCCs vs non-VCCs

A case study on *Apache HTTP Server* with 68 post-release vulnerabilities and 124 VCCs.



Size matters



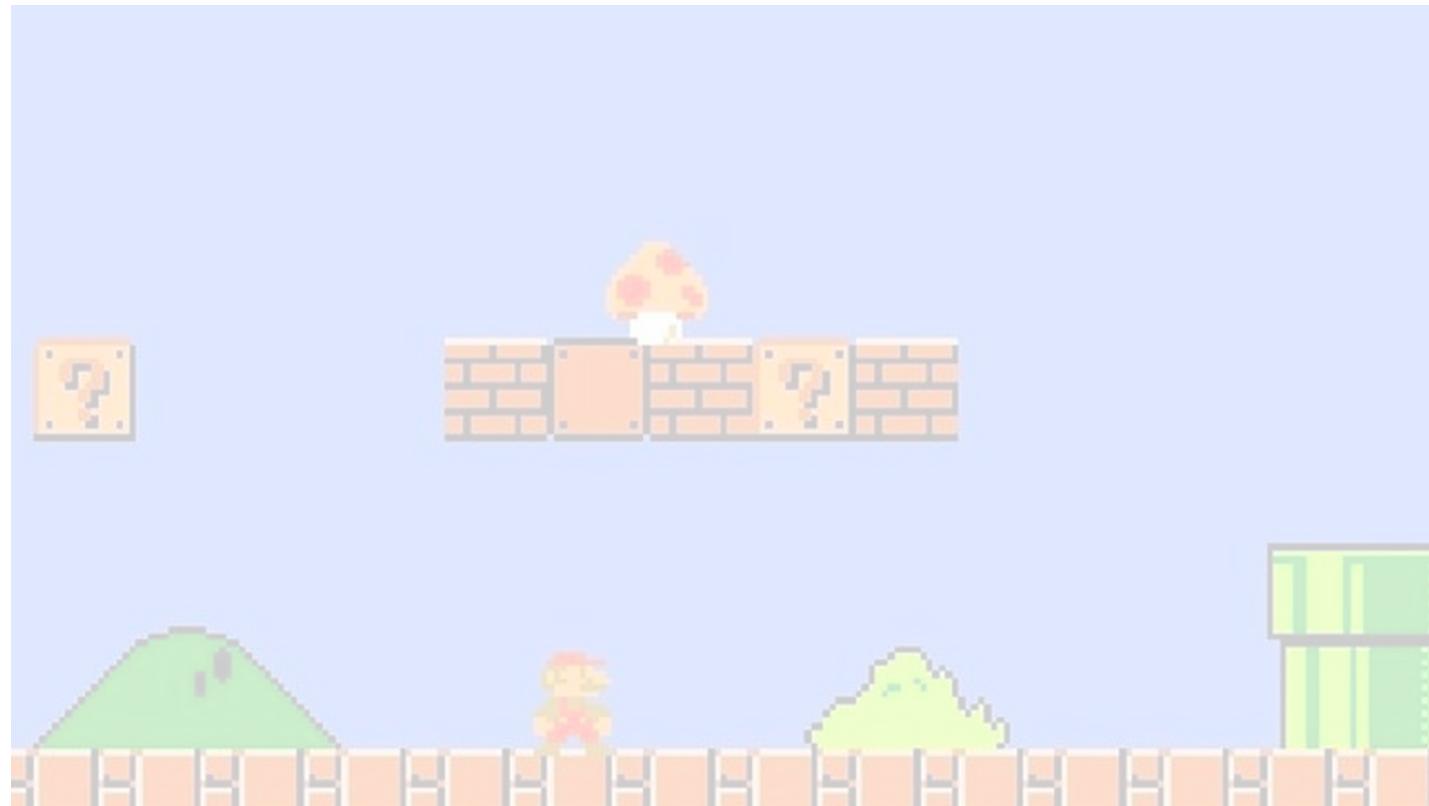
Don't step on
someone's toes

VCCs are made by **new authors** in **15% more cases** than non-VCCs.

Key Characteristics of VCCs

VCCs vs non-VCCs

A case study on *Apache HTTP Server* with 68 post-release vulnerabilities and 124 VCCs.



Size matters



Don't step on
someone's toes



A leopard CAN
change its spots

VCCs affect **existing files** in **87% of the cases** rather than new files.

Key Characteristics of VCCs

VCCs vs non-VCCs

A case study on *Apache HTTP Server* with 68 post-release vulnerabilities and 124 VCCs.

Large commits might increase the chance of contributing to a vulnerability.

Key Characteristics of VCCs

VCCs vs non-VCCs

A case study on *Apache HTTP Server* with 68 post-release vulnerabilities and 124 VCCs.

Large commits might increase the chance of contributing to a vulnerability.

Changing **other developers' code** might increase the chance of contributing to a vulnerability.

Key Characteristics of VCCs

VCCs vs non-VCCs

A case study on *Apache HTTP Server* with 68 post-release vulnerabilities and 124 VCCs.

Large commits might increase the chance of contributing to a vulnerability.

Changing **other developers' code** might increase the chance of contributing to a vulnerability.

Vulnerabilities are more likely to be added when **modifying existing files** rather than creating new files.

Yeah, cool.
How can we
mine them?

Mining VCCs: A First Approach

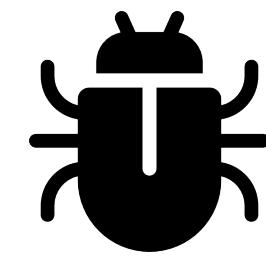
Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.

Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.



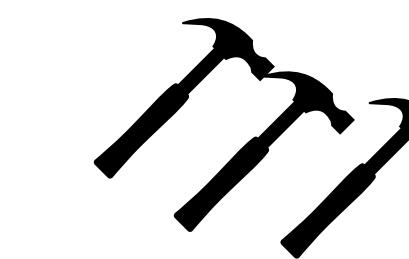
Post-release
vulnerability

Can be a known vulnerability from NVD
or another source, it is the same.

Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.

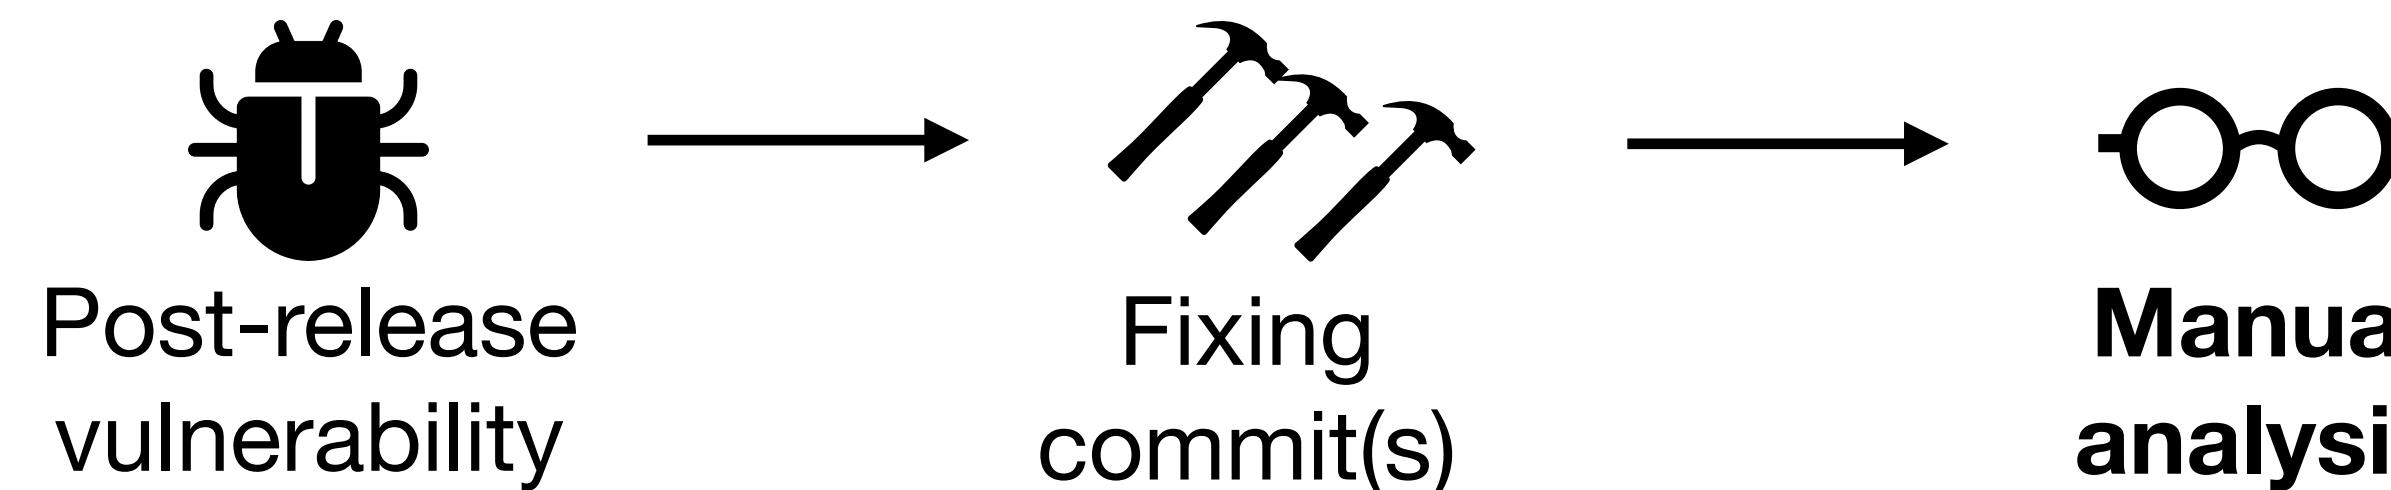


We assume the vulnerability is already mapped to its fixing commits.

Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

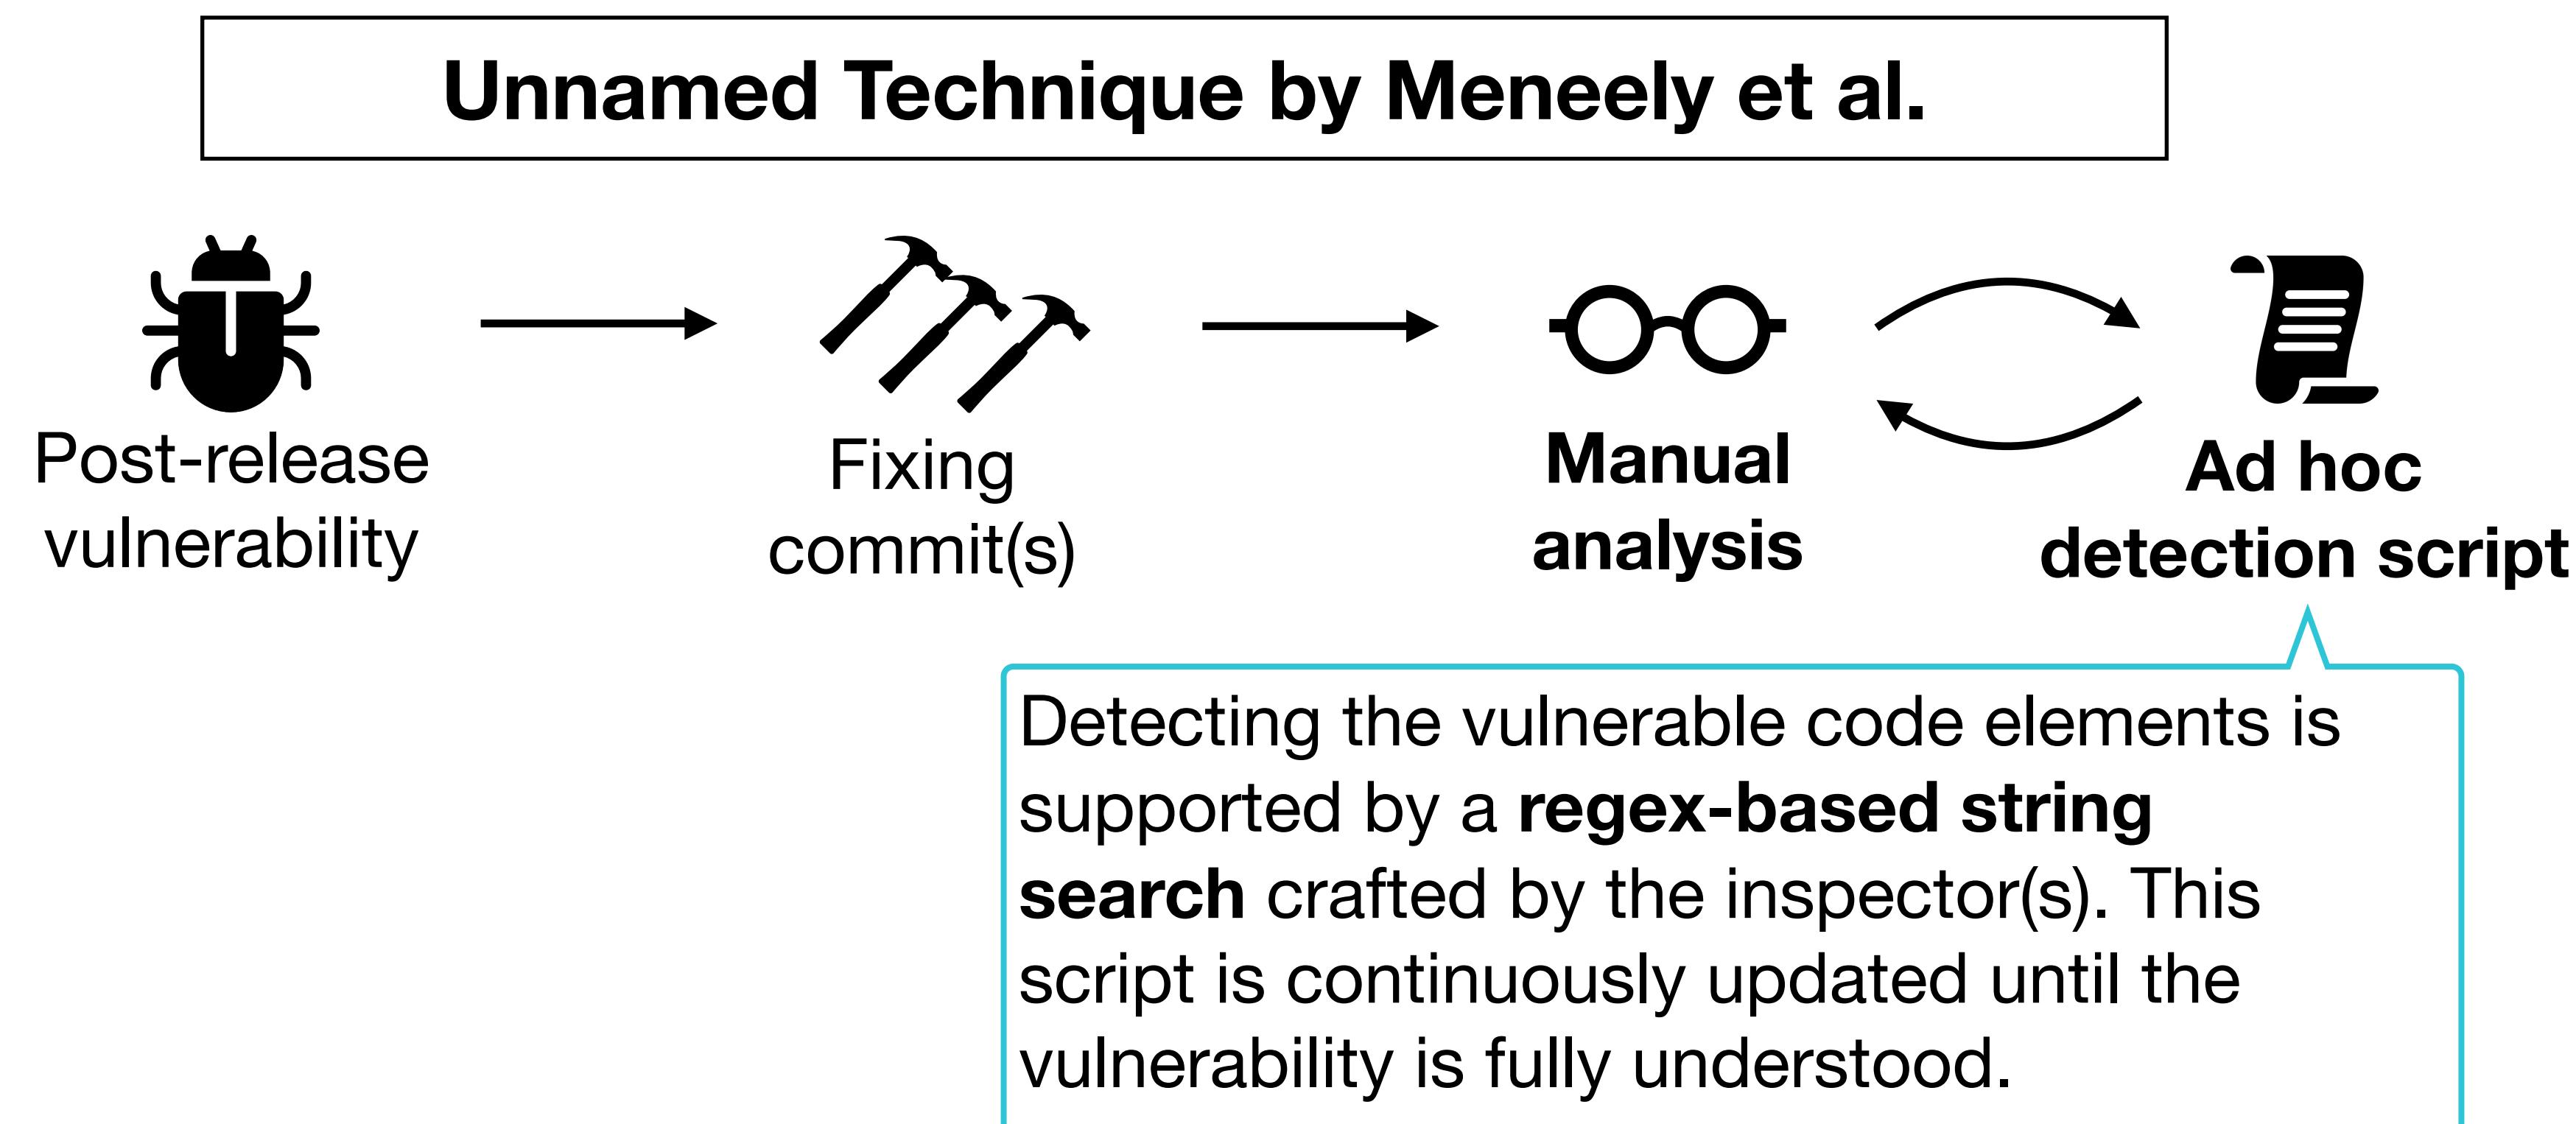
Unnamed Technique by Meneely et al.



One (or more) inspectors examine(s) the patch and its context to find the **vulnerable code elements** (statements). All the fixing commits are analyzed as one single big commit.

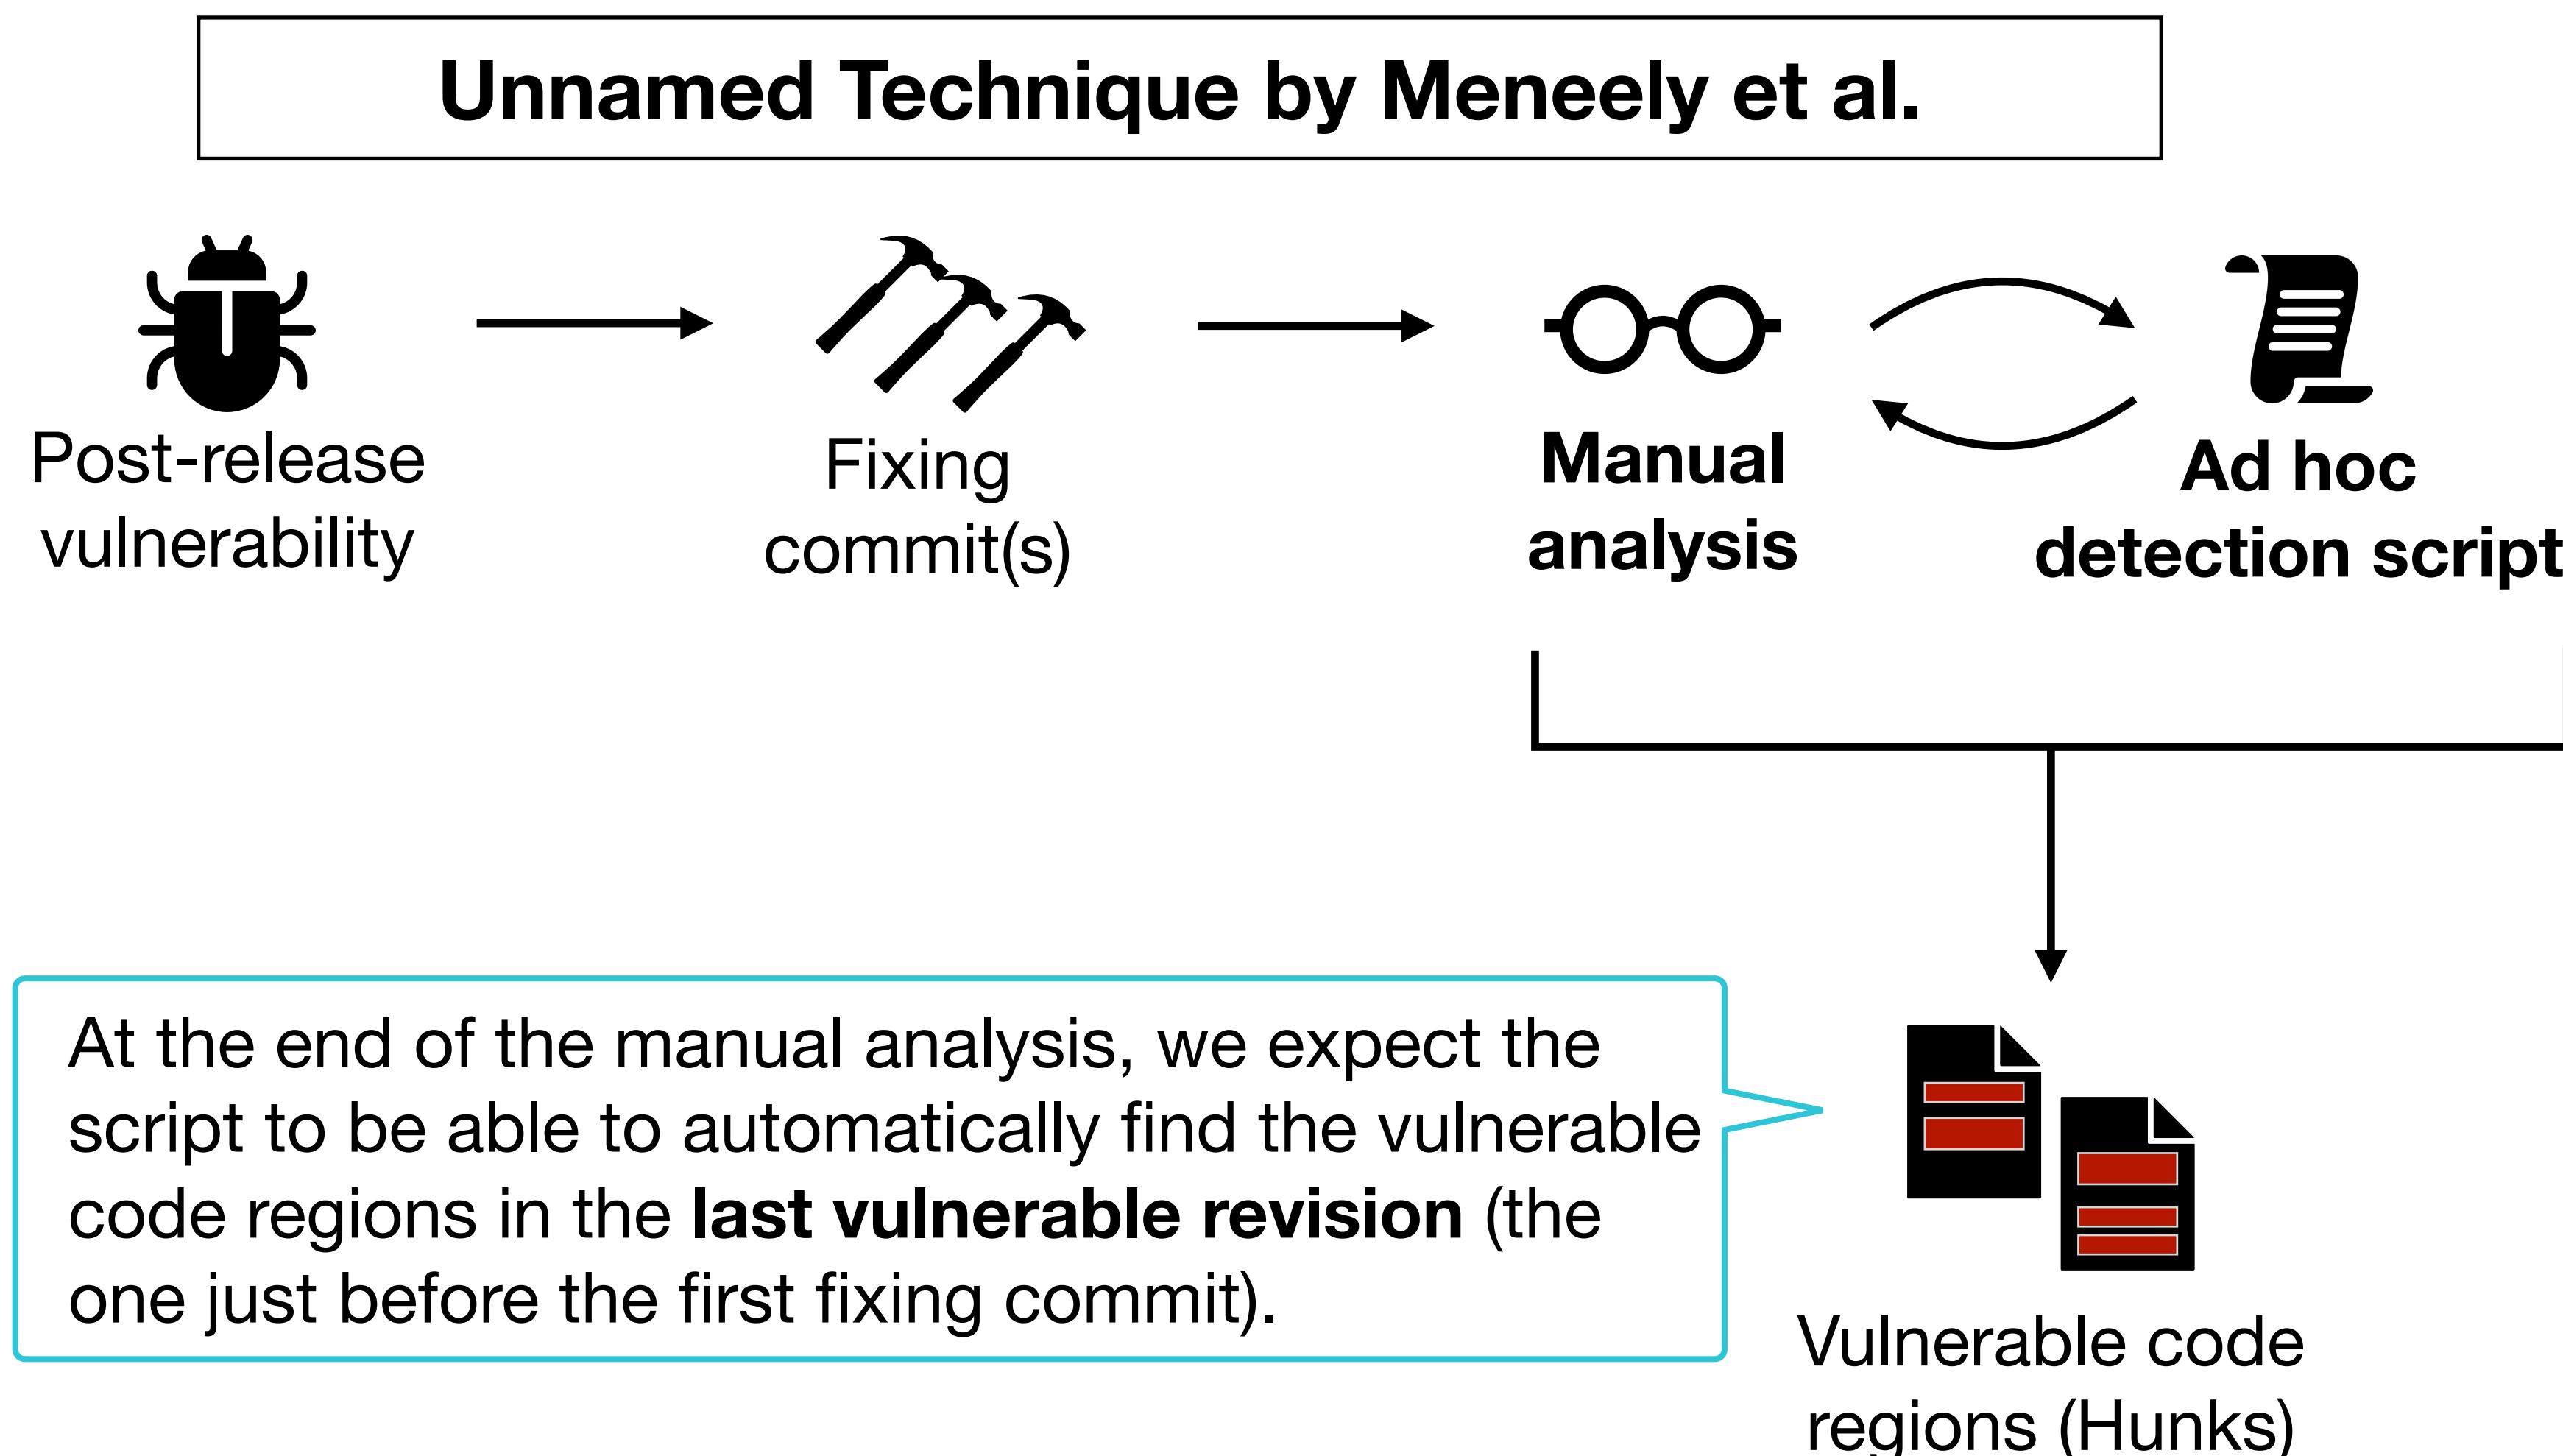
Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.



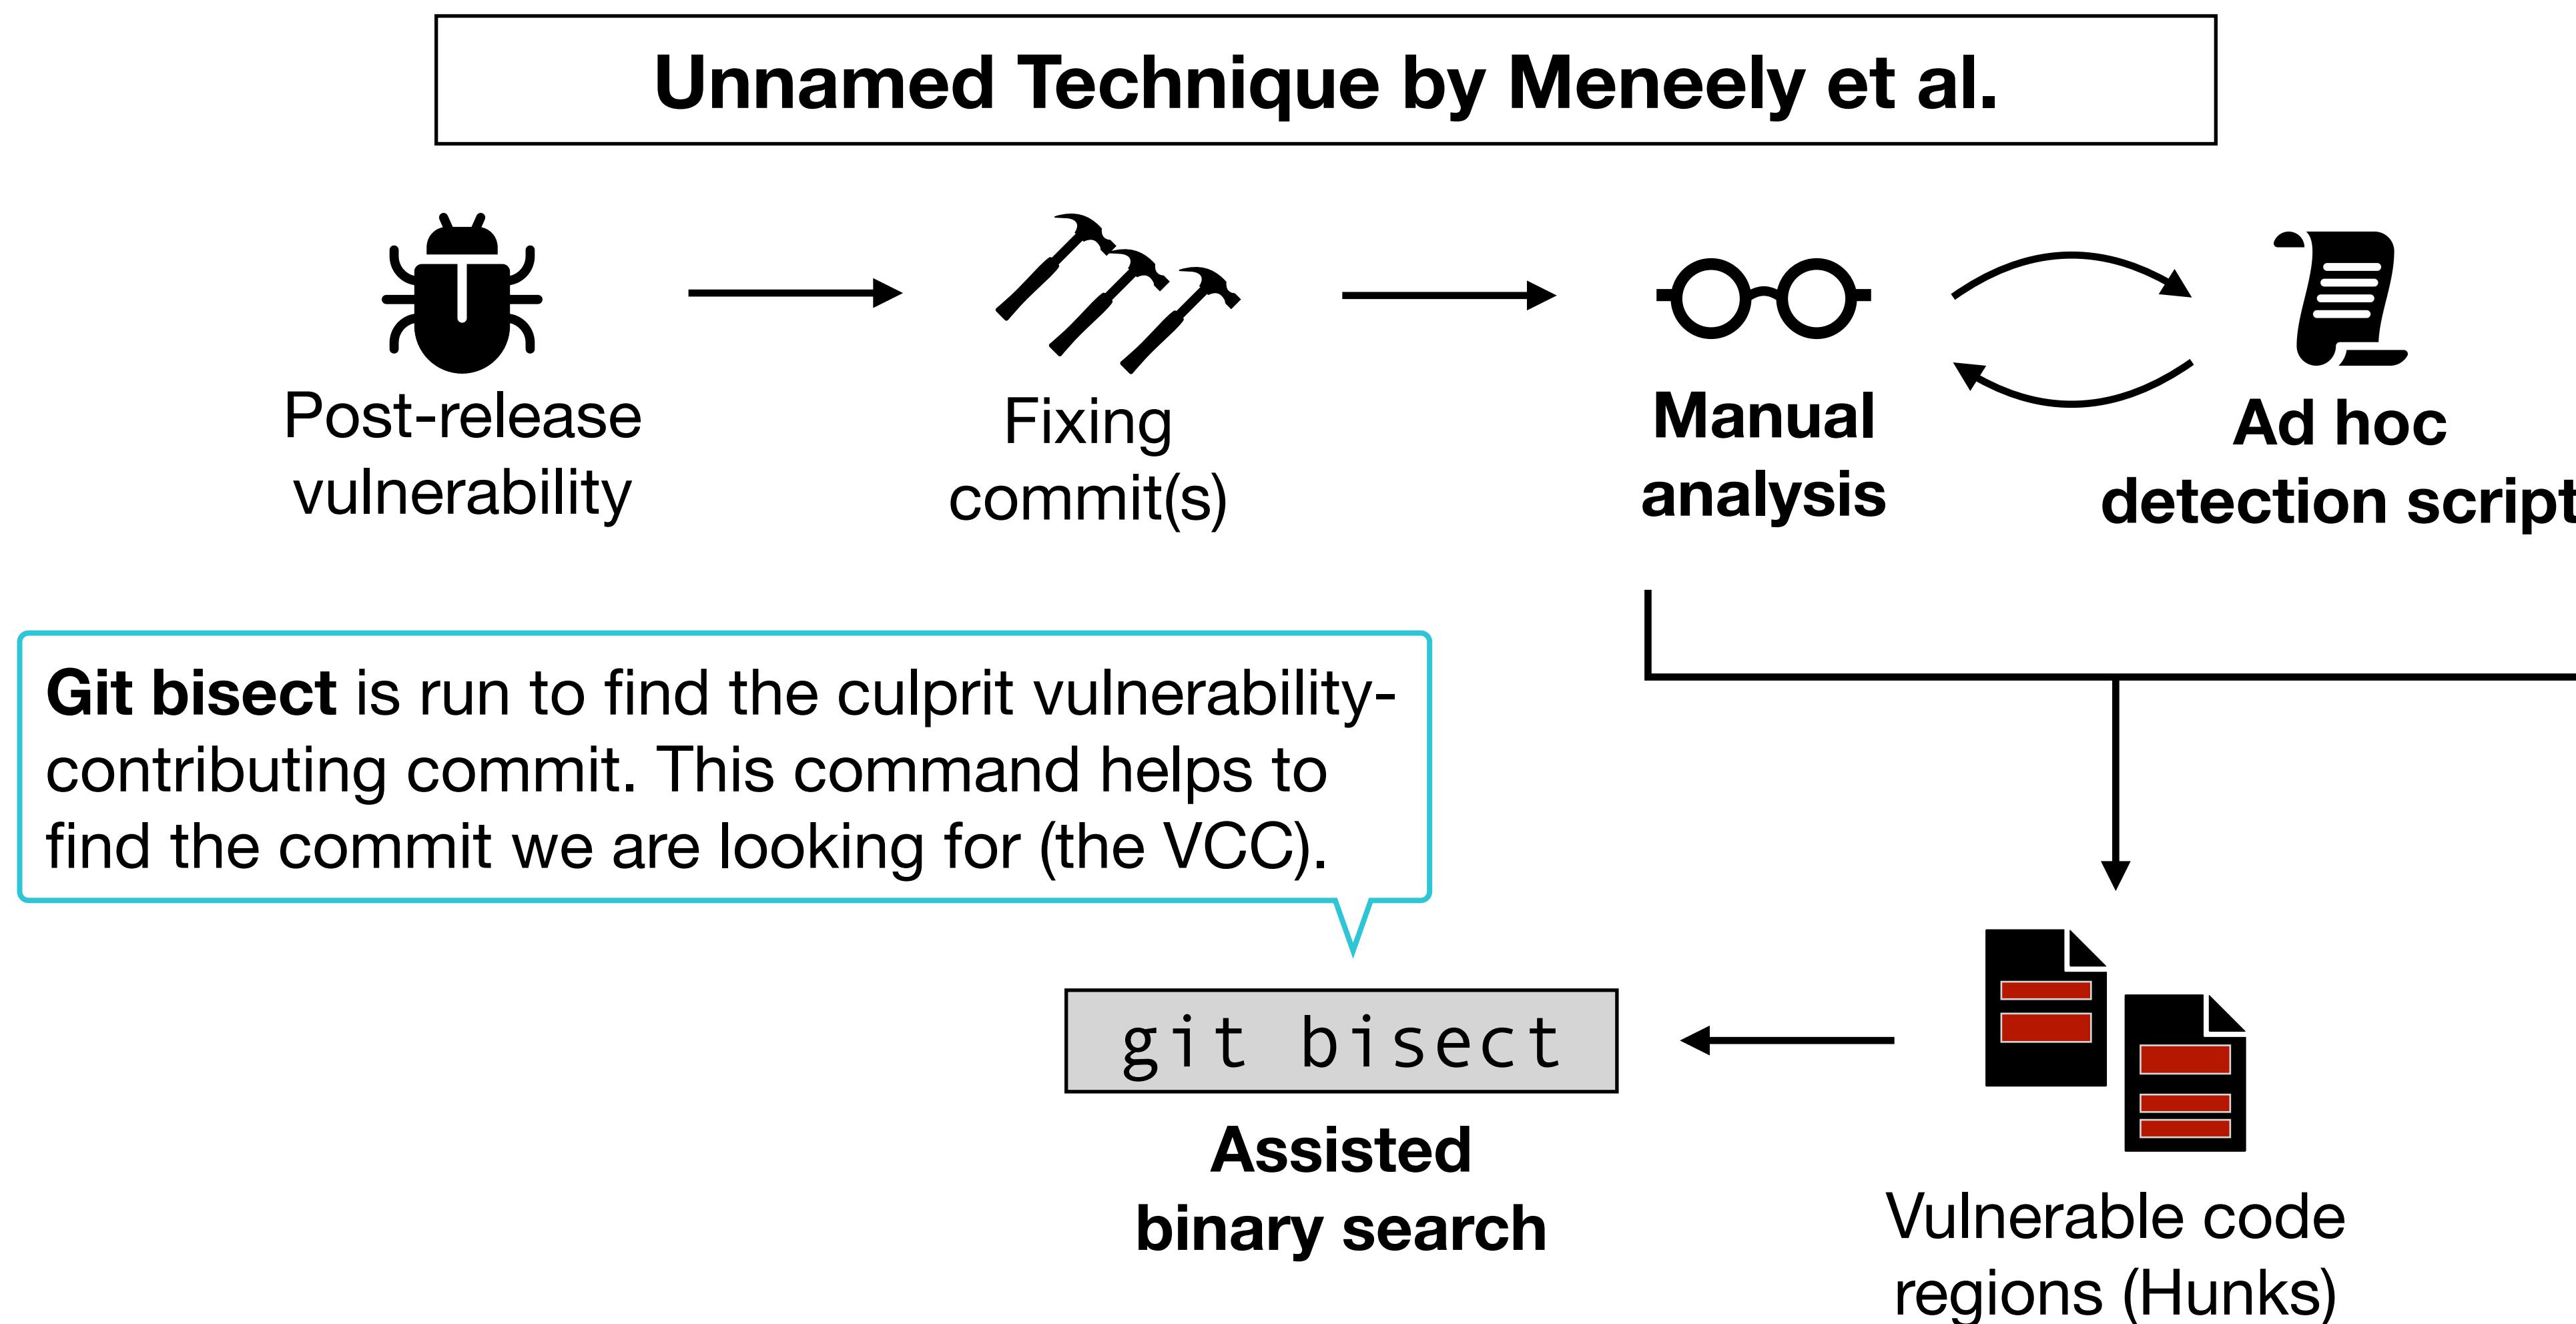
Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.



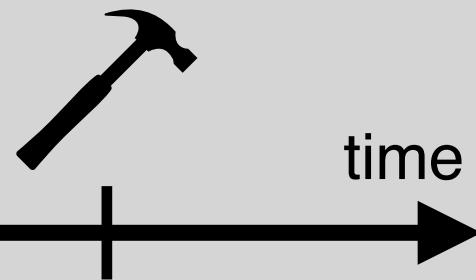
Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.

git bisect

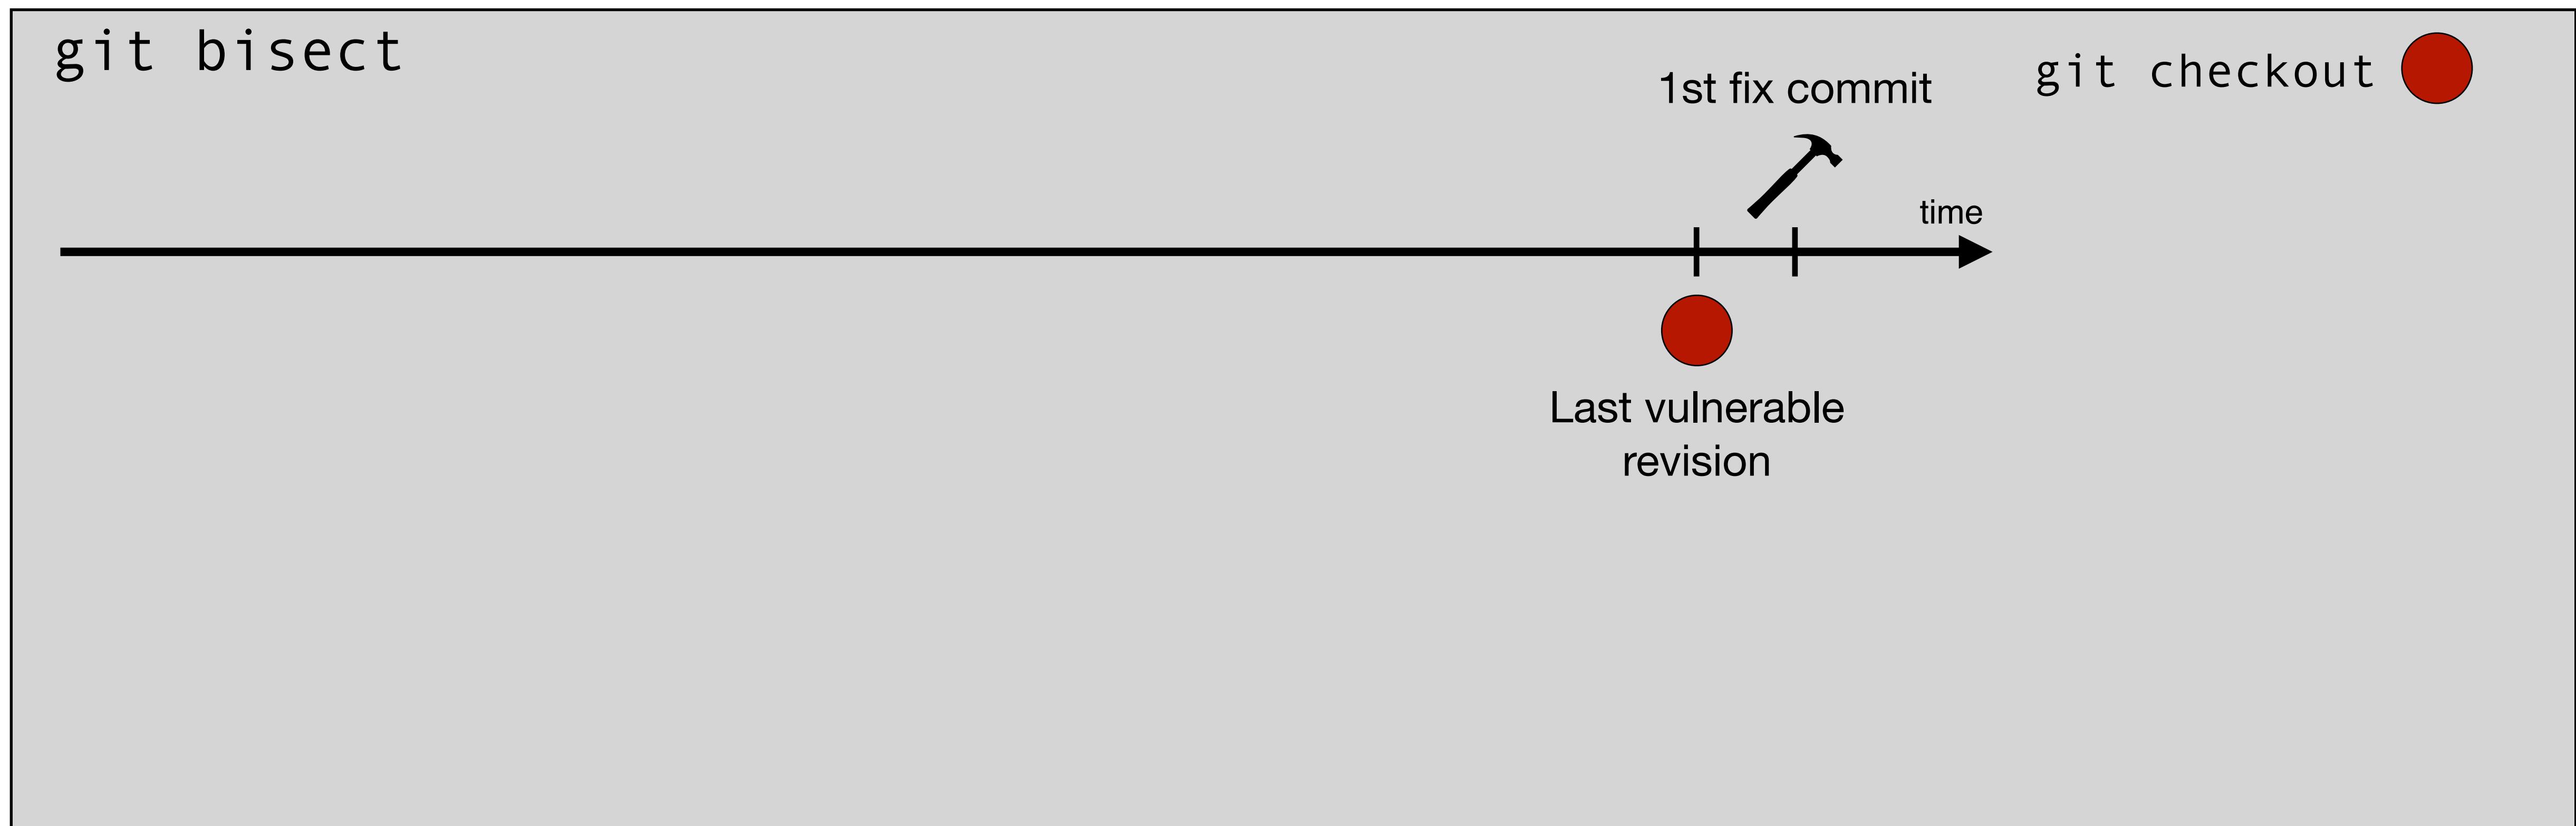
1st fix commit



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

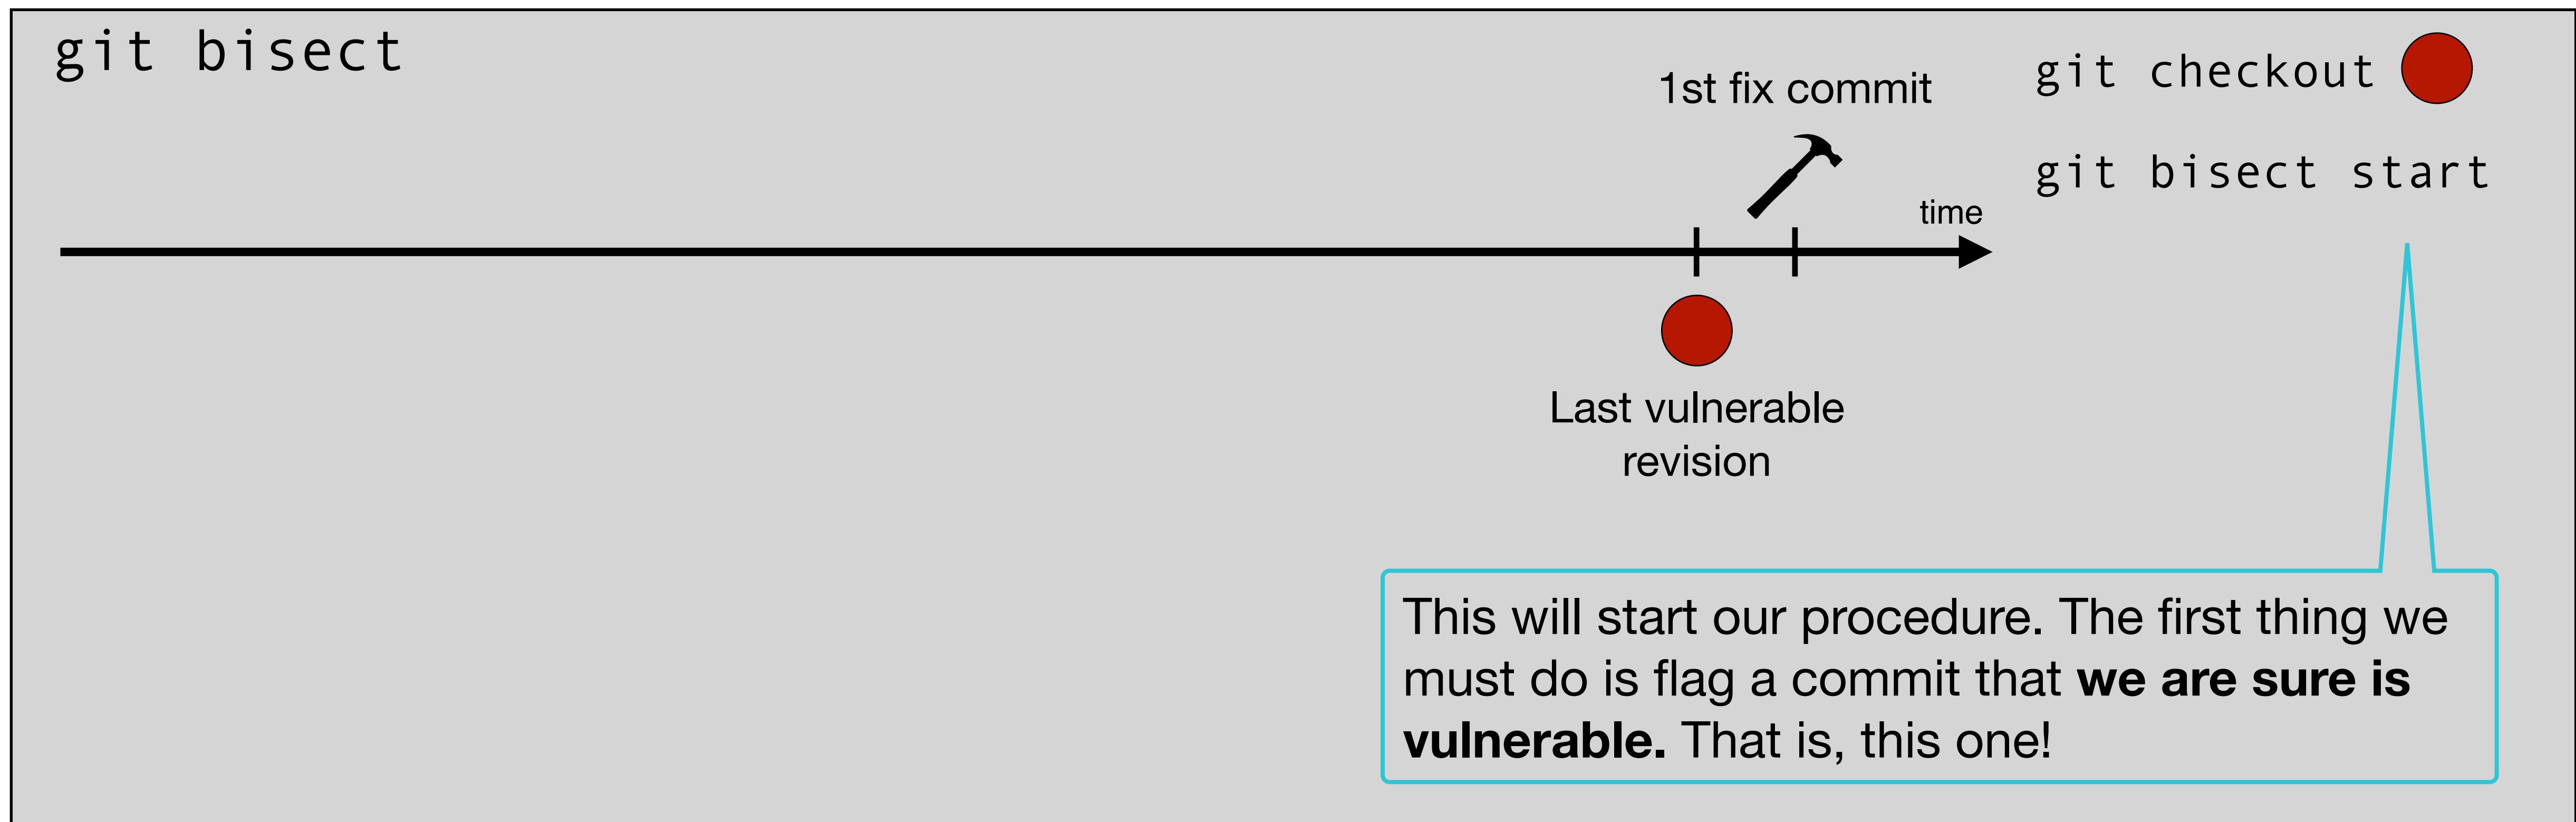
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

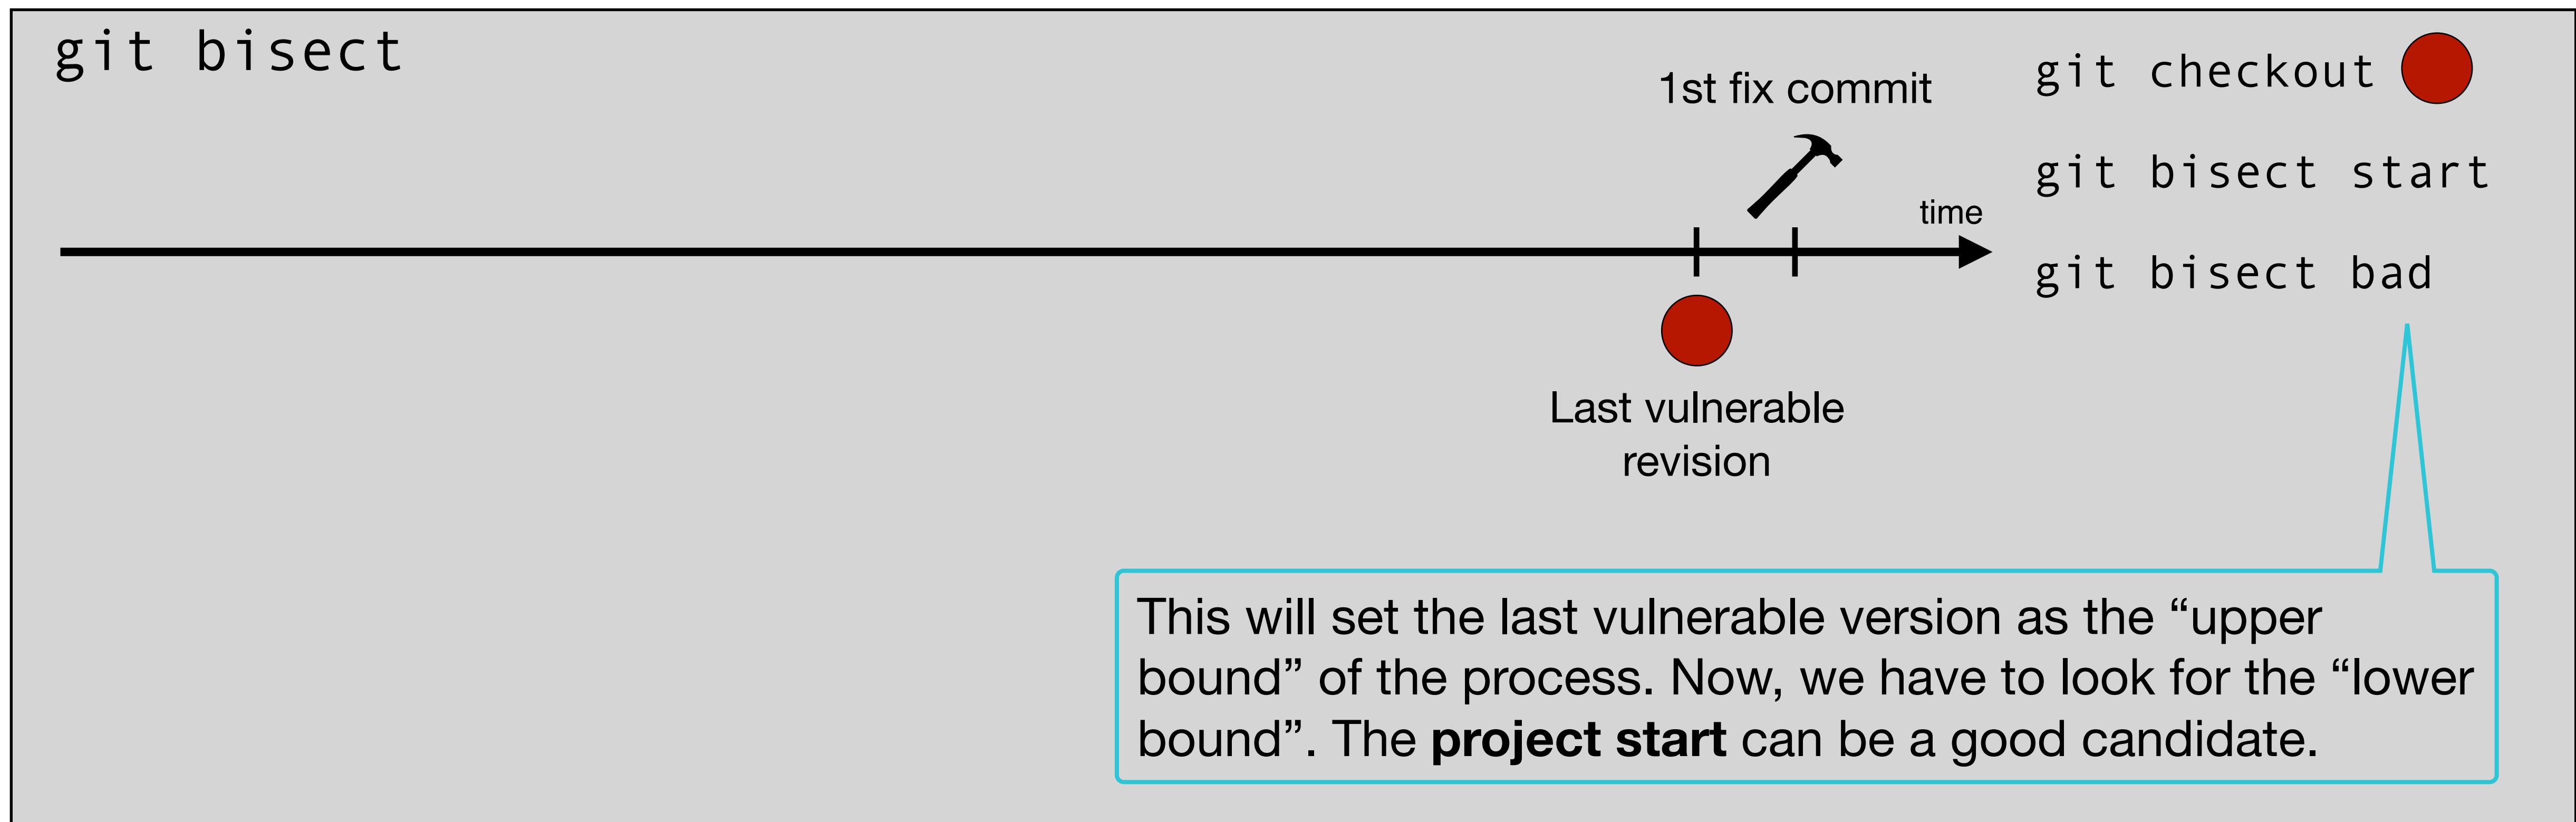
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

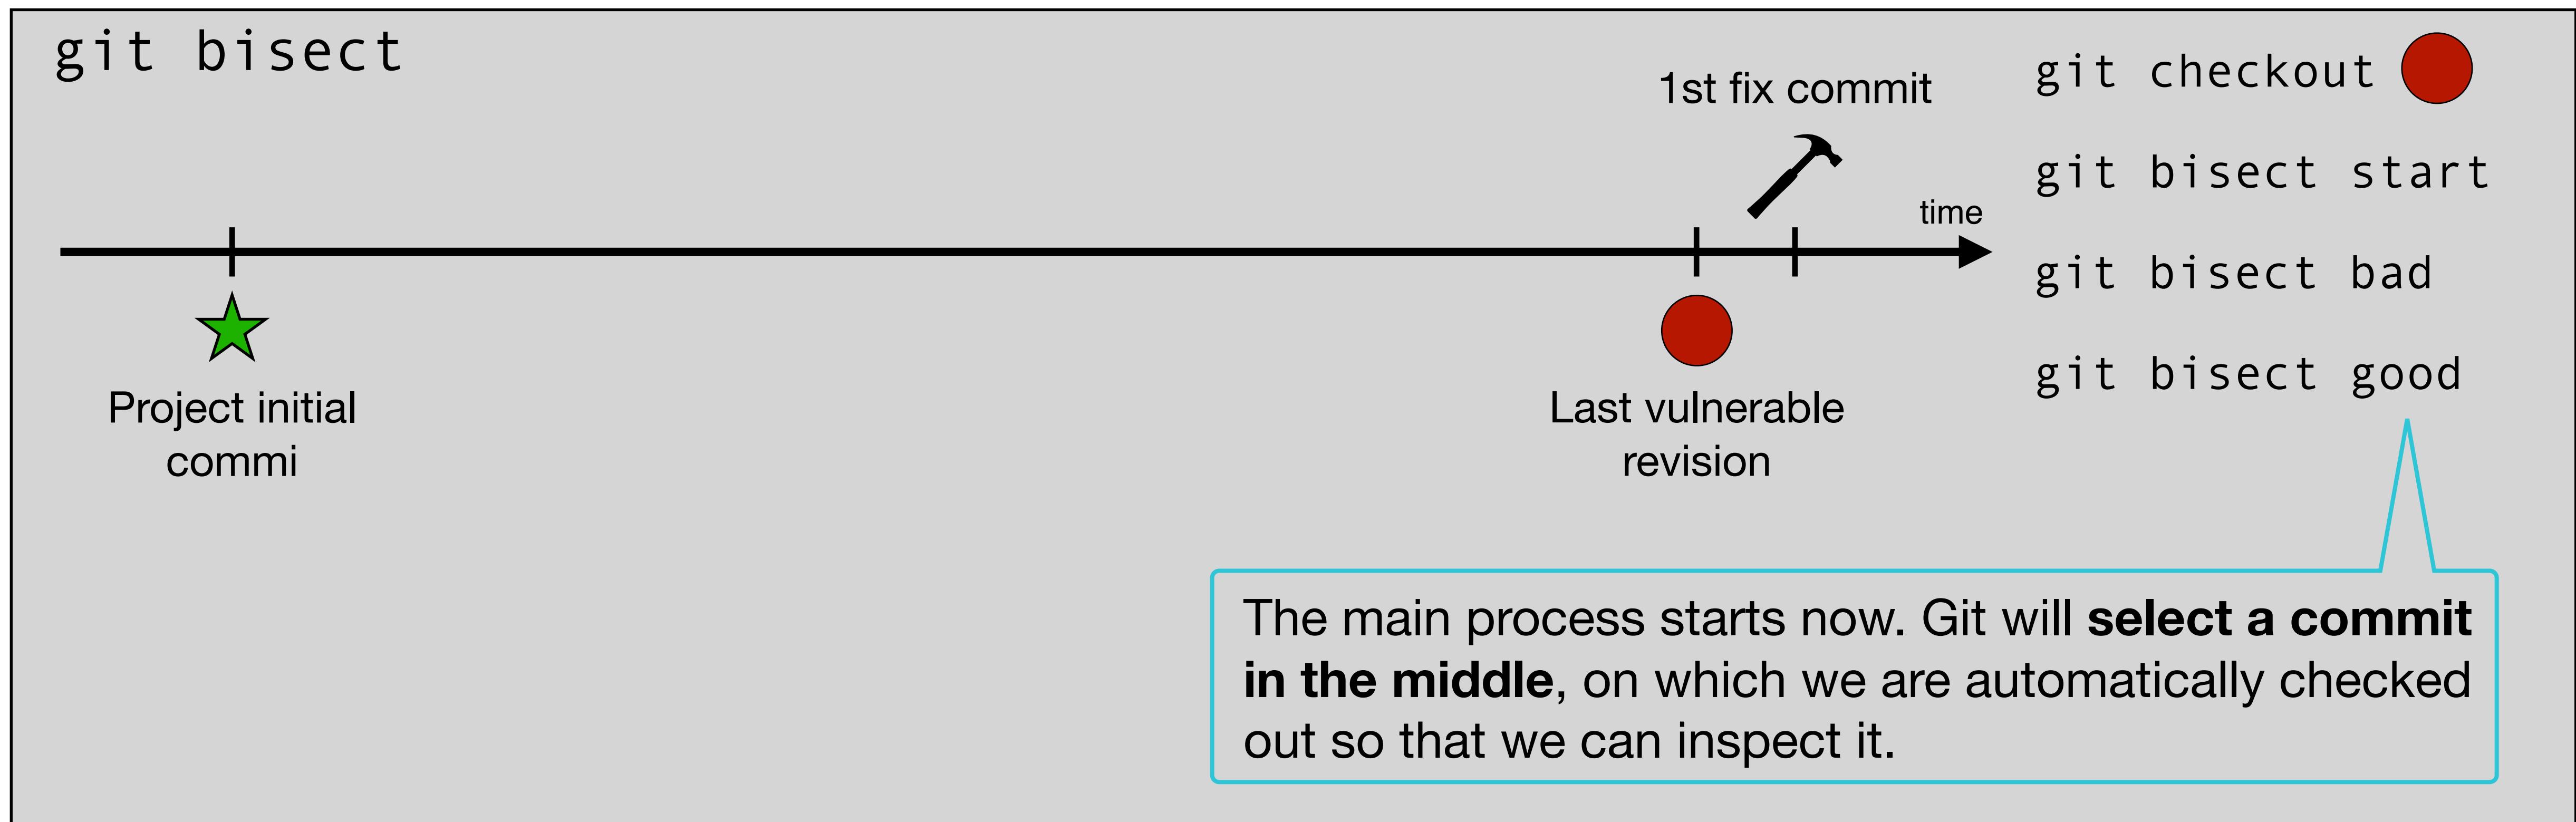
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

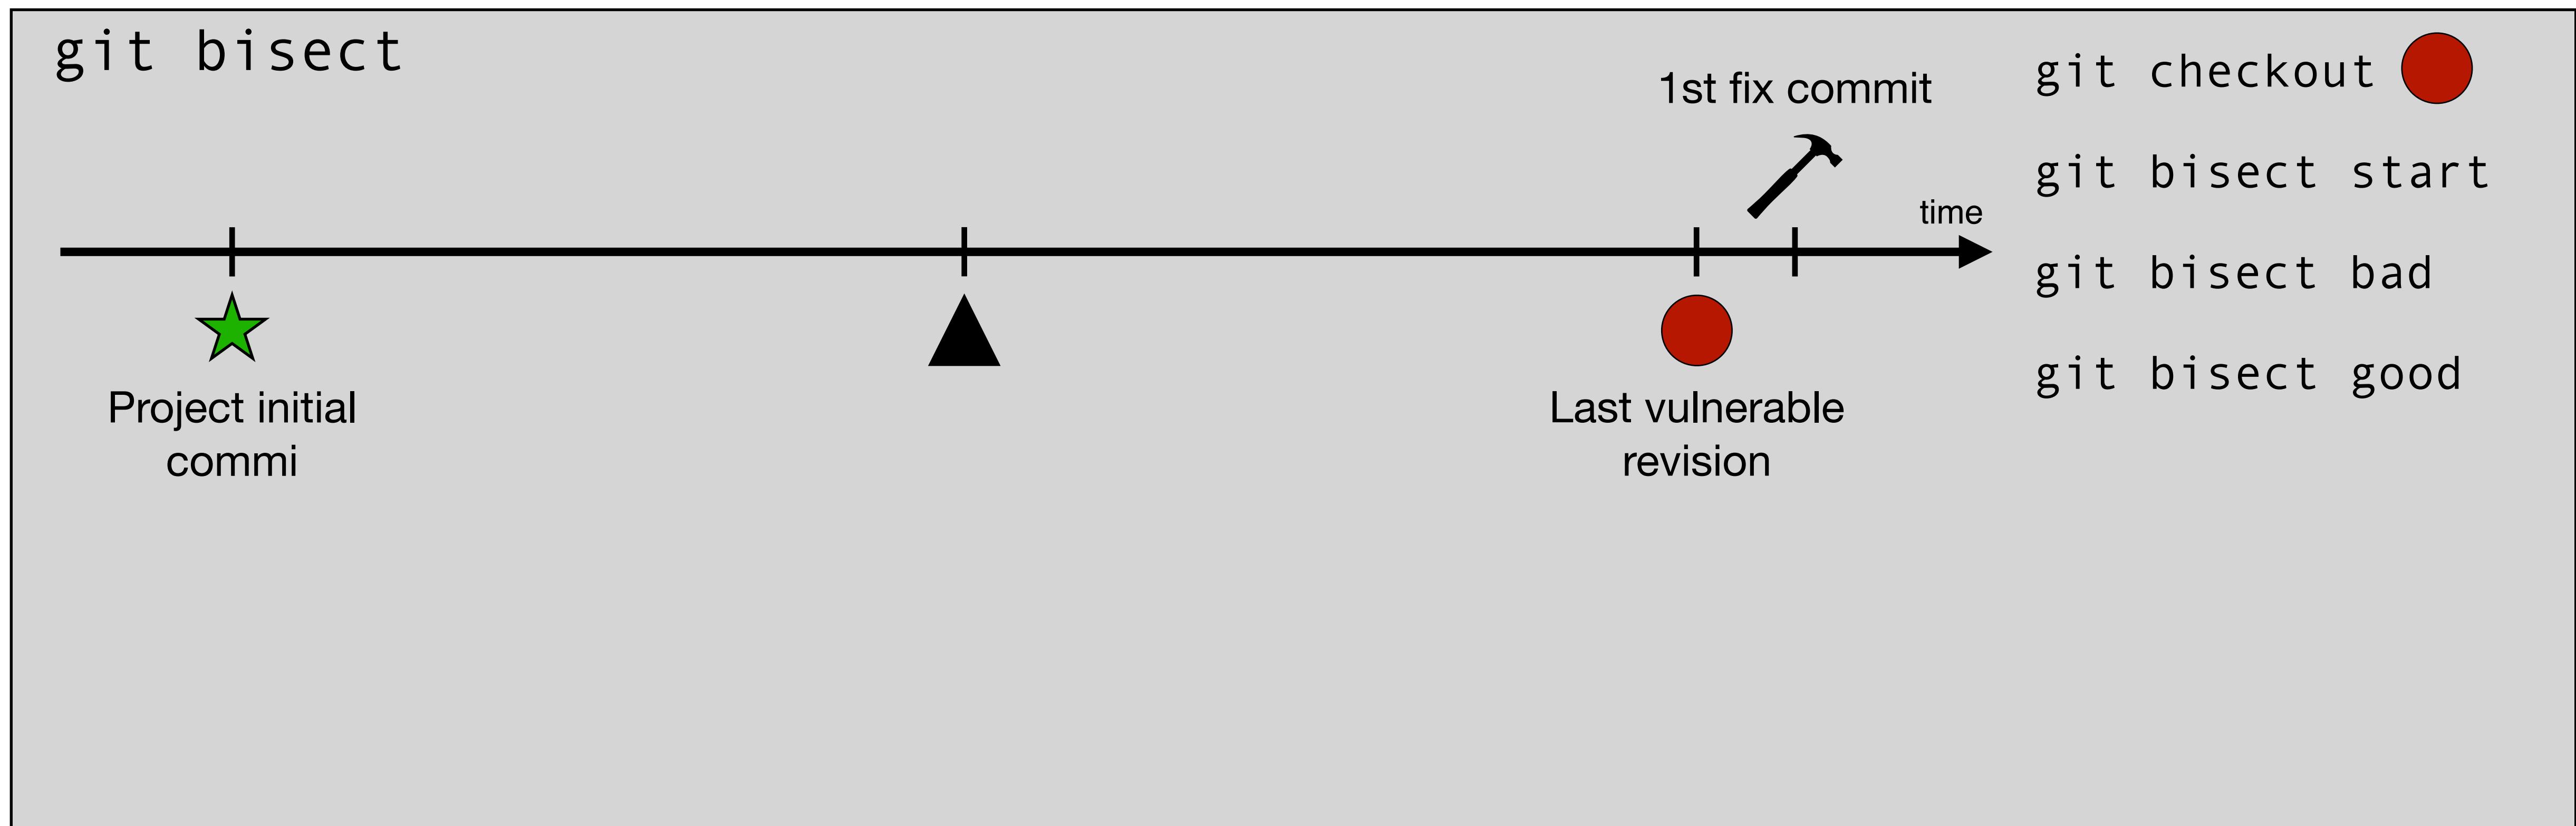
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

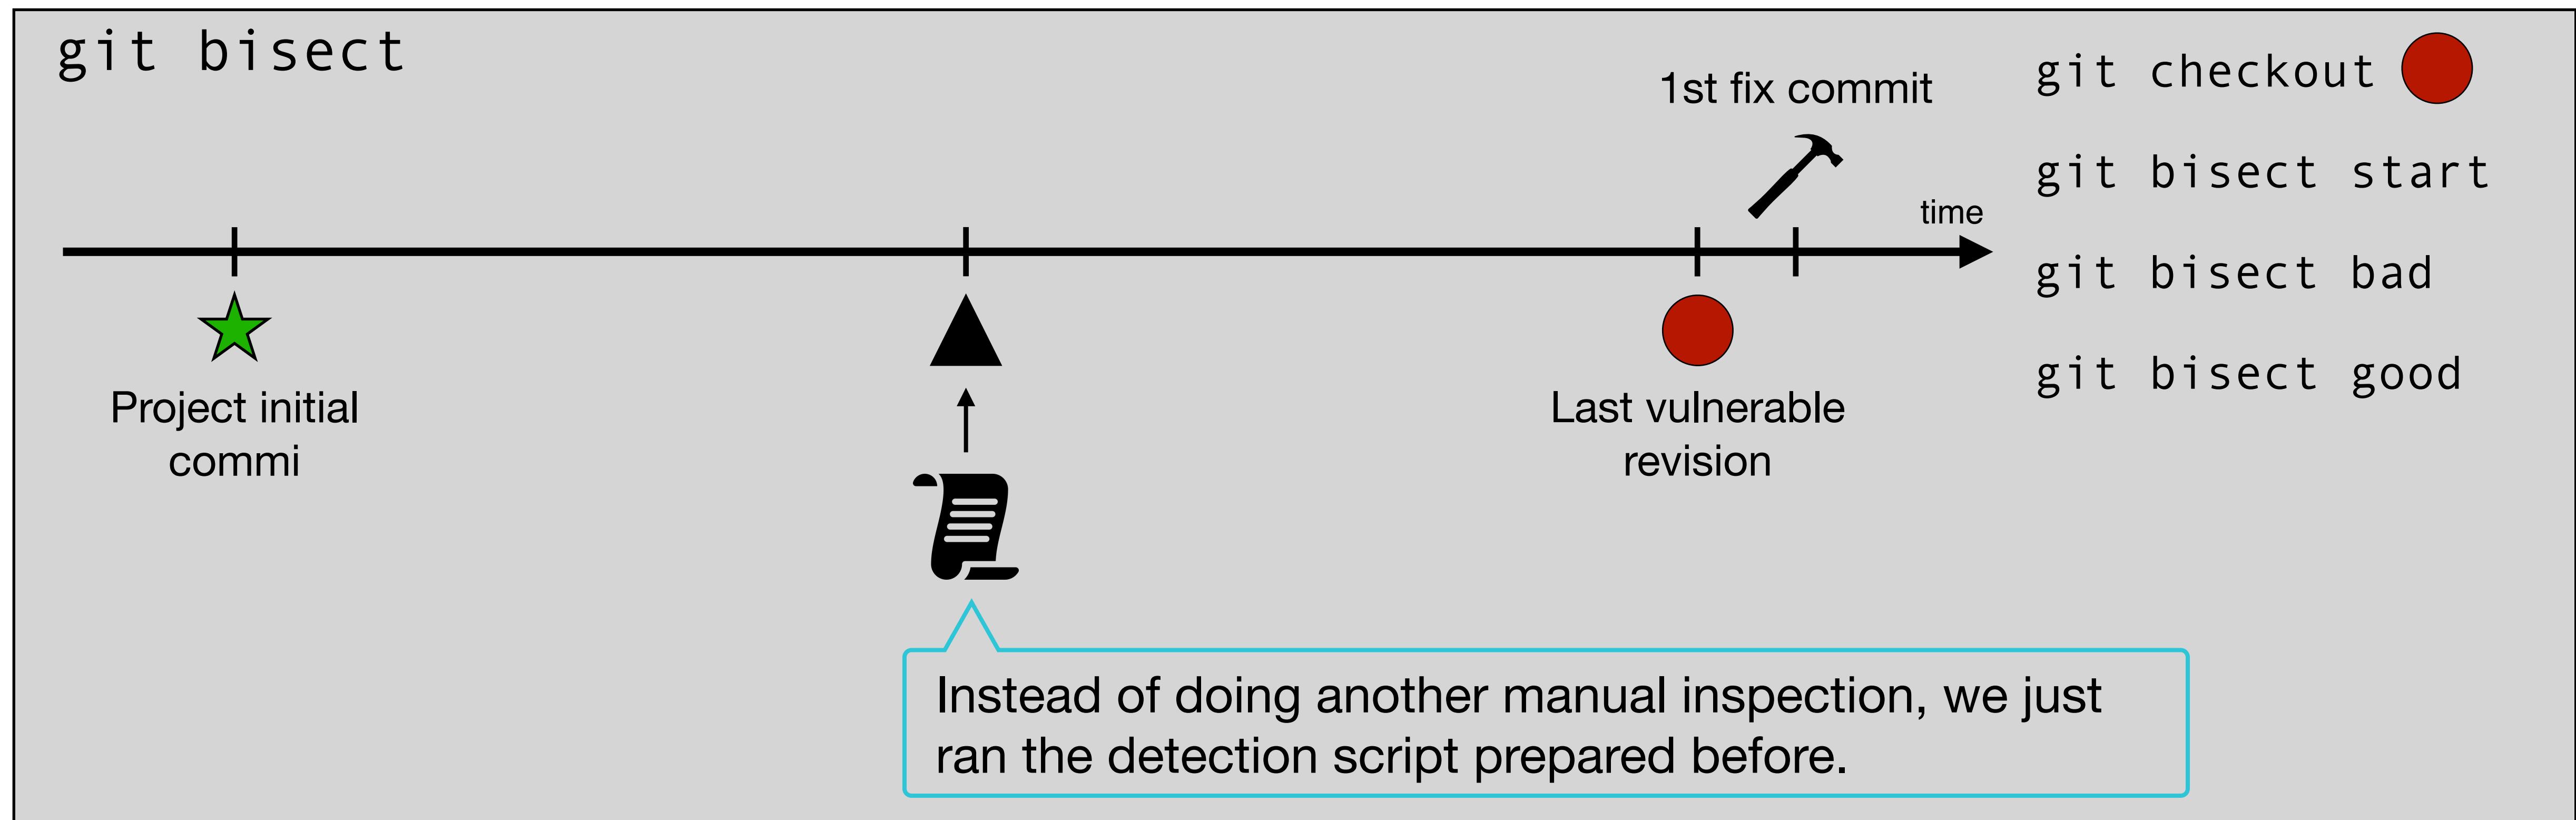
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

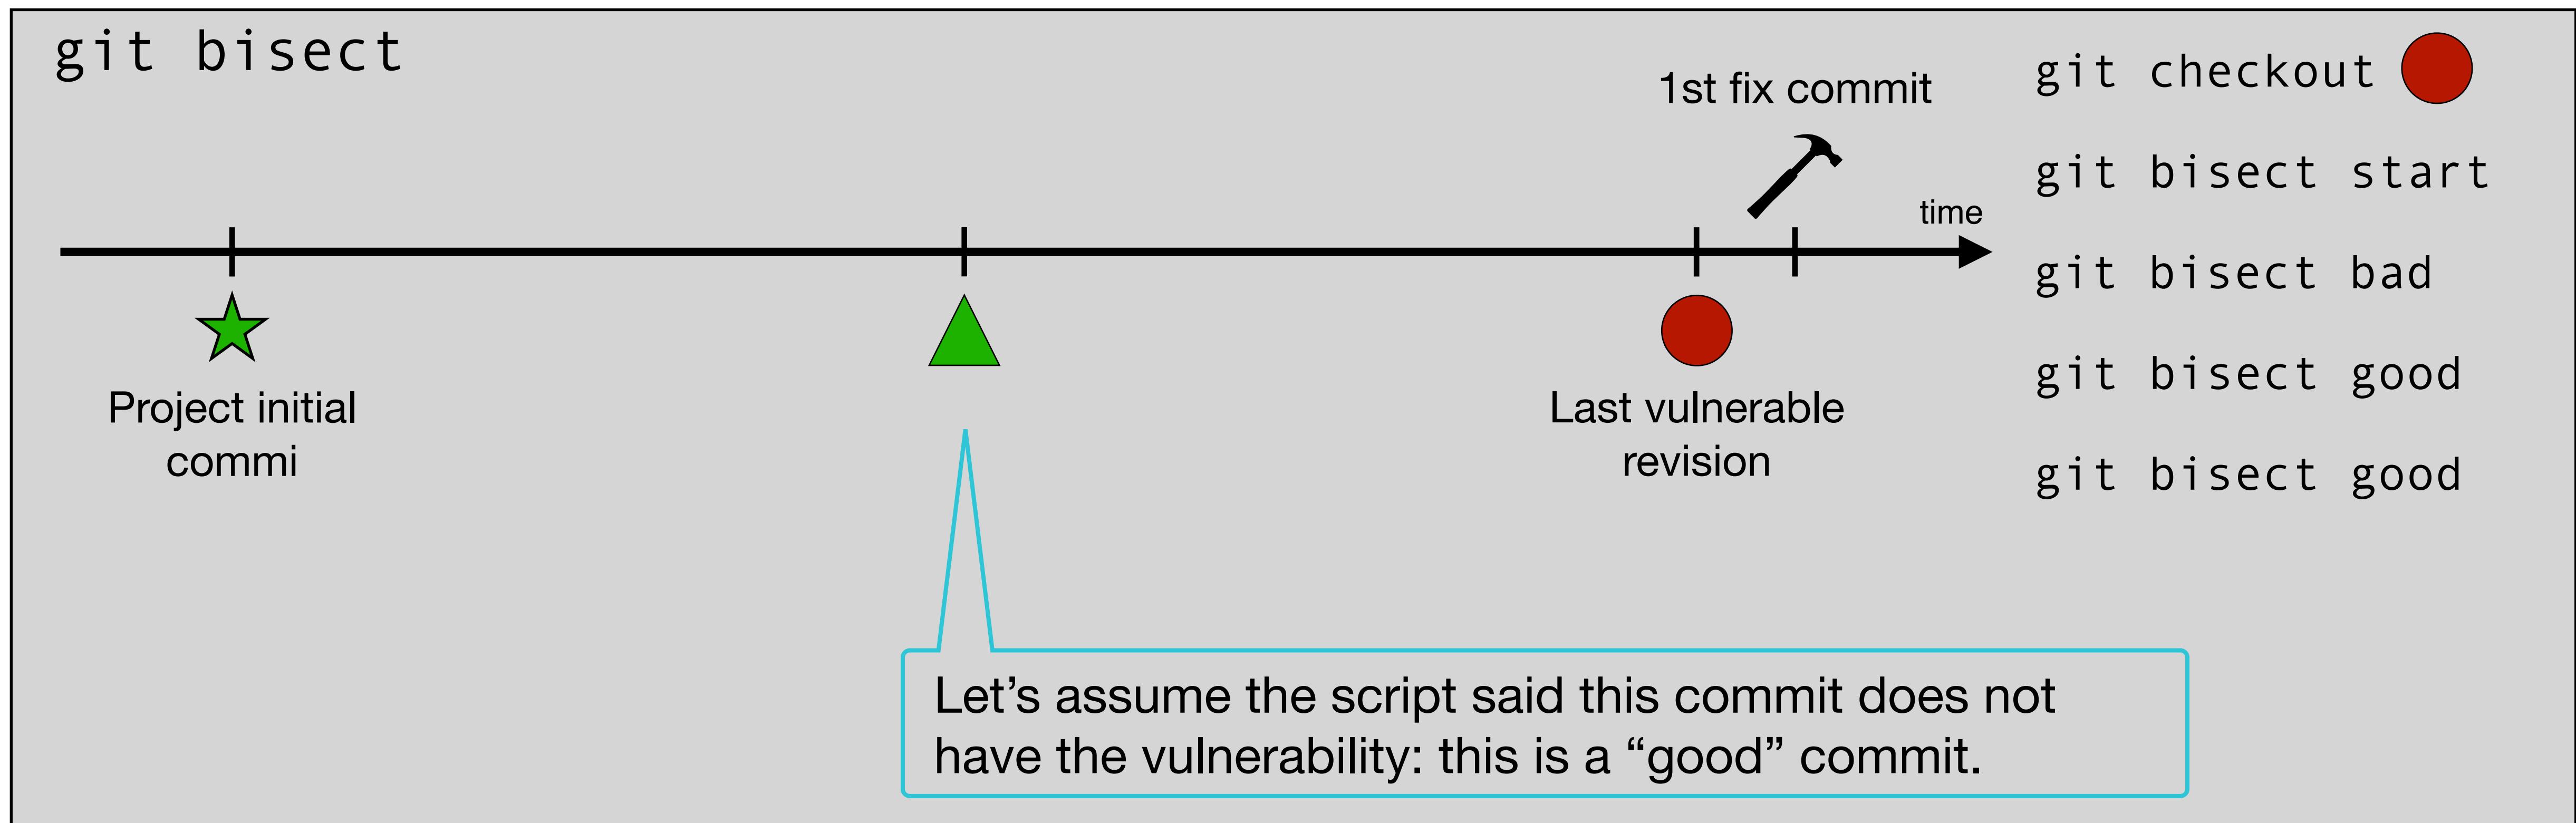
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

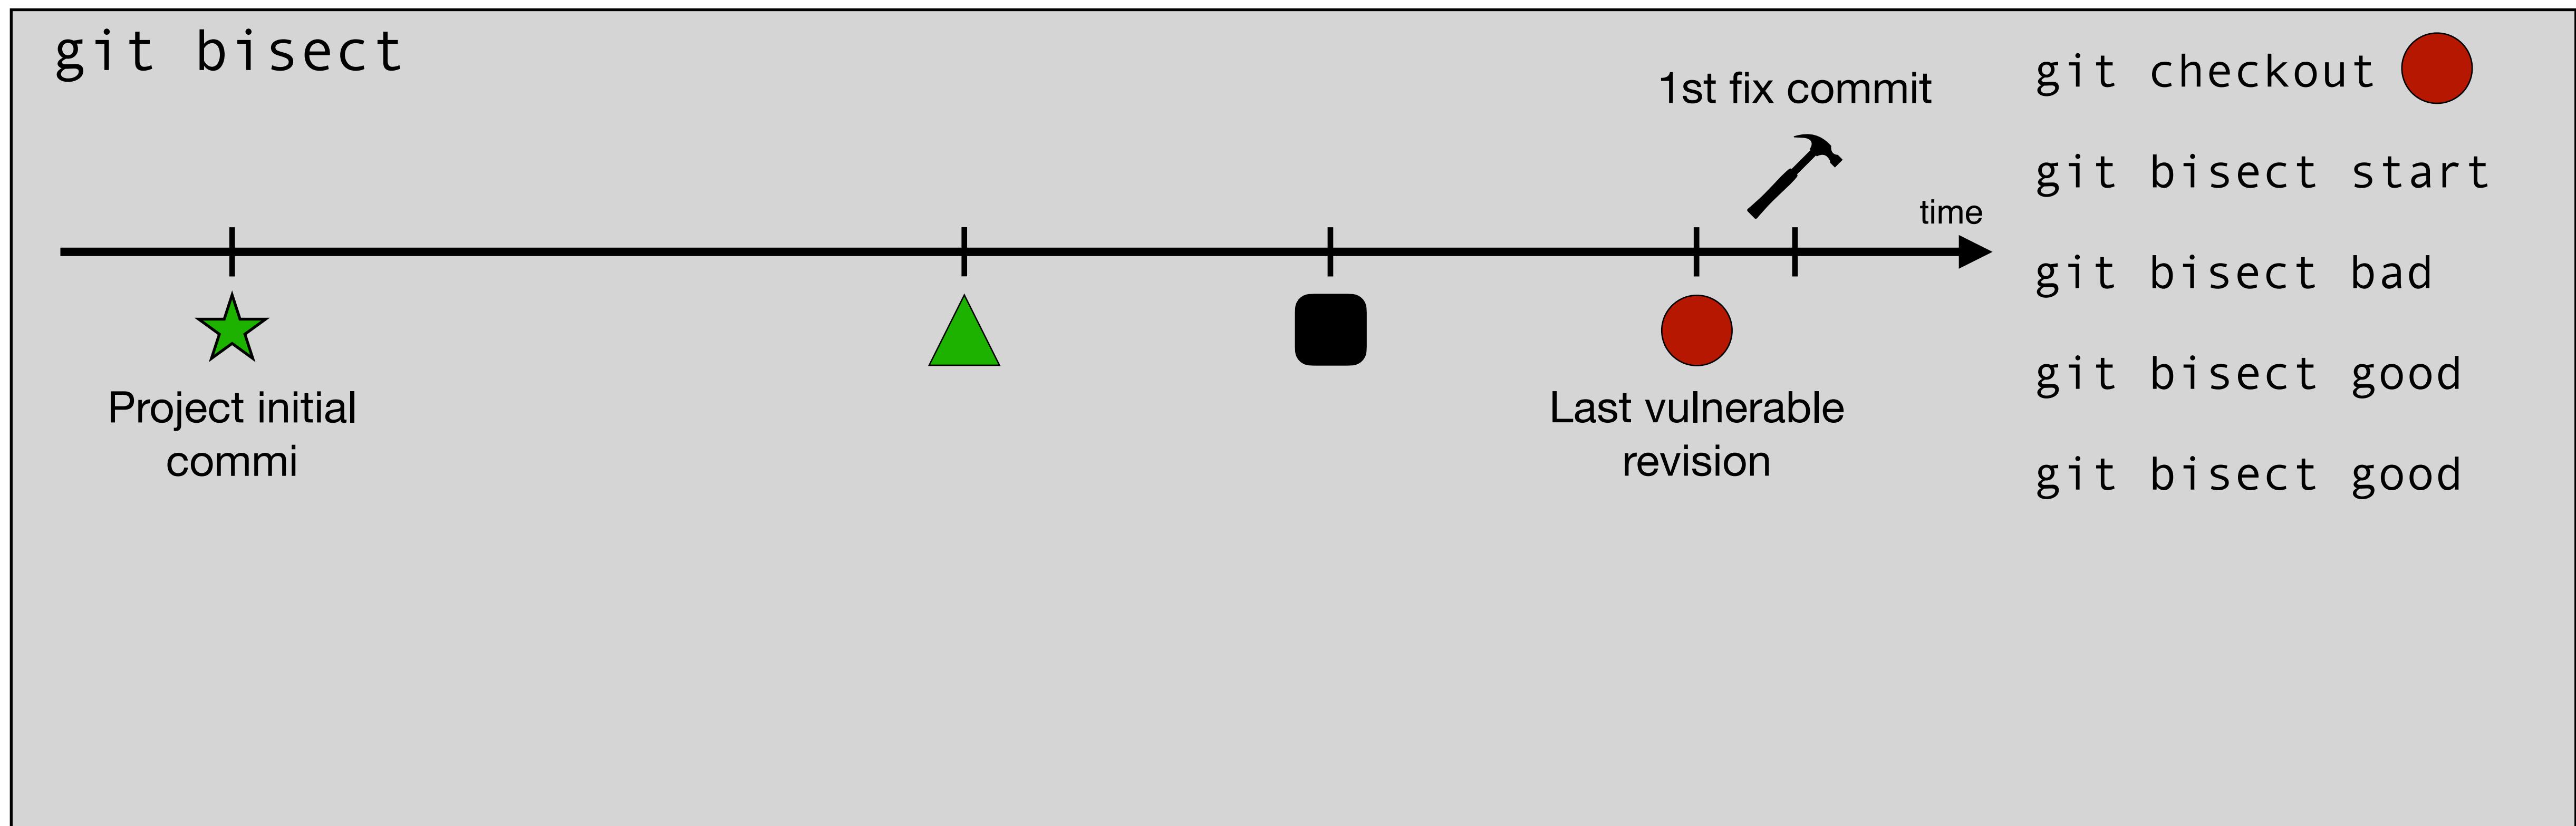
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

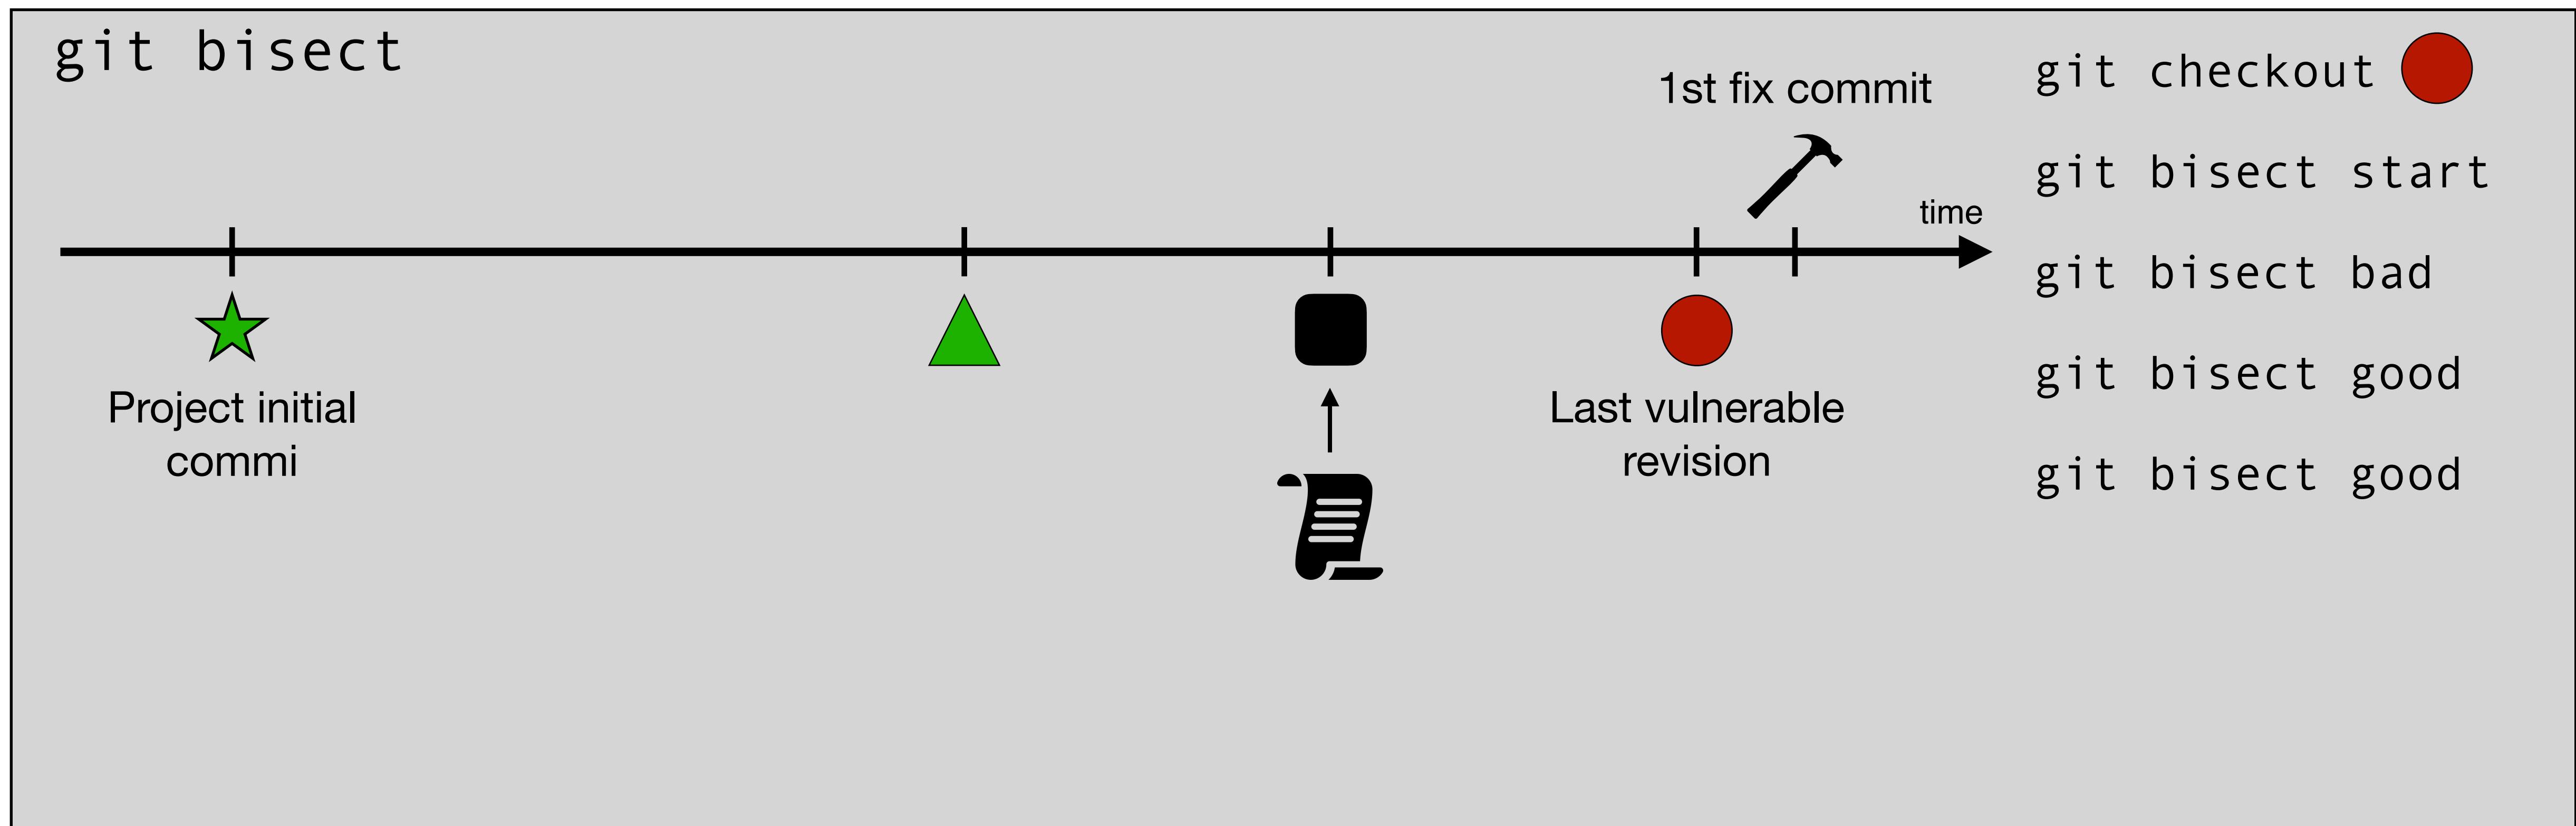
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

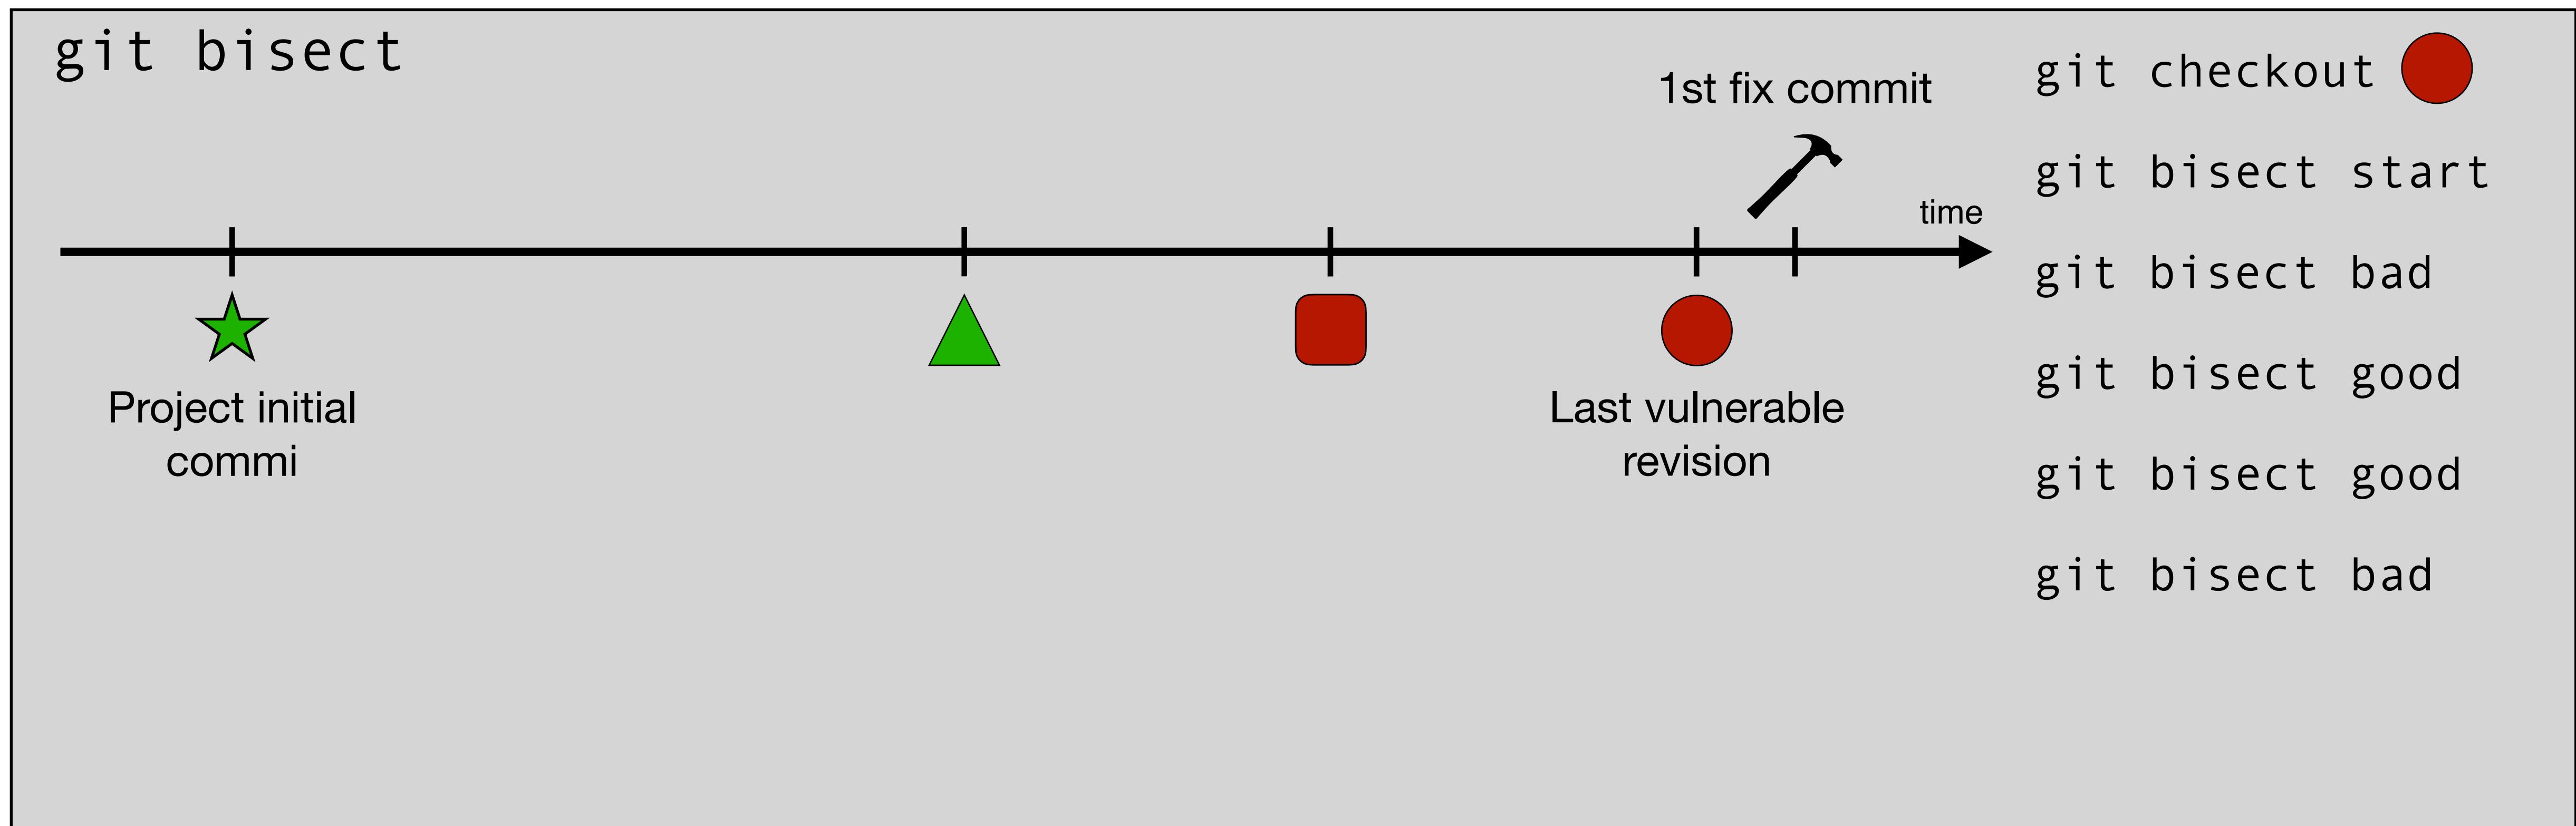
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

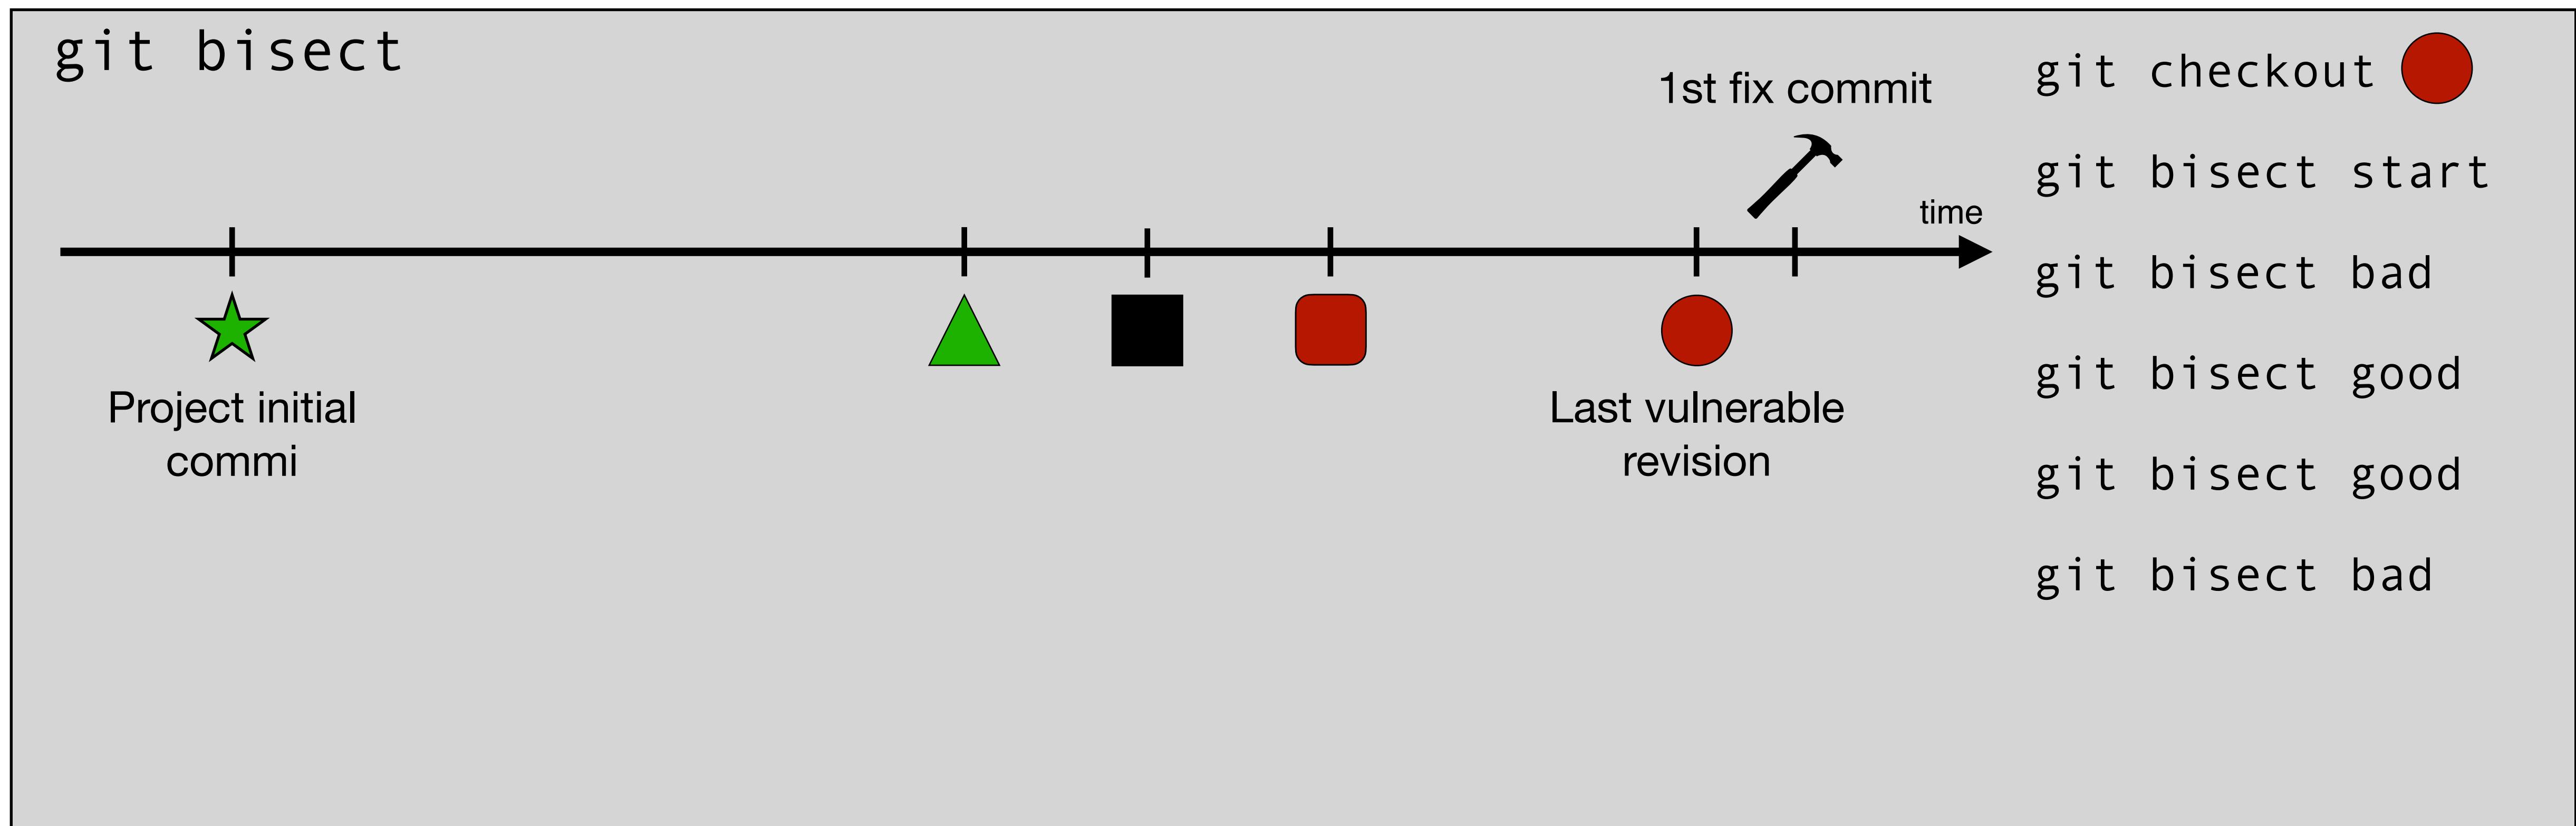
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

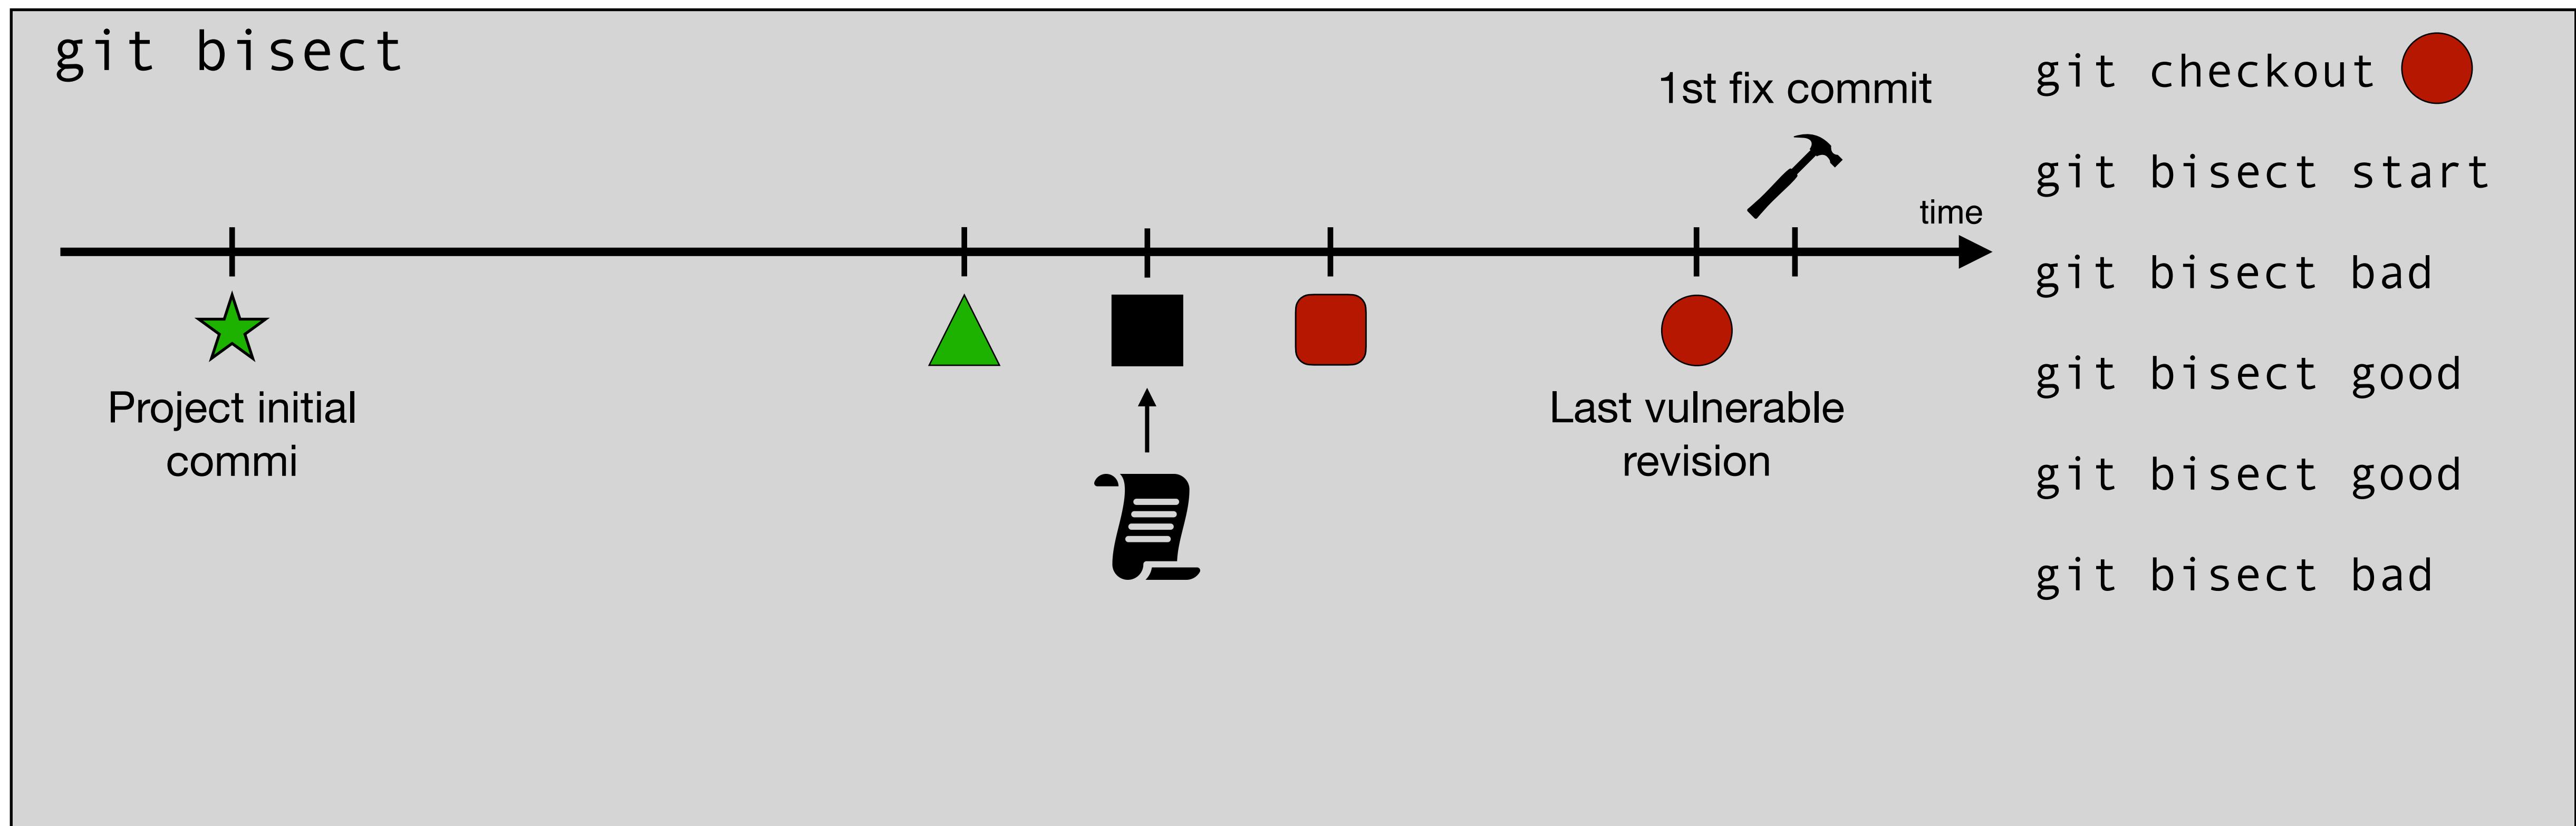
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

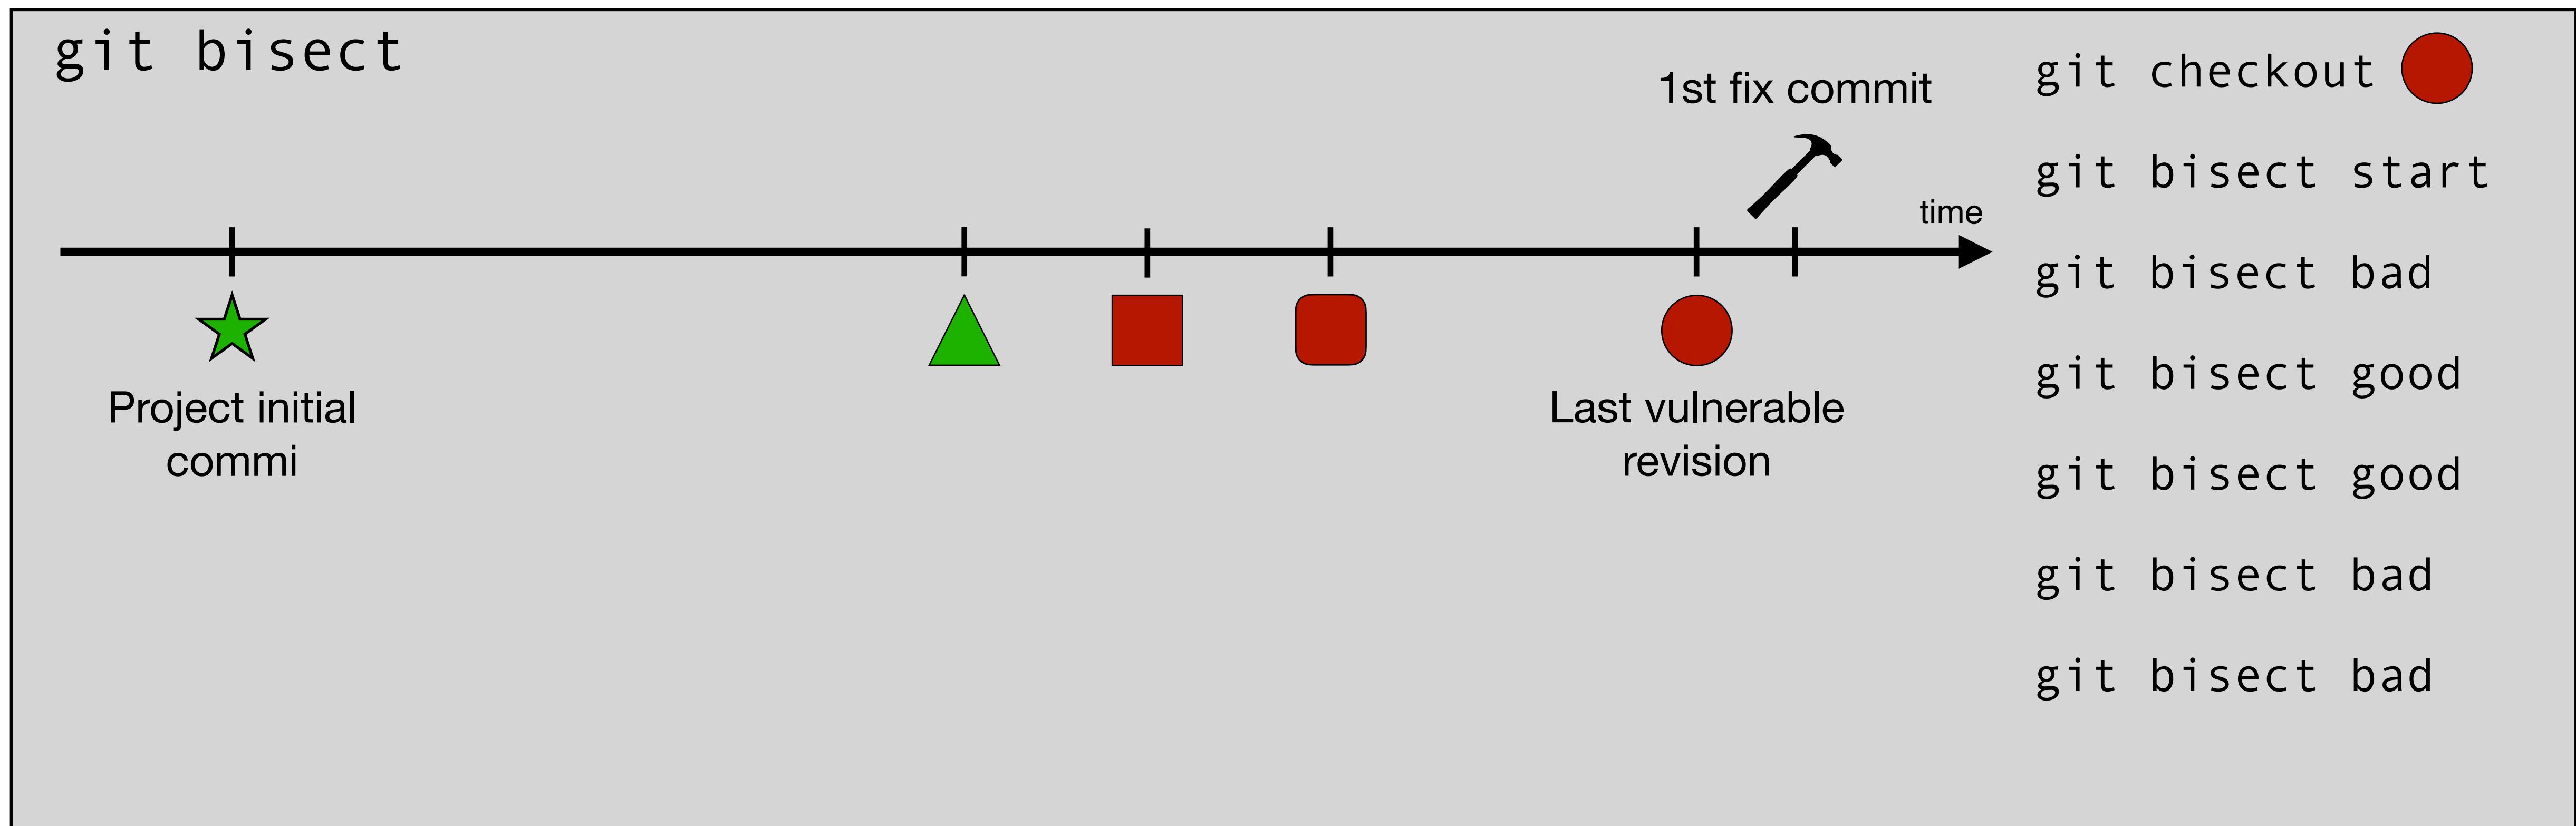
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

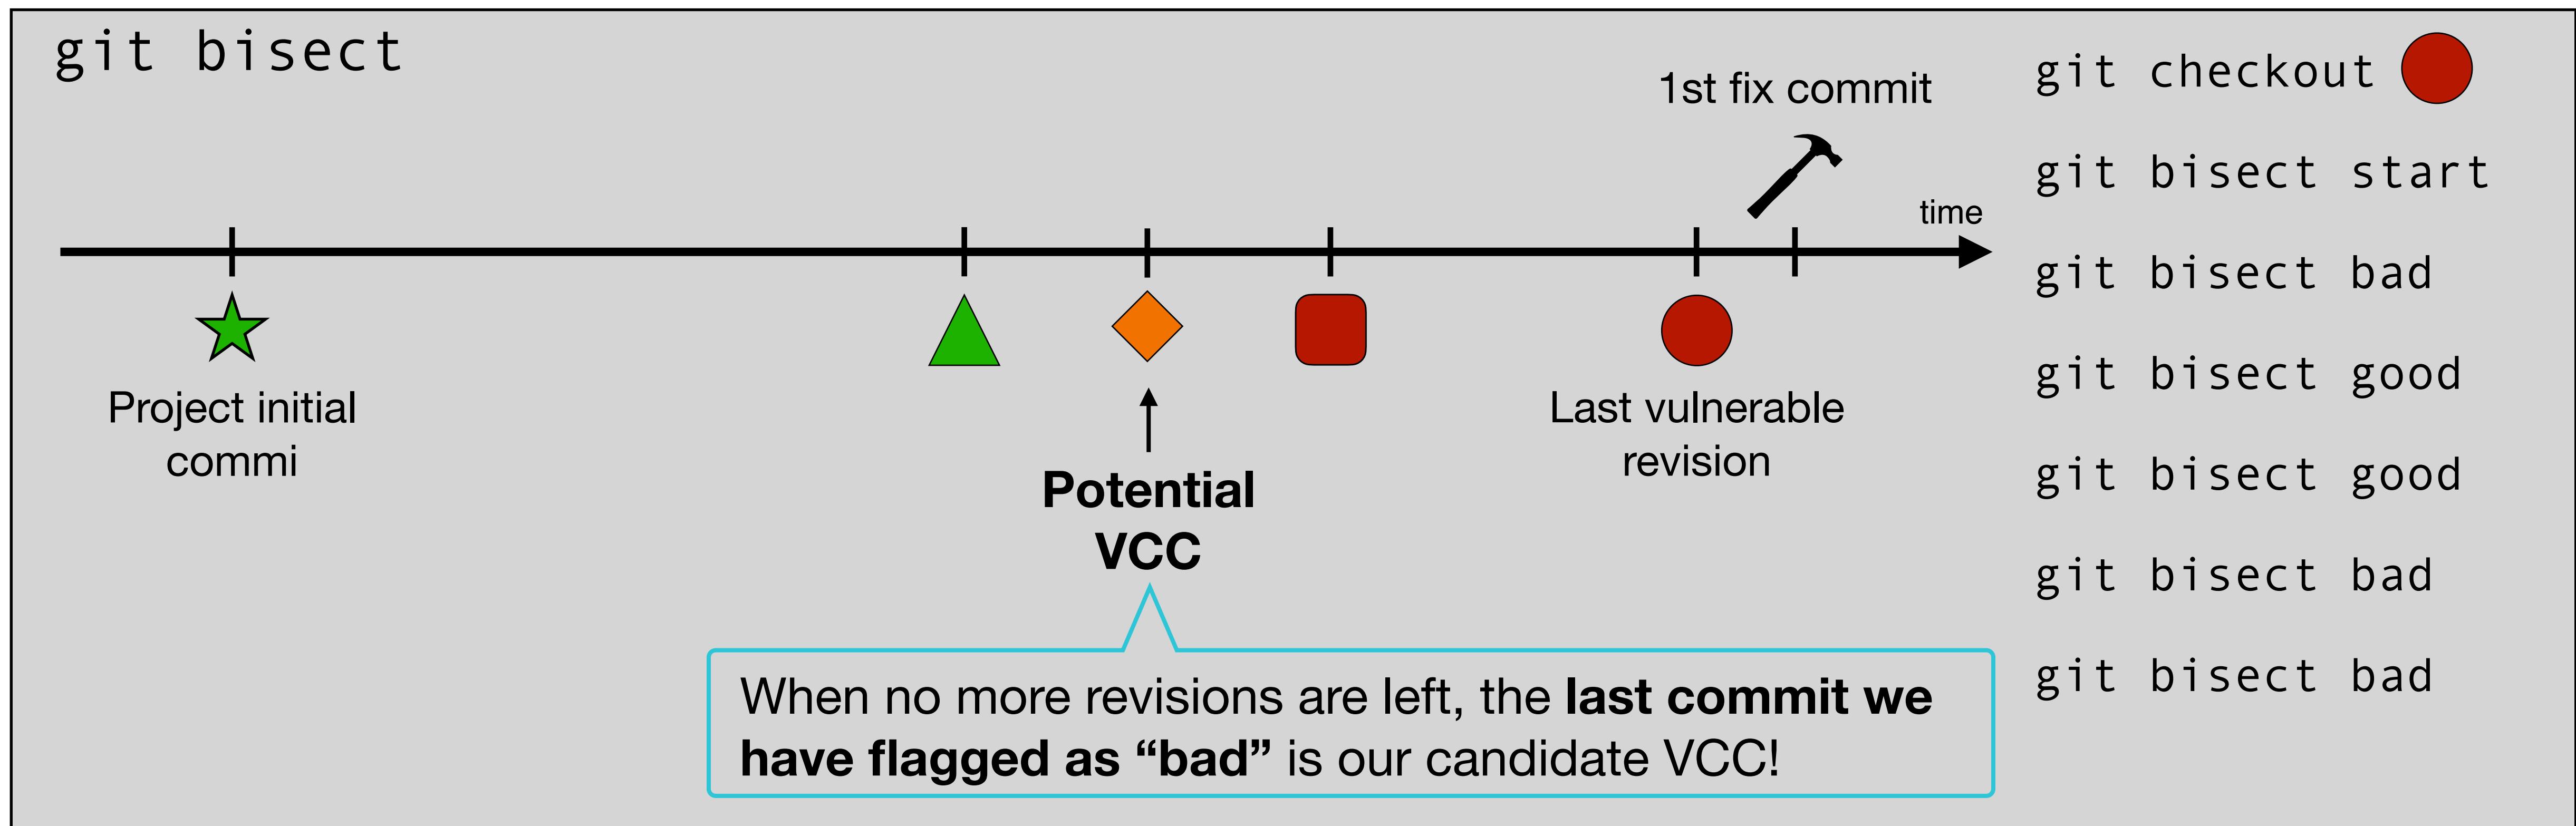
Unnamed Technique by Meneely et al.



Mining VCCs: A First Approach

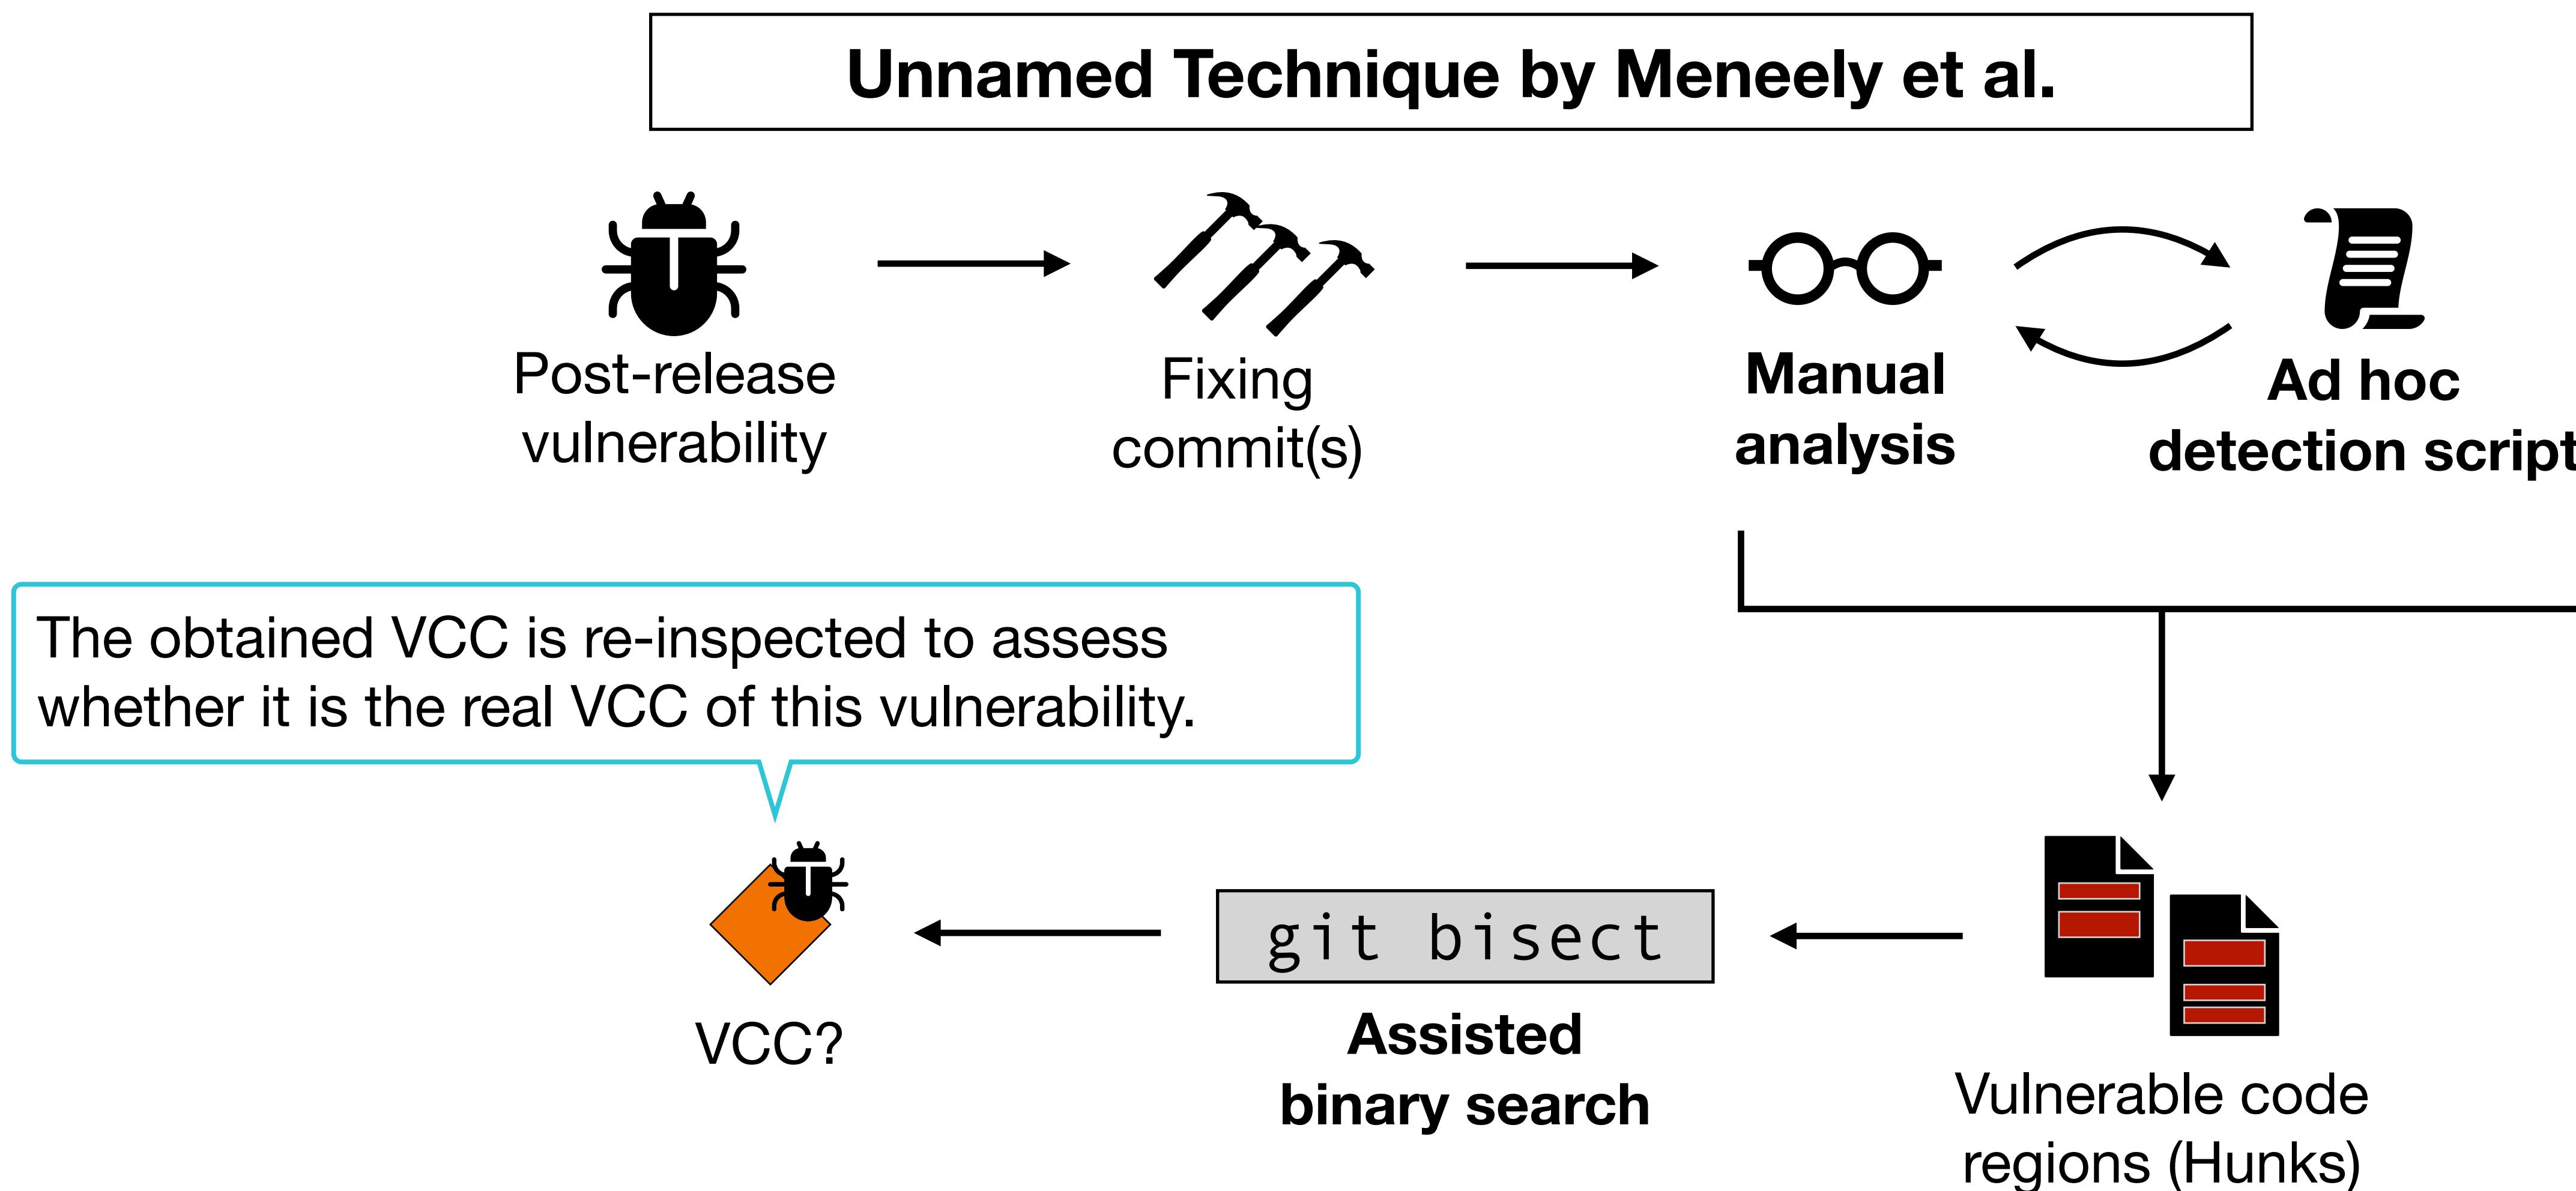
Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.



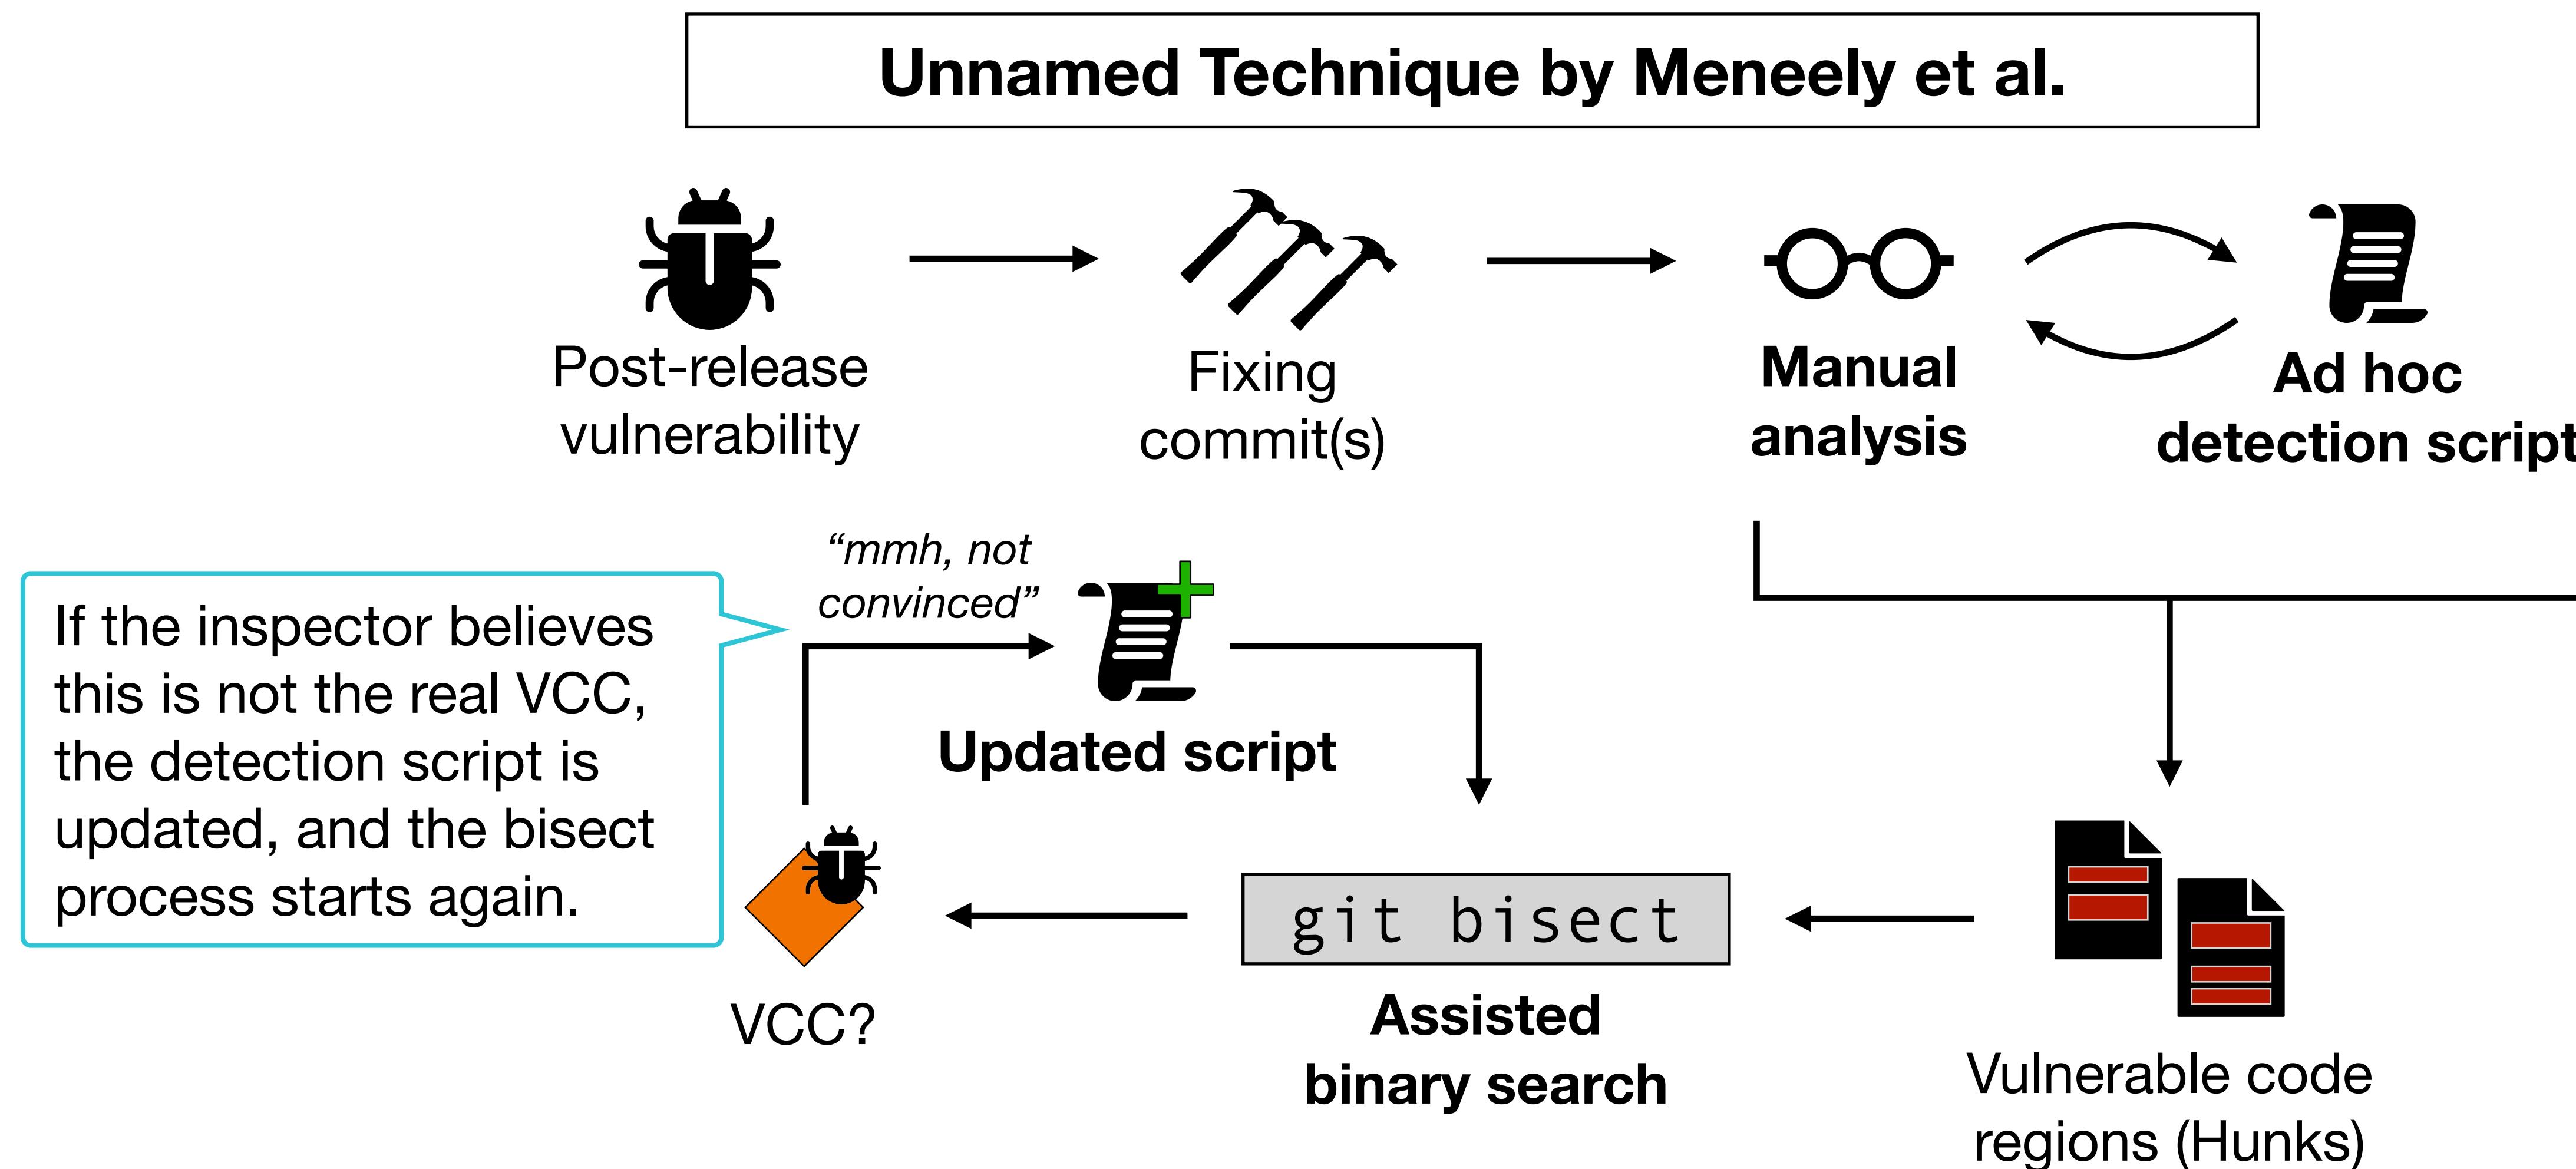
Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.



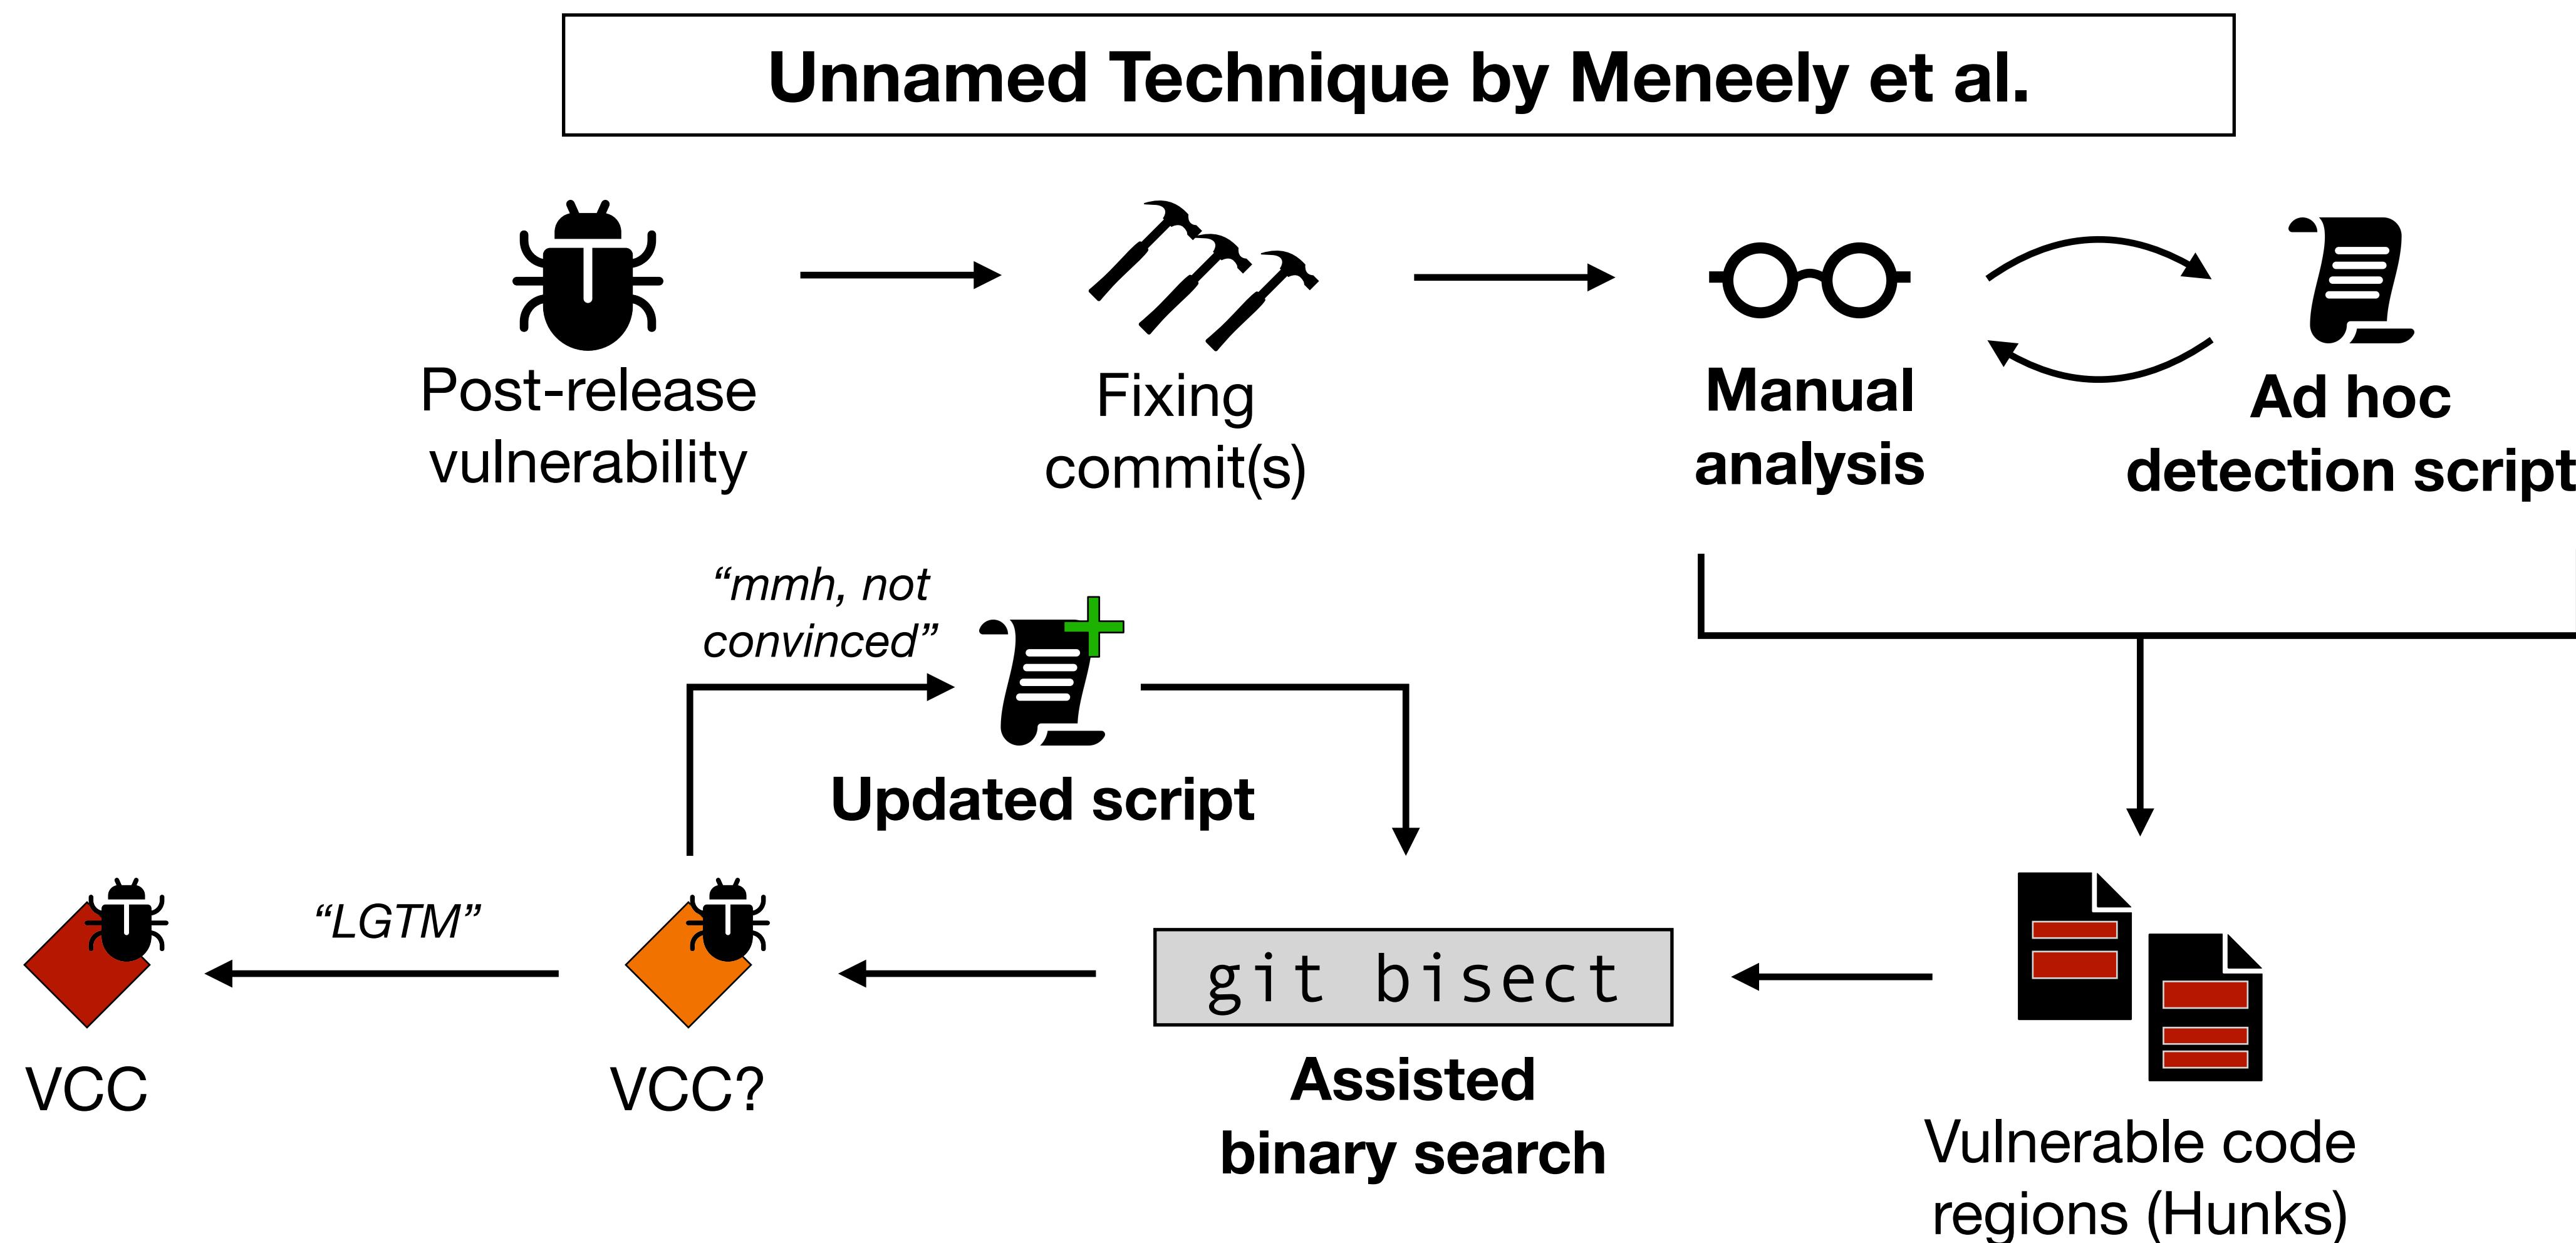
Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories



Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.



Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!

Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!

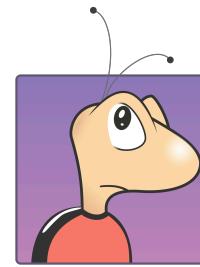
Śliwerski, Zimmermann, Zeller (SZZ)



Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!

Śliwerski, Zimmermann, Zeller (SZZ)



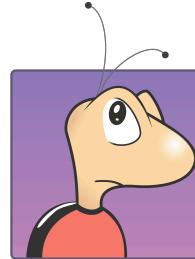
Project Bug
Tracker

The original approach relies on *Bugzilla*, but we can mine any bug tracker or similar database.

Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!

Śliwerski, Zimmermann, Zeller (SZZ)



Project Bug
Tracker



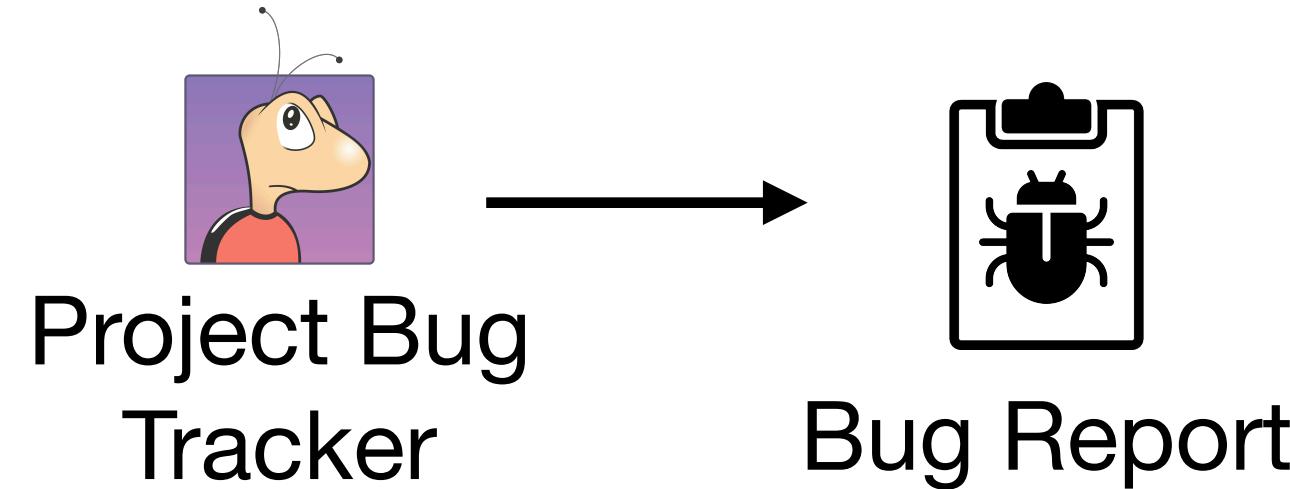
Project
History

The original approach relies on CVS (*Concurrent Versioning System*), but here we consider *git*.

Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!

Śliwerski, Zimmermann, Zeller (SZZ)



We pick a bug report for which we want to know its *bug-inducing commits* (BICs).

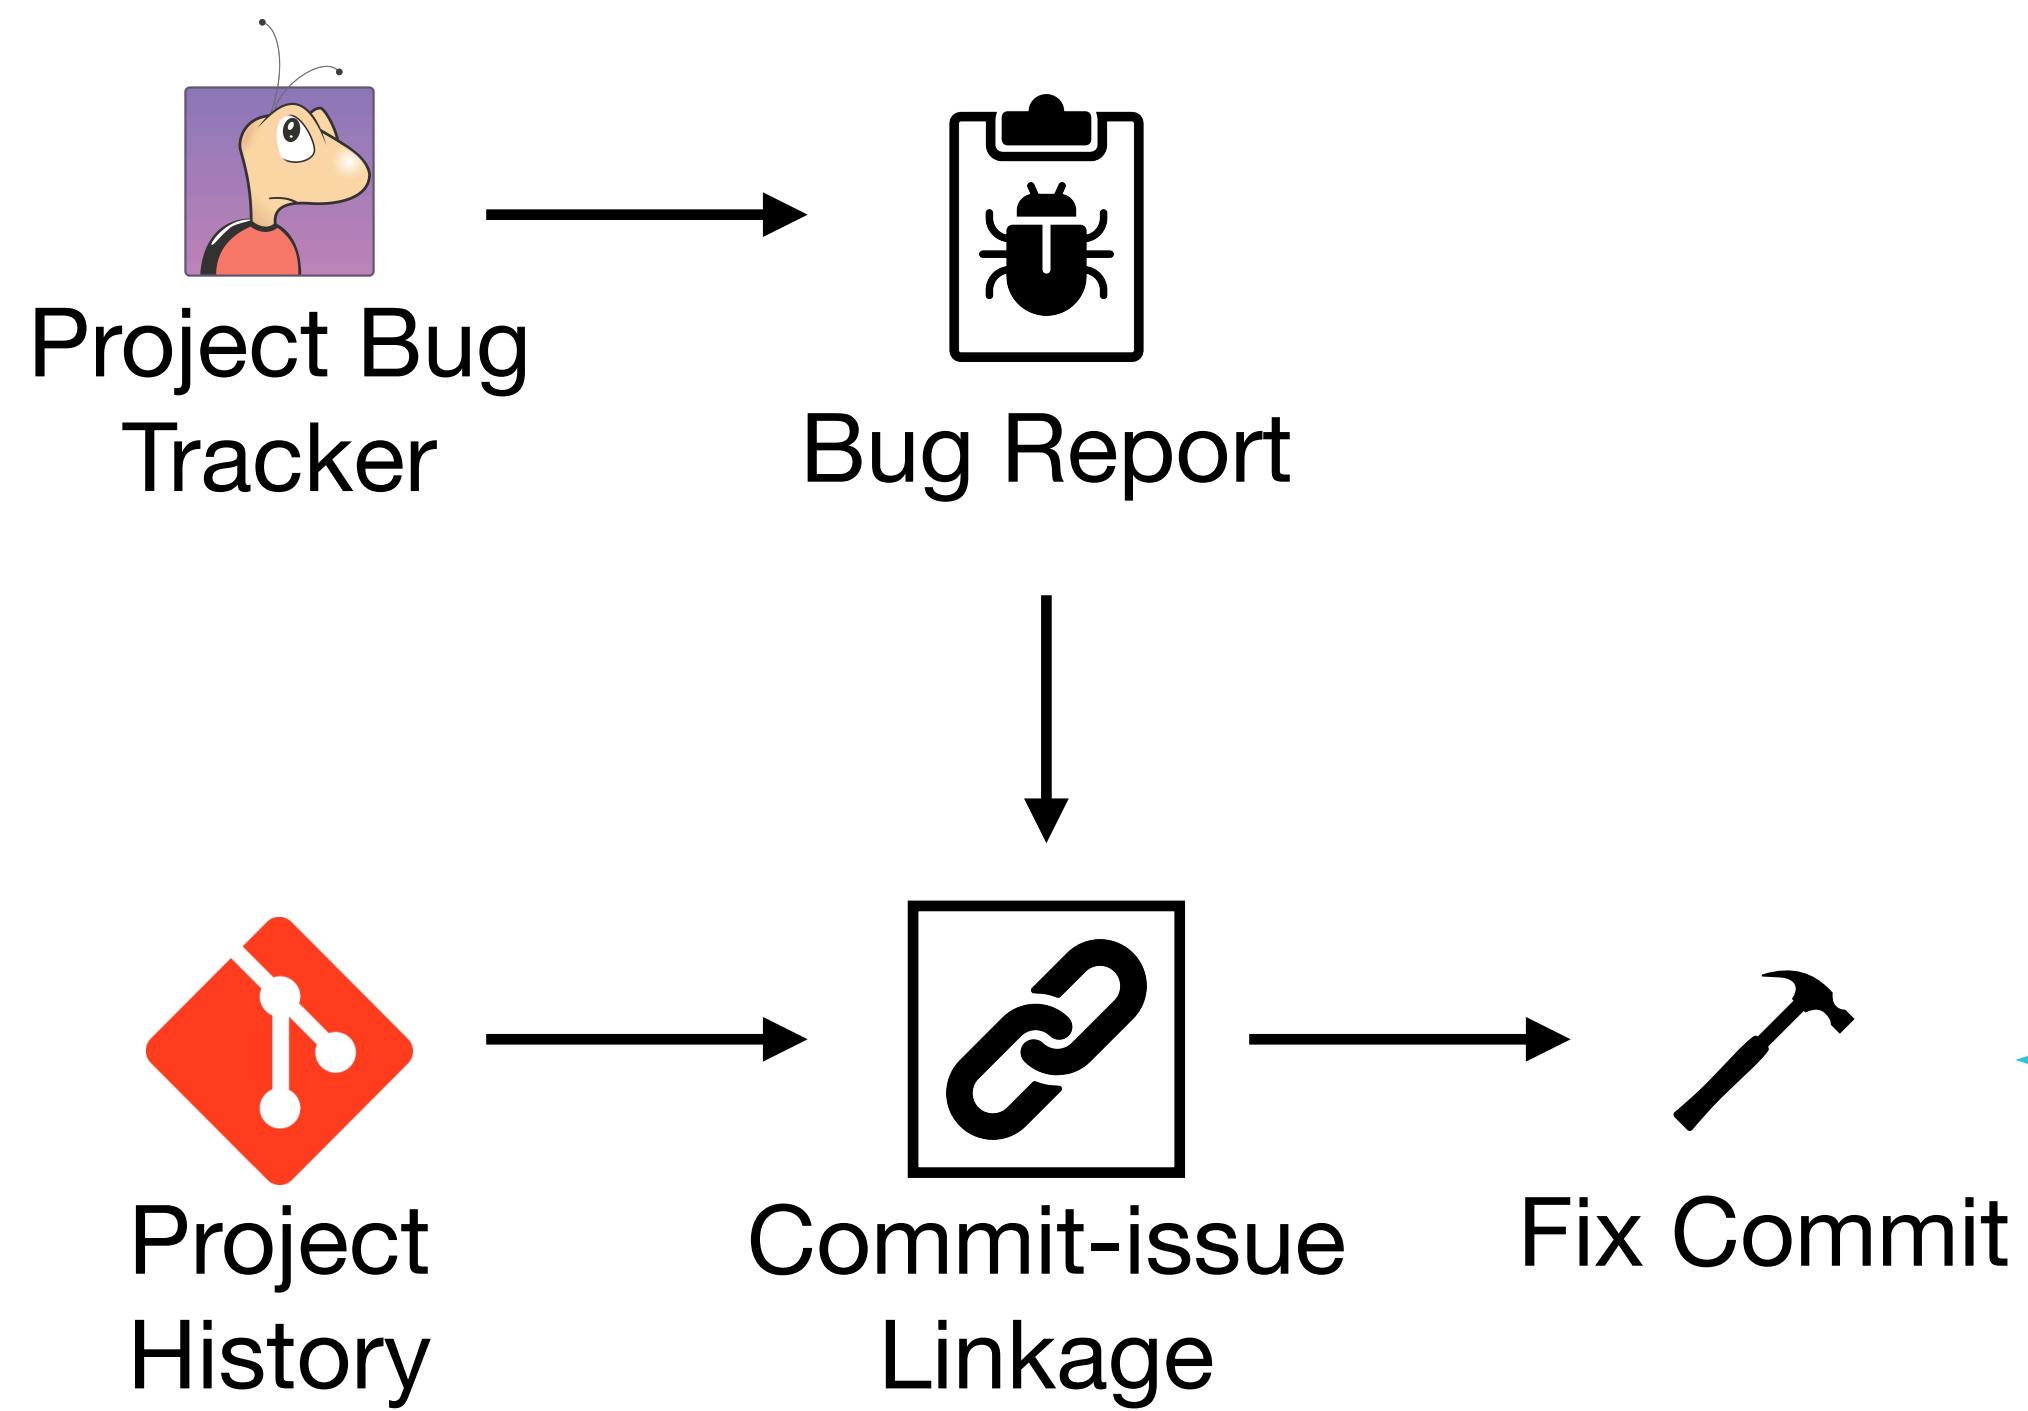


Project History

Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!

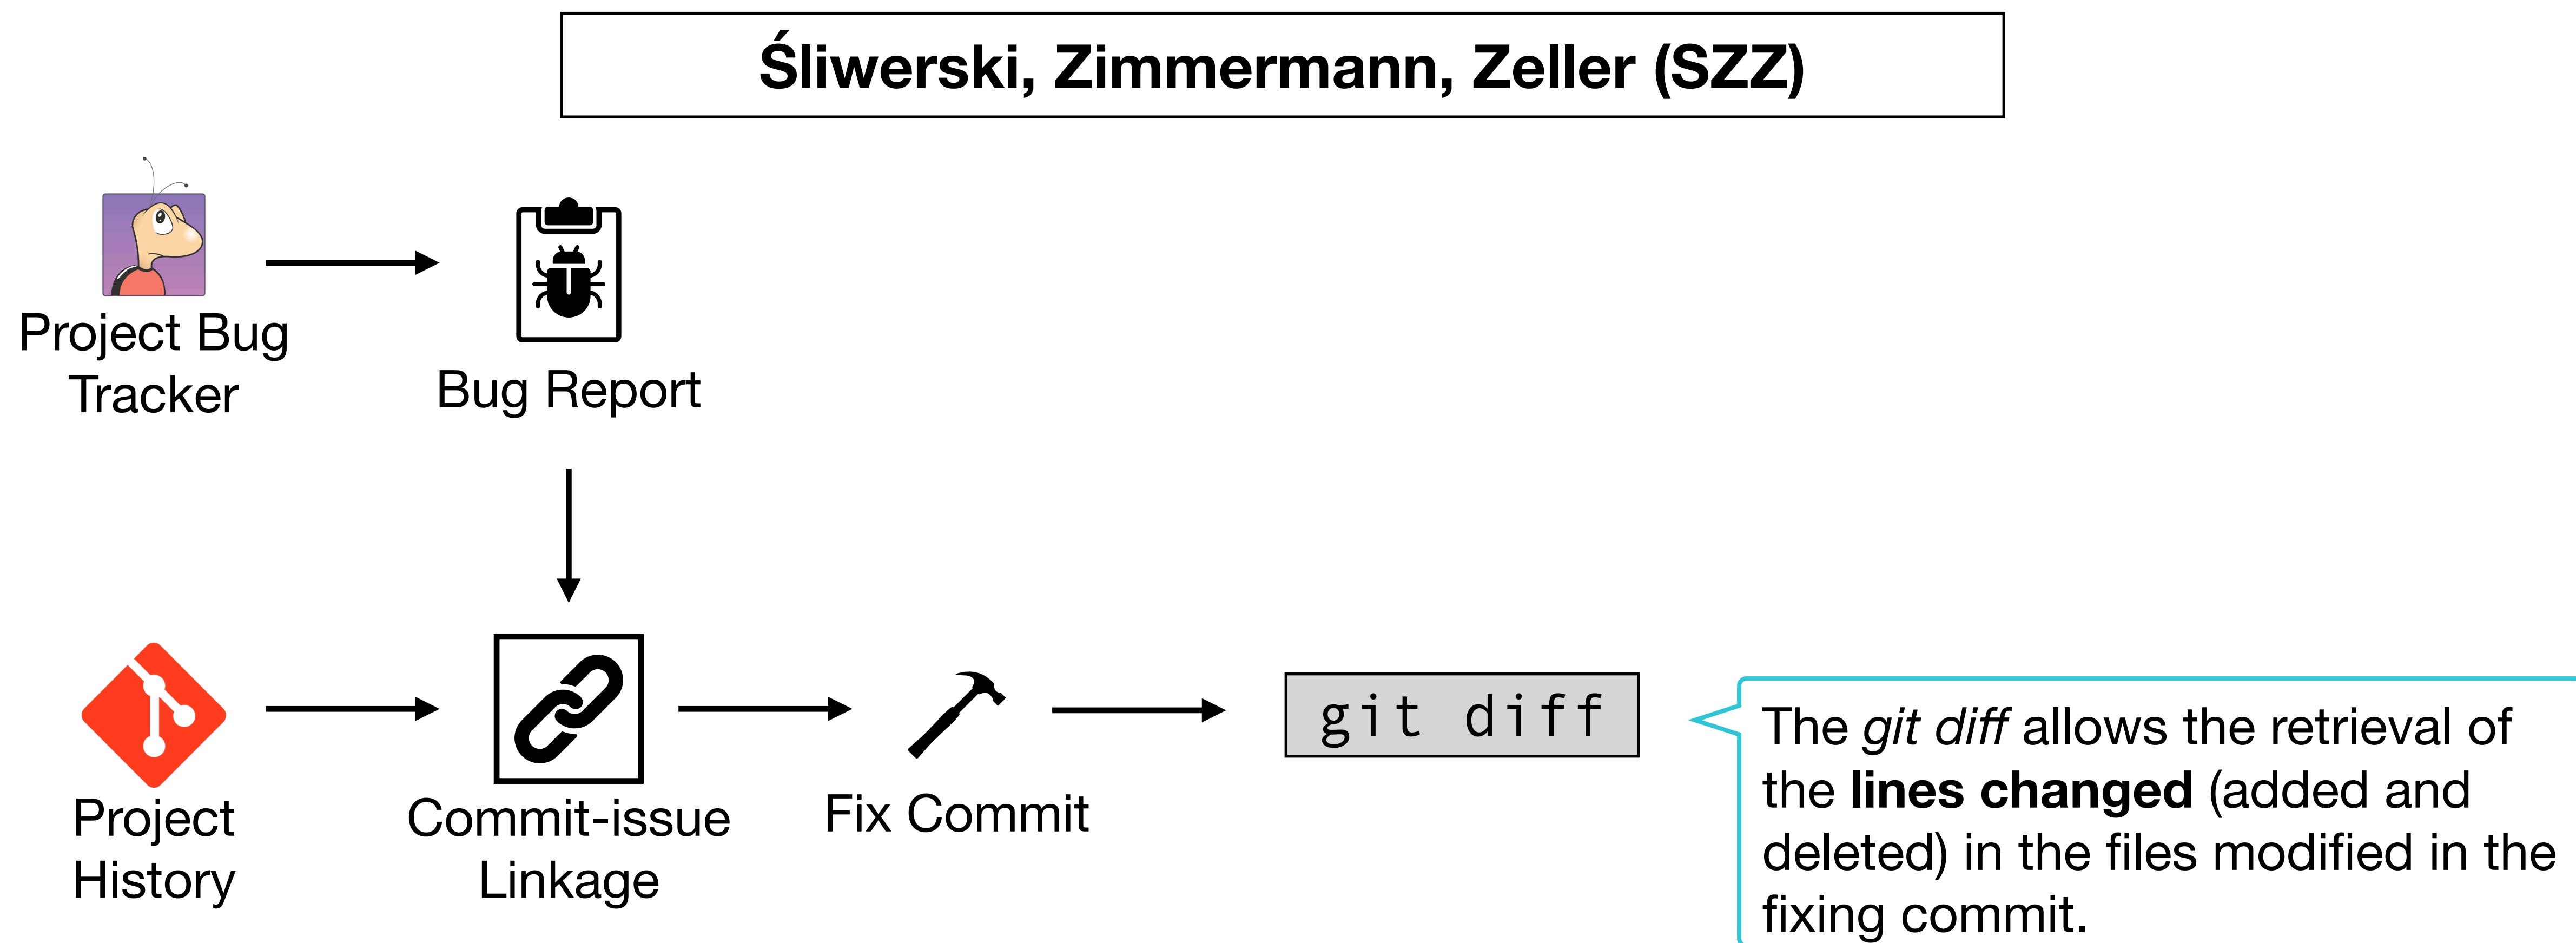
Śliwerski, Zimmermann, Zeller (SZZ)



We can run any commit-issue link algorithm we want. The original approach uses a **pattern-based search**, looking for the bug ID (a number) inside the commit messages. In any case, we just want the *bug-fixing commit*.

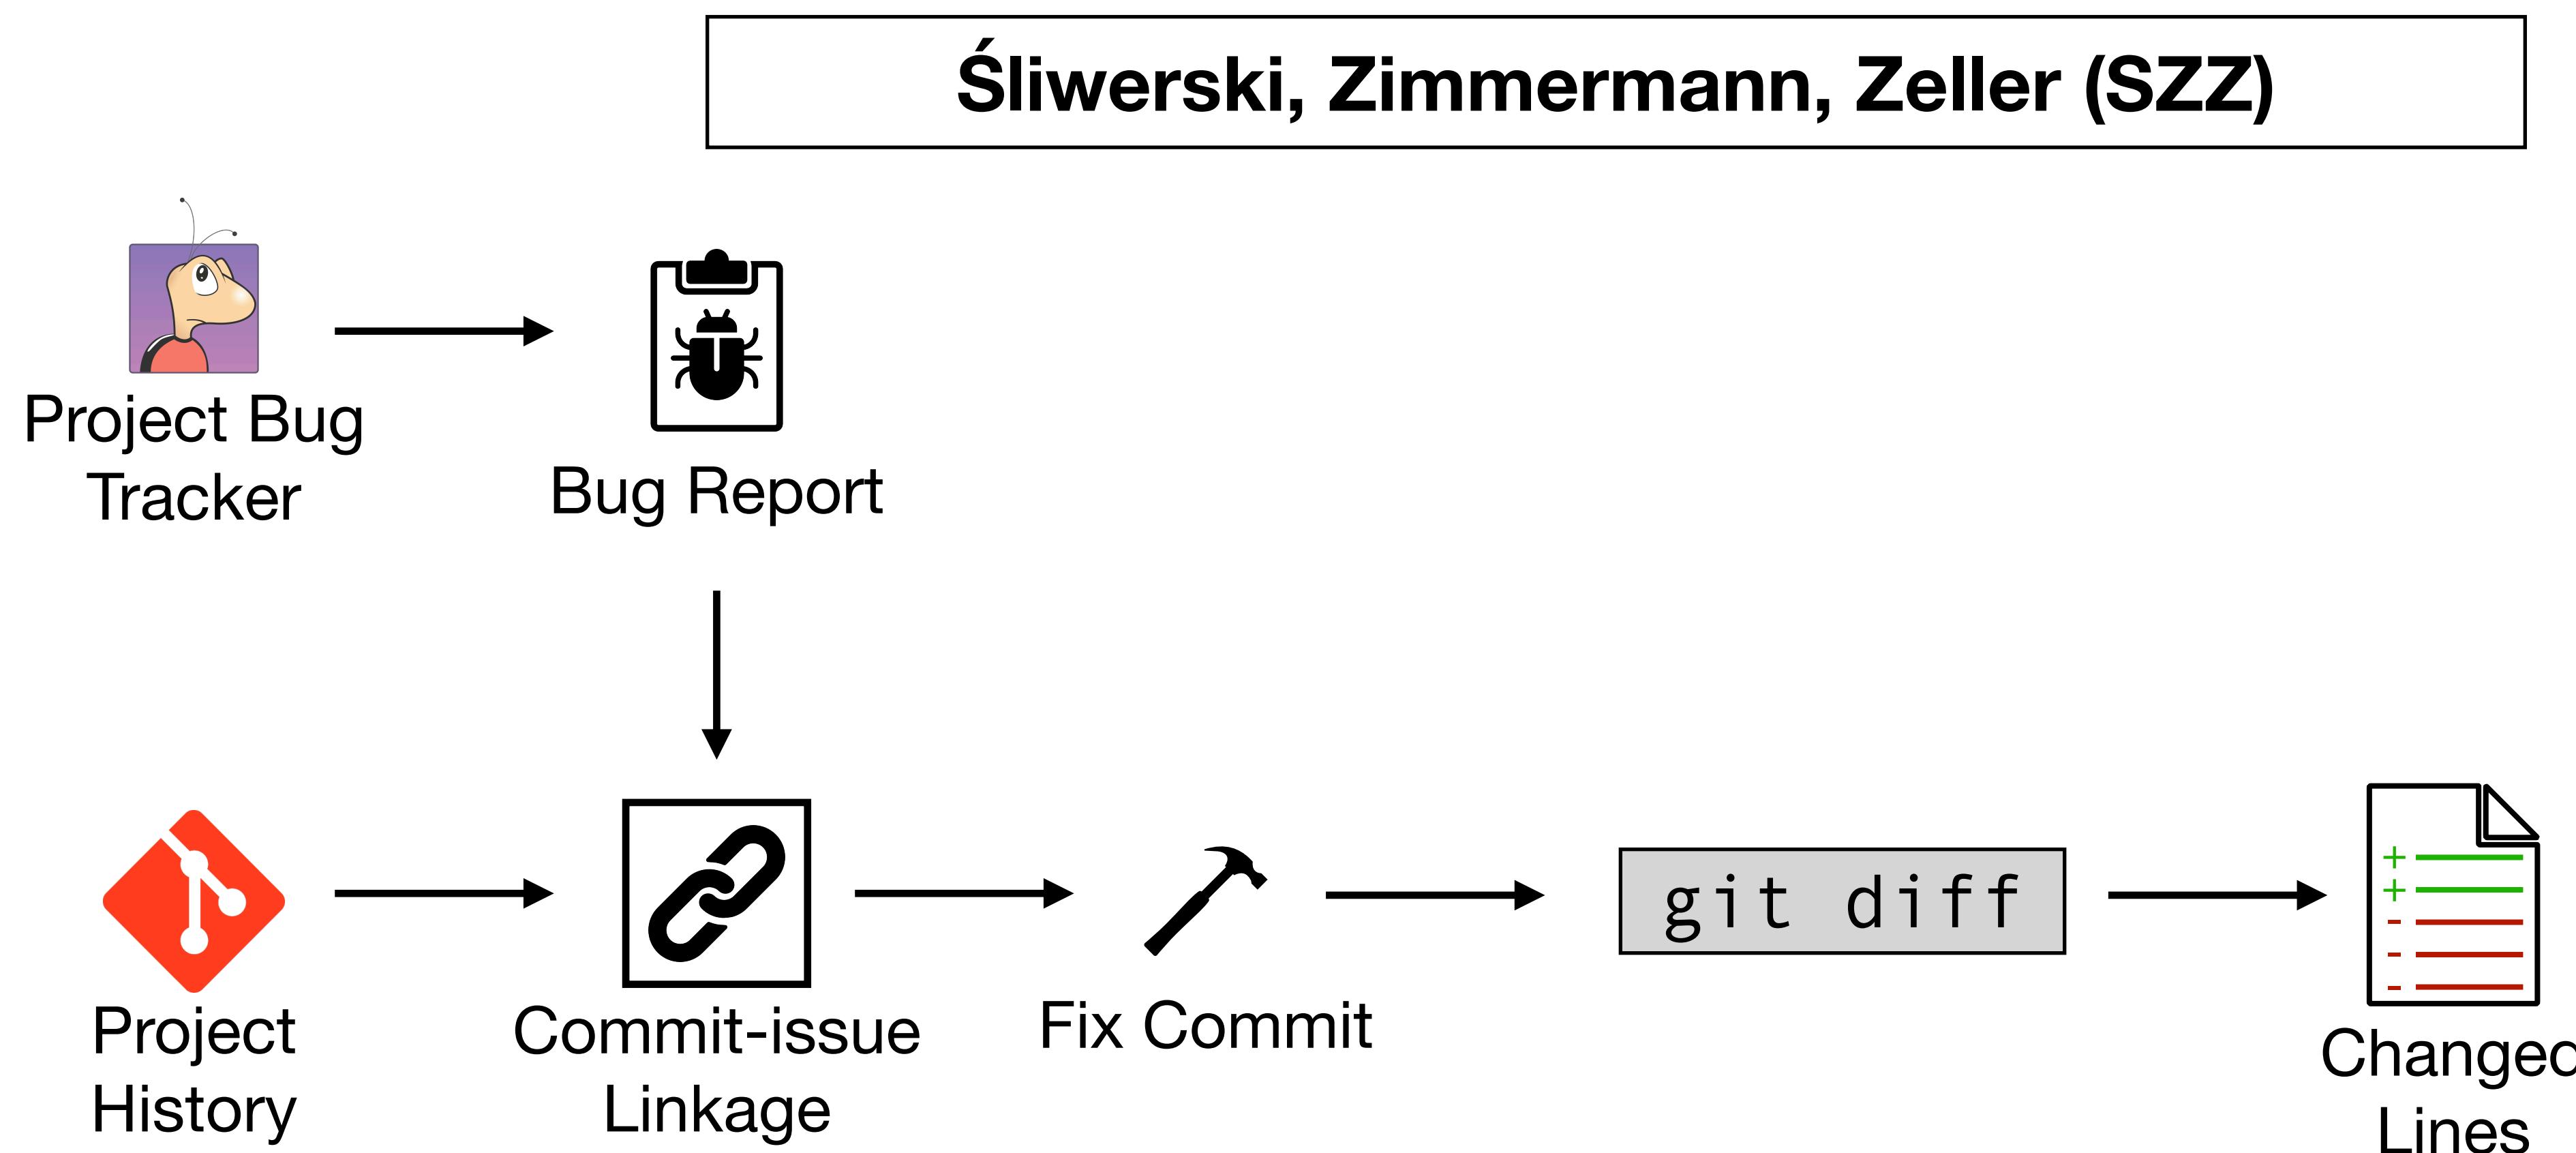
Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



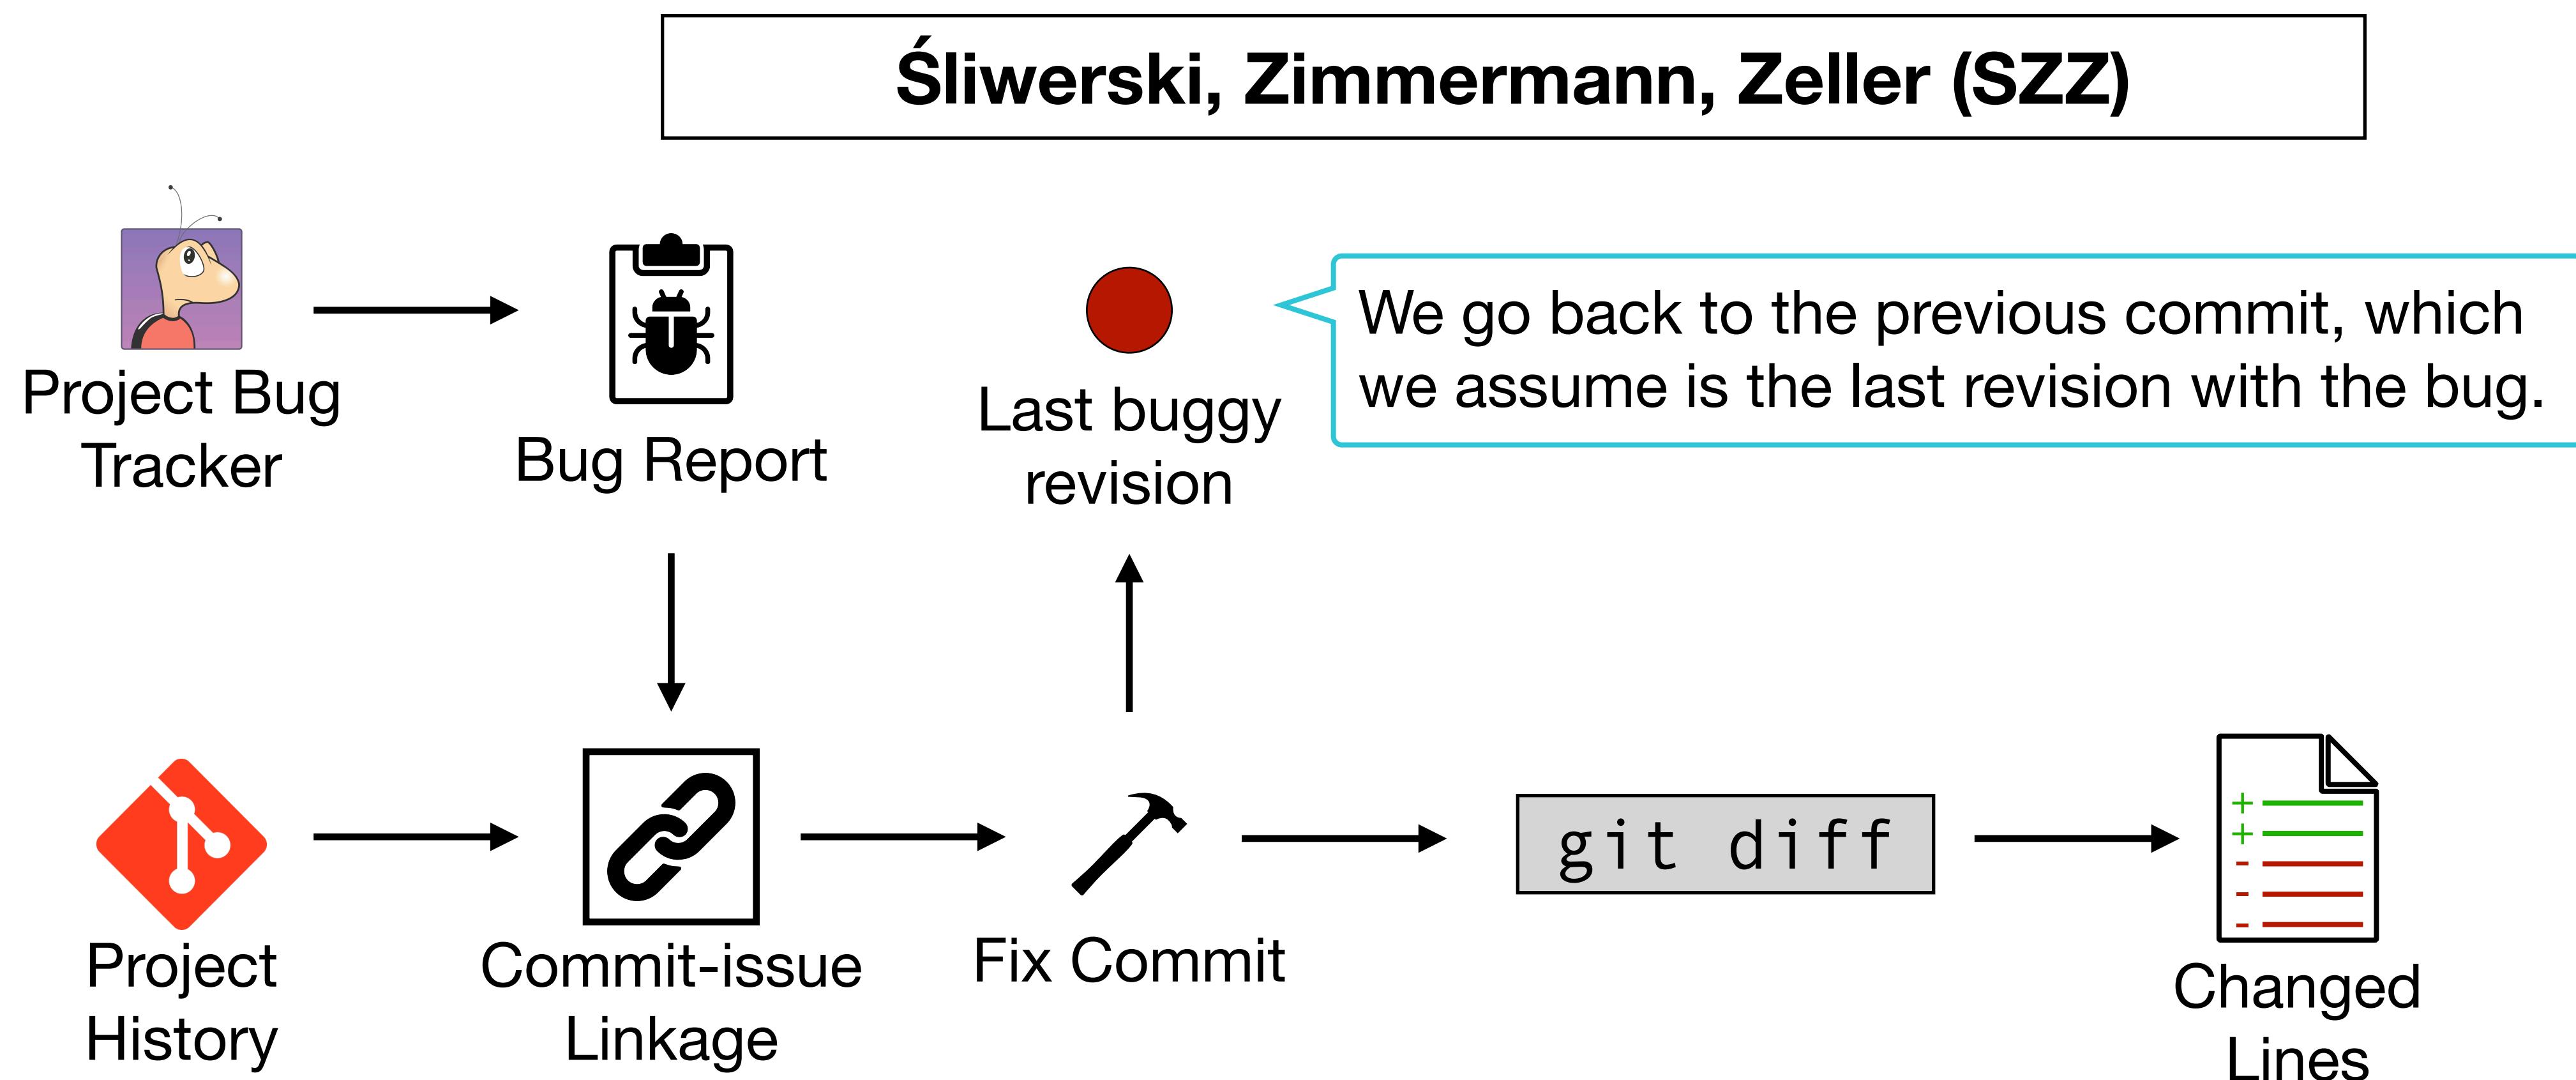
Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



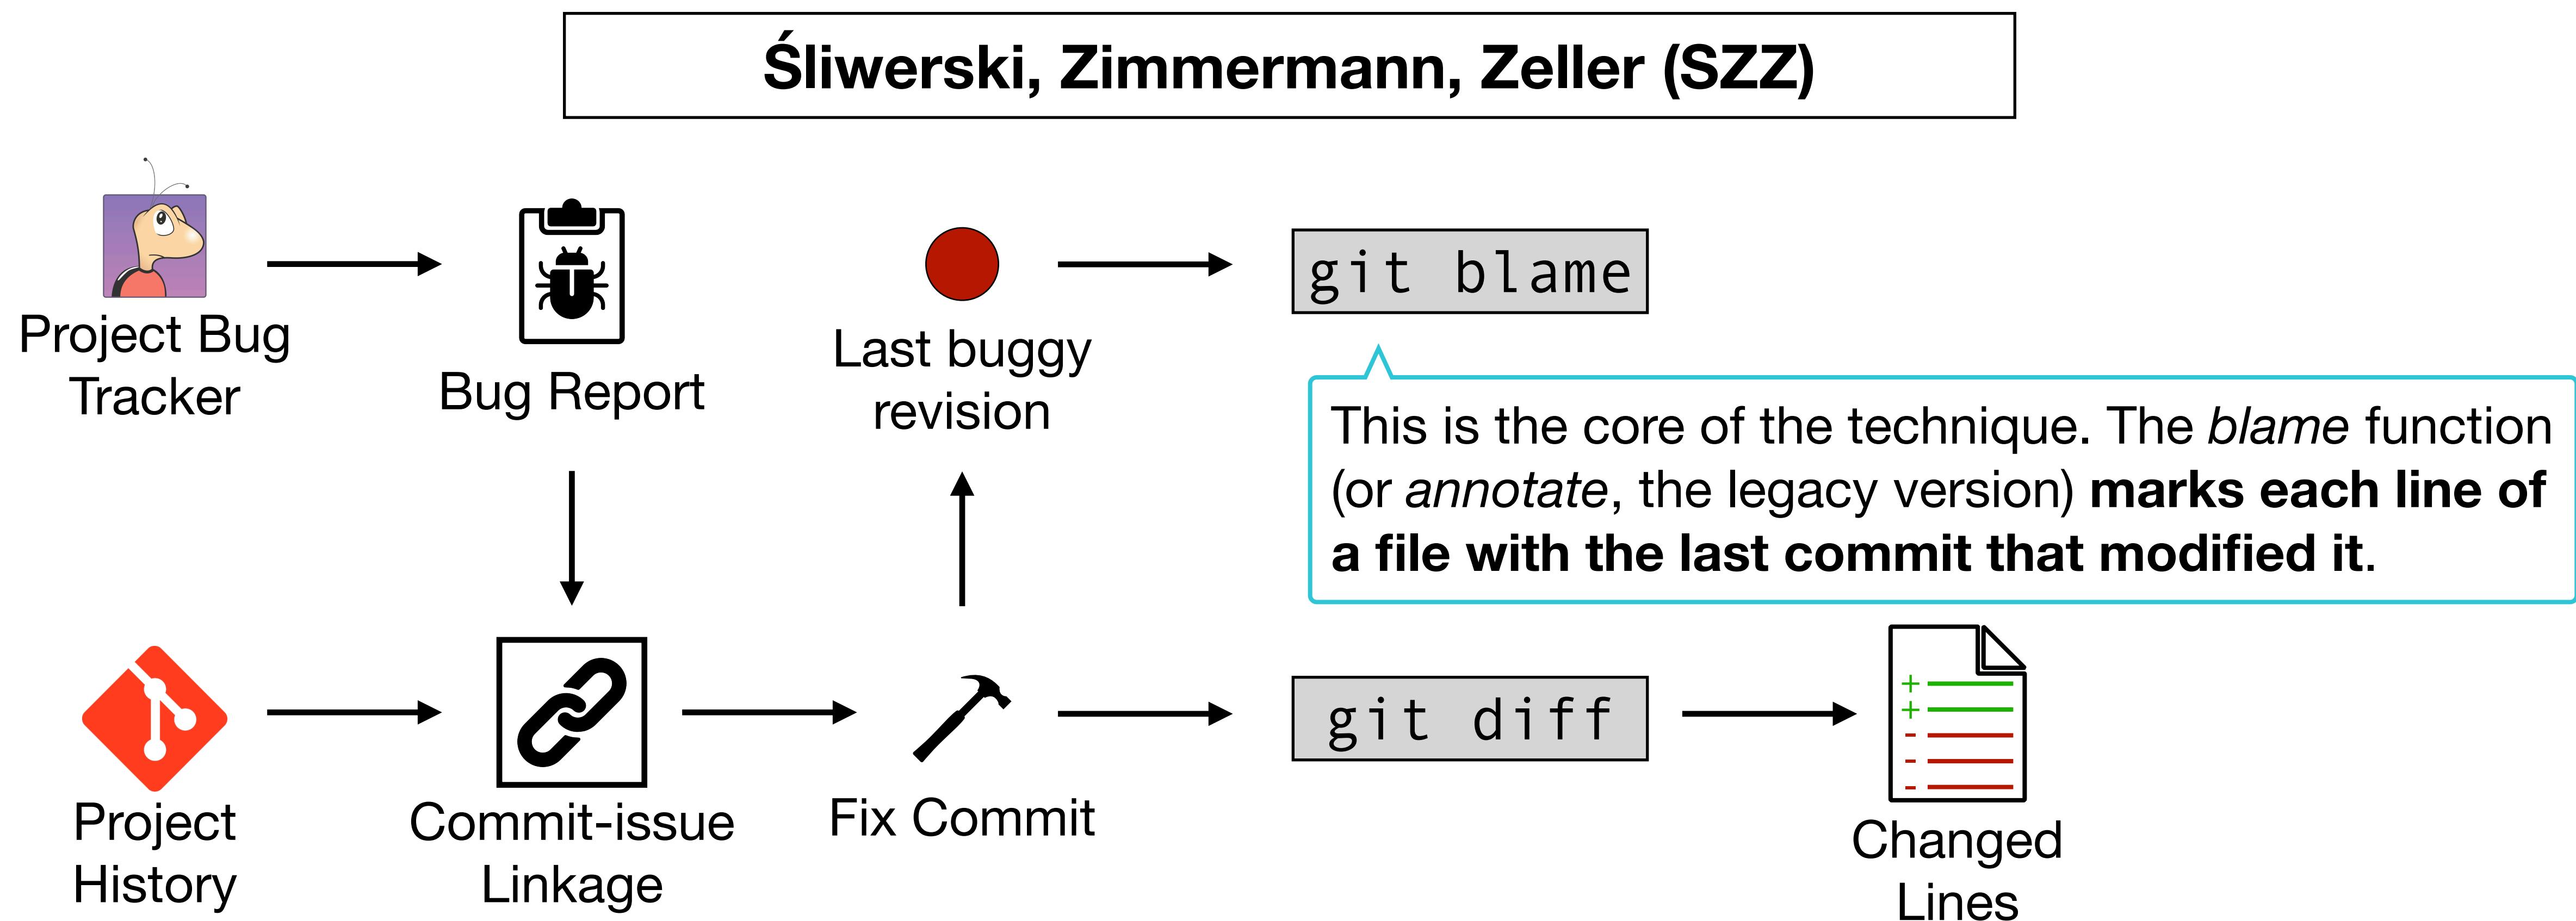
Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



MSR for Vulnerability Prediction — Mining VCCs

commons-csv / src / main / java / org / apache / commons / csv / CSVParser.java ↑ Top

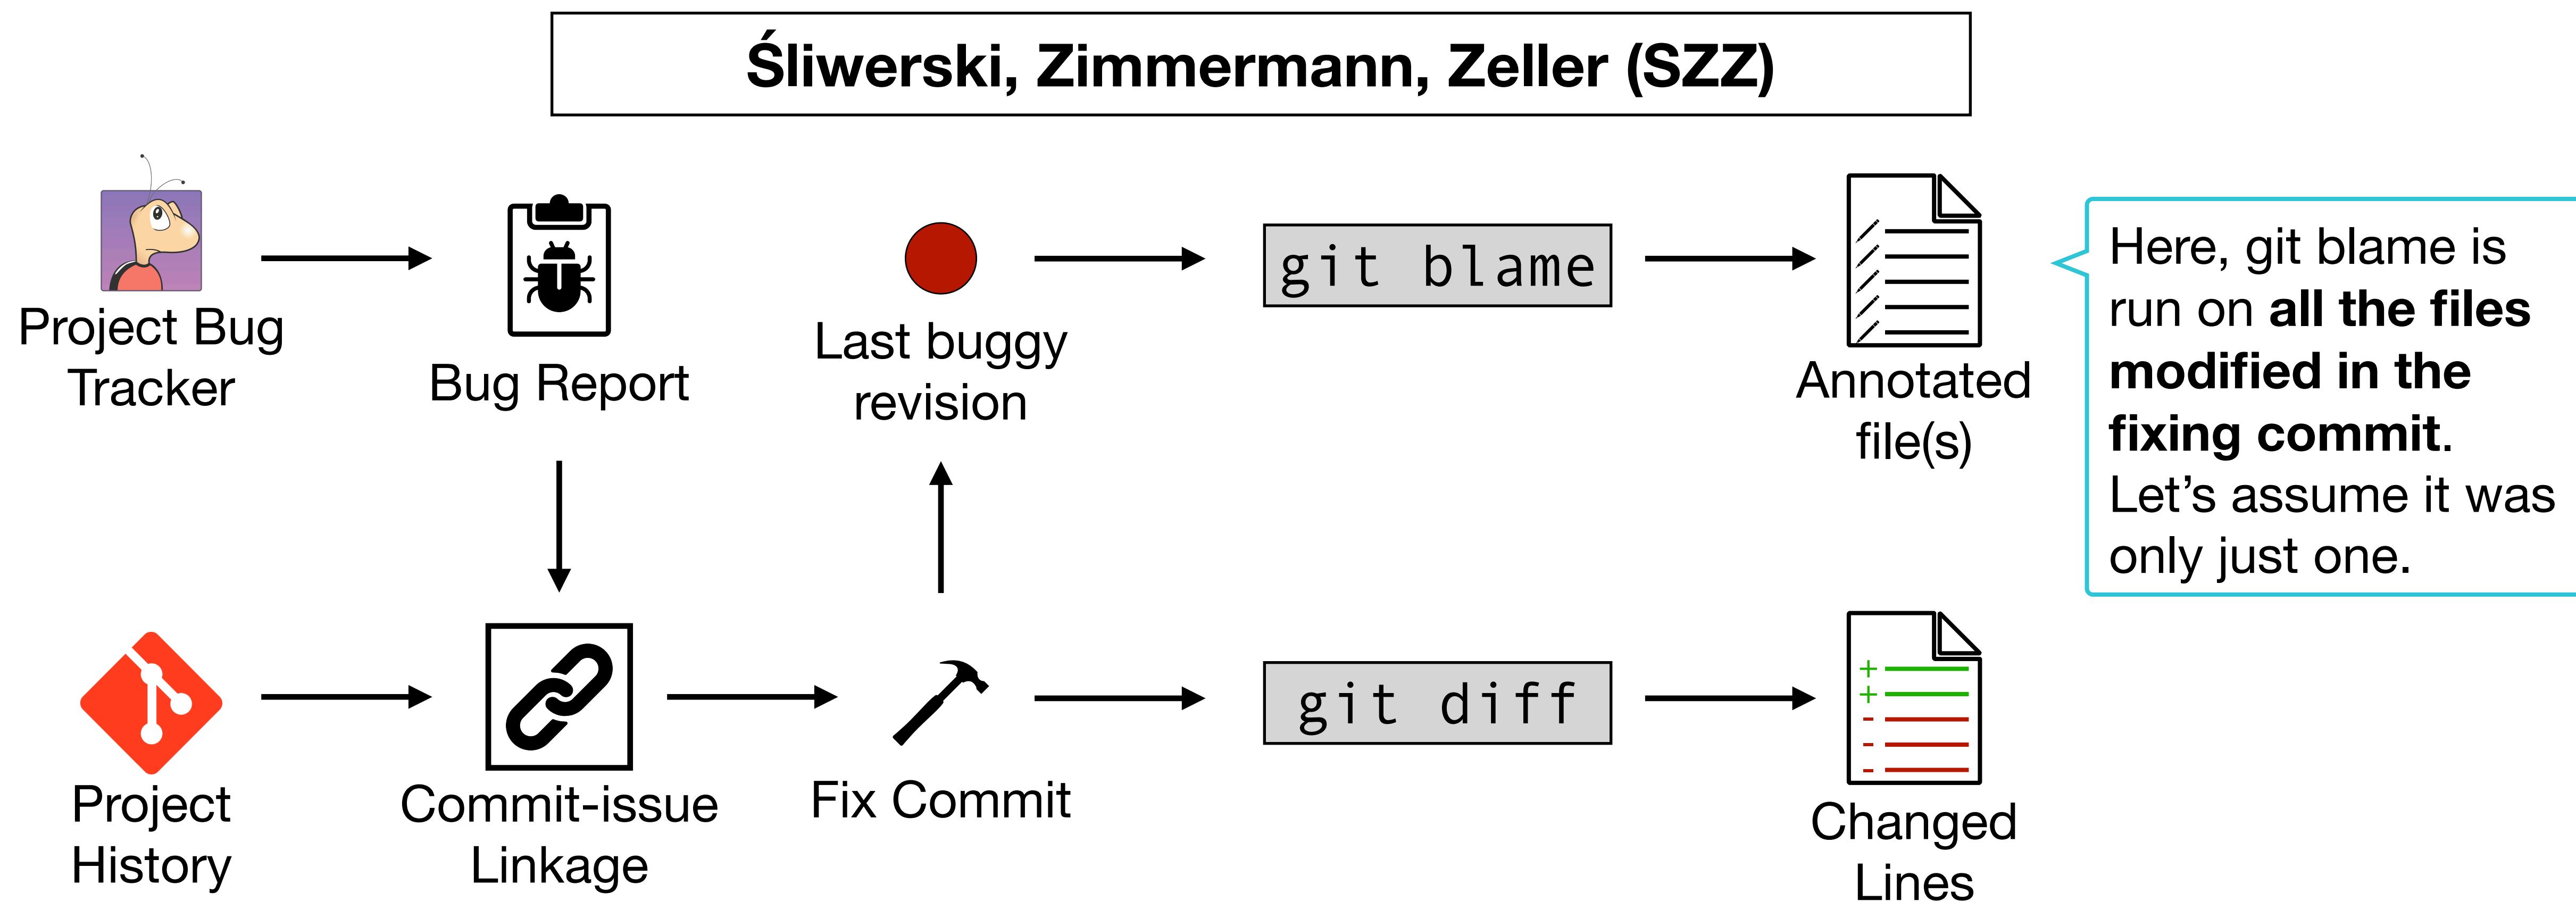
Code Blame 823 lines (759 loc) · 29.2 KB Raw Copy Download Edit More

Older Newer Contributors 10

Time Ago	Author	Commit Message	Line No.	Code
4 years ago		[CSV-239] Add CSVRecord.ge...	476	private Headers createHeaders() throws IOException {
4 years ago		[CSV-239] Cannot get headers ...	477	Map<String, Integer> hdrMap = null;
4 years ago		[CSV-239] Add CSVRecord.ge...	478	List<String> headerNames = null;
4 years ago		[CSV-239] Cannot get headers ...	479	final String[] formatHeader = this.format.getHeader();
			480	if (formatHeader != null) {
4 years ago		[CSV-239] Cannot get headers ...	481	hdrMap = createEmptyHeaderMap();
4 years ago		[CSV-239] Cannot get headers ...	482	String[] headerRecord = null;
			483	if (formatHeader.length == 0) {
			484	// read the header from the first line of the file
			485	final CSVRecord nextRecord = this.nextRecord();
			486	if (nextRecord != null) {
			487	headerRecord = nextRecord.values();
8 months ago		[CSV-304] Accessors for hea...	488	headerComment = nextRecord.getComment();
4 years ago		[CSV-239] Cannot get headers ...	489	}
			490	} else {
			491	if (this.format.getSkipHeaderRecord()) {
8 months ago		Guard against NPE in createH...	492	final CSVRecord nextRecord = this.nextRecord();
			493	if (nextRecord != null) {
			494	headerComment = nextRecord.getComment();
			495	}
4 years ago		[CSV-239] Cannot get headers ...	496	}
			497	headerRecord = formatHeader;
			498	}
			499	// build the name to index mappings
			500	if (headerRecord != null) {

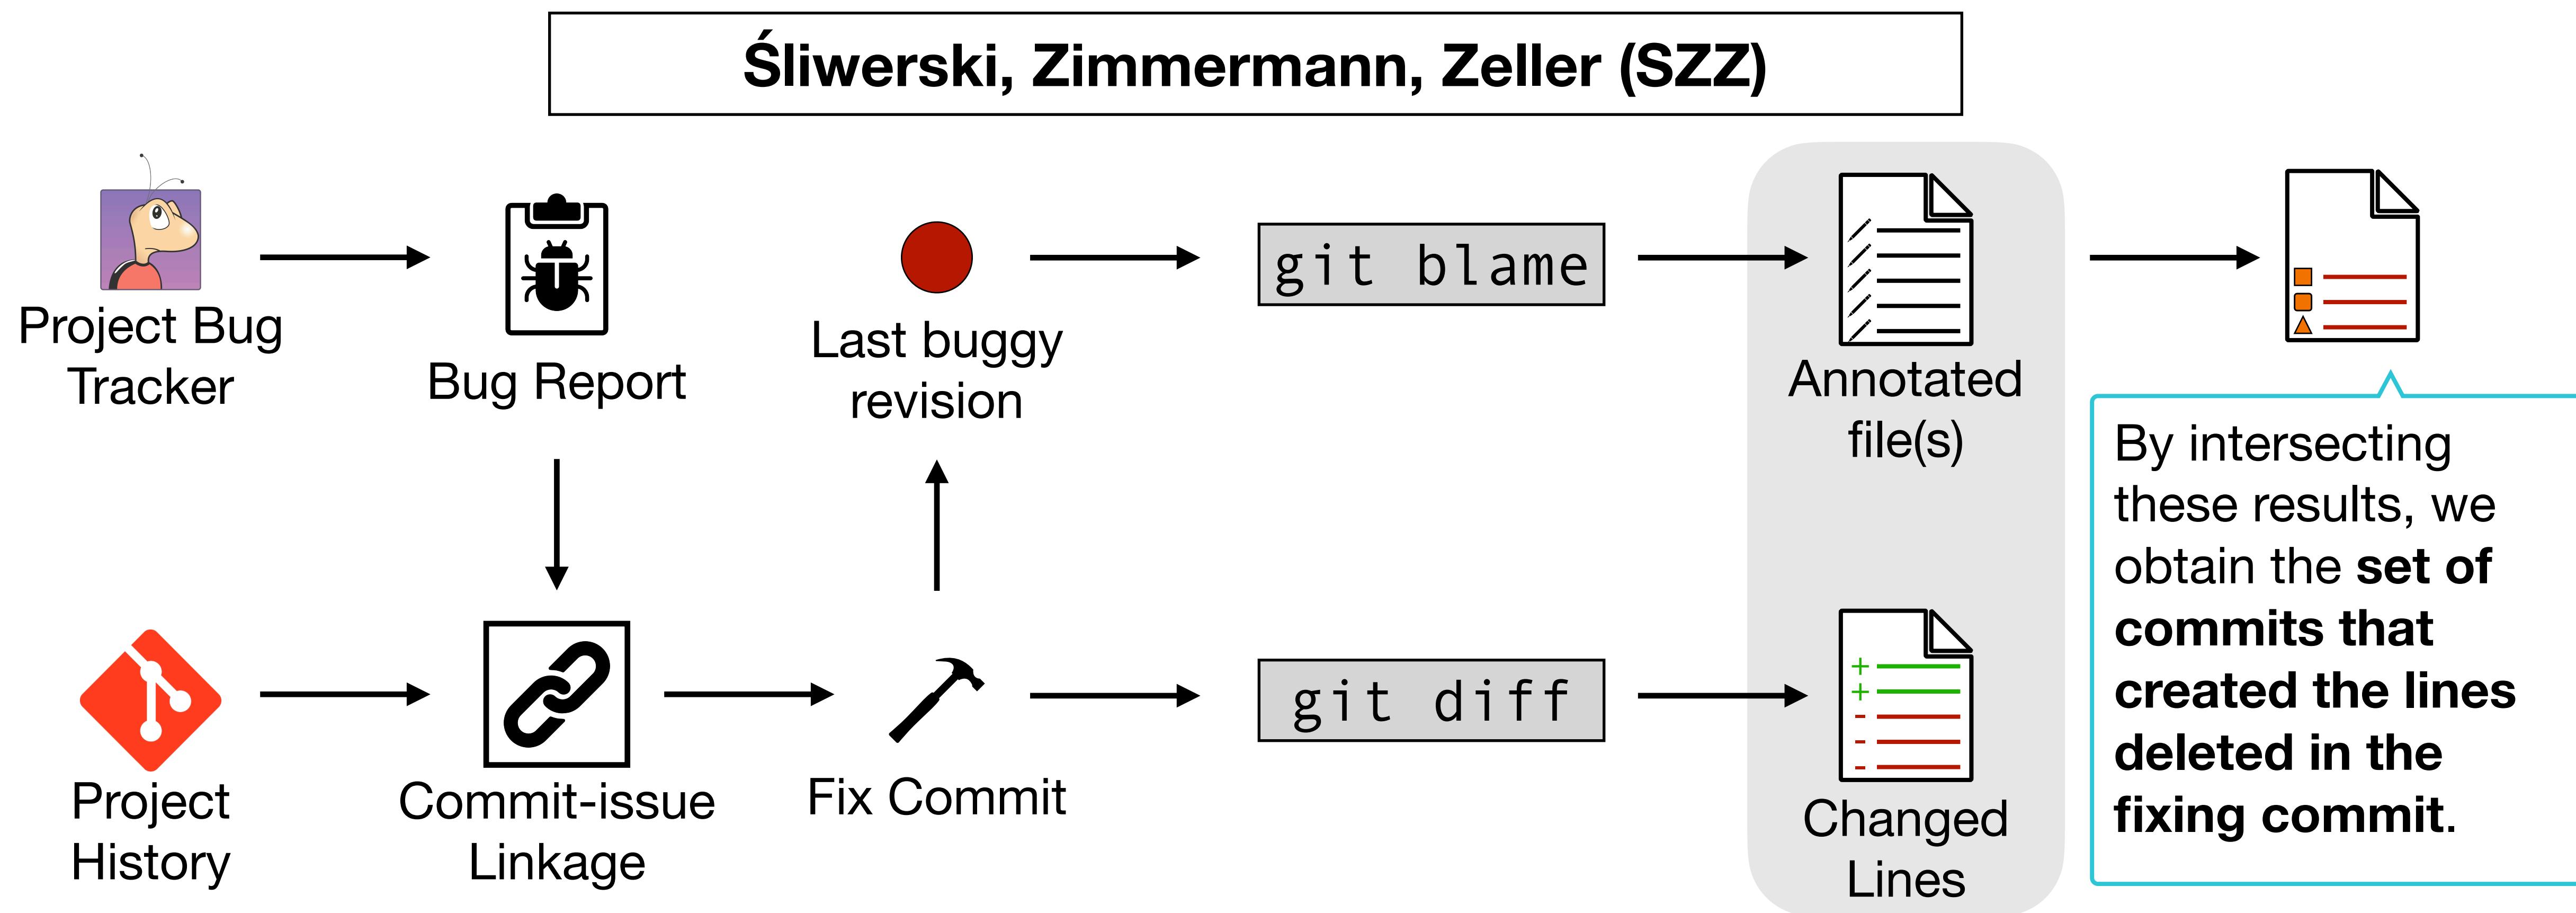
Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



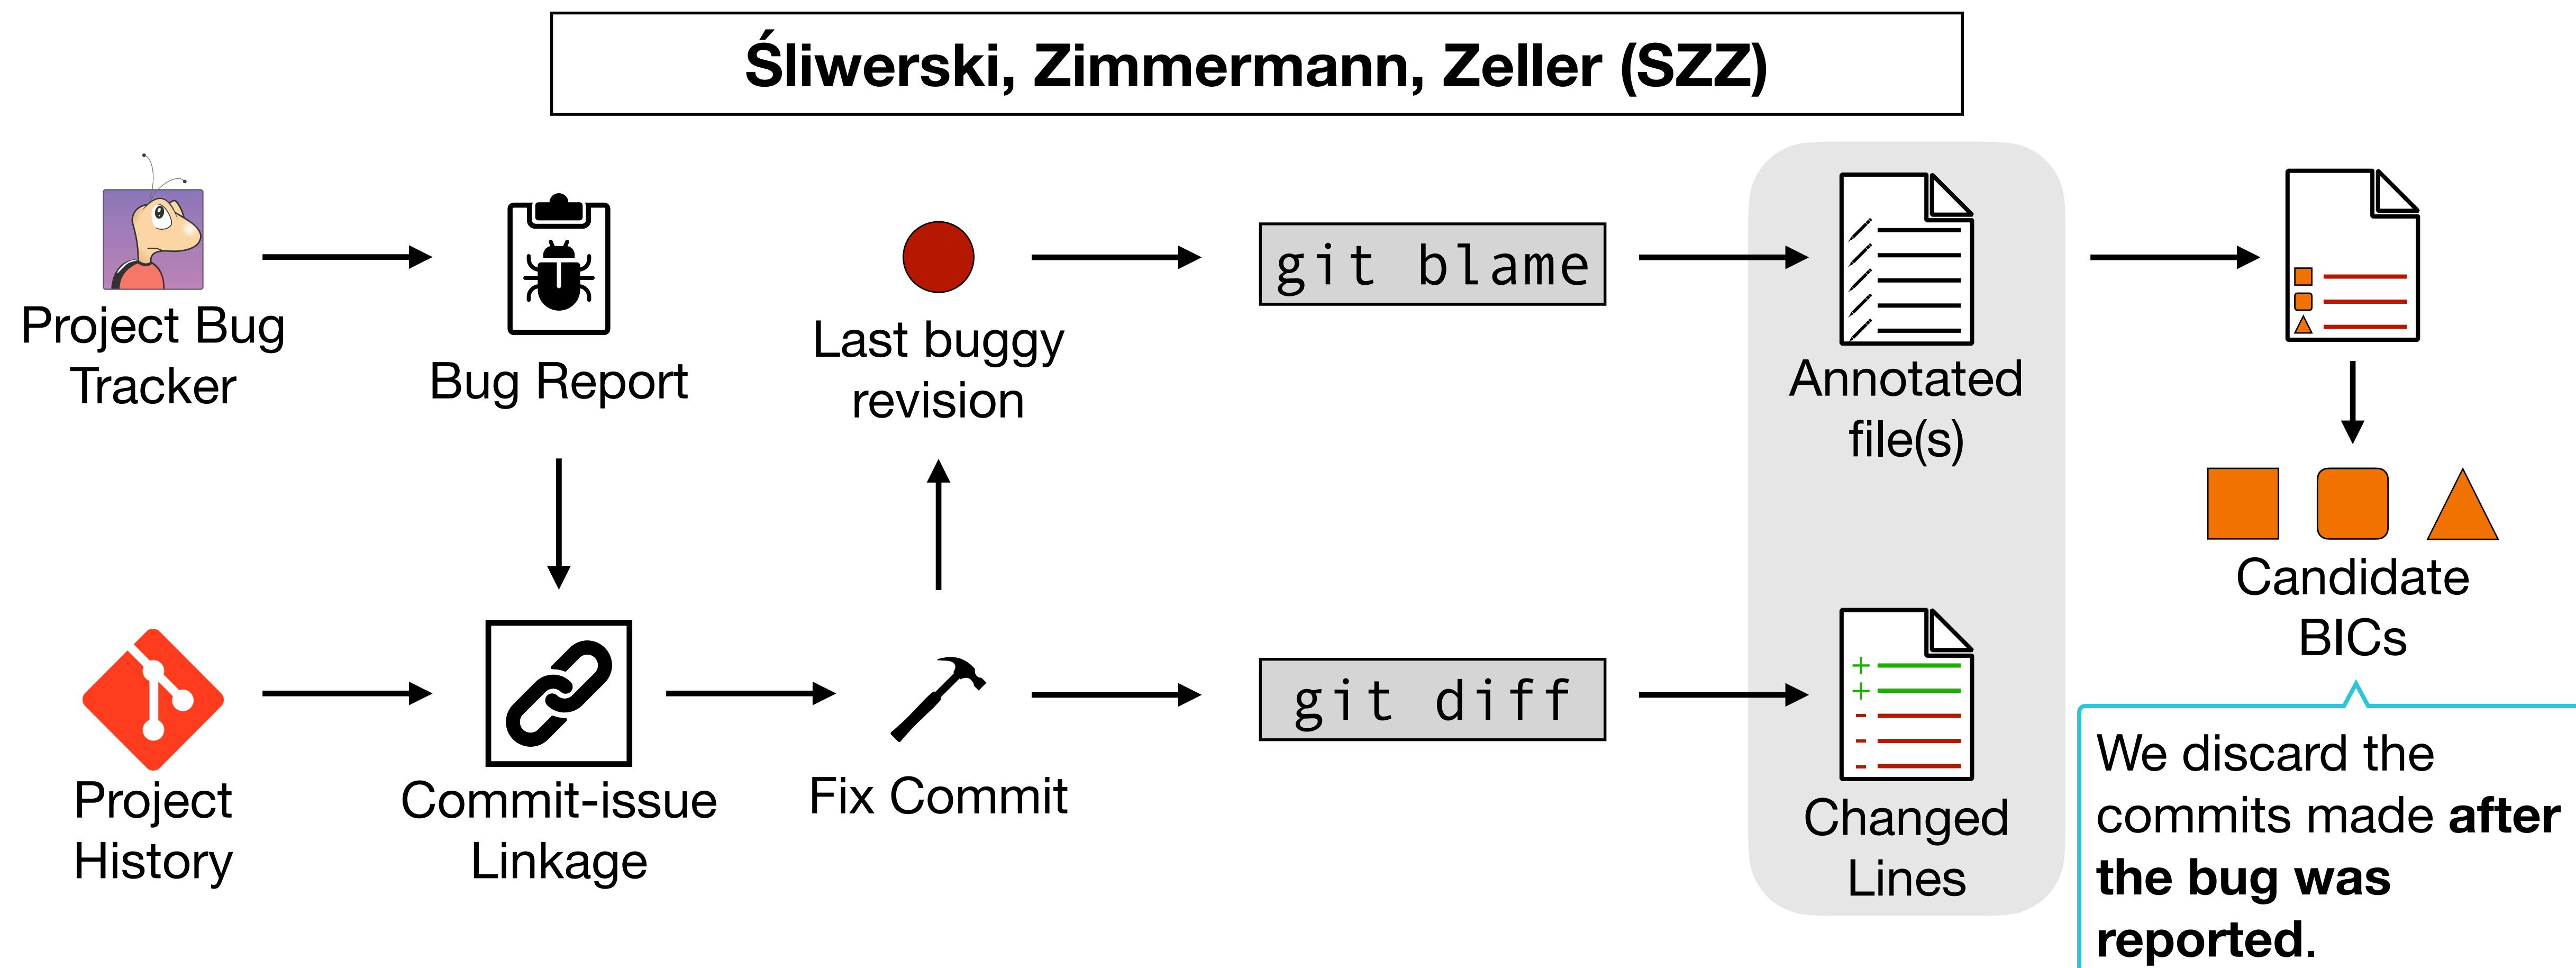
Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



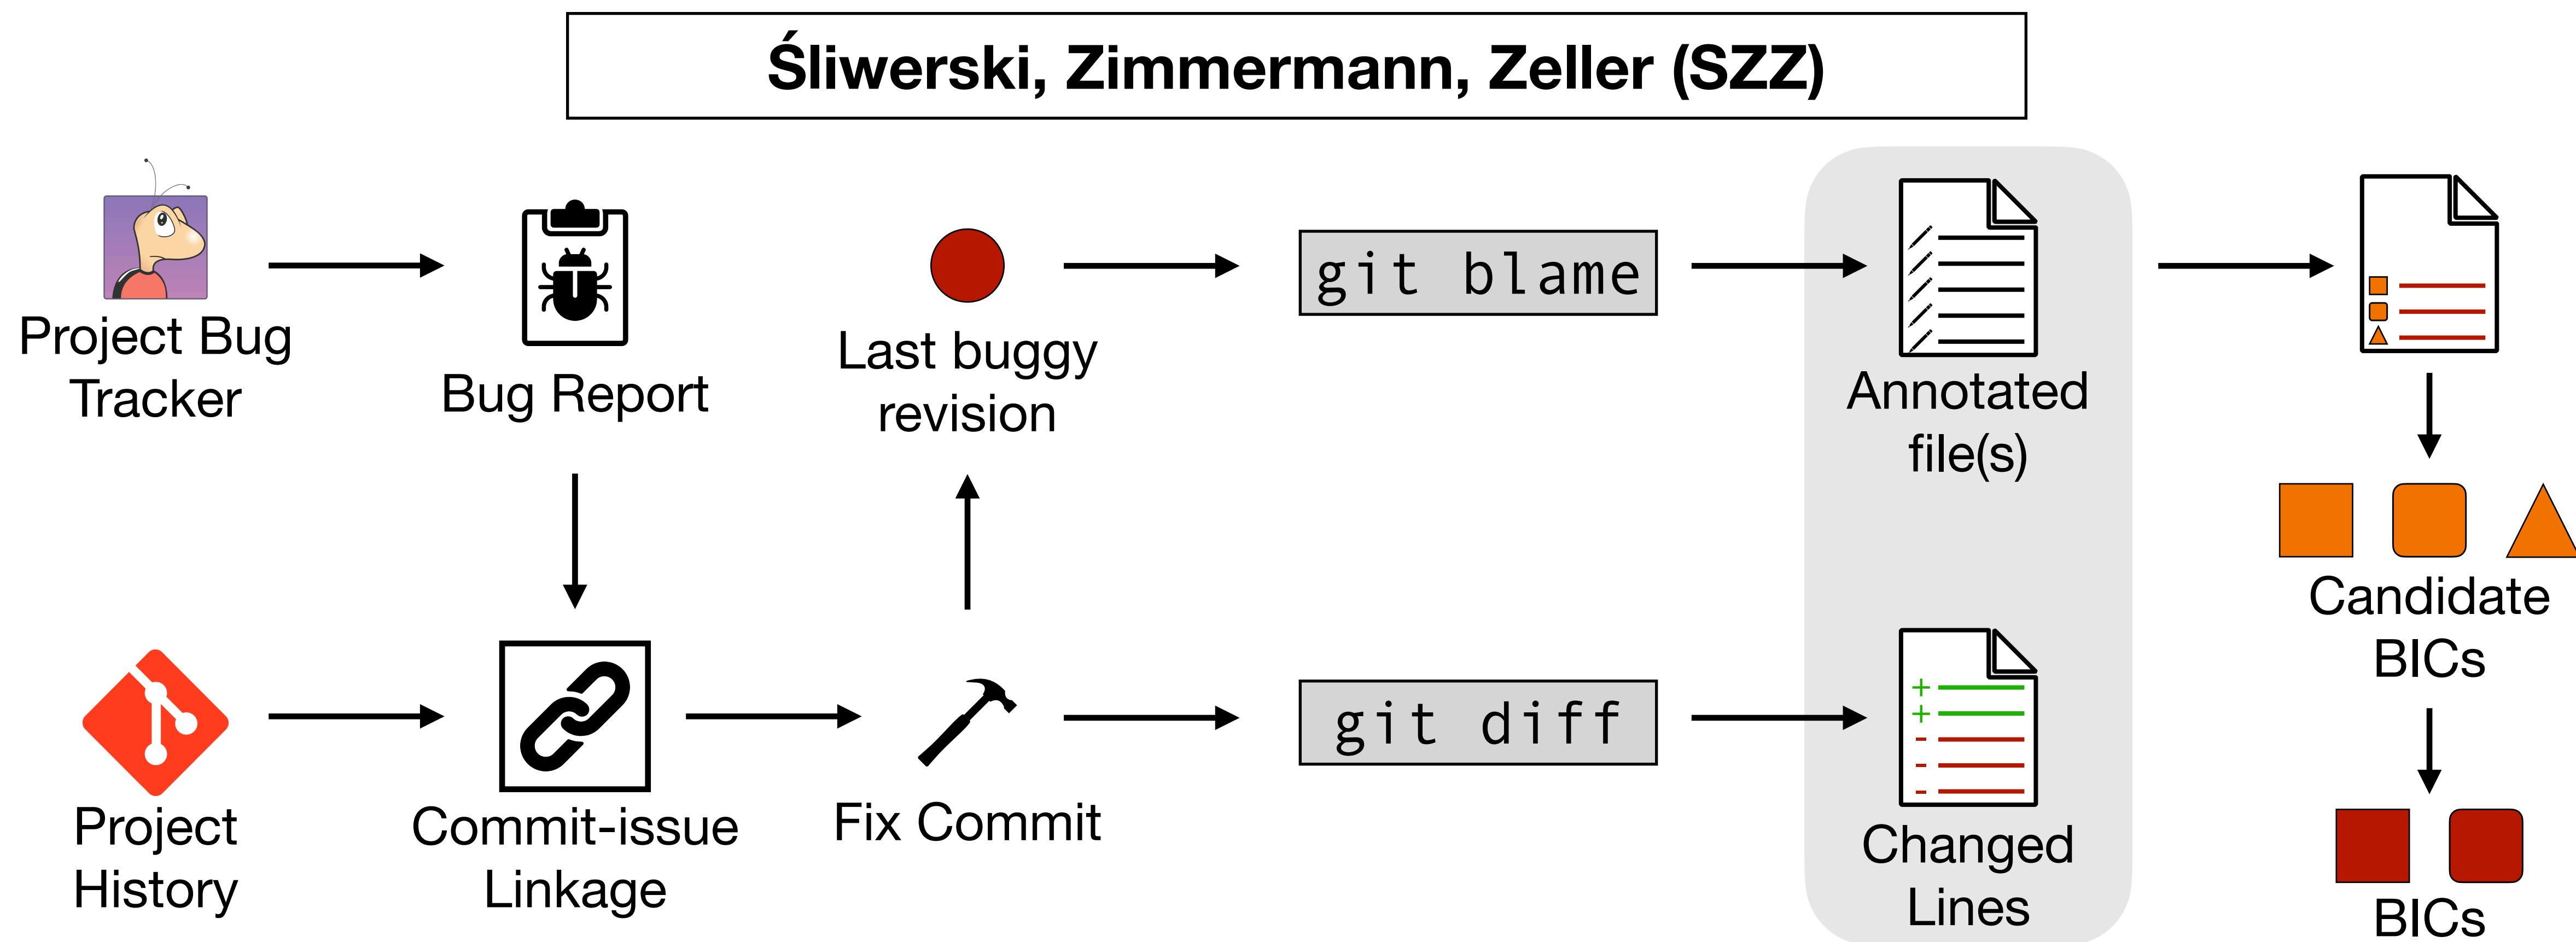
Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



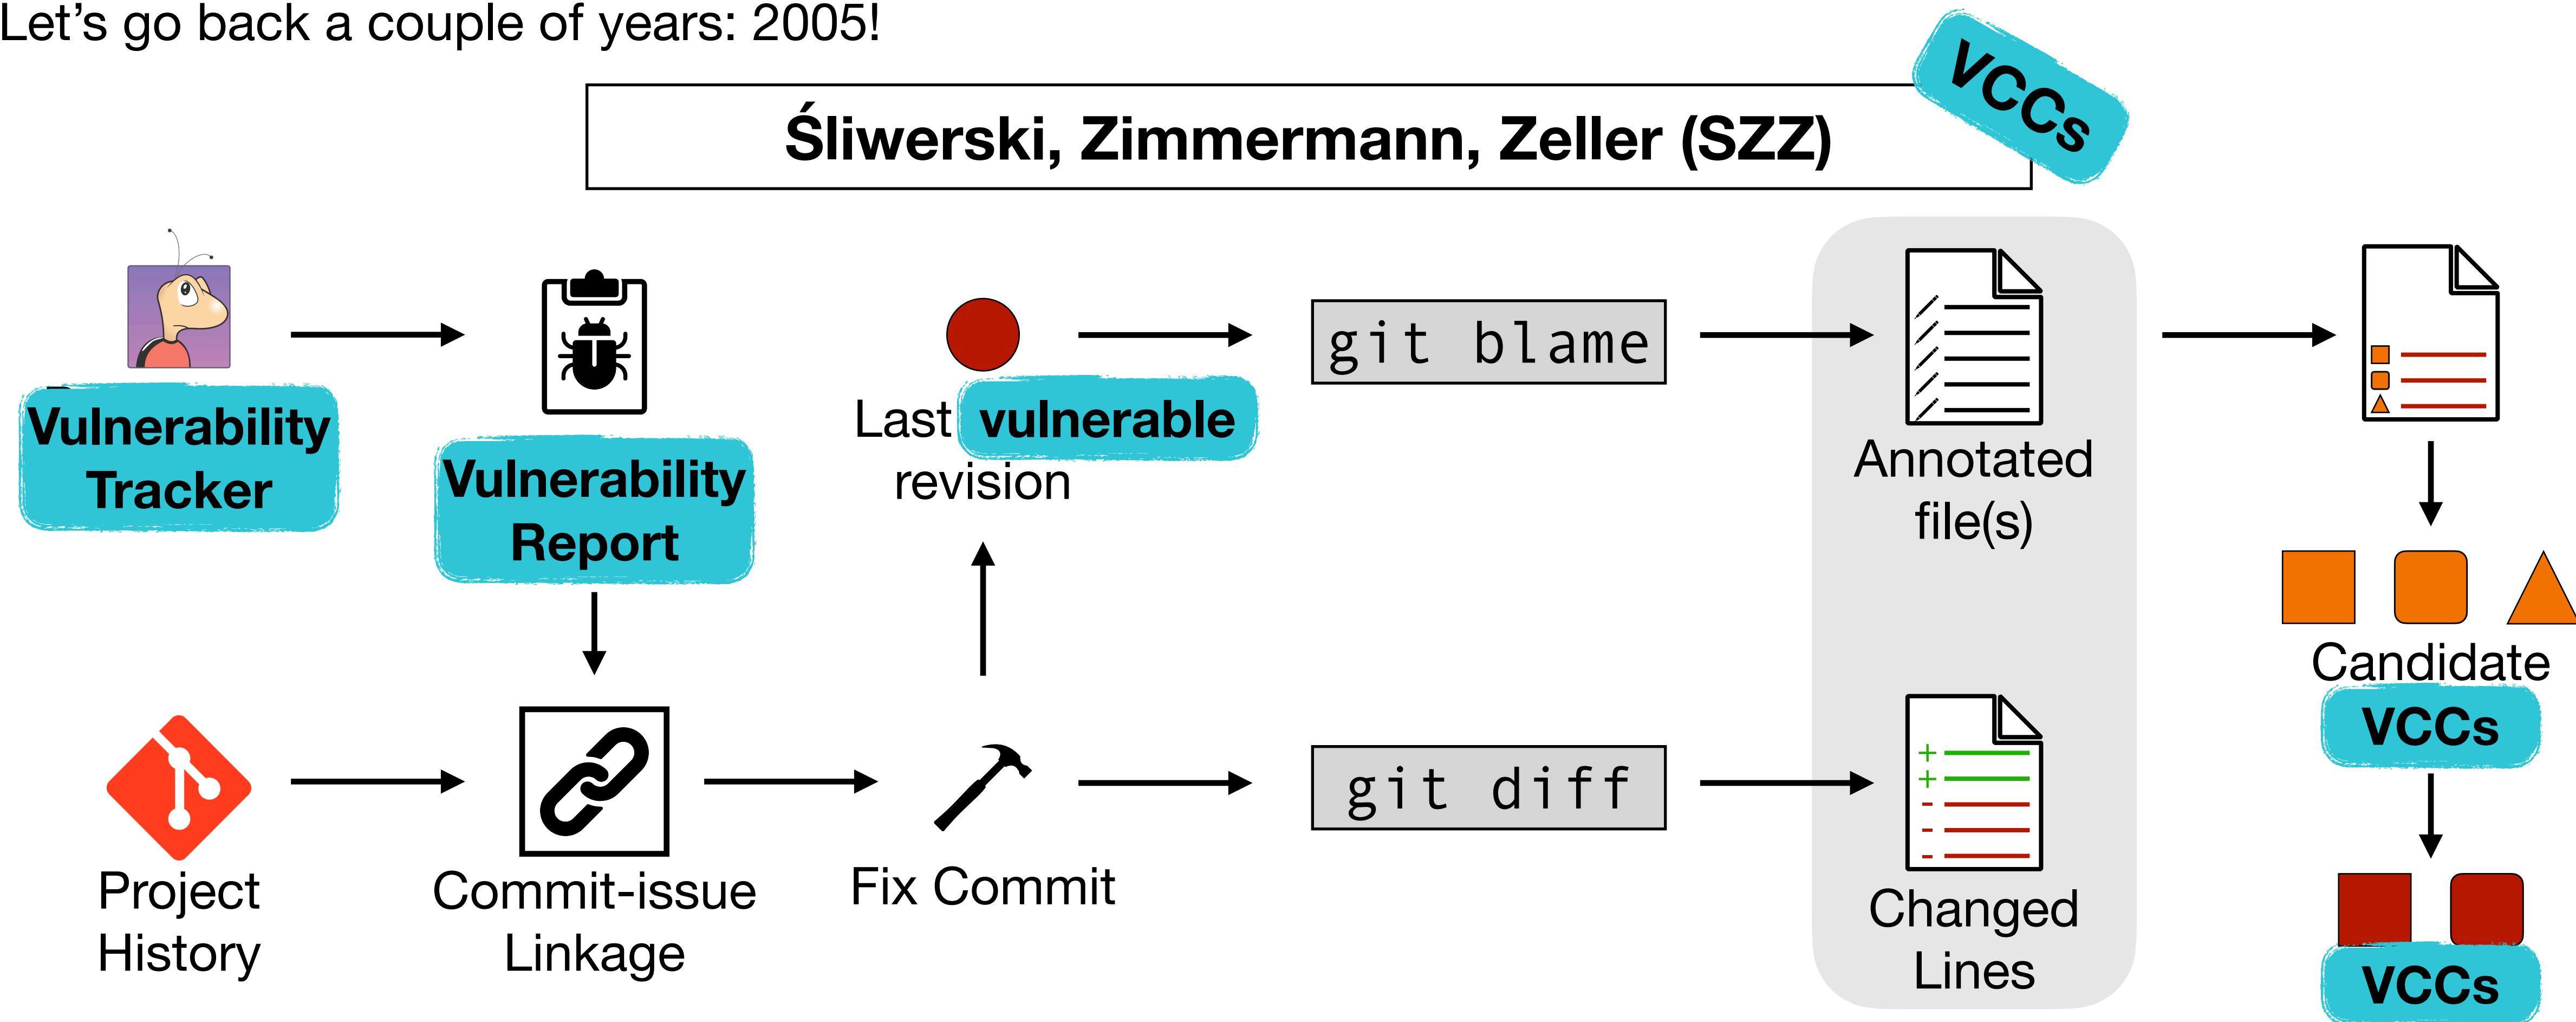
Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



Mining VCCs: Borrowing from the Bug World

Meneely et al.'s technique doesn't scale: it's manual and time-consuming. We need a fully-automated solution.
Let's go back a couple of years: 2005!



Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Comments and Blank Lines

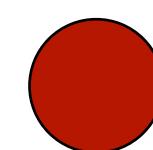
If the fixing commit also **modified an existing comment** or **removed a blank line**, the BICs (or VCCs) resulting from blaming these lines would be false positives: they made no real contribution to the bug.

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

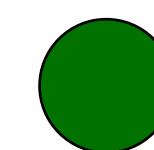
Comments and Blank Lines

If the fixing commit also **modified an existing comment** or **removed a blank line**, the BICs (or VCCs) resulting from blaming these lines would be false positives: they made no real contribution to the bug.



Last buggy/vulnerable

```
1: public void foo() {  
2:     // print report  
3:     if (report == null)  
4:     {  
5:         println(report);  
6:     }  
7: }
```



Fixed Revision

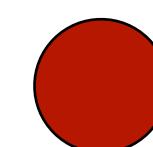
```
1: public void foo() {  
2:     // print out report  
3:     if (report != null)  
4:     {  
5:         println(report);  
6:     }  
7: }
```

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Comments and Blank Lines

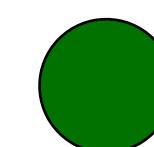
If the fixing commit also **modified an existing comment** or **removed a blank line**, the BICs (or VCCs) resulting from blaming these lines would be false positives: they made no real contribution to the bug.



Last buggy/vulnerable

Two lines changed,
one was just
deleted.

```
1: public void foo() {  
2:     // print report  
3:     if (report == null)  
4:     {  
5:         println(report);  
6:     }  
7: }
```



Fixed Revision

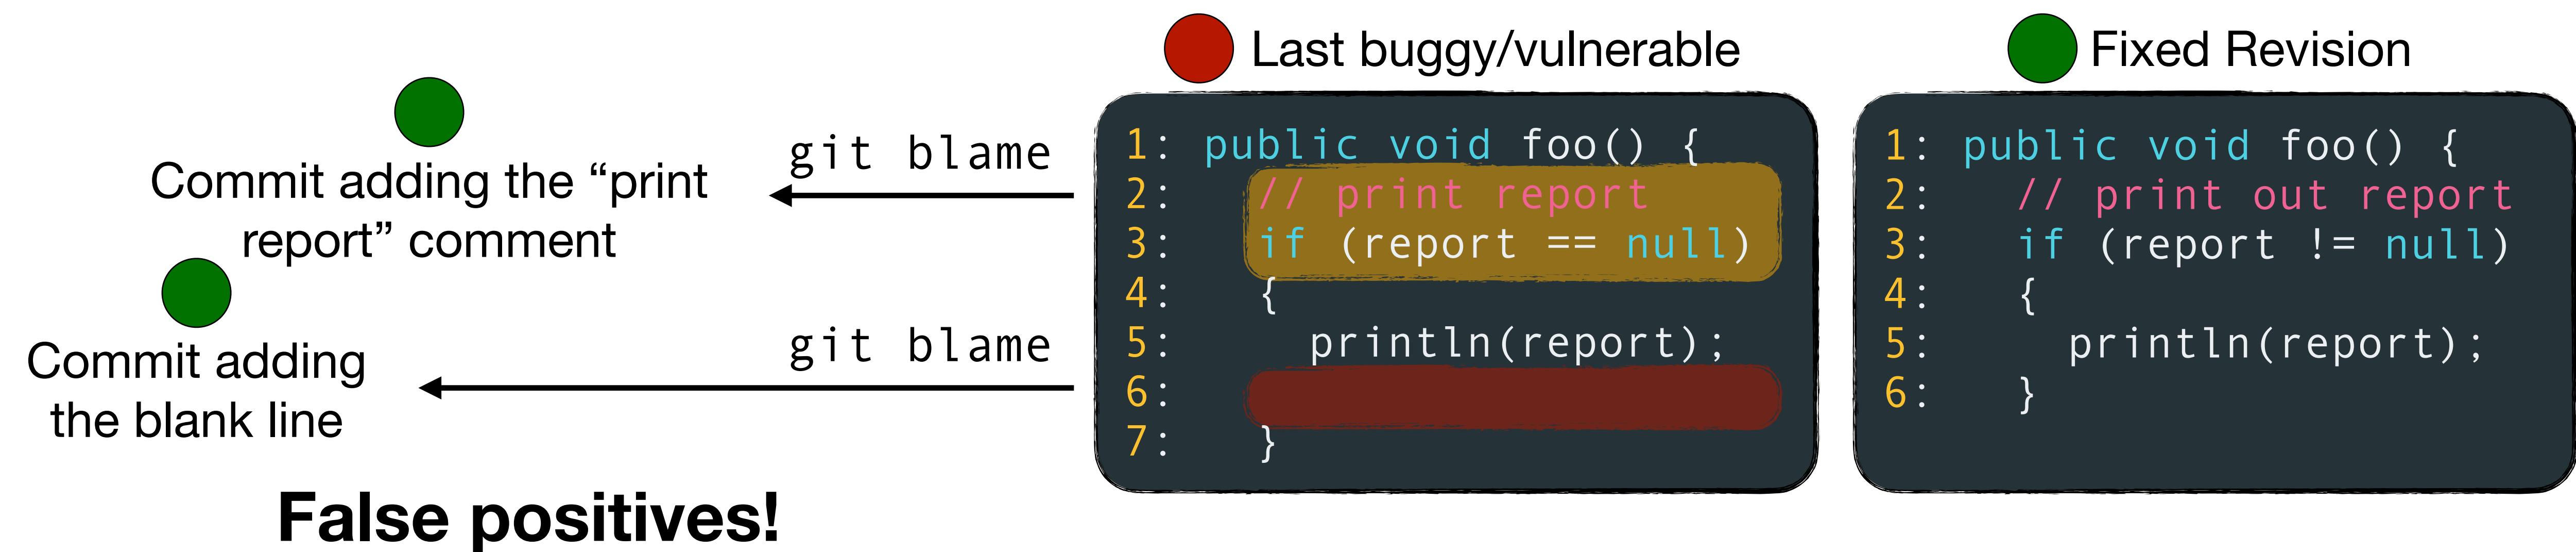
```
1: public void foo() {  
2:     // print out report  
3:     if (report != null)  
4:     {  
5:         println(report);  
6:     }  
7: }
```

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Comments and Blank Lines

If the fixing commit also **modified an existing comment** or **removed a blank line**, the BICs (or VCCs) resulting from blaming these lines would be false positives: they made no real contribution to the bug.



Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Format/Aesthetic Changes

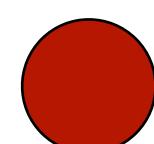
If the fixing commit **modified a line that underwent at least one format change after the bug was introduced**, the BICs (or VCCs) resulting from blaming these lines would be false positives, and the real BICs (VCCs) will be false negatives.

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

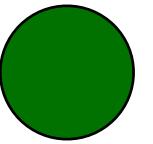
Format/Aesthetic Changes

If the fixing commit **modified a line that underwent at least one format change after the bug was introduced**, the BICs (or VCCs) resulting from blaming these lines would be false positives, and the real BICs (VCCs) will be false negatives.



Last buggy/vulnerable

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```



Fixed Revision

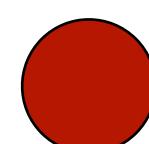
```
1: public void foo() {  
2:     if (folder != null)  
3:         return;
```

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

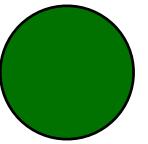
Format/Aesthetic Changes

If the fixing commit **modified a line that underwent at least one format change after the bug was introduced**, the BICs (or VCCs) resulting from blaming these lines would be false positives, and the real BICs (VCCs) will be false negatives.



Last buggy/vulnerable

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```



Fixed Revision

```
1: public void foo() {  
2:     if (folder != null)  
3:         return;
```

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Format/Aesthetic Changes

If the fixing commit **modified a line that underwent at least one format change after the bug was introduced**, the BICs (or VCCs) resulting from blaming these lines would be false positives, and the real BICs (VCCs) will be false negatives.

Revision B

```
1: public void foo() {  
2:     if (folder == null) return;
```

Last buggy/vulnerable

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```

Revision A

```
1: public void foo() {  
2:     if (folder != null) return;
```

Revision C

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Format/Aesthetic Changes

If the fixing commit **modified a line that underwent at least one format change after the bug was introduced**, the BICs (or VCCs) resulting from blaming these lines would be false positives, and the real BICs (VCCs) will be false negatives.

The commit that brought A to B is adding the bug/vulnerability!

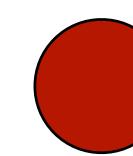
Revision A

```
1: public void foo() {  
2:     if (folder != null) return;
```



Revision B

```
1: public void foo() {  
2:     if (folder == null) return;
```



Last buggy/vulnerable

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```

Revision C

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

Format/Aesthetic Changes

If the fixing commit **modified a line that underwent at least one format change after the bug was introduced**, the BICs (or VCCs) resulting from blaming these lines would be false positives, and the real BICs (VCCs) will be false negatives.

C is the last commit that changed line 2 (false positive), shadowing B (false negative)!

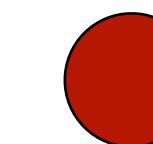
```
1: public void foo() {  
2:     if (folder != null) return;
```

Revision A



Revision B

```
1: public void foo() {  
2:     if (folder == null) return;
```



Last buggy/vulnerable

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```

Revision C

```
1: public void foo() {  
2:     if (folder == null)  
3:         return;
```

git blame

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

SZZ by Kim et al.

Automatic Identification of Bug-Introducing Changes

Sunghun Kim¹, Thomas Zimmermann², Kai Pan¹, E. James Whitehead, Jr.¹

¹*University of California,
Santa Cruz, CA, USA
{hunkim, pankai, ejw}@cs.ucsc.edu*

²*Saarland University,
Saarbrücken, Germany
tz@dcn.org*

Abstract

Bug-fixes are widely used for predicting bugs or finding risky parts of software. However, a bug-fix does not contain information about the change that initially introduced a bug. Such bug-introducing changes can help identify important properties of software bugs such as correlated factors or causalities. For example, they reveal which developers or what kinds of source code changes introduce more bugs. In contrast to bug-fixes that are relatively easy to obtain, the extraction of bug-introducing changes is challenging.

In this paper, we present algorithms to automatically and accurately identify bug-introducing changes. We remove false positives and false negatives by using annotation graphs, by ignoring non-semantic source code changes, and outlier fixes. Additionally, we validated that the fixes we used are true fixes by a manual inspection. Altogether, our algorithms can remove about 38%–51% of false positives and 14%–15% of false negatives compared to the previous algorithm. Finally, we show applications of bug-introducing changes that demonstrate their value for research.

1. Introduction

Today, software bugs remain a constant and costly fixture of industrial and open source software development. To manage the flow of bugs, software projects carefully control their changes using software configuration management (SCM) systems, capture bug reports using bug tracking software (such as Bugzilla), and then record which change in the SCM system fixes a specific bug in the change tracking system.

The progression of a single bug is as follows. A programmer makes a change to a software system, either to add new functionality, restructure the code, or to repair an existing bug. In the process of making this change, they inadvertently introduce a bug into the software. We call this a *bug-introducing change*, the modification in which a bug was injected into the software. At some later time, this bug manifests itself in some undesired external behavior, which is recorded in a bug tracking system. Subsequently, a developer modifies the project's source code, possibly changing multiple files, and repairs the bug. They commit this change to the SCM system,

permanently recording the change. As part of the commit, developers commonly (but not always) record in the SCM system change log the identifier of the bug report that was just fixed. We call this modification a *bug-fix change*.

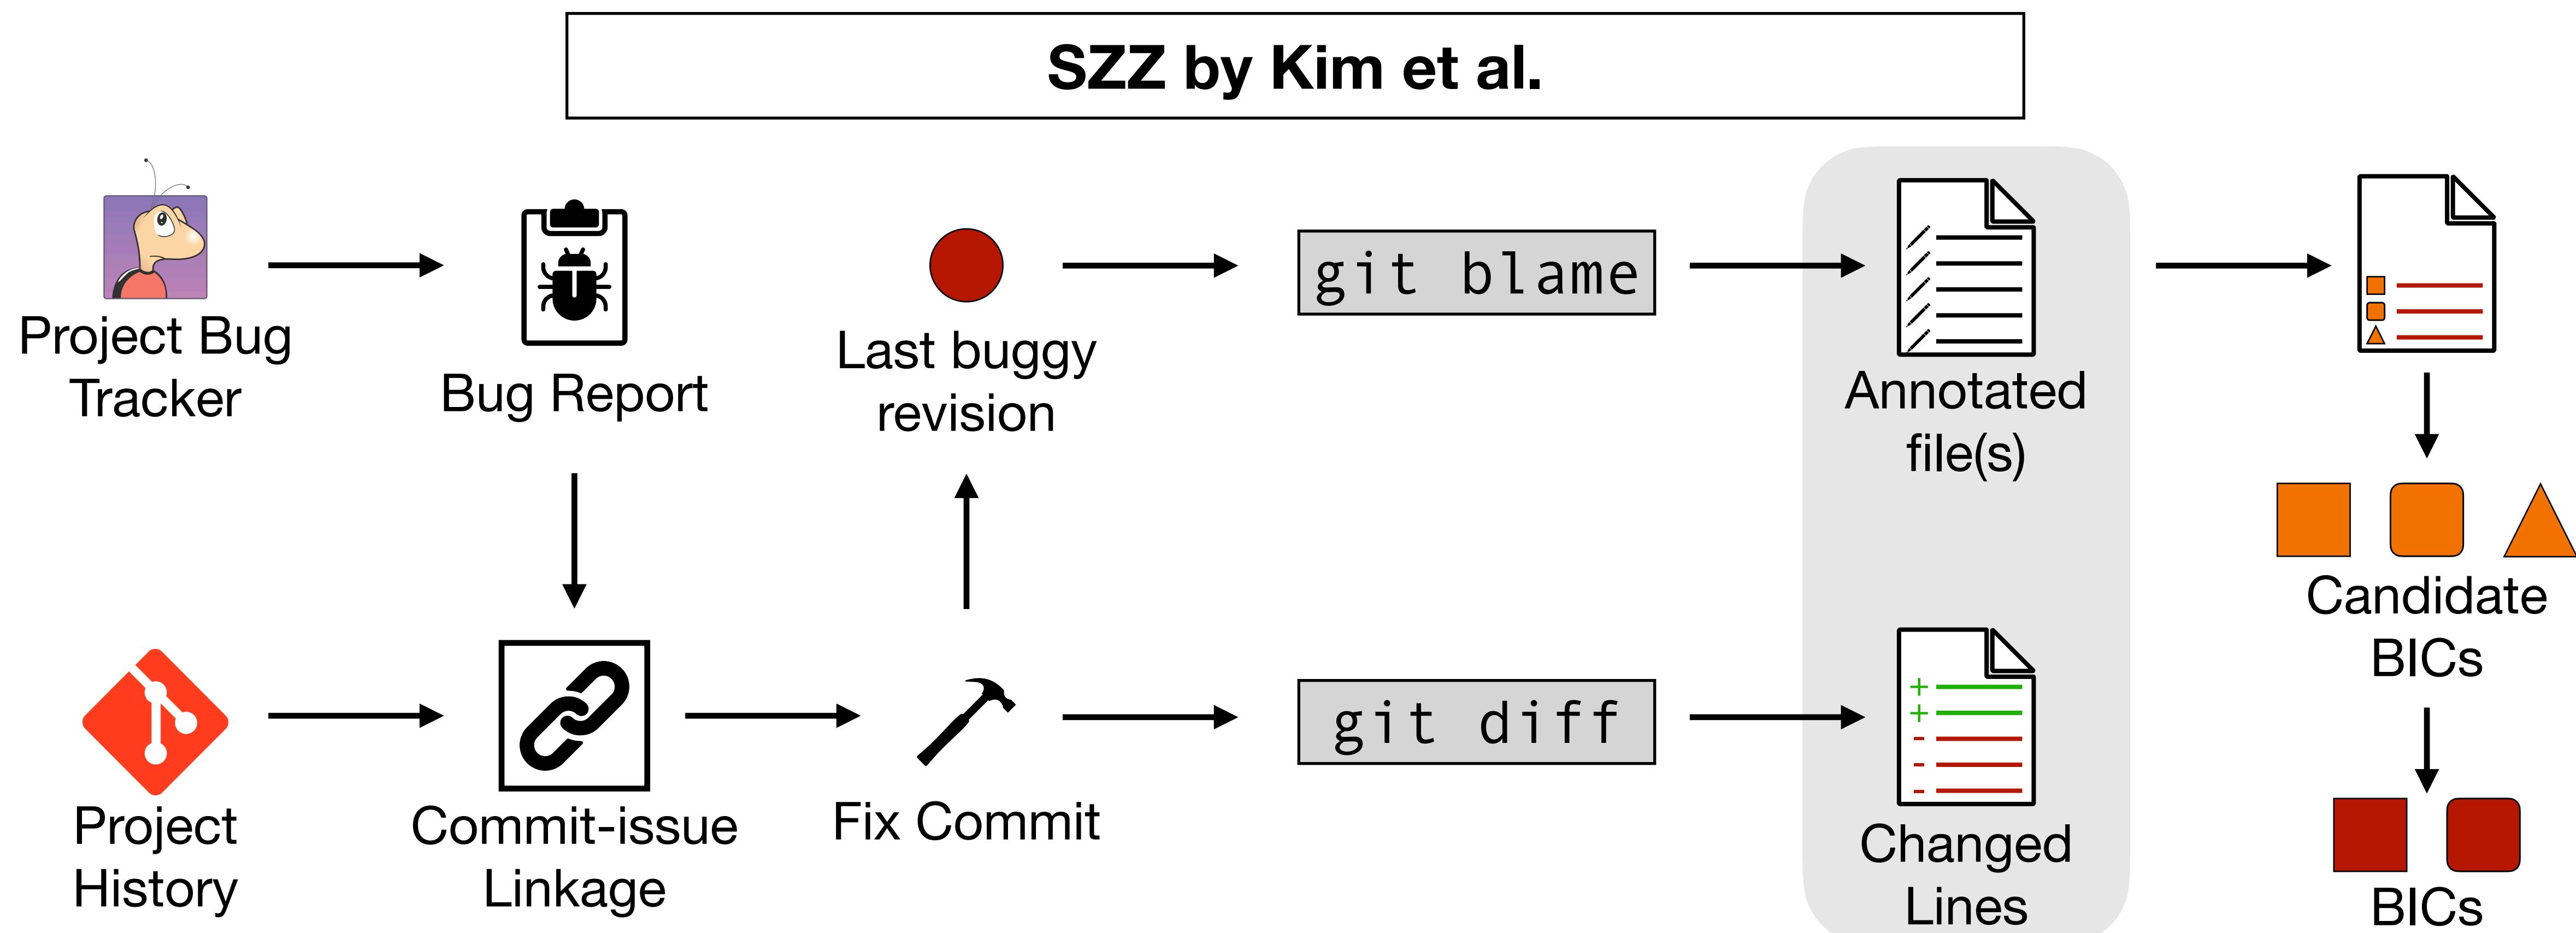
Software evolution research leverages the history of changes and bug reports that accretes over time in SCM systems and bug tracking systems to improve our understanding of how a project has grown. It offers the possibility that by examining the history of changes made to a software project, we might better understand patterns of bug introduction, and raise developer awareness that they are working on risky—that is, bug-prone—sections of a project. For example, if we can find rules that associate bug-introducing changes with certain source code change patterns (such as signature changes that involve parameter addition [11]), it may be possible to identify source code change patterns that are bug-prone.

Due to the widespread use of bug tracking and SCM systems, the most readily available data concerning bugs are the bug-fix changes. It is easy to mine an SCM repository to find those changes that have repaired a bug. To do so, one examines change log messages in two ways: searching for keywords such as “Fixed” or “Bug” [12] and searching for references to bug reports like “#42233” [2, 4, 16]. With bug-fix information, researchers can determine the *location* of a bug. This permits useful analysis, such as determining per-file bug counts, predicting bugs, finding risky parts of software [7, 13, 14], or visually revealing the relationship between bugs and software evolution [3].

The major problem with bug-fix data is that it sheds no light on *when* a bug was injected into the code and *who* injected it. The person fixing a bug is often not the person who first made the bug, and the bug-fix must, by definition, occur after the bug was first injected. Bug-fix data also provides imprecise data on *where* a bug occurred. Since functions and methods change their names over time, the fact that a fix was made to function “foo” does not mean the function still had that name when the bug was injected; it could have been named “bar” then. In order to deeply understand the phenomena surrounding the introduction of bugs into code, such as correlated factors and causalities, we need access to the actual moment and point the bug was introduced. This is tricky, and the focus of our paper.

Mining VCCs: Borrowing from the Bug World

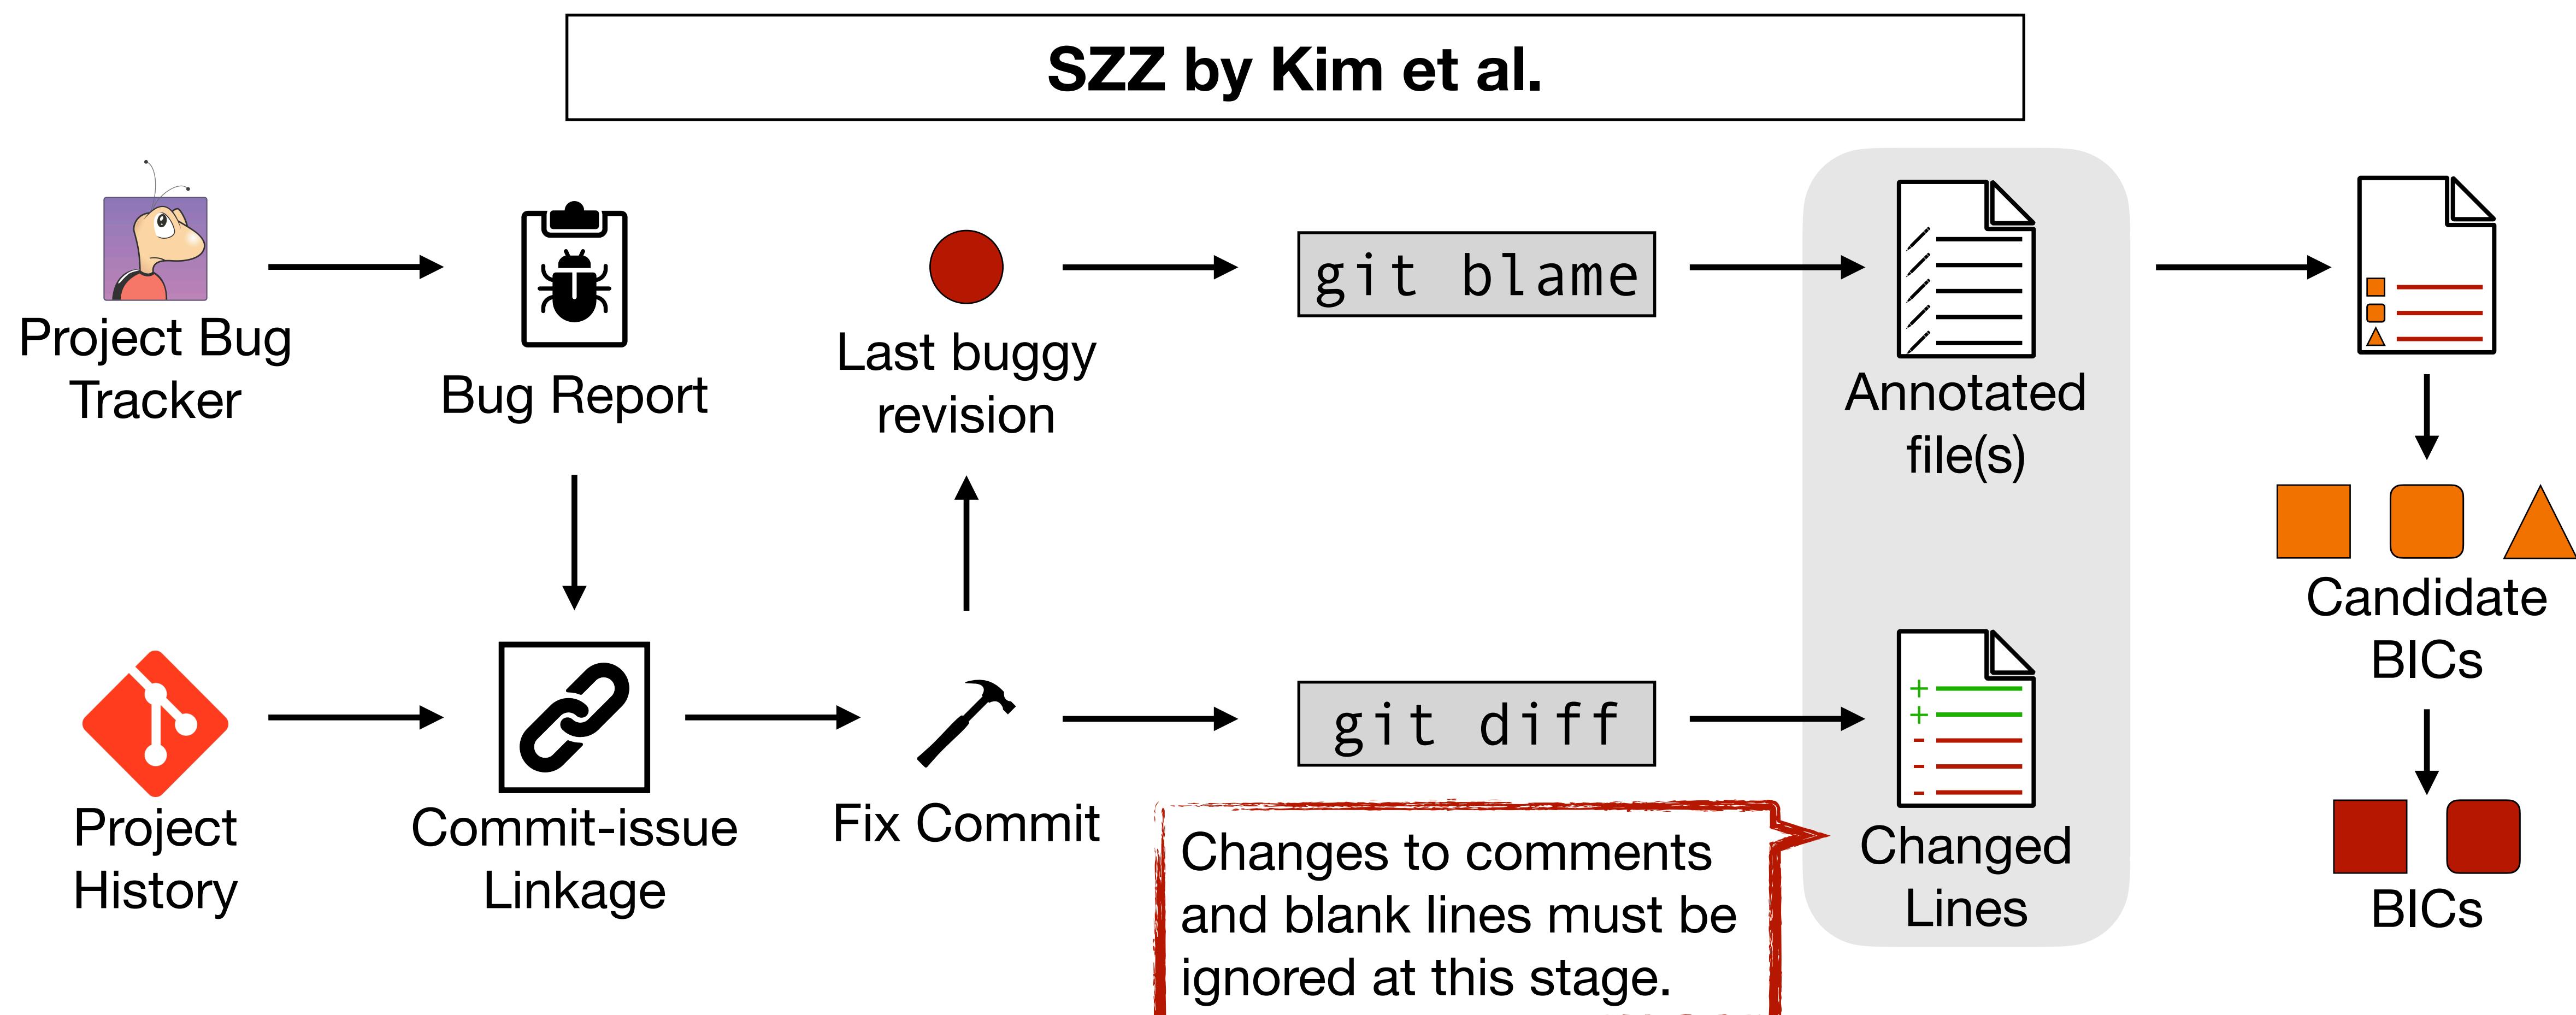
The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.



Let's go back to the original SZZ...

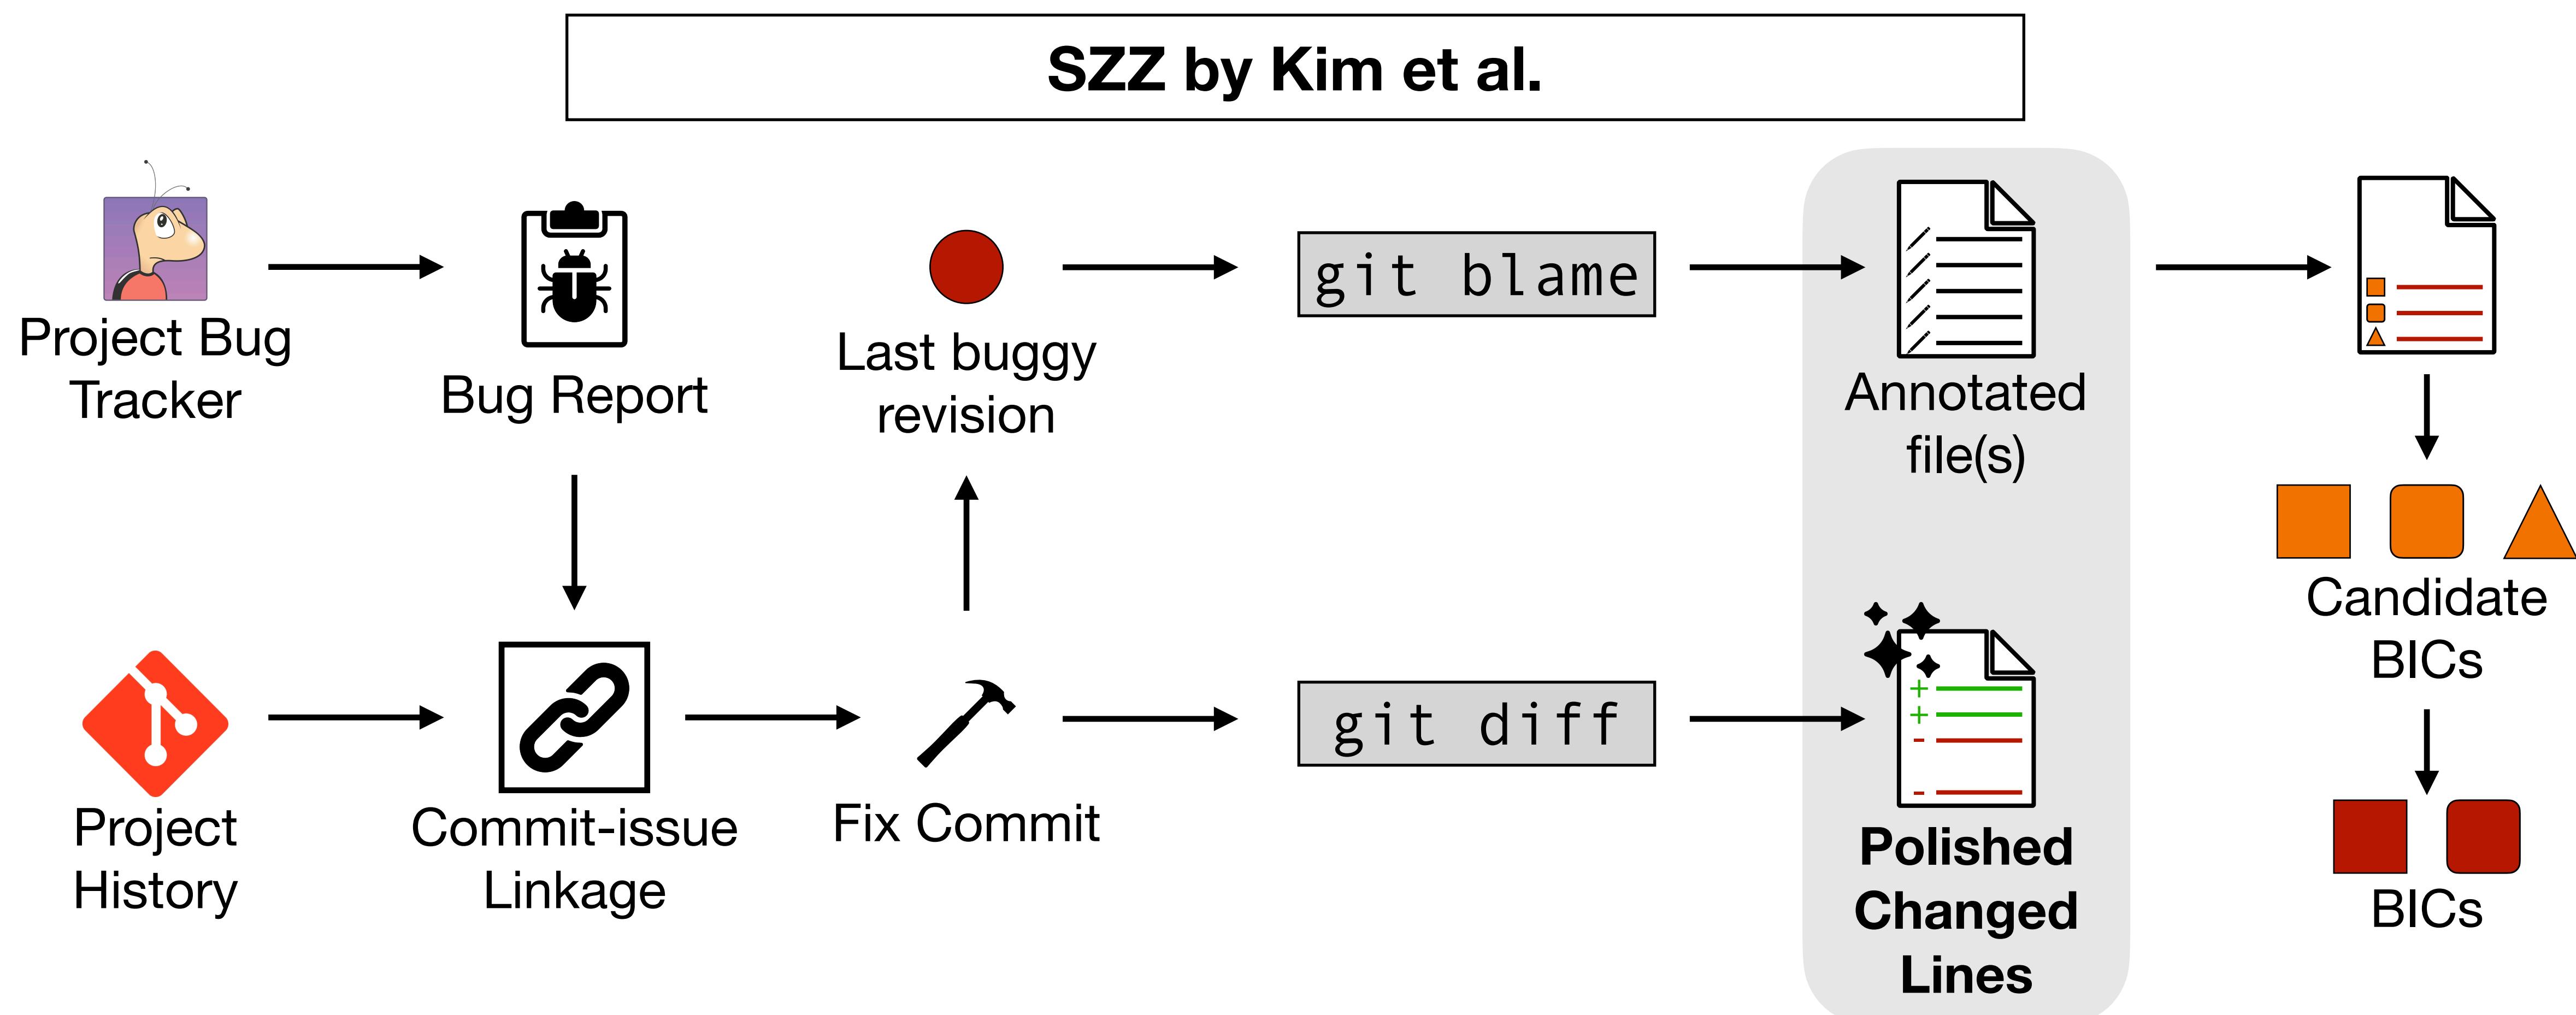
Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.



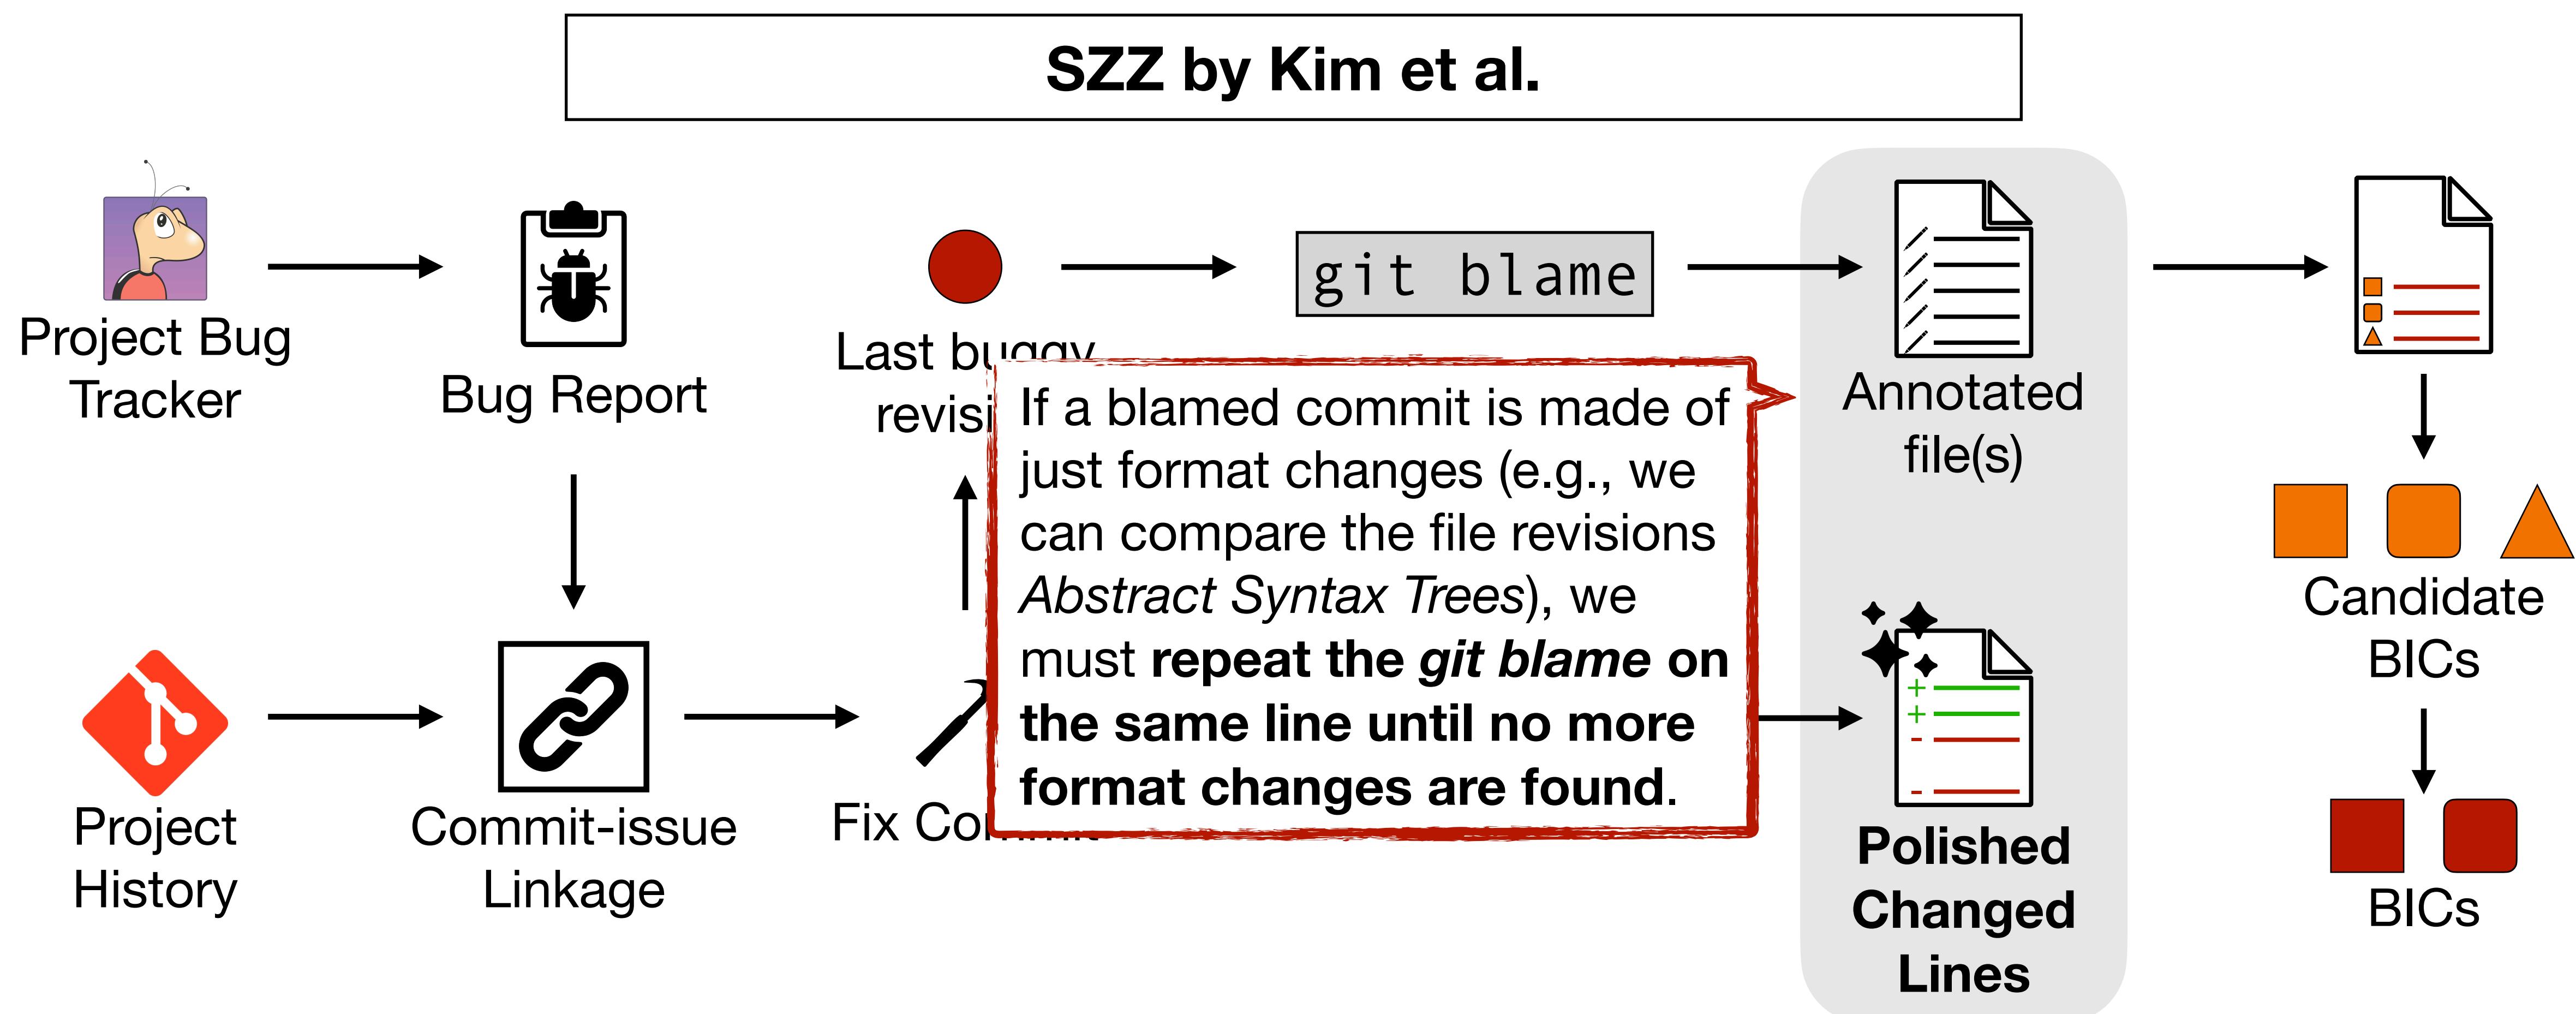
Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.



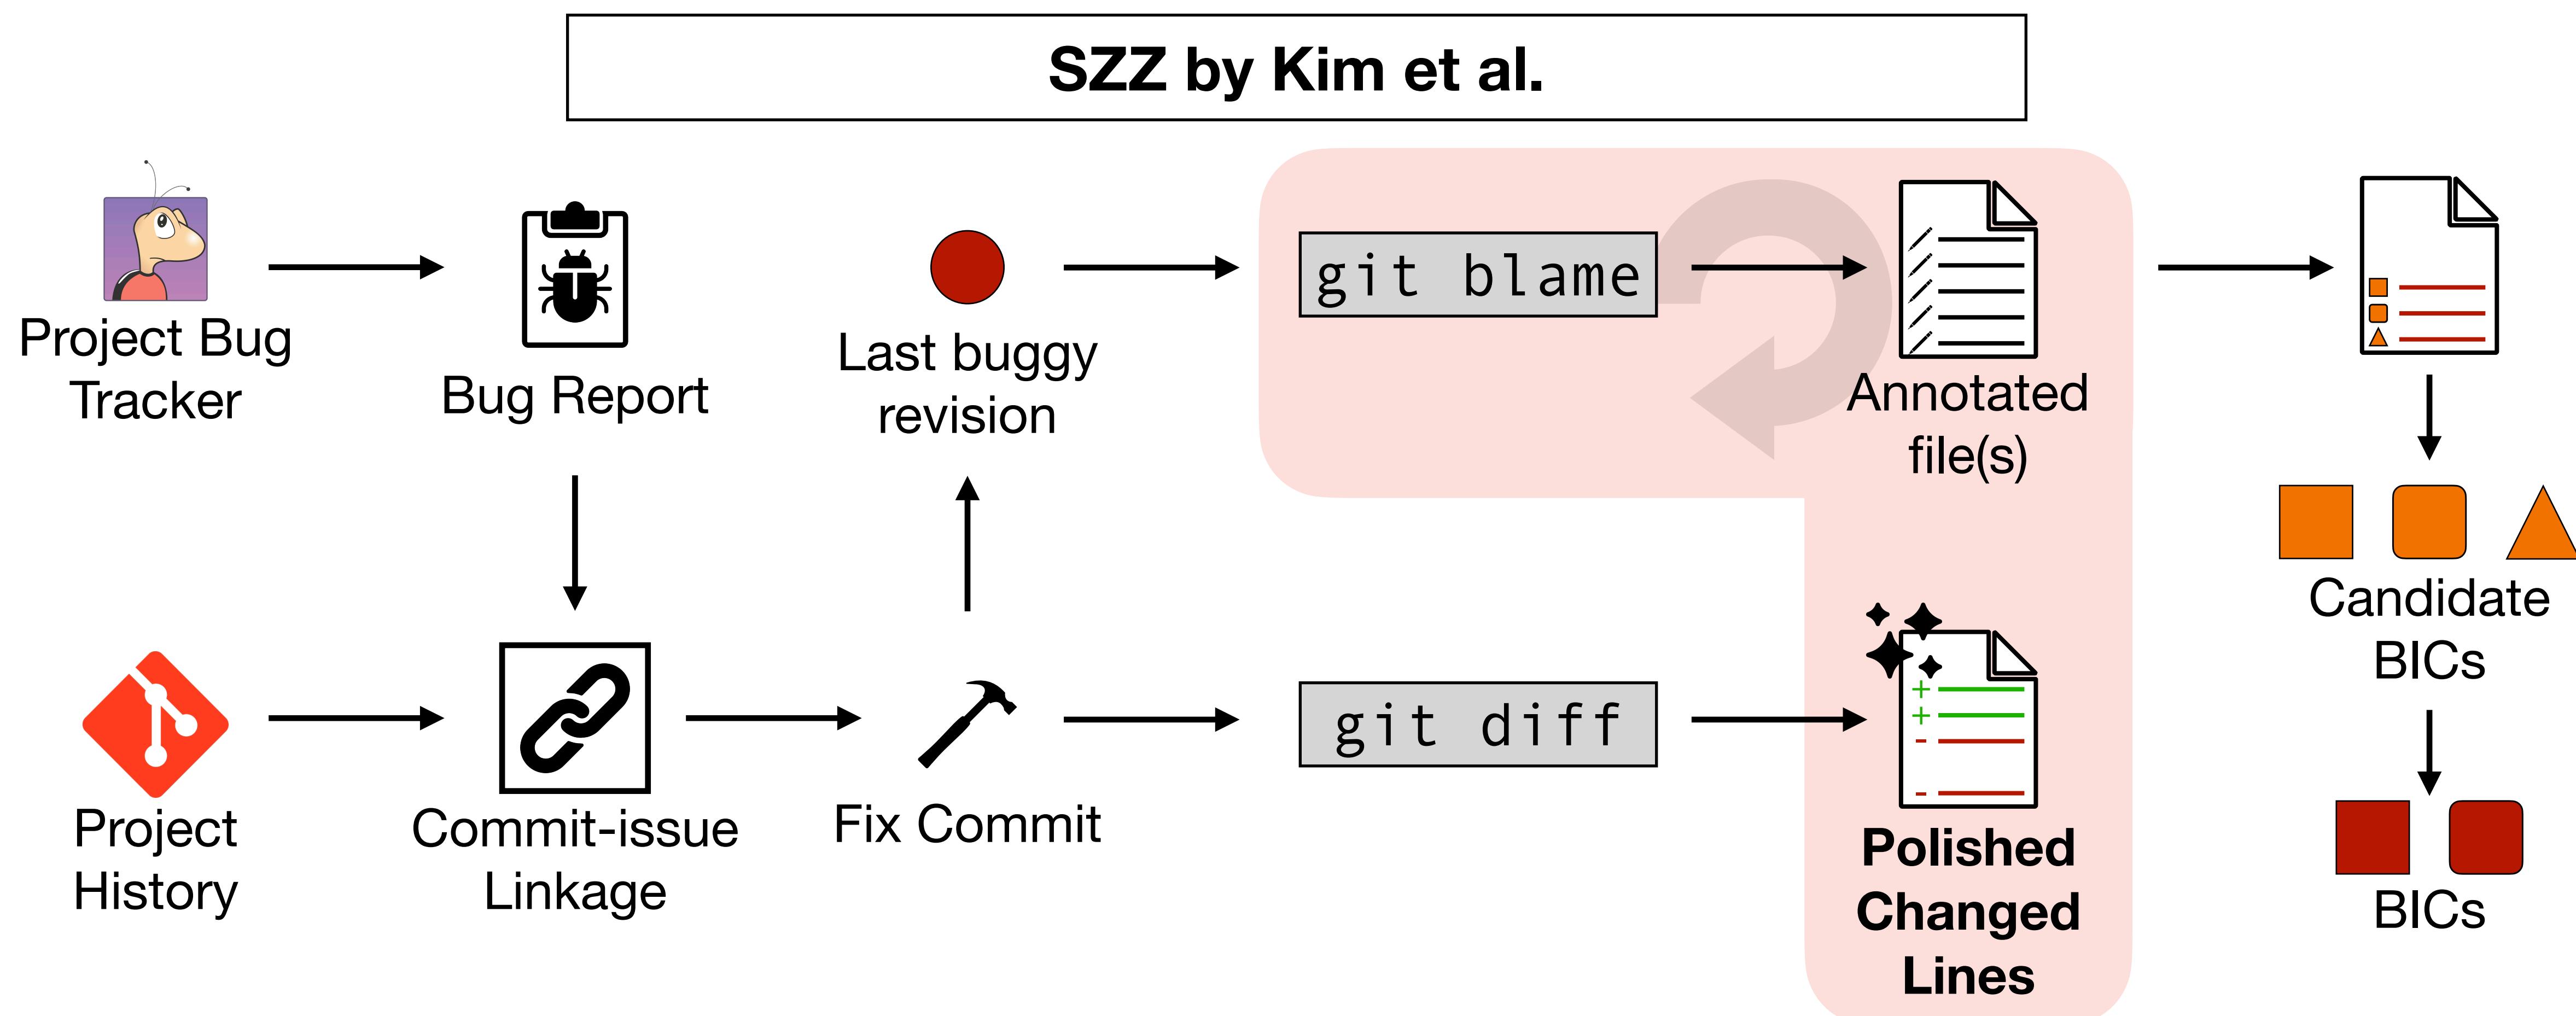
Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.



Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.



Mining VCCs: Borrowing from the Bug World

This is surely a good improvement, but there are still some more problems...

Mining VCCs: Borrowing from the Bug World

This is surely a good improvement, but there are still some more problems...

Meta-changes

The set of candidate BICs/VCCs might be made of commits that do not really modify the source code, e.g., **merge commits**, which incorporate commits from one branch into another.

Mining VCCs: Borrowing from the Bug World

This is surely a good improvement, but there are still some more problems...

Meta-changes

The set of candidate BICs/VCCs might be made of commits that do not really modify the source code, e.g., **merge commits**, which incorporate commits from one branch into another.

SZZ by da Costa et al.

A Framework for Evaluating the Results of the SZZ Approach for Identifying Bug-Introducing Changes

Daniel Alencar da Costa, Shane McIntosh, Weiyi Shang, Uirá Kulesza, Roberta Coelho, and Ahmed E. Hassan

Abstract—The approach proposed by Śliwerski, Zimmermann, and Zeller (SZZ) for identifying bug-introducing changes is at the foundation of several research areas within the software engineering discipline. Despite the foundational role of SZZ, little effort has been made to evaluate its results. Such an evaluation is a challenging task because the ground truth is not readily available. By acknowledging such challenges, we propose a framework to evaluate the results of alternative SZZ implementations. The framework evaluates the following criteria: (1) the earliest bug appearance, (2) the future impact of changes, and (3) the realism of bug introduction. We use the proposed framework to evaluate five SZZ implementations using data from ten open source projects. We find that previously proposed improvements to SZZ tend to inflate the number of incorrectly identified bug-introducing changes. We also find that a single bug-introducing change may be blamed for introducing hundreds of future bugs. Furthermore, we find that SZZ implementations report that at least 46 percent of the bugs are caused by bug-introducing changes that are years apart from one another. Such results suggest that current SZZ implementations still lack mechanisms to accurately identify bug-introducing changes. Our proposed framework provides a systematic mean for evaluating the data that is generated by a given SZZ implementation.

Index Terms—SZZ, evaluation framework, bug detection, software repository mining

1 INTRODUCTION

SOFTWARE bugs are costly to fix [1]. For instance, a recent study suggests that developers spend approximately half of their time fixing bugs [2]. Hence, reducing the required time and effort to fix bugs is an alluring research problem with plenty of potential for industrial impact.

After a bug has been reported, a key task is to identify the root cause of the bug such that a team can learn from its mistakes. Hence, researchers have developed several approaches to identify prior bug-introducing changes, and to use such knowledge to avoid future bugs [3], [4], [5], [6], [7], [8], [9], [10].

A popular approach to identify bug-introducing changes was proposed by Śliwerski, Zimmermann, and Zeller (‘‘SZZ’’ for short) [9], [11]. The SZZ approach first looks for

bug-fixing changes by searching for the recorded bug ID in change logs. Once these bug-fixing changes are identified, SZZ analyzes the lines of code that were changed to fix the bug. Finally, SZZ traces back through the code history to find when the changed code was introduced (i.e., the supposed bug-introducing change(s)).

Two lines of prior work highlight the foundational role of SZZ in software engineering (SE) research. The first line includes studies of how bugs are introduced [9], [10], [12], [13], [14], [15], [16], [17], [18], [19], [20], [21], [22]. For example, by studying the bug-introducing changes that are identified by SZZ, researchers are able to correlate characteristics of code changes (e.g., time of day that a change is recorded [9]) with the introduction of bugs. The second line of prior work includes studies that leverage the knowledge of prior bug-introducing changes in order to avoid the introduction of such changes in the future. For example, one way to avoid the introduction of bugs is to perform *just-in-time* (JIT) quality assurance, i.e., to build models that predict if a change is likely to be a bug-introducing change before integrating such a change into a project’s code base. [6], [8], [23], [24], [25].

Despite the foundational role of SZZ, the current evaluations of SZZ-generated data (the indicated bug-introducing changes) are limited. When evaluating the results of SZZ implementations, prior work relies heavily on manual analysis [9], [11], [26], [27]. Since it is infeasible to analyze all of the SZZ results by hand, prior studies select a small sample for analysis. While the prior manual analyses yield valuable insights, the domain experts (e.g., developers or testers) were not consulted. These experts can better judge

- D.A. da Costa, U. Kulesza, and R. Coelho are with the Department of Informatics and Applied Mathematics (DIMAp), Federal University of Rio Grande do Norte, Natal-RN 59078-970, Brazil.
E-mail: dcosta@dimap.ufrn.br; uira, roberta@dimap.ufrn.br.
- S. McIntosh is with the Department of Electrical and Computer Engineering, McGill University, Montreal, QC H3A 0C4, Canada.
E-mail: shane.mcintosh@mcgill.ca.
- W. Shang is with the Department of Computer Science and Software Engineering, Concordia University, Montreal, QC H4B 1R6, Canada.
E-mail: shang@cs.concordia.ca.
- U. Kulesza is with the Software Analysis and Intelligence Lab (SAIL), School of Computing, Queen’s University, Kingston, ON K7L 3N6, Canada.
E-mail: ahmed@cs.queensu.ca.

Manuscript received 4 Sept. 2015; revised 13 Sept. 2016; accepted 30 Sept. 2016. Date of publication 10 Oct. 2016; date of current version 24 July 2017. Recommended for acceptance by A. Zeller.
For information on obtaining reprints of this article, please send e-mail to: reprints@iee.org, and reference the Digital Object Identifier below.
Digital Object Identifier no. 10.1109/TSE.2016.2616306

0098-5589 © 2016 IEEE. Personal use is permitted, but republication/redistribution requires IEEE permission.
See http://www.ieee.org/publications_standards/publications/rights/index.html for more information.

Authorized licensed use limited to: Universita degli Studi di Salerno. Downloaded on May 09, 2023 at 14:55:22 UTC from IEEE Xplore. Restrictions apply.

Mining VCCs: Borrowing from the Bug World

This is surely a good improvement, but there are still some more problems...

Meta-changes

The set of candidate BICs/VCCs might be made of commits that do not really modify the source code, e.g., **merge commits**, which incorporate commits from one branch into another.

Basically, it's a variant of the SZZ by Kim et al. that **ignores merge commits while traversing the history with the repeated *git blames*.**

SZZ by da Costa et al.

A Framework for Evaluating the Results of the SZZ Approach for Identifying Bug-Introducing Changes

Daniel Alencar da Costa, Shane McIntosh, Weiyi Shang, Uirá Kulesza, Roberta Coelho, and Ahmed E. Hassan

Abstract—The approach proposed by Śliwerski, Zimmermann, and Zeller (SZZ) for identifying bug-introducing changes is at the foundation of several research areas within the software engineering discipline. Despite the foundational role of SZZ, little effort has been made to evaluate its results. Such an evaluation is a challenging task because the ground truth is not readily available. By acknowledging such challenges, we propose a framework to evaluate the results of alternative SZZ implementations. The framework evaluates the following criteria: (1) the earliest bug appearance, (2) the future impact of changes, and (3) the realism of bug introduction. We use the proposed framework to evaluate five SZZ implementations using data from ten open source projects. We find that previously proposed improvements to SZZ tend to inflate the number of incorrectly identified bug-introducing changes. We also find that a single bug-introducing change may be blamed for introducing hundreds of future bugs. Furthermore, we find that SZZ implementations report that at least 46 percent of the bugs are caused by bug-introducing changes that are years apart from one another. Such results suggest that current SZZ implementations still lack mechanisms to accurately identify bug-introducing changes. Our proposed framework provides a systematic mean for evaluating the data that is generated by a given SZZ implementation.

Index Terms—SZZ, evaluation framework, bug detection, software repository mining

1 INTRODUCTION

SOFTWARE bugs are costly to fix [1]. For instance, a recent study suggests that developers spend approximately half of their time fixing bugs [2]. Hence, reducing the required time and effort to fix bugs is an alluring research problem with plenty of potential for industrial impact.

After a bug has been reported, a key task is to identify the root cause of the bug such that a team can learn from its mistakes. Hence, researchers have developed several approaches to identify prior bug-introducing changes, and to use such knowledge to avoid future bugs [3], [4], [5], [6], [7], [8], [9], [10].

A popular approach to identify bug-introducing changes was proposed by Śliwerski, Zimmermann, and Zeller (“SZZ” for short) [9], [11]. The SZZ approach first looks for

bug-fixing changes by searching for the recorded bug ID in change logs. Once these bug-fixing changes are identified, SZZ analyzes the lines of code that were changed to fix the bug. Finally, SZZ traces back through the code history to find when the changed code was introduced (i.e., the supposed bug-introducing change(s)).

Two lines of prior work highlight the foundational role of SZZ in software engineering (SE) research. The first line includes studies of how bugs are introduced [9], [10], [12], [13], [14], [15], [16], [17], [18], [19], [20], [21], [22]. For example, by studying the bug-introducing changes that are identified by SZZ, researchers are able to correlate characteristics of code changes (e.g., time of day that a change is recorded [9]) with the introduction of bugs. The second line of prior work includes studies that leverage the knowledge of prior bug-introducing changes in order to avoid the introduction of such changes in the future. For example, one way to avoid the introduction of bugs is to perform *just-in-time* (JIT) quality assurance, i.e., to build models that predict if a change is likely to be a bug-introducing change before integrating such a change into a project’s code base. [6], [8], [23], [24], [25].

Despite the foundational role of SZZ, the current evaluations of SZZ-generated data (the indicated bug-introducing changes) are limited. When evaluating the results of SZZ implementations, prior work relies heavily on manual analysis [9], [11], [26], [27]. Since it is infeasible to analyze all of the SZZ results by hand, prior studies select a small sample for analysis. While the prior manual analyses yield valuable insights, the domain experts (e.g., developers or testers) were not consulted. These experts can better judge

- D.A. da Costa, U. Kulesza, and R. Coelho are with the Department of Informatics and Applied Mathematics (DIMAp), Federal University of Rio Grande do Norte, Natal-RN 59078-970, Brazil.
E-mail: dcosta@dimap.ufrn.br; uira@dimap.ufrn.br;
- S. McIntosh is with the Department of Electrical and Computer Engineering, McGill University, Montreal, QC H3A 0C4, Canada.
E-mail: shane.mcintosh@mail.mcgill.ca;
- W. Shang is with the Department of Computer Science and Software Engineering, Concordia University, Montreal, QC H4B 1R6, Canada.
E-mail: shang@cs.concordia.ca;
- A.E. Hassan is with the Software Analysis and Intelligence Lab (SAIL), School of Computing, Queen’s University, Kingston, ON K7L 3N6, Canada.
E-mail: ahmed@cs.queensu.ca.

Manuscript received 4 Sept. 2015; revised 13 Sept. 2016; accepted 30 Sept. 2016. Date of publication 10 Oct. 2016; date of current version 24 July 2017. Recommended for acceptance by A. Zeller.
For information on obtaining reprints of this article, please send e-mail to reprints@ieee.org, and reference the Digital Object Identifier below.
Digital Object Identifier no. 10.1109/TSE.2016.2616306

0989-5895 © 2016 IEEE. Personal use is permitted, but republication/redistribution requires IEEE permission.
See http://www.ieee.org/publications_standards/publications/rights/index.html for more information.

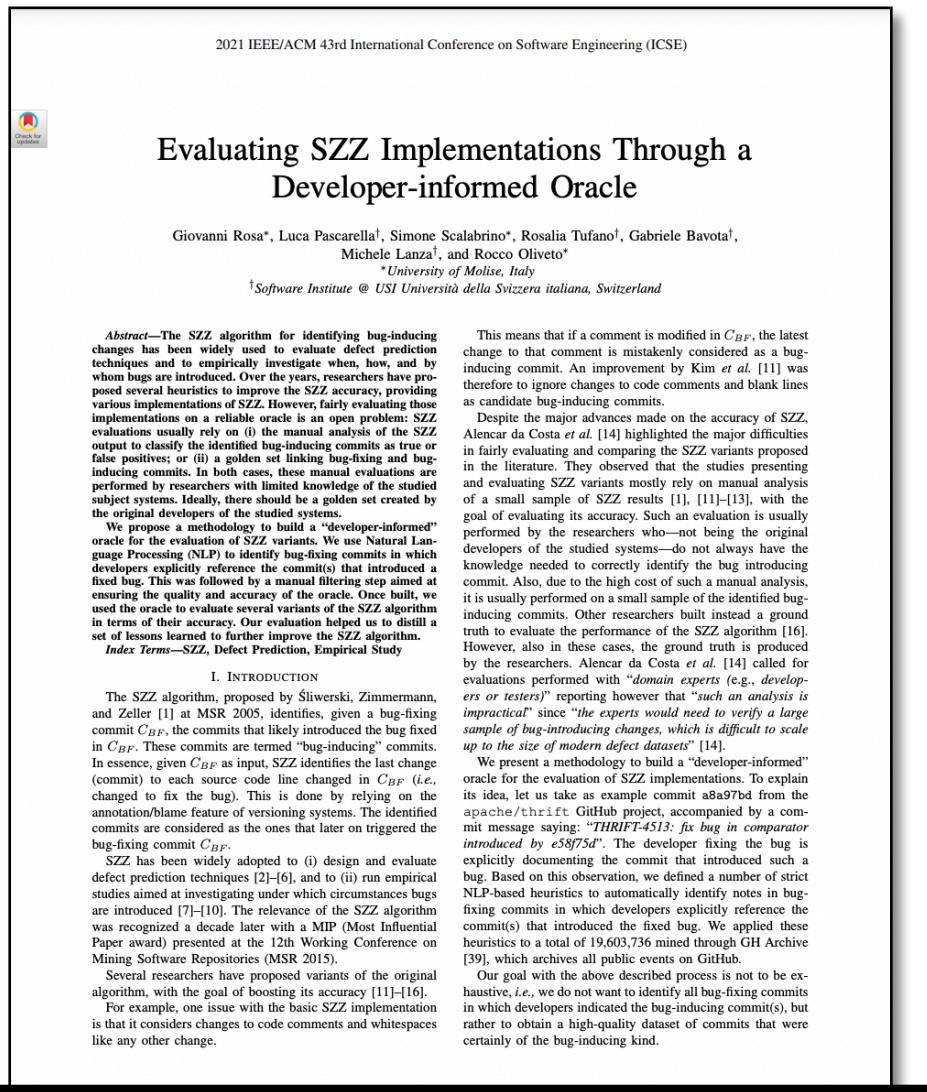
Authorized licensed use limited to: Universita degli Studi di Salerno. Downloaded on May 09, 2023 at 14:55:22 UTC from IEEE Xplore. Restrictions apply.

Mining VCCs: Borrowing from the Bug World

Many SZZ variants have been proposed over the years. It is difficult to remember them all or understand which is better. Luckily, some studies put things in order.

Mining VCCs: Borrowing from the Bug World

Many SZZ variants have been proposed over the years. It is difficult to remember them all or understand which is better. Luckily, some studies put things in order.



Rosa et al.

Comparison of nine SZZ variants on 123 OSS projects.

da Costa et al.

Comparison of five SZZ variants on ten OSS projects.



Rodríguez-Pérez et al.

Comparison of four SZZ variants on two OSS projects.

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

**Do Bugs Foreshadow Vulnerabilities?
A Study of the Chromium Project**

Felipe Camilo, Andrew Menseley, and Meiyappan Nagappan
Department of Software Engineering
Rochester Institute of Technology,
134 Lomb Memorial Drive
Rochester, NY, USA 14603
1-585-475-829
[\(f6c7162.mnxse.mnxse\)@rit.edu](mailto:(f6c7162.mnxse.mnxse)@rit.edu)

Abstract—In developers face ever-increasing pressure to engineer secure software, researchers are building an understanding of security-sensitive bugs (i.e. vulnerabilities). Research into mining software repositories has greatly increased our understanding of software bugs. However, vulnerabilities are conceptually different than traditional bugs. Vulnerabilities represent an abuse of functionality as opposed to wrong or insufficient functionality compared to what is expected. In this study, we conducted an empirical analysis of pre-release and post-release vulnerabilities in the Chromium project. We found 374 pre-release and 76 post-release vulnerabilities. Our results show that the relationship between bugs and vulnerabilities is weak. We found that 76 post-release vulnerabilities are associated with 10 pre-release vulnerabilities. While we found no association between bugs and vulnerabilities, we did find an association between the number of features, SLOC, and number of pre-release vulnerabilities. We also found that the relationship between bugs and vulnerabilities is stronger than any of the non-security bug categories. In a separate analysis, we found that the files with the highest density of bugs also had the highest density of highest vulnerability density. These results indicate that bugs and vulnerabilities are empirically distinct groups, warranting the need for more research targeting vulnerabilities specifically.

Keywords:—Software engineering, security, mining, software repositories, vulnerability, bug, developer behavior, open source, empirical study.

Computers & Security 99 (2020) 102067

Available online at www.sciencedirect.com
ScienceDirect
journal homepage: www.elsevier.com/locate/cose

Investigating the vulnerability fixing process in OSS projects: Peculiarities and challenges

Gerardo Canfora, Andrea Di Sorbo^a, Sara Forootani, Antonio Pirozzi,
Corrado Aaron Visaggio

Department of Engineering, University of Sannio, Italy

ABSTRACT
Although vulnerabilities can be considered and treated as bugs, they present numerous peculiarities compared to other types of bugs (canonical bugs in the remainder of the paper). A vulnerability adds functionality to a system, as it allows an adversary to misuse or abuse the system. In contrast, a canonical bug is a bug that is introduced by the execution of a requirement, and thus degrades the functionality of the system. This difference is important in the fixing process of vulnerabilities. By mining the repositories of 6 open source projects, we characterize the differences in the fixing process between vulnerabilities and canonical bugs, highlighting the peculiarities and the challenges for future research. Results clearly demonstrate that (i) re-assignments (the changes observed in canonical bugs) are required for fixing the developers able to handle vulnerability-related bugs, (ii) developers' security-related skills should be profiled, to improve the efficiency of the security bug assignment tasks, and, consequently, reduce the re-assignments, and (iii) vulnerabilities require more effort, contributors and time to define the fixing strategy but smaller time to fix than canonical bugs.

© 2020 Elsevier Ltd. All rights reserved.

Keywords:—Security, bug fixing, Process improvement, Software maintenance and evolution, Bug management, Empirical study

Information and Software Technology 142 (2022) 106745
Contents lists available at ScienceDirect
Information and Software Technology
journal homepage: www.elsevier.com/locate/infsof

Patchworking: Exploring the code changes induced by vulnerability fixing activities

Gerardo Canfora^a, Andrea Di Sorbo^a, Sara Forootani^{a,*}, Matias Martinez^b, Corrado A. Visaggio^b

^a Department of Engineering, University of Sannio, Italy
^b Université Polytechnique Hauts-de-France, France

ARTICLE INFO
Keywords: Software vulnerabilities, Software maintenance, Empirical study

ABSTRACT
Context: Identifying and repairing vulnerable code is a critical software maintenance task. Change impact analysis plays an important role during software maintenance, as it helps software maintainers to figure out the potential effects of a change before it is applied. However, while the software engineering community has extensively studied techniques and tools for performing impact analysis of change requests, there are no approaches for identifying the code changes induced by vulnerability fixing activities.
Objective: We hypothesize that similar vulnerabilities may present similar strategies for patching. More specifically, our work aims at understanding whether the class of the vulnerability to fix may determine the type of changes induced by the repair.
Method: To verify our conjecture, in this paper, we examine 524 security patches applied to vulnerabilities belonging to ten different weaknesses categories and extracted from 98 different open-source projects written in Java.
Results: We obtain empirical evidence that vulnerabilities of the same type are often resolved by applying similar code changes. This finding provides a new perspective for vulnerability fixing activities.
Conclusion: On the one hand, our findings open the way to better management of software maintenance activities when dealing with software vulnerabilities. Indeed, vulnerability class information could be exploited to better predict how the code will be affected by the fix, how the structure of the code will change (i.e., completed, copied, deleted, and well change), and the effort required for the fix. On the other hand, our results can be leveraged for improving automated strategies supported by developers when they have to deal with security flaws.

1. Introduction
Software maintenance involves critical tasks, as fixing code defects and developing the software according to the emerging needs of users [1]. In the context of mining analysis and impact analysis after “identifying the potential consequences of a change”, or estimating what “kind of change is needed to accomplish a change” [2]. In particular, change impact analysis plays an important role in software maintenance, as it allows identifying and analyzing the code changes that are often required to implement a change [3]. According to the definition provided by the MITRE corporation, a security vulnerability is a “flaw in a software, firmware, hardware, or service component resulting from a weakness that can be exploited, causing an unauthorized access or other damage to the system or its users” [4]. The objective of this research is to improve our fundamental understanding of vulnerabilities by empirically evaluating the conceptual differences between bugs and vulnerabilities.

The objective of this research is to improve our fundamental understanding of vulnerabilities by empirically evaluating the conceptual differences between bugs and vulnerabilities. We conducted an empirical study on the Chromium open source project (a.k.a. Google Chrome). We collected code reviews, post-release vulnerabilities, version control data, and bug data over six years. In the project, we conducted regression analysis to understand the relationship between bugs and vulnerabilities. Developers are reminders as recent as Shellshock and Heartbleed are reminders that small mistakes can lead to widespread problems. To engineer secure software, developers need a scientific and rigorous understanding of how to detect and prevent vulnerabilities.

We can build an understanding of vulnerabilities by viewing them as security-sensitive bugs. That is, a vulnerability can be defined as a “software defect that violates an implicit or explicit security policy” [1]. Recently, mining software repositories has greatly increased our understanding of software quality via empirical study of bugs. Researchers have provided a myriad of metrics, prediction models, hypothesis

* Corresponding author.
Email addresses: [\(G. Canfora\)](mailto:canfora@unisannio.it), [\(A. Di Sorbo\)](mailto:dsorbo@unisannio.it), [\(S. Forootani\)](mailto:forootani@unisannio.it), [\(M. Martinez\)](mailto:matias.martinez@uphf.fr), [\(C.A. Visaggio\).](mailto:visaggio@unisannio.it)

¹ <https://ccn.mitre.org/about/terminology.html>.

² <https://www.cvedetails.com/vulnerability-by-type.php>.

³ <https://nvd.nist.gov/vuln/detail/CVE-2018-1000656>.

⁴ <https://doi.org/10.1016/j.infsof.2021.106745>.

Camilo et al.

Comparison of pre-release bugs and post-release vulnerabilities in Chromium.

Canfora et al.

Comparison of bug and vulnerability fixing commits in six OSS projects.

Canfora et al.

In-depth analysis of the changes made in vulnerability fixing commits in 98 Java projects.

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ **We need other VCC-specific techniques!**

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

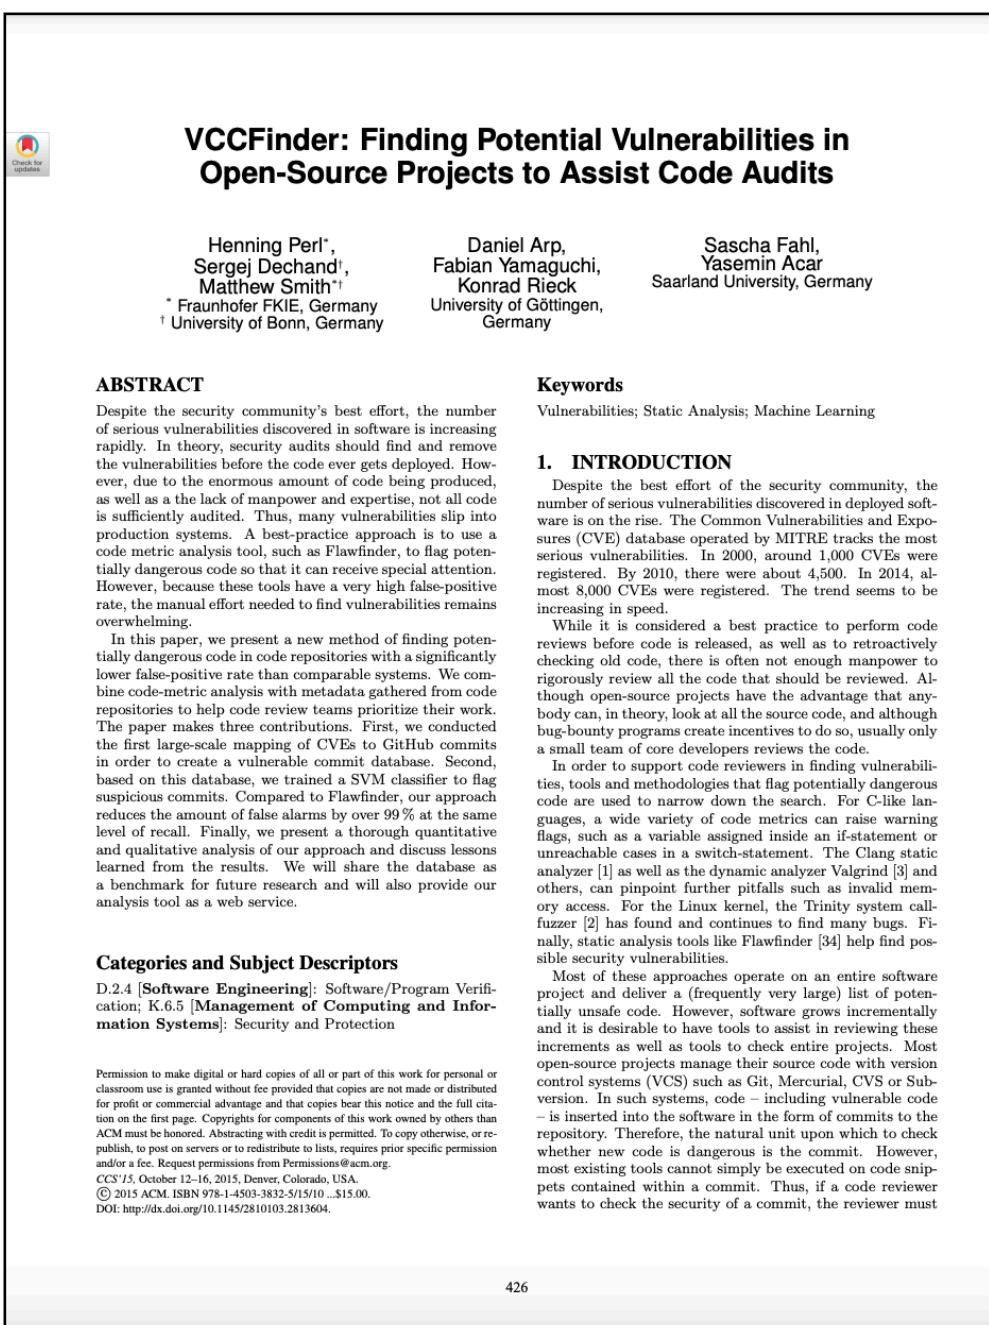
V-SZZ by Bao et al.

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.



SZZ by Yang et al.

A modified version of the original SZZ but:

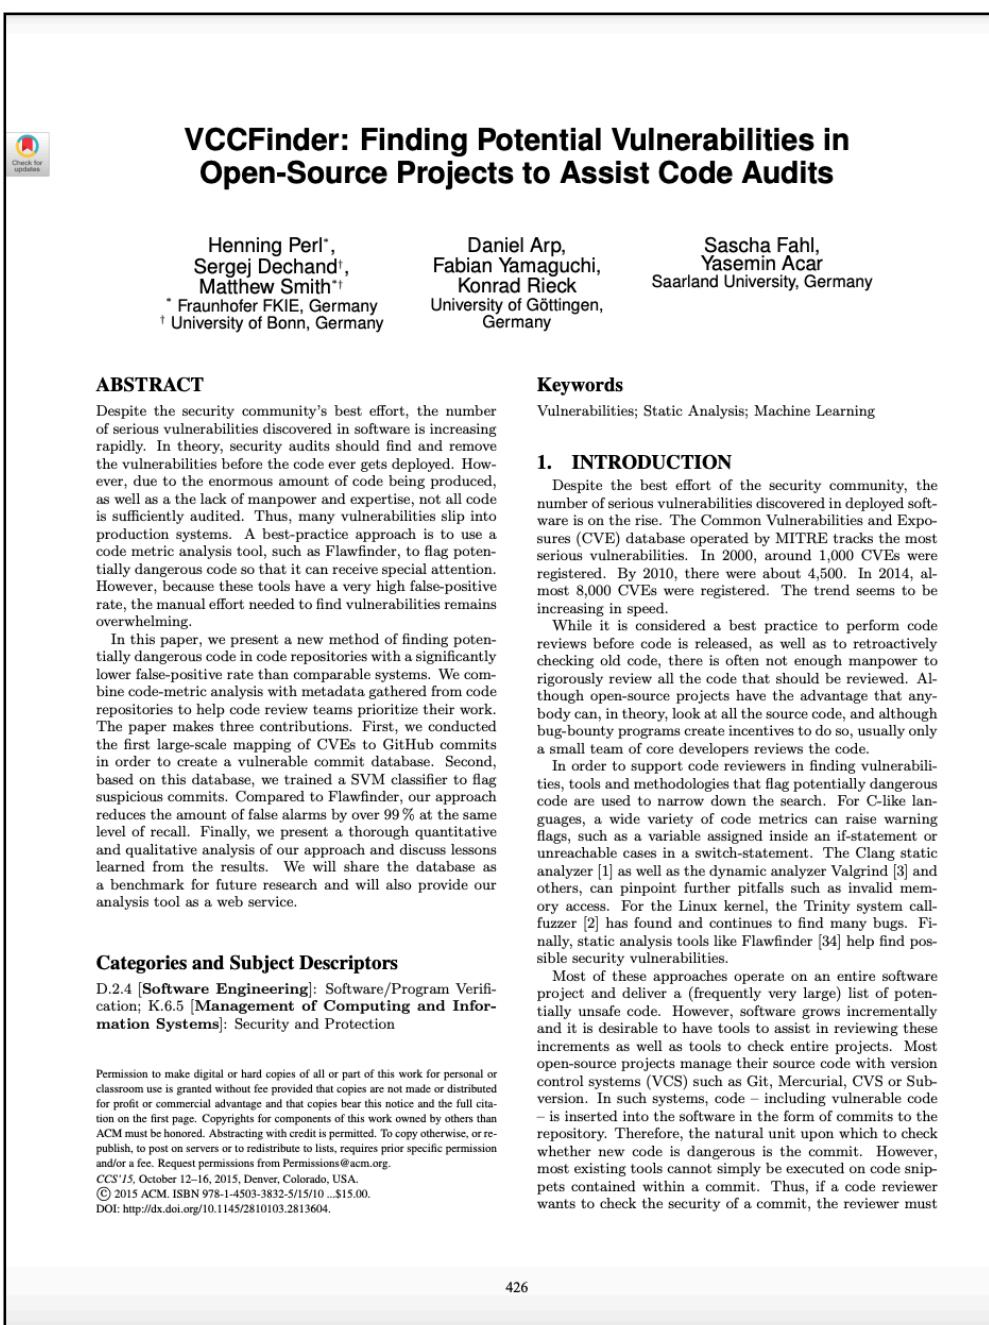
V-SZZ by Bao et al.

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.



SZZ by Yang et al.

V-SZZ by Bao et al.

A modified version of the original SZZ but:

👉 Documentation files (e.g., README) are ignored.

Mining VCCs: Ad hoc Approaches

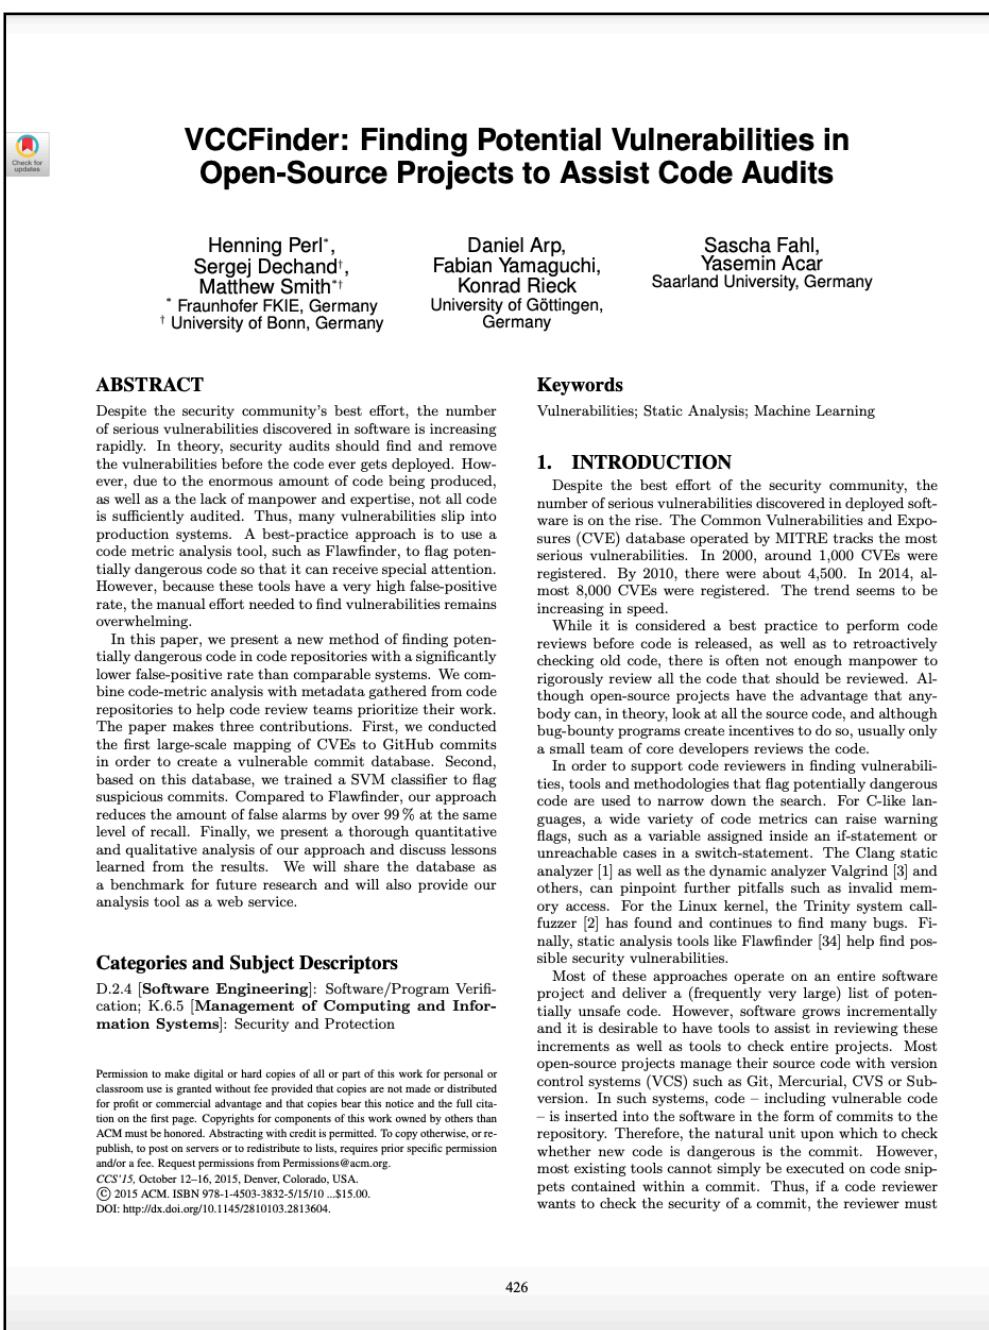
Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

V-SZZ by Bao et al.



A modified version of the **original SZZ** but:

- ➡ Documentation files (e.g., README) are ignored.
- ➡ In addition to the blames on the deleted lines, this variant also considers the blames on **the lines around the block of new lines**.

Mining VCCs: Ad hoc Approaches

Okay but
Indeed, th

SZZ by

Blamed ←

Blamed ←

```
int main(int argc, char* argv[]) {  
    char buff[65], *temp;  
    temp = argv[1] ? argv[1] : "";  
    if (argc > 0 && strlen(argv[1]) > 64)  
        strcpy(buff, temp);  
    printf("%s", "bye");  
}
```

SZZ by

v-SZZ

Rationale

Some vulnerabilities are fixed by adding missing checks, e.g., an *if* added before reading from a buffer. Hence, the **context** around the new code blocks might be responsible for the vulnerability.

Mining VCCs: Ad hoc Approaches

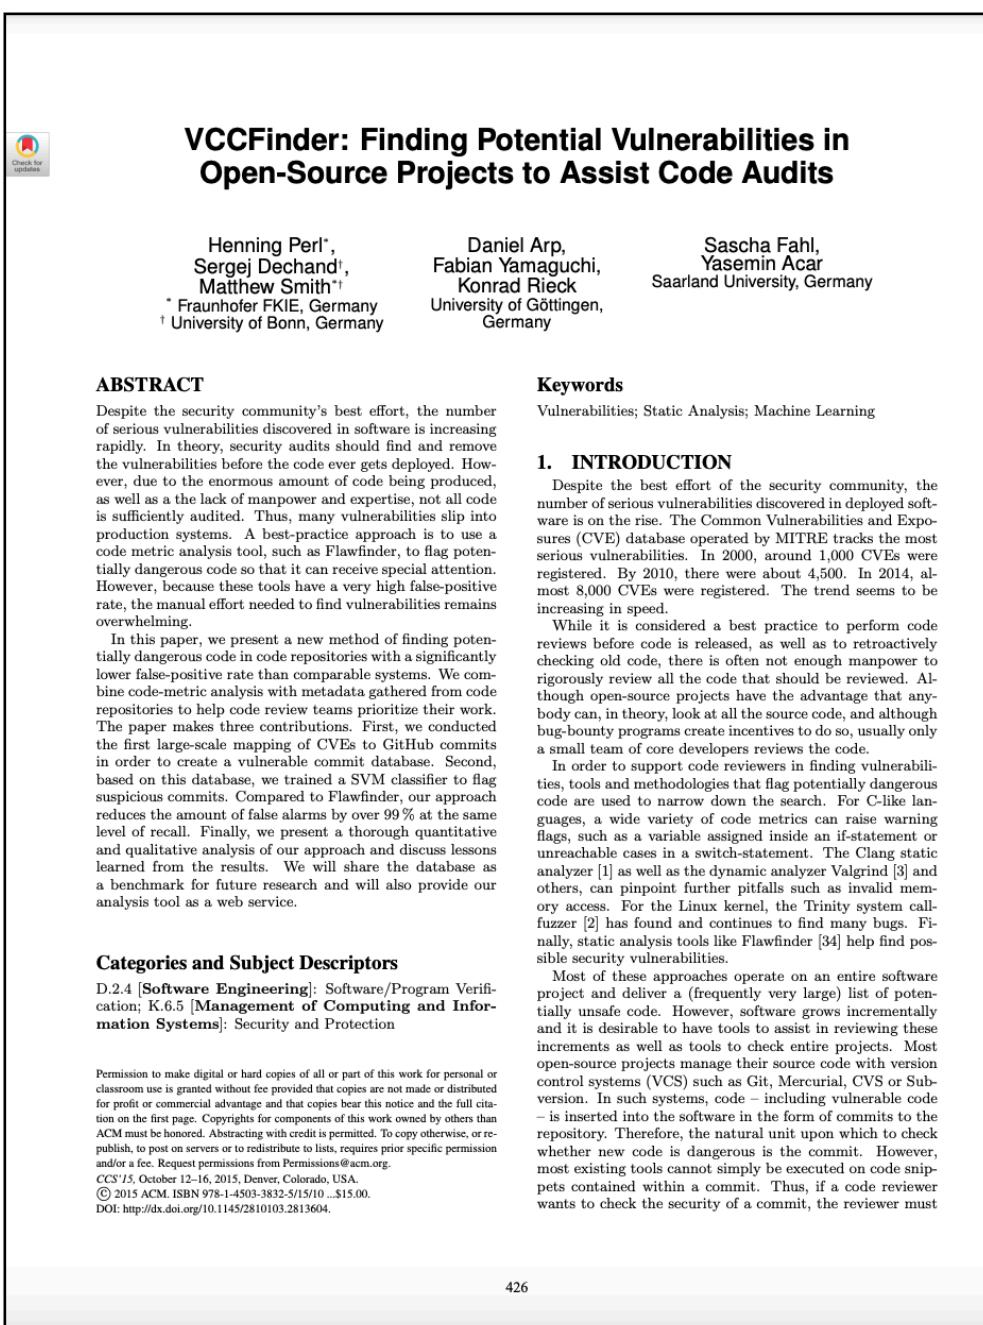
Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

V-SZZ by Bao et al.



A modified version of the **original SZZ** but:

- Documentation files (e.g., README) are ignored.
- In addition to the blames on the deleted lines, this variant also considers the blames on **the lines around the block of new lines**.
- It returns **only the most blamed commit**. In case of a tie, all the commits with the top score are returned (ex aequo).

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

VulDigger: A Just-in-Time and Cost-Aware Tool for Digging Vulnerability-Contributing Changes

Limin Yang*, Xiangxue Li^{†*} and Yu Yu[‡]

^{*}Department of Computer Science and Technology, East China Normal University, Shanghai, China
[†]Westone Cryptologic Research Center, Beijing, China; National Engineering Laboratory for Wireless Security, XUPT
[‡]Department of Computer Science and Engineering, Shanghai Jiao Tong University, Shanghai, China
Email: lmyang@stu.ecnu.edu.cn, xxli@cs.ecnu.edu.cn, yuyu@yuyu.hk

Abstract—It has been widely adopted to minimize the maintenance cost by predicting potential vulnerabilities before code audits in academia and industry. Most previous research focused to mine file-level vulnerability prediction models, which are coarse grained and may suffer from cost-prohibitive and impractical security testing activities. In this paper, we focus on a cost-aware vulnerability prediction model and present a just-in-time change-level code audit tool called VulDigger to dig sensitive ones from a pool of code changes. Our contributions benefit from the case study of Mozilla Firefox by constructing a large-scale vulnerability-contributing changes (VCCs) dataset in a semi-automatic fashion. VulDigger is a change-level tool and with the means of established and new metrics derived from both software defect prediction and vulnerability prediction. Consequently, the precision of such tool is extremely promising (i.e., 92%) for an effort-aware software team. We also examine the return on investment by comparing a random predictor against most skeptical changes with fewer efforts to inspect. Our findings suggest that such model is capable of pinpointing 31% of all VCCs with only 20% of the effort it would take to audit all changes (i.e., 55% better than random predictor). Our outputs can assist as an early step of continuous security inspections as it provides immediate feedback once developers submit changes to their code base.

1. INTRODUCTION

Code audits and security testing have been cost-prohibitive processes since most people today don't test software until it gets into the deployment phase of its life cycle and such practice has been proved ineffective to locate vulnerabilities or security bugs. Considering the disastrous consequence an exploited vulnerability could cause, e.g., Heartbleed [1], and to reduce the inspection effort, researchers have proposed a multitude of vulnerability prediction models for assisting and prioritizing code audits [2], [3], [4], [5].

Most of these studies focus on predicting vulnerable-prone modules (i.e., files or components) and can be beneficial in some contexts. However, these predictions are generally made too late. One of the suggested solutions is to apply security testing on each phase of the development cycle. Early detection is desirable as the later it gets into testing, the higher the cost and the more difficult it is to validate whether the fix is correct. Therefore, some researchers introduced change-level prediction methods and concentrated on predictions of vulnerability-contributing commits/changes (VCCs) [6]. Similar to the field of software defect prediction,

advantages of change-level predictions are [7]: (1) Prediction is suitable for code snippets and thus smaller regions of code needs inspection instead of huge files/components. (2) Developers that are responsible for VCCs can easily be traced and they can assist security experts or fix the security bugs by themselves with all design decisions fresh in their minds. (3) Predictions are made early and just-in-time as immediate feedback is given once a change is submitted to the code base. **Challenges for change-level predictions**—To our best knowledge, there are two main challenges for change-level vulnerability predictions except [6], we attribute it to the following two challenges:

- *The lack of a ground-truth dataset.* It's arduous to determine which code changes that indeed induced a vulnerability due to the multiplicity of code changes. Therefore, building a VCC ground-truth dataset is challenging and requires considerable human effort.
- *The disorderly structure of code changes.* Code changes could not retain the original structure and integrity like files or components, hence many established measures (e.g., code complexity, coupling, and cohesion) and commercial analysis tools (e.g., Understand C++) are not directly applicable.

Perl et al. [6] analyzed 66 open-source projects in GitHub and presented a database with 640 VCCs. Compared to Flawfinder [8], their results reduced false positives with significantly. However, they didn't consider the actual effort in code audits as some VCCs are huge (i.e., with thousands of lines of modifications).

Our contributions. We therefore build a cost-aware change-level vulnerability prediction model based on the code churn of a change. Through the study of Mozilla Firefox project — one of the most vulnerable open-source projects and has been the target in a plethora of vulnerability studies yet most of them focus on file/component-level predictions, our contributions can be outlined as follows:

- We present a change-level code review tool – VulDigger, to flag suspicious code changes immediately on the time of submitting by deriving features from software defect prediction models along with some new metrics (e.g., the maximum changes has been made in the past for files modified in a change). The precision of such tool is extremely promising (i.e., 92%) for a cost-aware software team.

V-SZZ by Bao et al.

A modified version of the SZZ by Perl et al. but:

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

VulDigger: A Just-in-Time and Cost-Aware Tool for Digging Vulnerability-Contributing Changes

Limin Yang*, Xiangxue Li[†]* and Yu Yu[‡]
^{*}Department of Computer Science and Technology, East China Normal University, Shanghai, China
[†]Westone Cryptologic Research Center, Beijing, China; National Engineering Laboratory for Wireless Security, XUPT
[‡]Department of Computer Science and Engineering, Shanghai Jiao Tong University, Shanghai, China
Email: lmyang@stu.ecnu.edu.cn, xxli@cs.ecnu.edu.cn, yuyu@sjtu.edu.cn

Abstract—It has been widely adopted to minimize the maintenance cost by predicting potential vulnerabilities before code audits in academia and industry. Most previous research focused on file/component-level vulnerability prediction models, which are coarse-grained and may suffer from cost-prohibitive and impractical security testing activities. In this paper, we focus on a cost-aware vulnerability prediction model and present a just-in-time change-level code review tool called VulDigger to dig sensitive ones from a pool of code changes. Our contributions benefit from the case study of Mozilla Firefox by constructing a large-scale vulnerability-contributing changes (VCCs) dataset in a semi-automatic fashion. This dataset is a challenging dataset with a mixture of established and new metrics derived from both software defect prediction and vulnerability prediction. Consequently, the precision of such tool is extremely promising (i.e., 92%) for an effort-aware software team. We also examine the return on investment by comparing a random predictor against most deployed changes with fewer lines to inspect. Our findings suggest that such model is capable of pinpointing 31% of all VCCs with only 20% of the effort it would take to audit all changes (i.e., 55% better than random predictor). Our outputs can assist as an early step of continuous security inspections as it provides immediate feedback once developers submit changes to their code base.

INTRODUCTION

Code audits and security testing have been cost-prohibitive processes since most people today don't test software until it gets into the deployment phase of its life cycle and such practice has been proved ineffective to locate vulnerabilities or security bugs. Considering the disastrous consequence an exploited vulnerability could cause, e.g., Heartbleed [1], and to reduce the inspection effort, researchers have proposed a multitude of vulnerability prediction models for assisting and prioritizing code audits [2], [3], [4], [5].

Most of these studies focus on predicting vulnerable-prone modules (i.e., files or components) and can be beneficial in some contexts. However, these predictions are generally made too late. One of the suggested solutions is to apply security testing on each phase of the development cycle. Early detection is desirable as the later it gets into testing, the higher the cost and the more difficult it is to validate whether the changes are safe. Therefore, some researchers introduced change-level prediction methods and concentrated on predictions of vulnerability-contributing commits/changes (VCCs) [6]. Similar to the field of software defect prediction,

978-1-5090-5019-2/17/\$31.00 ©2017 IEEE
Authorized licensed use limited to: Università degli Studi di Salerno. Downloaded on May 10, 2023 at 07:51:56 UTC from IEEE Xplore. Restrictions apply.

SZZ by Yang et al.

V-SZZ by Bao et al.

A modified version of the SZZ by Perl et al. but:

👉 Test and non-C/C++ files are ignored. Changes to comments, empty lines, and whitespaces are ignored as well.

Mining VCCs: Ad hoc Approaches

Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

V-SZZ by Bao et al.

VulDigger: A Just-in-Time and Cost-Aware Tool for Digging Vulnerability-Contributing Changes

Limin Yang*, Xiangxue Li[†]* and Yu Yu[‡]

^{*}Department of Computer Science and Technology, East China Normal University, Shanghai, China
[†]Westone Cryptologic Research Center, Beijing, China; National Engineering Laboratory for Wireless Security, XUPT
[‡]Department of Computer Science and Engineering, Shanghai Jiao Tong University, Shanghai, China
Email: lmyang@stu.ecnu.edu.cn, xxli@cs.ecnu.edu.cn, yuyu@yuyu.hk

Abstract—It has been widely adopted to minimize the maintenance cost by predicting potential vulnerabilities before code audits in academia and industry. Most previous research focused to identify file-level vulnerabilities. Predictions models are coarse grained and may suffer from cost-prohibitive and impractical security testing activities. In this paper, we focus on a cost-aware vulnerability prediction model and present a just-in-time change-level code review tool called VulDigger to dig sensitive ones from a pool of code changes. Our contributions benefit from the case study of Mozilla Firefox by constructing a large-scale vulnerability-contributing changes (VCCs) dataset in a semi-automatic fashion. The dataset is a change-level dataset with millions of established and new metrics derived from both software defect prediction and vulnerability prediction. Consequently, the precision of such tool is extremely promising (i.e., 92%) for an effort-aware software team. We also examine the return on investment of using a random predictor to locate most skeptical changes with fewer efforts to inspect. Our findings suggest that such model is capable of pinpointing 31% of all VCCs with only 20% of the effort it would take to audit all changes (i.e., 55% better than random predictor). Our outputs can assist an early step of continuous security inspections as it provides immediate feedback once developers submit changes to their code base.

INTRODUCTION

Code audits and security testing have been cost-prohibitive processes since most people today don't test software until it gets into the deployment phase of its life cycle and such practice has been proved ineffective to locate vulnerabilities or security bugs. Considering the disastrous consequence an exploited vulnerability could cause, e.g., Heartbleed [1], and to reduce the inspection effort, researchers have proposed a multitude of vulnerability prediction models for assisting and prioritizing code audits [2], [3], [4], [5].

Most of these studies focus on predicting vulnerable-prone modules (i.e., files or components) and can be beneficial in some contexts. However, these predictions are generally made too late. One of the suggested solutions is to apply security testing on each phase of the development cycle. Early detection is desirable as the later it gets into testing, the higher the cost and the more difficult it is to validate whether the changes are safe. Therefore, some researchers introduced change-level prediction methods and concentrated on predictions of vulnerability-contributing commits/changes (VCCs) [6]. Similar to the field of software defect prediction,

advantages of change-level predictions are [7]: (1) Prediction is suitable for code snippets and thus smaller regions of code needs inspection instead of huge files/components. (2) Developers that are responsible for VCCs can easily be traced and they can assist security experts or fix the security bugs by themselves with all design decisions fresh in their minds. (3) Predictions are made early and just-in-time as immediate feedback is given once a change is submitted to the code base. Challenges for change-level predictions are [7]: To our best knowledge, there is no ground-truth change-level vulnerability predictions except [6], we attribute it to the following two challenges:

- *The lack of a ground-truth dataset.* It's arduous to determine which code changes that indeed induced a vulnerability due to the multiplicity of code changes. Therefore, building a VCC ground-truth dataset is challenging and requires considerable human effort.
- *The disorderly structure of code changes.* Code changes could not retain the original structure and integrity like files or components, hence many established measures (e.g., code complexity, coupling, and cohesion) and commercial analysis tools (e.g., Understand C++) are not directly applicable.

Perl et al. [6] analyzed 66 open-source projects in GitHub and presented a dataset with 640 VCCs. Compared to Flawfinder [8], their results reduced false positives with significantly. However, they didn't consider the actual effort in code audits as some VCCs are huge (i.e., with thousands of lines of modifications).

Our contributions. We therefore build a cost-aware change-level vulnerability prediction model based on the code churn of a change. Through the study of Mozilla Firefox project – one of the most vulnerable open-source projects and has been the target in a plethora of vulnerability studies yet most of them focus on file/component-level predictions, our contributions can be outlined as follows:

- We present a change-level code review tool – VulDigger, to flag suspicious code changes immediately on the time of submitting by deriving features from software defect prediction models along with some new metrics (e.g., the maximum changes has been made in the past for files modified in a change). The precision of such tool is extremely promising (i.e., 92%) for a cost-aware software team.

A modified version of the SZZ by Perl et al. but:

- Test and non-C/C++ files are ignored. Changes to comments, empty lines, and whitespaces are ignored as well.
- For each new line added, the blame around this line is considered **only if it contains a C/C++ keyword or a function call**.

Mining VCCs: Ad hoc Approaches

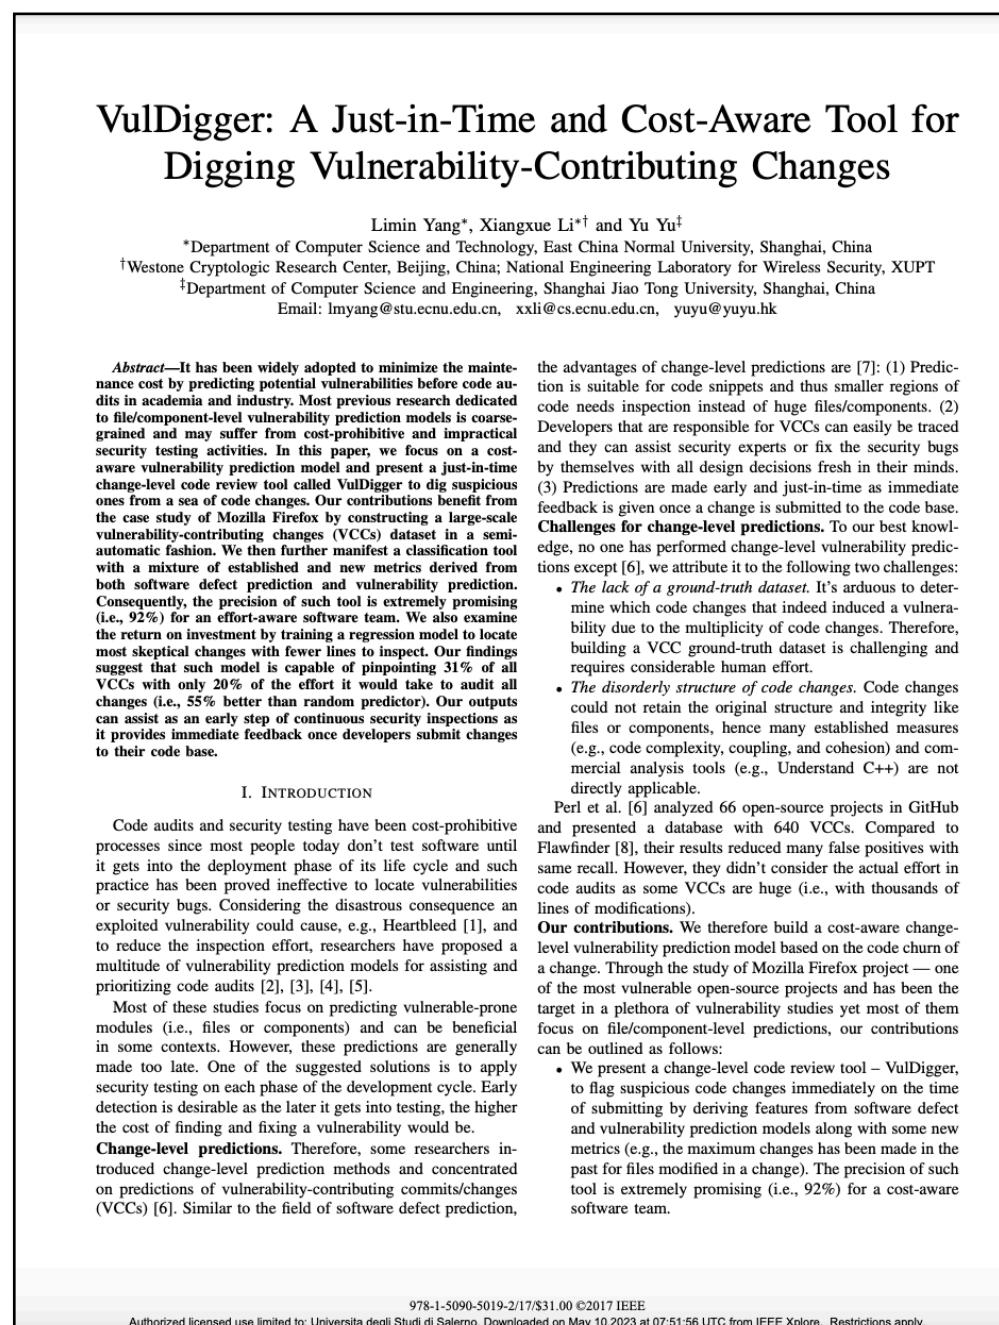
Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

v-SZZ by Bao et al.



A modified version of the SZZ by Perl et al. but:

- 👉 Test and non-C/C++ files are ignored. Changes to comments, empty lines, and whitespaces are ignored as well.
- 👉 For each new line added, the blame around this line is considered **only if it contains a C/C++ keyword or a function call**.
- 👉 Unlike the Perl et al. variant, it considers the blames around blocks of new lines **only if they do not contain new functions**.

Mining VCCs: Ad hoc Approaches

Okay but
Indeed, th

SZZ by

Blamed ←
Blamed ←
NOT blamed ←

```
int main(int argc, char* argv[]) {  
    char buff[65], *temp;  
    temp = argv[1] ? argv[1] : "";  
    if (argc > 0 && my_len(argv[1]) > 64)  
        strcpy(buff, temp);  
    printf("%s", "bye");  
}  
int my_len(char* buff) {  
    return strlen(buff);  
}
```

Rationale

Functions can be added anywhere in the file. Hence, the local context does not always involve meaningful parts.

Mining VCCs: Ad hoc Approaches

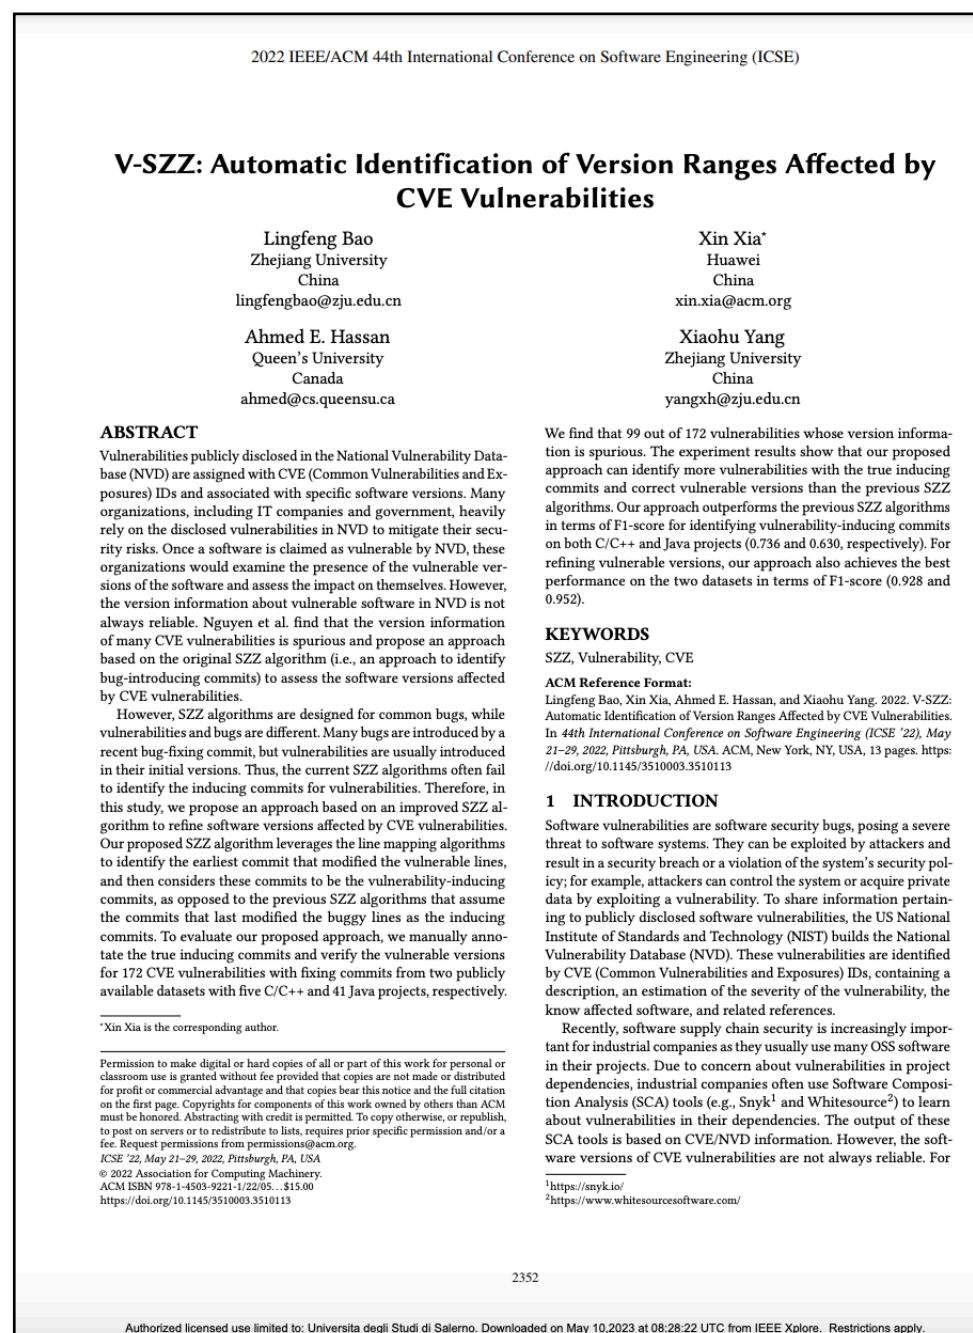
Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

V-SZZ by Bao et al.



A modified version of the SZZ by Kim et al. but:

Mining VCCs: Ad hoc Approaches

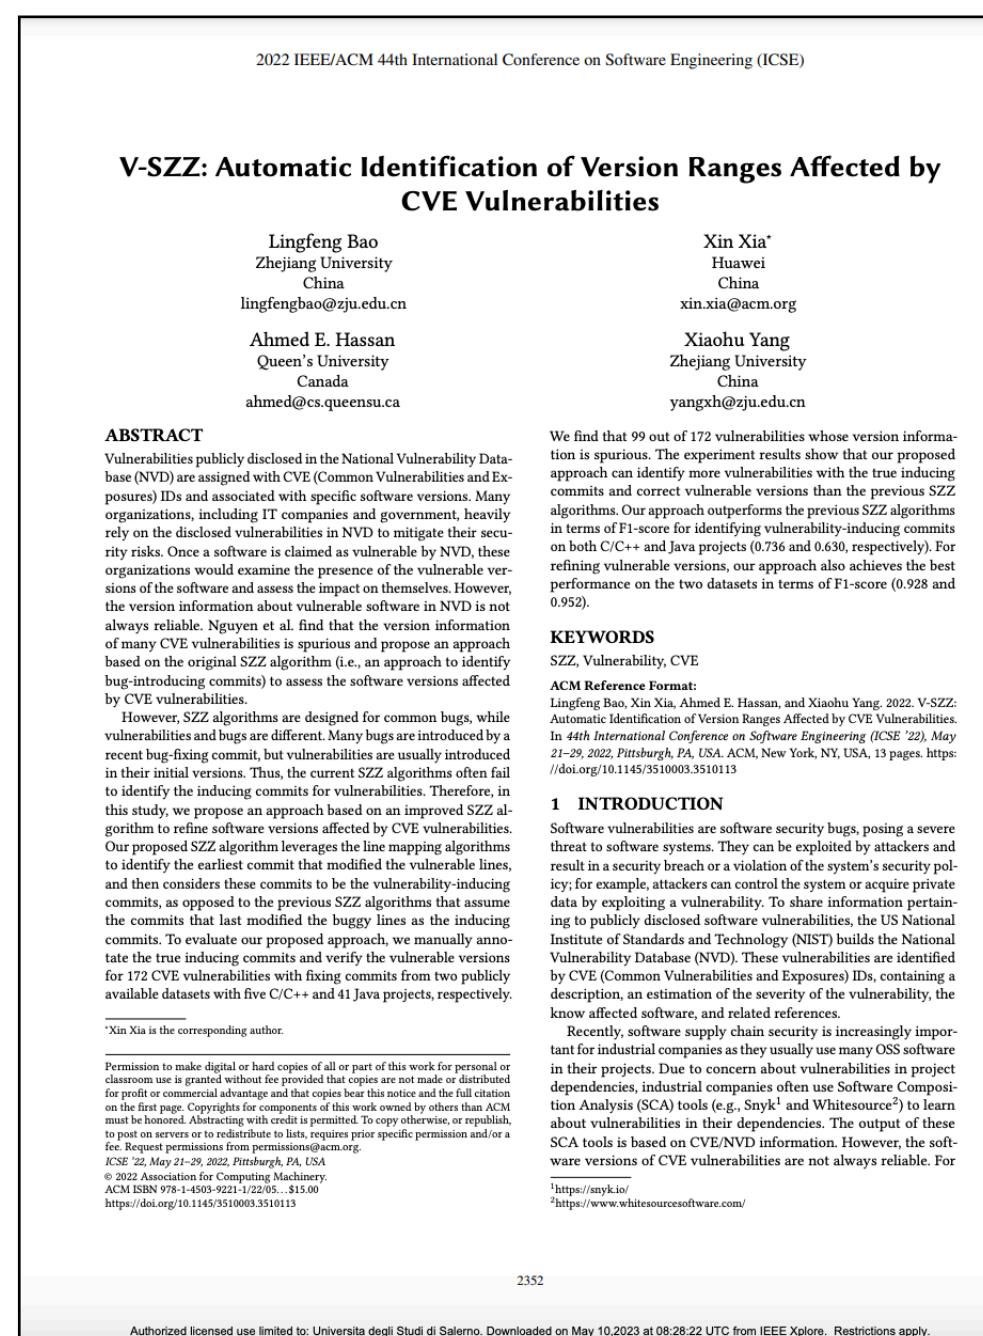
Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

V-SZZ by Bao et al.



A modified version of the **SZZ** by Kim et al. but:

👉 The git blame is **repeated beyond format changes** until reaching the commits that created the blamed lines. This approach is supported by both AST and string similarity matching.

Mining VCCs: Ad hoc Approaches

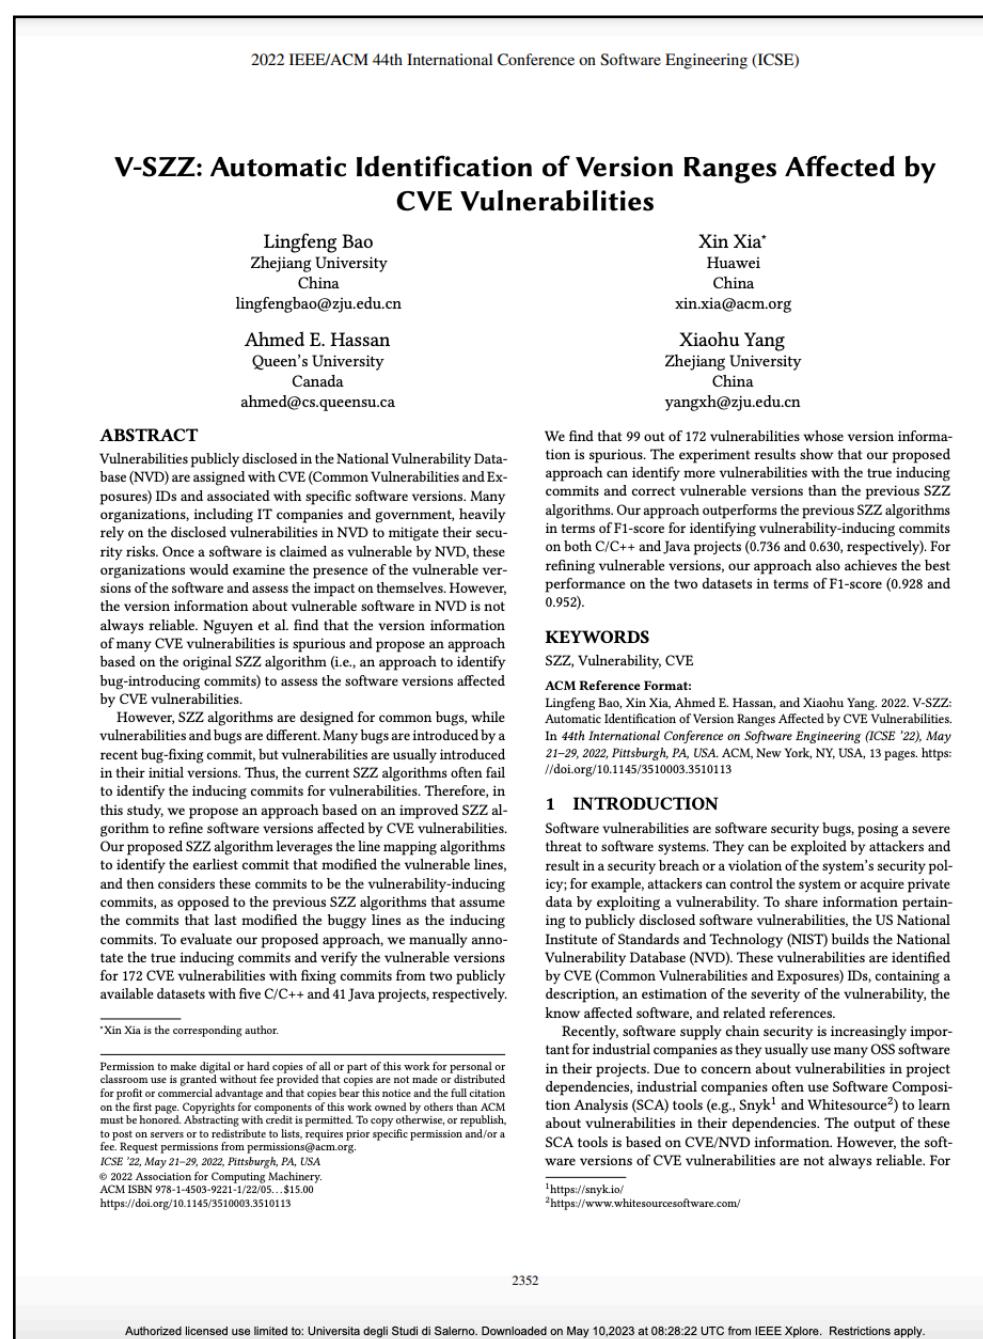
Okay but reusing the algorithms meant for bugs does not work well for VCCs.
Indeed, there are studies explaining how bugs and vulnerabilities differ.

→ We need other VCC-specific techniques!

SZZ by Perl et al.

SZZ by Yang et al.

V-SZZ by Bao et al.



A modified version of the **SZZ** by Kim et al. but:

👉 The git blame is **repeated beyond format changes** until reaching the commits that created the blamed lines. This approach is supported by both AST and string similarity matching.

Rationale

According to certain studies, many vulnerabilities are *foundational*, i.e., introduced early in the project, even before the first release.

Are we sure
they work?

Performance Indicators

How can we be sure VCC mining algorithms work as expected? We want our algorithm to minimize:

Performance Indicators

How can we be sure VCC mining algorithms work as expected? We want our algorithm to minimize:

False Positives

The algorithm returned a commit that was not a real VCC.

False Negatives

The algorithm did not return one (or more) real VCC.

Performance Indicators

How can we be sure VCC mining algorithms work as expected? We want our algorithm to minimize:

False Positives

The algorithm returned a commit that was not a real VCC.

False Negatives

The algorithm did not return one (or more) real VCC.

From the *Information Retrieval* world, we commonly use these metrics to evaluate such approaches:

Performance Indicators

How can we be sure VCC mining algorithms work as expected? We want our algorithm to minimize:

False Positives

The algorithm returned a commit that was not a real VCC.

False Negatives

The algorithm did not return one (or more) real VCC.

From the *Information Retrieval* world, we commonly use these metrics to evaluate such approaches:

$$Precision = \frac{|correct \cap identified|}{|identified|}$$

“Among the found VCCs, how many are correct?”

Performance Indicators

How can we be sure VCC mining algorithms work as expected? We want our algorithm to minimize:

False Positives

The algorithm returned a commit that was not a real VCC.

False Negatives

The algorithm did not return one (or more) real VCC.

From the *Information Retrieval* world, we commonly use these metrics to evaluate such approaches:

$$Precision = \frac{|correct \cap identified|}{|identified|}$$

“Among the found VCCs, how many are correct?”

$$Recall = \frac{|correct \cap identified|}{|correct|}$$

“Among the correct VCCs, how many did I find?”

Performance Indicators

How can we be sure VCC mining algorithms work as expected? We want our algorithm to minimize:

False Positives

The algorithm returned a commit that was not a real VCC.

False Negatives

The algorithm did not return one (or more) real VCC.

From the *Information Retrieval* world, we commonly use these metrics to evaluate such approaches:

$$Precision = \frac{|correct \cap identified|}{|identified|}$$

“Among the found VCCs, how many are correct?”

$$Recall = \frac{|correct \cap identified|}{|correct|}$$

“Among the correct VCCs, how many did I find?”

$$F - measure = \frac{2}{\frac{1}{Precision} + \frac{1}{Recall}}$$

“Trade-off between precision and recall”

Performance Indicators

How can we be sure VCC mining algorithms work as expected? We want our algorithm to minimize:

False Positives

The algorithm returned a commit that was not a real VCC.

False Negatives

The algorithm did not return one (or more) real VCC.

From the *Information Retrieval* world, we commonly use these metrics to evaluate such approaches:

$$Precision = \frac{|correct \cap identified|}{|identified|}$$

$$Recall = \frac{|correct \cap identified|}{|correct|}$$

$$F - measure = \frac{2}{\frac{1}{Precision} + \frac{1}{Recall}}$$

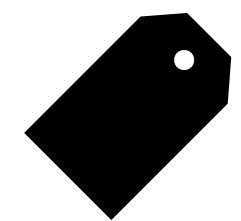
**But how do we determine
this “correct” set?**

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:

Building the Ground Truth

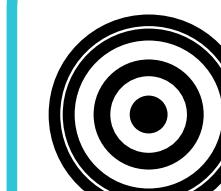
We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:



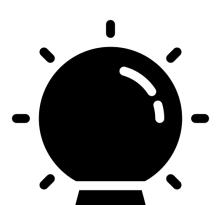
Exhaustive Labeling



Bisect-driven Labeling



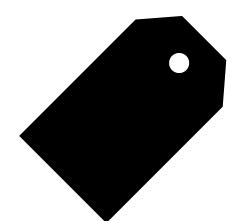
Precision Assessment



Developer-informed Oracle

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:



Exhaustive Labeling



Bisect-driven Labeling

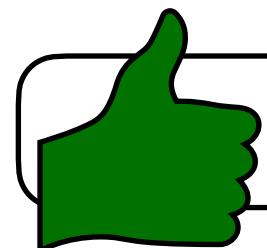


Precision Assessment



Developer-informed Oracle

For each vulnerability, we manually inspect all the commits in the project and assess whether it is a VCC. Complete but time-consuming.



Recommended when...

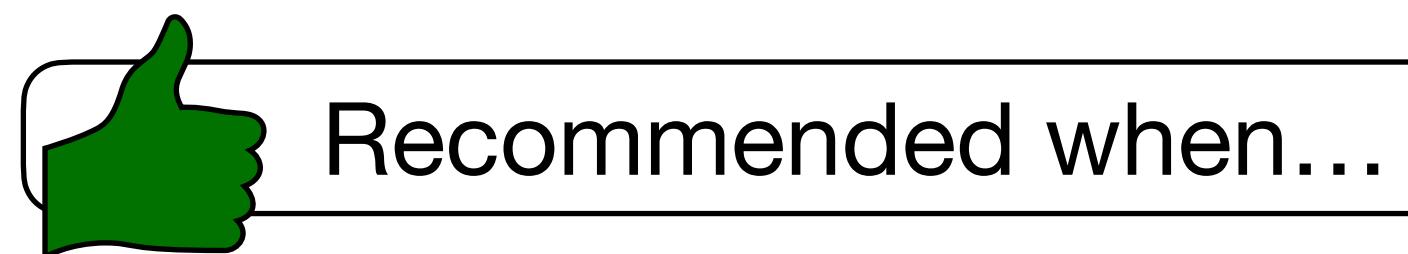
we want to be exhaustive (!) or just want to analyze a few vulnerabilities.

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:



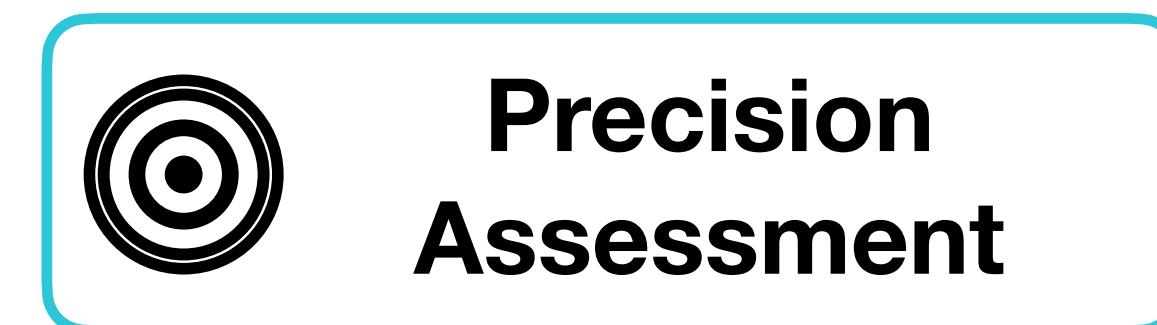
For each vulnerability, we run git bisect until we find at least one VCC. Inspired by the Meneely et al. mining technique. Less complete but faster, reducing the workload by a logarithmic factor.



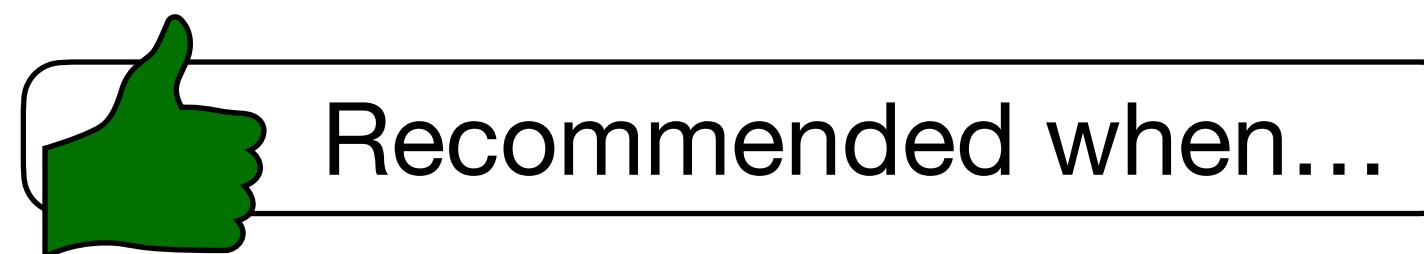
we don't need a complete *correct* set,
and we have time to inspect.

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:



For each commit flagged as VCC by the algorithm, we inspect it to assess whether it is a real VCC. This will not produce the *correct* set, but only *correctly identified*. Hence, we are not aware of the “missed” VCCs.



Recommended when...

we are only interested in assessing the precision.

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:



For each vulnerability, we process the fixing commit message to retrieve mentions of the culprit commit(s). Developers sometimes explicitly indicate the commit where the vulnerability was introduced. This method has a fully automated part based on NLP/text mining and an (optional) manual assessment part.



we don't need a complete *correct* set and, we want developers' experience.

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:



Exhaustive
Labeling



Bisect-driven
Labeling



Precision
Assessment



Developer-
informed Oracle

EXAMPLE

CVE-2011-5321 (NULL pointer dereference) in Linux Kernel was fixed in commit c290f835 by just adding a single line of code.

“

TTY: drop driver reference in tty_open fail path

When tty_driver_lookup_tty fails in tty_open, we forget to drop a reference to the tty driver. This was added by commit 4a2b5fd (Move tty lookup/reopen to caller). [...]

```

1869: if (!tty) {
1870:     /* check whether we're reopening an existing tty */
1871:     tty = tty_driver_lookup_tty(driver, inode, index);
1872:     if (IS_ERR(tty)) {
1874:         tty_unlock();
1875:         mutex_unlock(&tty_mutex);
1876:         tty_driver_kref_put(driver);
1877:         return PTR_ERR(tty);
1878:     }
1879: }
```

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the “standard” for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:



EXAMPLE

CVE-2011-5321 (NULL pointer dereference) in Linux Kernel was fixed in commit `c290f835` by just adding a single line of code.

“
*TTY: drop driver reference in tty_open fail path
When tty_driver_lookup_tty fails in tty_open, we forget to drop a reference to the tty driver. This was added by commit 4a2b5fd (Move tty lookup/reopen to caller). [...]*

According to the developer who fixed this vulnerability, this is a VCC (which involuntarily introduced the vulnerability while refactoring some code).

**How can I
use them?**

Available Tools

OpenSZZ

Command-line tool written in Java implementing the **standard SZZ**, analyzing GitHub repositories and Jira issues.

SZZUnleashed

Collection of Python and Java scripts implementing the **SZZ by Williams and Spacco** (not seen).

PyDriller

Python library for repository mining, including an implementation of **SZZ by Kim et al.**

Archeogit

Command-line tool written in Python implementing the **SZZ by Perl et al.**

V-SZZ

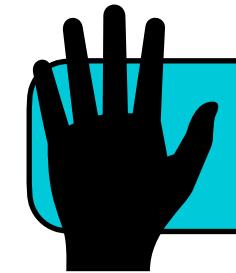
Collection of Python scripts replicating **V-SZZ by Bao et al.**

PySZZ

Collection of Python scripts implementing **several SZZ variants** with a uniform interface.

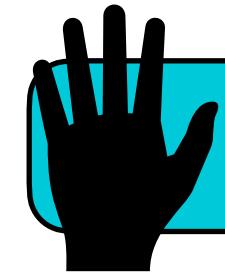
Isn't there
something
ready to use?

Available Datasets



Curated

Available Datasets

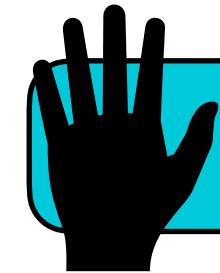


Curated

Vulnerability
History Project

Database of curated histories of 2,677 vulnerabilities of eight open-source projects. Built by **class assignments in a Master's degree course held at RIT**.

Available Datasets



Curated

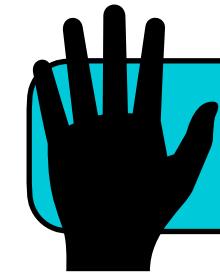
Vulnerability
History Project

Database of curated histories of 2,677 vulnerabilities of eight open-source projects. Built by **class assignments in a Master's degree course held at RIT**.

Java
VCC Dataset

Dataset of 100 VCCs of 71 known vulnerabilities affecting popular Java projects. Built by **manually analyzing the history aided by blames on fixing commits**.

Available Datasets



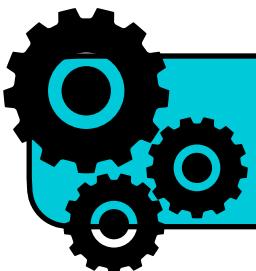
Curated

Vulnerability
History Project

Database of curated histories of 2,677 vulnerabilities of eight open-source projects. Built by **class assignments in a Master's degree course held at RIT**.

Java
VCC Dataset

Dataset of 100 VCCs of 71 known vulnerabilities affecting popular Java projects. Built by **manually analyzing the history aided by blames on fixing commits**.



Mined

Available Datasets



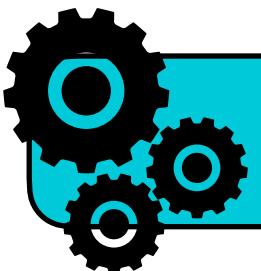
Curated

Vulnerability
History Project

Database of curated histories of 2,677 vulnerabilities of eight open-source projects. Built by **class assignments in a Master's degree course held at RIT**.

Java
VCC Dataset

Dataset of 100 VCCs of 71 known vulnerabilities affecting popular Java projects. Built by **manually analyzing the history aided by blames on fixing commits**.

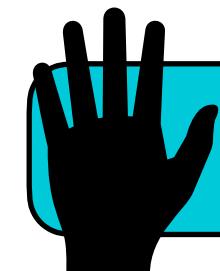


Mined

Secret Life
Dataset

Dataset of 12,256 VCCs of 3,663 vulnerabilities affecting 1,096 open-source projects. Built by **running an SZZ variant by Iannone et al.** (not seen).

Available Datasets



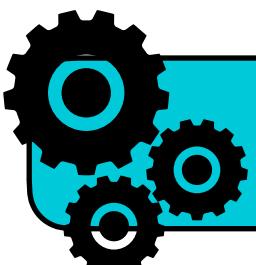
Curated

Vulnerability History Project

Database of curated histories of 2,677 vulnerabilities of eight open-source projects. Built by **class assignments in a Master's degree course held at RIT.**

Java VCC Dataset

Dataset of 100 VCCs of 71 known vulnerabilities affecting popular Java projects. Built by **manually analyzing the history aided by blames on fixing commits.**



Mined

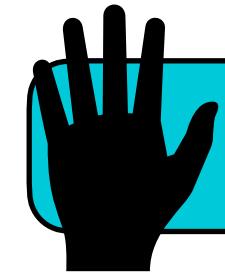
Secret Life Dataset

Dataset of 12,256 VCCs of 3,663 vulnerabilities affecting 1,096 open-source projects. Built by **running an SZZ variant by Iannone et al.** (not seen).

FrontEndART Dataset

Dataset of ~700 VCCs of 564 vulnerabilities affecting 198 Java projects. Built by **running an SZZ variant by Aladics et al.**

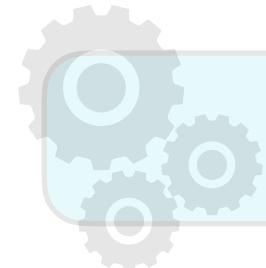
Available Datasets



Curated

Vulnerability
History Project

Database of curated histories of 2,677 vulnerabilities of eight open-source projects. Built by **class assignments in a Master's degree course held at RIT.**



Secret Life
Dataset

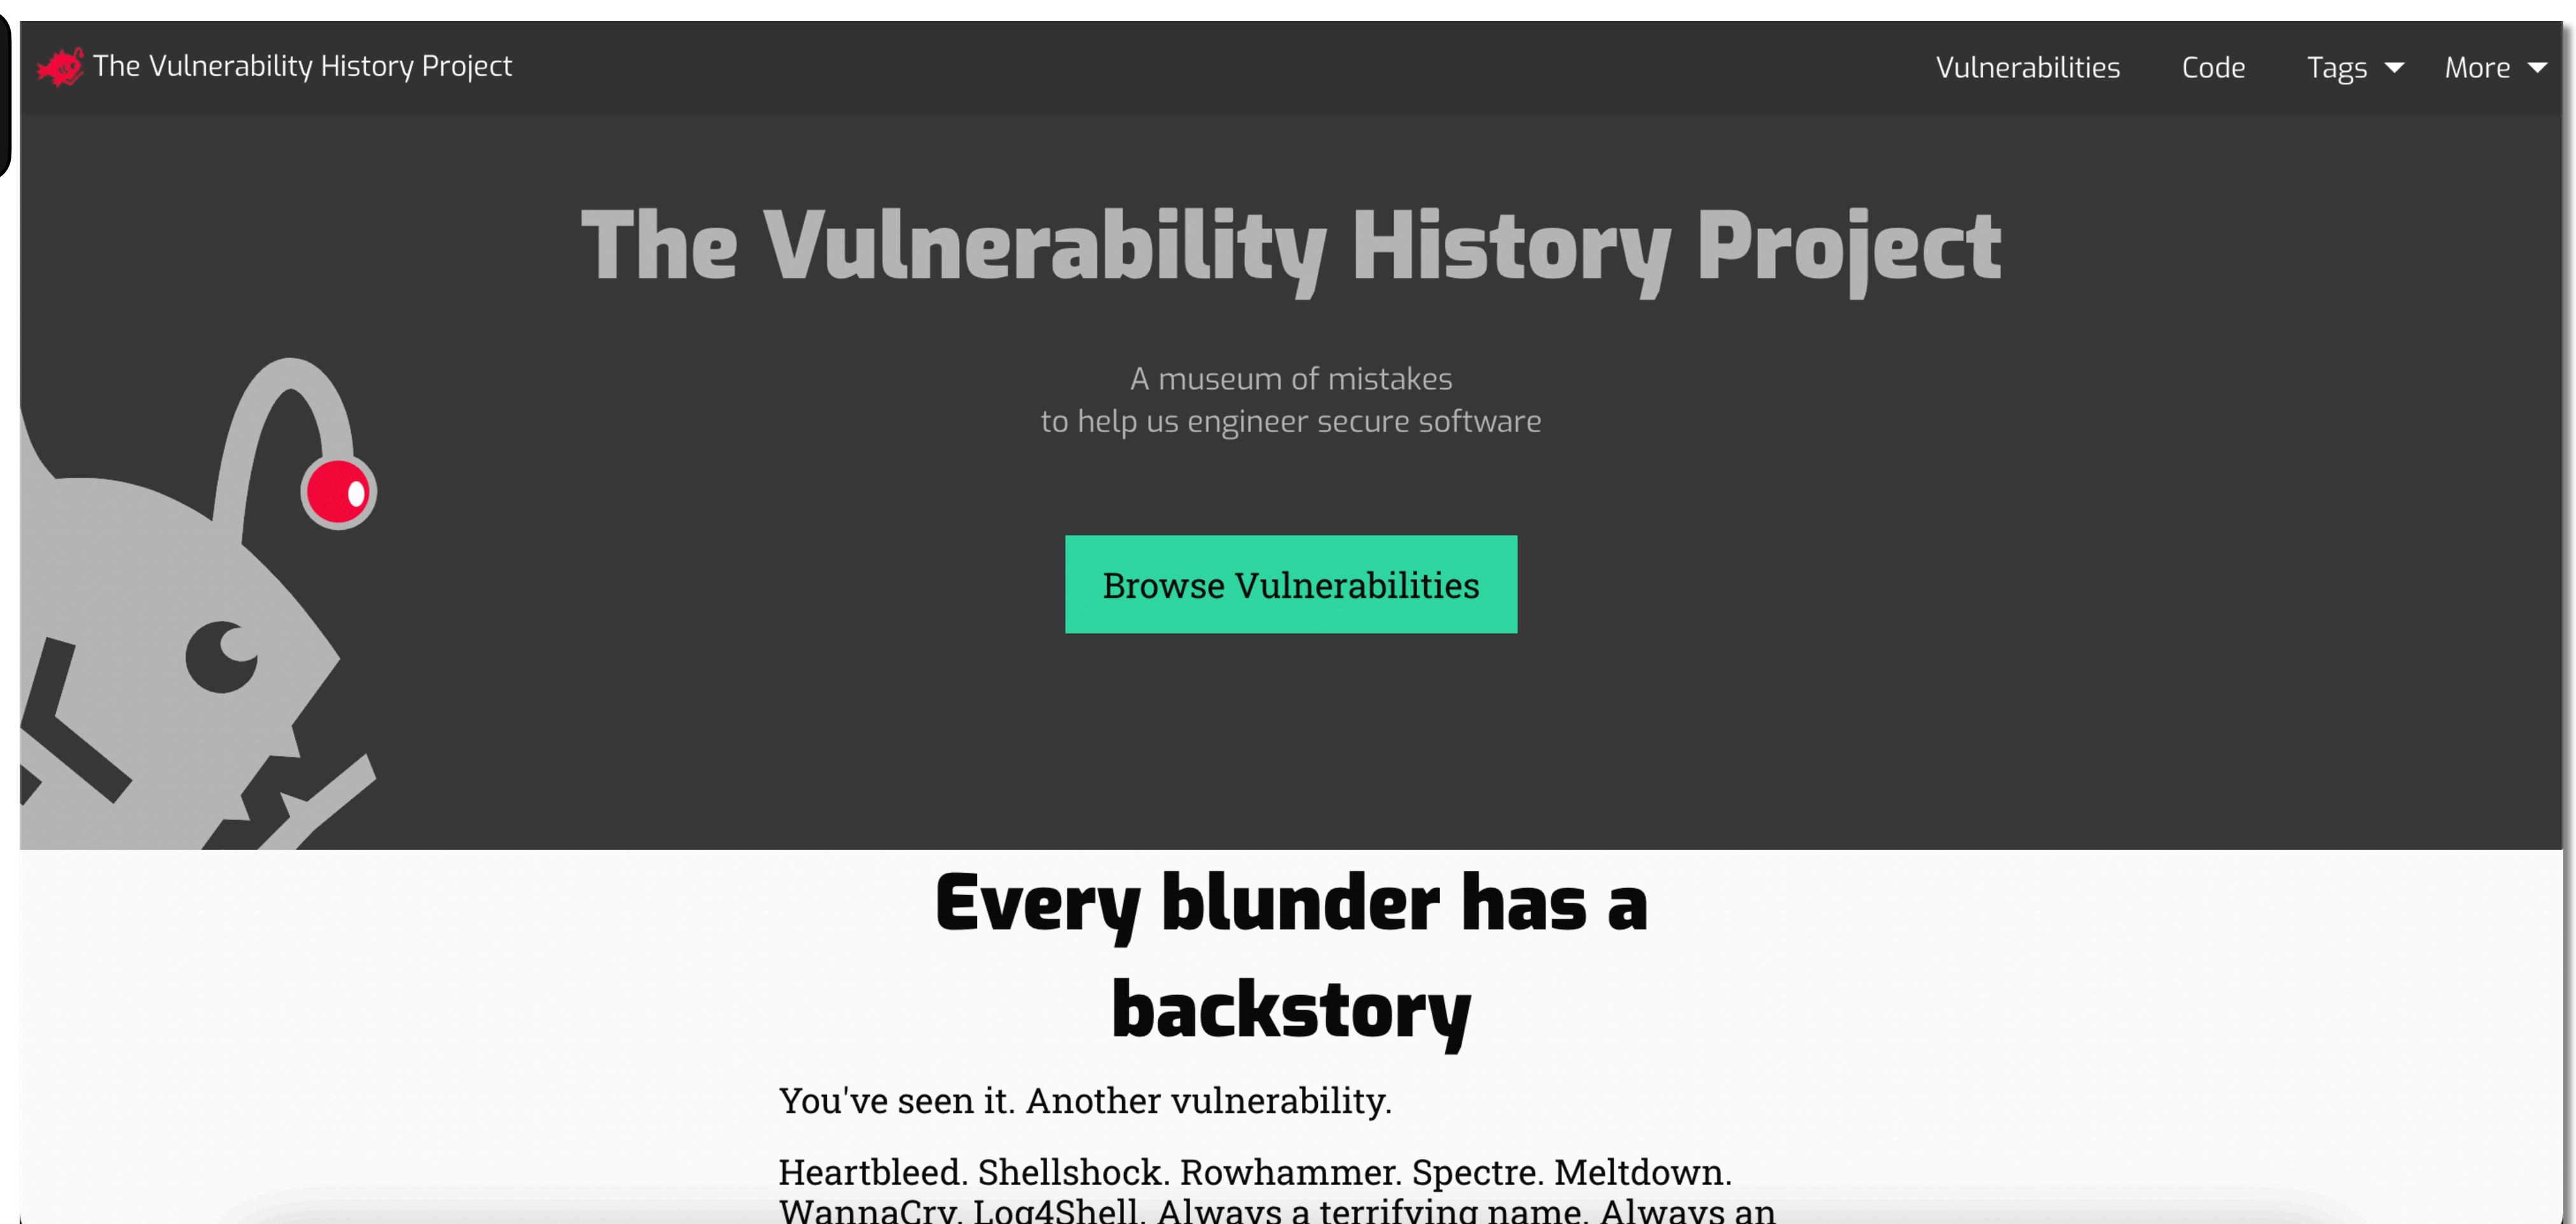
Dataset of 12,256 VCCs of 3,663 vulnerabilities affecting 1,096 open-source projects. Built by **running an SZZ variant by Iannone et al.** (not seen).

Provenance
Dataset

Dataset of ~700 VCCs of 564 vulnerabilities affecting 198 Java projects. Built by **running an SZZ variant by Aladics et al.**

Available Datasets

Vulnerability
History Project



The screenshot shows the homepage of The Vulnerability History Project. At the top left is a black button with white text: "Vulnerability History Project". To its right is the project's logo, a red cartoon mouse with a single red eye, followed by the text "The Vulnerability History Project". On the far right of the header are navigation links: "Vulnerabilities", "Code", "Tags ▾", and "More ▾". The main title "The Vulnerability History Project" is centered in large, light gray letters. Below it is a subtitle: "A museum of mistakes to help us engineer secure software". A green button labeled "Browse Vulnerabilities" is positioned below the subtitle. In the bottom right corner of the main dark area, there is a large, stylized text quote: "Every blunder has a backstory". At the very bottom of the page, in a smaller white section, is the text: "You've seen it. Another vulnerability. Heartbleed. Shellshock. Rowhammer. Spectre. Meltdown. WannaCry. Log4Shell. Always a terrifying name. Always an".

The Vulnerability History Project

A museum of mistakes
to help us engineer secure software

Browse Vulnerabilities

Every blunder has a
backstory

You've seen it. Another vulnerability.
Heartbleed. Shellshock. Rowhammer. Spectre. Meltdown.
WannaCry. Log4Shell. Always a terrifying name. Always an

Available Datasets

The Vulnerability History Project

Vulnerabilities Code Tags ▾ More ▾

CVE-2017-12615

A horizontal timeline from 2007 to 2017 showing commit activity. Yellow dots represent commits, with a higher density in 2017. A red dot is visible in 2007. A yellow cat icon is positioned above the timeline.

2007 2008 2009 2010 2011 2012 2013 2014 2015 2016 2017

Vulnerability-contributing commit for CVE-2017-12615:
Phase 1: Setting eol and mime types
July 20th, 2006

changes new developer refactors
same directory fix vcc

Origin to Fix

In Apache Tomcat on Windows, an attacker could upload a JSP (JavaServer Page, essentially a web page with java code) that would be later executed by the server. This worked if a "/" was added at the end of the file extension.

CWE-650: Trusting HTTP Permission Methods on the Server Side Language: Java Lesson: Code Refactors
Lesson: Distrust Input Lesson: Least Privilege Lesson: Too Many Cooks Lifetime: 5+ years
Project: Tomcat Tomcat subsystem: resources VCC

8 Upvotes

Mistakes Made Tag Notes Fix VCC Curator Notes Curate Articles

Built by running a CI variant by Radouan El Amraoui

Available Datasets

Vulnerability History Project

 Mistakes Made  Tag Notes  Fix  VCC  Curator Notes  Curate  Articles

 **CWE-650: Trusting HTTP Permission Methods on the Server Side**

[Learn more about CWE-650: Trusting HTTP Permission Methods on the Server Side.](#)

 **Lesson: Distrust Input**
File input was missing some sanitization, as using a "/" would allow the malicious file to go through to the server.
[Learn more about Lesson: Distrust Input.](#)

 **Lifetime: 5+ years**
4039.3 days, or 11.1 years
[Learn more about Lifetime: 5+ years.](#)

 **Language: Java**

[Learn more about Language: Java.](#)

 **Lesson: Least Privilege**
They should be checking auth privileges at all times and not let the file where the vulnerability is to impact the rest of the program.
[Learn more about Lesson: Least Privilege.](#)

 **Project: Tomcat**
[Learn more about Project: Tomcat.](#)

 **Lesson: Code Refactors**
129 refactors took place during the vulnerability.
[Learn more about Lesson: Code Refactors.](#)

 **Lesson: Too Many Cooks**
64 different developers made commits to the files fixed for this vulnerability.
[Learn more about Lesson: Too Many Cooks.](#)

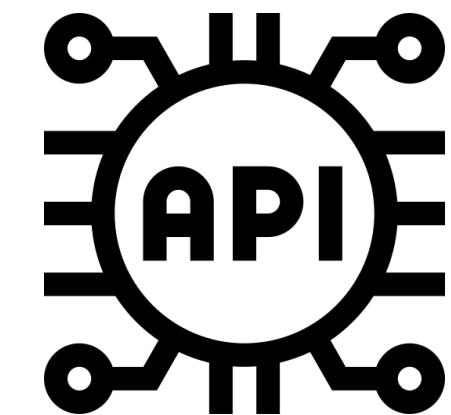
 **Tomcat subsystem: resources**
[Learn more about Tomcat subsystem: resources.](#)

 VCC



Available Datasets

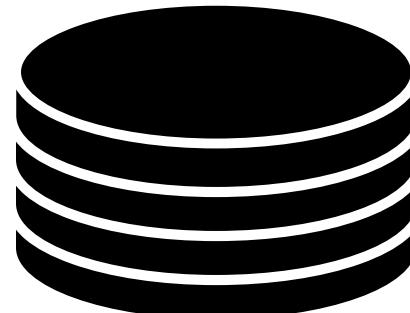
Vulnerability
History Project



RESTful API

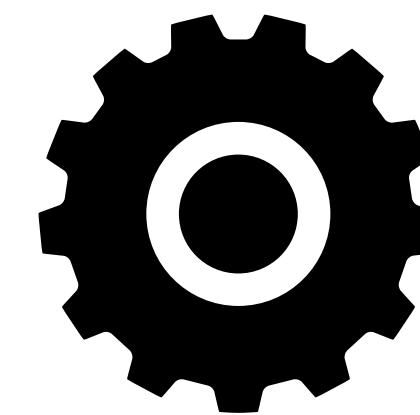
Retrieving data with simple
HTTP requests.

VHP can be mined in several ways



Raw Data

The list of vulnerabilities is
available in a repository of its
organization in GitHub.



Ad Hoc Tool

The organization in GitHub offers
a dedicated command-line tool.

I'm
hungry!

Wrap up

Wrap up

MSR for Vulnerability Prediction – Vulnerability-contributing Commits

Key Characteristics of VCCs

VCCs vs non-VCCs
A case study on Apache HTTP Server with 68 post-release vulnerabilities and 124 VCCs.

- Large commits might increase the chance of contributing to a vulnerability.
- Changing other developers' code might increase the chance of contributing to a vulnerability.
- Vulnerabilities are more likely to be added when modifying existing files rather than creating new files.

A. Meneely et al., "When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits," 2013 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement, Baltimore, MD, USA, 2013, pp. 65-74, doi: 10.1109/ESEM.2013.19.

Definition & Characteristics of VCCs

Wrap up

MSR for Vulnerability Prediction – Vulnerability-contributing Commits

Key Characteristics of VCCs

VCCs vs non-VCCs
A case study on Apache HTTP Server with 68 post-release vulnerabilities and 124 VCCs.

- Large commits might increase the chance of contributing to a vulnerability.
- Changing other developers' code might increase the chance of contributing to a vulnerability.
- Vulnerabilities are more likely to be added when modifying existing files rather than creating new files.

A. Meneely et al., "When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits," 2013 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement, Baltimore, MD, USA, 2013, pp. 65-74, doi: 10.1109/ESEM.2013.19.

Definition & Characteristics of VCCs

MSR for Vulnerability Prediction – Mining VCCs

Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.

```
graph LR
    A[Post-release vulnerability] --> B[Fixing commit(s)]
    B --> C[Manual analysis]
    B --> D[Ad hoc detection script]
    C --> E[VCC?]
    C --> F[VCC]
    D --> G[Updated script]
    G --> H[git bisect]
    G --> I[Assisted binary search]
    H --> J[VCC?]
    I --> K[VCC]
    E -- "mmh, not convinced" --> G
    J -- "LGTM" --> F
```

Meneely et al. technique (git bisect)

Wrap up

MSR for Vulnerability Prediction – Vulnerability-contributing Commits

Key Characteristics of VCCs

VCCs vs non-VCCs
A case study on Apache HTTP Server with 68 post-release vulnerabilities and 124 VCCs.

- Large commits might increase the chance of contributing to a vulnerability.
- Changing other developers' code might increase the chance of contributing to a vulnerability.
- Vulnerabilities are more likely to be added when modifying existing files rather than creating new files.

A. Meneely et al., "When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits," 2013 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement, Baltimore, MD, USA, 2013, pp. 65-74, doi: 10.1109/ESEM.2013.19.

Definition & Characteristics of VCCs

MSR for Vulnerability Prediction – Mining VCCs

Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.

```

graph LR
    A[Post-release vulnerability] --> B[Fixing commit(s)]
    B --> C[Manual analysis]
    B --> D[Ad hoc detection script]
    C --> E[Updated script]
    D --> F[Vulnerable code regions]
    E --> G[VCC?]
    F --> G
    G -- "VCC" --> H[VCC]
    G -- "git bisect" --> I[Assisted binary search]
    I --> J[Vulnerable code regions]
    J --> K[VCC?]
    K -- "VCC?" --> L[VCC]
  
```

Meneely et al. technique (git bisect)

MSR for Vulnerability Prediction – Mining VCCs

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

SZZ by Kim et al.

```

graph TD
    A[Project Bug Tracker] --> B[Bug Report]
    B --> C[git blame]
    C --> D[Annotated file(s)]
    E[Project History] --> F[Commit-issue Linkage]
    F --> G[Fix Commit]
    F --> H[git diff]
    H --> I[Candidate BICs]
    H --> J[Polished Changed Lines]
    J --> K[BICs]
    L((Last buggy revision))
  
```

SZZ algorithm and variants (git blame)

Wrap up

MSR for Vulnerability Prediction – Vulnerability-contributing Commits

Key Characteristics of VCCs

VCCs vs non-VCCs
A case study on Apache HTTP Server with 68 post-release vulnerabilities and 124 VCCs.

- Large commits might increase the chance of contributing to a vulnerability.
- Changing other developers' code might increase the chance of contributing to a vulnerability.
- Vulnerabilities are more likely to be added when modifying existing files rather than creating new files.

A. Meneely et al., "When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits," 2013 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement, Baltimore, MD, USA, 2013, pp. 65-74, doi: 10.1109/ESEM.2013.19.

Definition & Characteristics of VCCs

MSR for Vulnerability Prediction – Mining VCCs

Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.

Meneely et al. technique (git bisect)

MSR for Vulnerability Prediction – Mining VCCs

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

SZZ by Kim et al.

SZZ algorithm and variants (git blame)

MSR for Vulnerability Prediction – Validating VCC Mining Algorithms

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the "standard" for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:

- Exhaustive Labeling
- Bisect-driven Labeling
- Precision Assessment
- Developer-informed Oracle

For each vulnerability, we process the fixing commit message to retrieve mentions of the culprit commit(s). Developers sometimes explicitly indicate the commit where the vulnerability was introduced. This method has a fully automated part based on NLP/text mining and an (optional) manual assessment part.

Recommended when... we don't need a complete correct set and, we want developers' experience.

G. Rosa, L. Pascarella, S. Scalabrino, R. Tufano, G. Bayota, M. Lanza, and R. Oliveto. 2021. Evaluating SZZ Implementations Through a Developer-informed Oracle. In Proceedings of the 43rd International Conference on Software Engineering (ICSE '21). IEEE Press, 436–447. https://doi.org/10.1109/ICSE43902.2021.00049

Performance Metrics & Ground Truth

Wrap up

MSR for Vulnerability Prediction – Vulnerability-contributing Commits

Key Characteristics of VCCs

VCCs vs non-VCCs
A case study on Apache HTTP Server with 68 post-release vulnerabilities and 124 VCCs.

- Large commits might increase the chance of contributing to a vulnerability.
- Changing other developers' code might increase the chance of contributing to a vulnerability.
- Vulnerabilities are more likely to be added when modifying existing files rather than creating new files.

A. Meneely et al., "When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits," 2013 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement, Baltimore, MD, USA, 2013, pp. 65-74, doi: 10.1109/ESEM.2013.19.

Definition & Characteristics of VCCs

MSR for Vulnerability Prediction – Mining VCCs

Mining VCCs: A First Approach

Now let's see how we can retrieve VCCs from project histories.

Unnamed Technique by Meneely et al.

```

graph LR
    A[Post-release vulnerability] --> B[Fixing commit(s)]
    B --> C[Manual analysis]
    B --> D[Ad hoc detection script]
    C --> E[VCC?]
    C --> F[git bisect]
    D --> G[Vulnerable code regions (Hunks)]
    E -- "LGTM" --> H[VCC]
    F -- "mmh, not convinced" --> E
    F --> I[Assisted binary search]
    I --> E
  
```

Meneely et al. technique (git bisect)

MSR for Vulnerability Prediction – Mining VCCs

Mining VCCs: Borrowing from the Bug World

The SZZ algorithm is quite intuitive, but, despite its simplicity, it has been a revolution in the MSR world. Yet, all that glitters is not gold: it has some problems.

SZZ by Kim et al.

```

graph TD
    A[Project Bug Tracker] --> B[Bug Report]
    B --> C[Last buggy revision]
    C --> D[git blame]
    D --> E[Annotated file(s)]
    F[Project History] --> G[Commit-issue Linkage]
    G --> H[Fix Commit]
    H --> I[git diff]
    I --> J[Polished Changed Lines]
    J --> K[Candidate BICs]
    K --> L[BICs]
    E --> M[Candidate BICs]
  
```

SZZ algorithm and variants (git blame)

MSR for Vulnerability Prediction – Validating VCC Mining Algorithms

Building the Ground Truth

We need to build a **ground truth** (a.k.a. *golden set*) that is the "standard" for evaluating the algorithms. In other words, a dataset of true VCCs and non-VCCs. We can employ some methods:

- Exhaustive Labeling
- Bisect-driven Labeling
- Precision Assessment
- Developer-informed Oracle

For each vulnerability, we process the fixing commit message to retrieve mentions of the culprit commit(s). Developers sometimes explicitly indicate the commit where the vulnerability was introduced. This method has a fully automated part based on NLP/text mining and an (optional) manual assessment part.

Recommended when... we don't need a complete correct set and, we want developers' experience.

G. Rosa, L. Pascarella, S. Scalabrino, R. Tufano, G. Bayota, M. Lanza, and R. Oliveto. 2021. Evaluating SZZ Implementations Through a Developer-informed Oracle. In Proceedings of the 43rd International Conference on Software Engineering (ICSE '21). IEEE Press, 436–447. https://doi.org/10.1109/ICSE43902.2021.00049

Performance Metrics & Ground Truth

MSR for Vulnerability Prediction – Tools and Datasets for VCC Mining

Available Datasets

Curated	Mined
Vulnerability History Project Database of curated histories of 2,677 vulnerabilities of eight open-source projects. Built by class assignments in a Master's degree course held at RIT.	Java VCC Dataset Dataset of 100 VCCs of 71 known vulnerabilities affecting popular Java projects. Built by manually analyzing the history aided by blames on fixing commits.
Secret Life Dataset Dataset of 12,256 VCCs of 3,663 vulnerabilities affecting 1,096 open-source projects. Built by running an SZZ variant by Iannone et al.	FrontEndART Dataset Dataset of ~700 VCCs of 564 vulnerabilities affecting 198 Java projects. Built by running an SZZ variant by Aladics et al.

Available Tools & Datasets

(Some) Open Challenges

Non-code-related Vulnerabilities

Not all vulnerabilities are caused by coding mistakes. Some of them are caused by improper configurations or, even worse, design issues.

(Some) Open Challenges

Non-code-related Vulnerabilities

Not all vulnerabilities are caused by coding mistakes. Some of them are caused by improper configurations or, even worse, design issues.

Tangled Changes

Not all fixing commits are focused on fixing the vulnerability: other collateral activities may be done.

(Some) Open Challenges

Non-code-related Vulnerabilities

Not all vulnerabilities are caused by coding mistakes. Some of them are caused by improper configurations or, even worse, design issues.

Tangled Changes

Not all fixing commits are focused on fixing the vulnerability: other collateral activities may be done.

Irrelevant Changes

Not all lines changed are directly related to the vulnerability, e.g., addition/removal of import statements, parameters reordering, etc.

(Some) Open Challenges

Non-code-related Vulnerabilities

Not all vulnerabilities are caused by coding mistakes. Some of them are caused by improper configurations or, even worse, design issues.

Tangled Changes

Not all fixing commits are focused on fixing the vulnerability: other collateral activities may be done.

Irrelevant Changes

Not all lines changed are directly related to the vulnerability, e.g., addition/removal of import statements, parameters reordering, etc.

Migrated Repositories

Many “old” projects were migrated from another VCS (e.g., svn to git), so their history might be incomplete (e.g., the initial commit is enormous).



References

Articles (1/2)

[Meneely et al.] **When a Patch Goes Bad: Exploring the Properties of Vulnerability-Contributing Commits:** <https://ieeexplore.ieee.org/document/6681339>

[Śliwerski et al.] **When do changes induce fixes?:** <https://dl.acm.org/doi/10.1145/1082983.1083147>

[Kim et al.] **Automatic Identification of Bug-Introducing Changes:** <https://ieeexplore.ieee.org/document/4019564>

[da Costa et al.] **A Framework for Evaluating the Results of the SZZ Approach for Identifying Bug-Introducing Changes:** <https://ieeexplore.ieee.org/document/7588121>

[Rodríguez-Pérez et al.] **How bugs are born: a model to identify how bugs are introduced in software components:** <https://link.springer.com/article/10.1007/s10664-019-09781-y>

[Rosa et al.] **Evaluating SZZ Implementations Through a Developer-informed Oracle:** <https://dl.acm.org/doi/10.1109/ICSE43902.2021.00049>

[Camilo et al.] **Do Bugs Foreshadow Vulnerabilities? A Study of the Chromium Project:** <https://ieeexplore.ieee.org/document/7180086>

References

Articles (2/2)

[Canfora et al.] **Investigating the vulnerability fixing process in OSS projects: Peculiarities and challenges:** <https://www.sciencedirect.com/science/article/abs/pii/S0167404820303400>

[Canfora et al.] **Patchworking: Exploring the code changes induced by vulnerability fixing activities:** <https://www.sciencedirect.com/science/article/abs/pii/S0950584921001932>

[Perl et al.] **VCCFinder: Finding Potential Vulnerabilities in Open-Source Projects to Assist Code Audits:** <https://dl.acm.org/doi/10.1145/2810103.2813604>

[Yang et al.] **VulDigger: A Just-in-Time and Cost-Aware Tool for Digging Vulnerability-Contributing Changes:** <https://ieeexplore.ieee.org/document/8254428>

[Iannone et al.] **The Secret Life of Software Vulnerabilities: A Large-Scale Empirical Study:** <https://ieeexplore.ieee.org/document/9672730>

[Bao et al.] **V-SZZ: Automatic Identification of Version Ranges Affected by CVE Vulnerabilities:** <https://ieeexplore.ieee.org/document/9794006>



MSR for Vulnerability Prediction

Mining Vulnerability-Contributing Commits

Emanuele Iannone

SeSa Lab @ University of Salerno, Italy

eiannone@unisa.it